Toward a Methodology for AI Architecture Evaluation: Comparing Soar and CLIPS

Scott Wallace & John Laird
May, 1999

Architectures, Tasks and Capabilities

- What capabilities are required by the task?
- What capabilities are supported by the architecture?

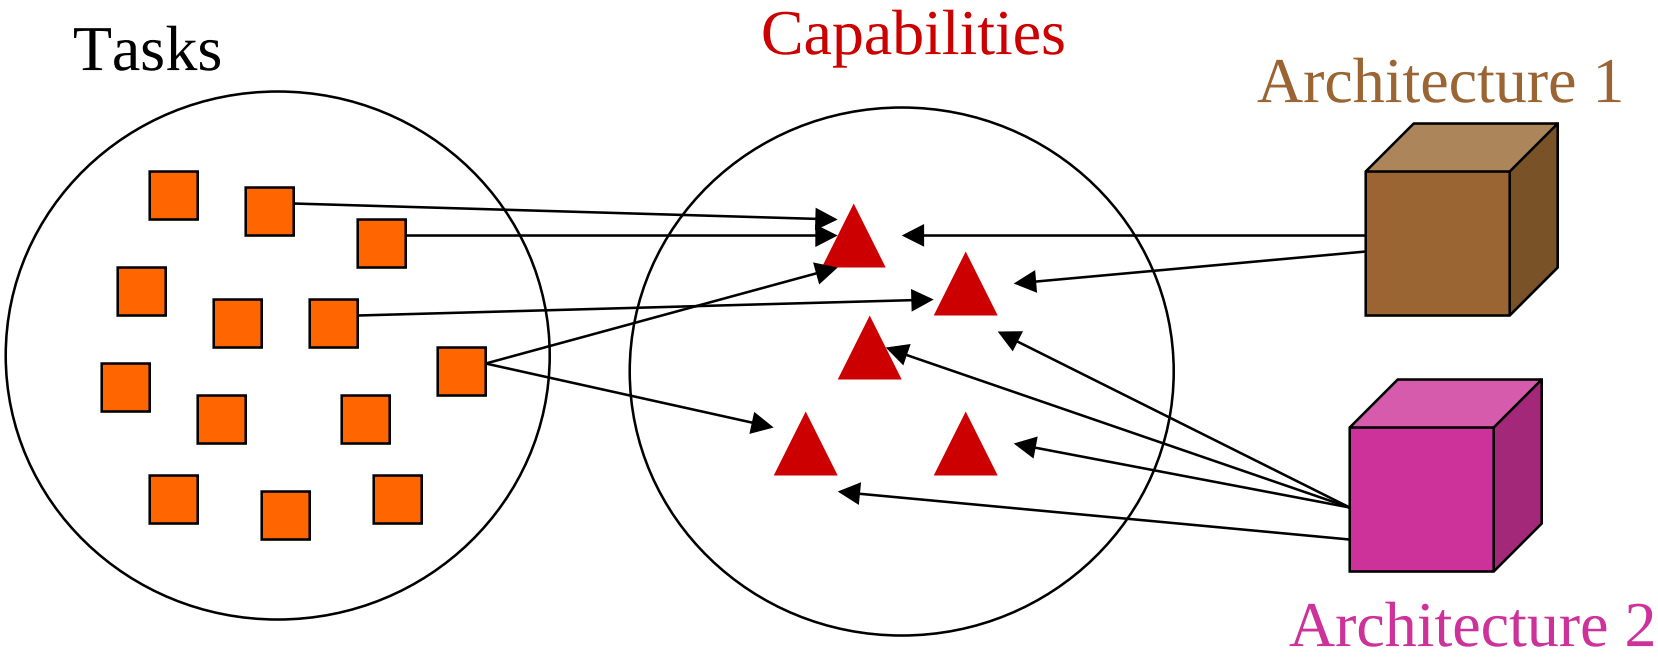

Result

- Performance profiles of architectures and tasks

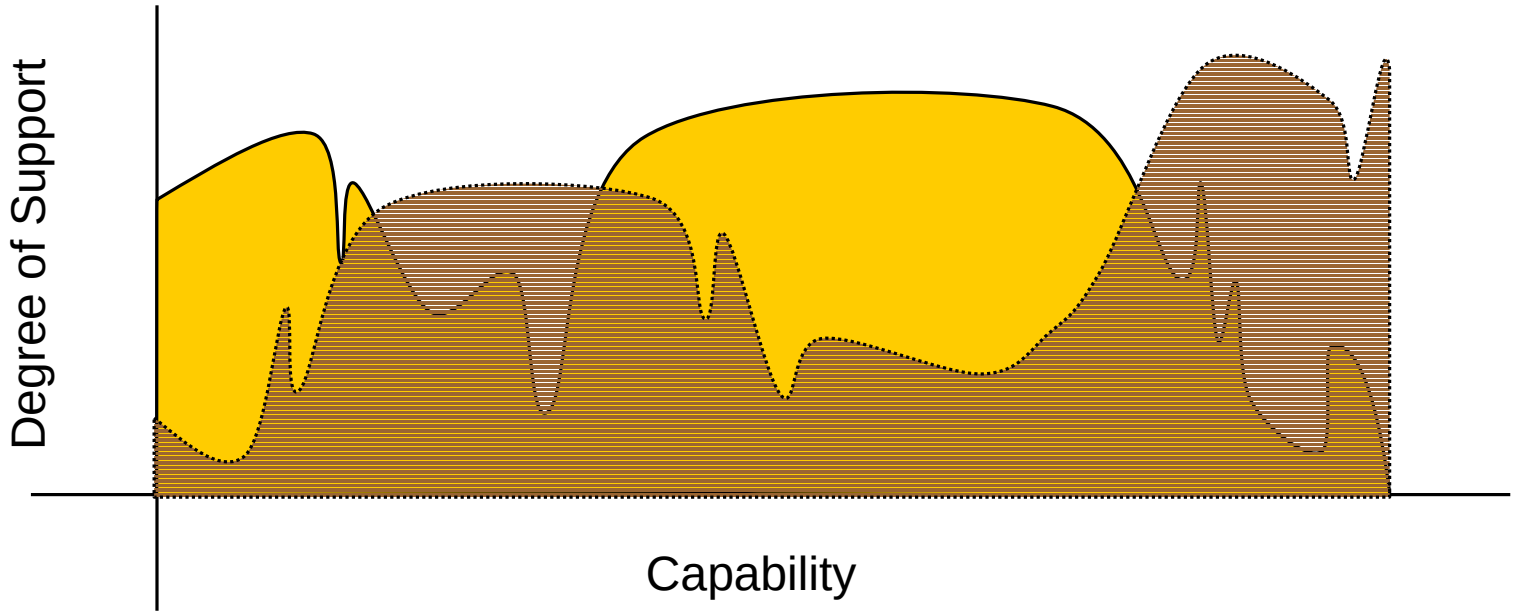

Methodology

- Degree of support:
  - Runtime performance
  - Number of rules
- Create comparable agents to evaluate a range of capabilities
- Maintain knowledge consistency
  - Detailed specification of agent's internal and external behavior
- Ensure agents are exposed to same stimuli
- Examine multiple architectural mechanisms
  - Create an array of agents which exhibit identical external behavior, but differ internally

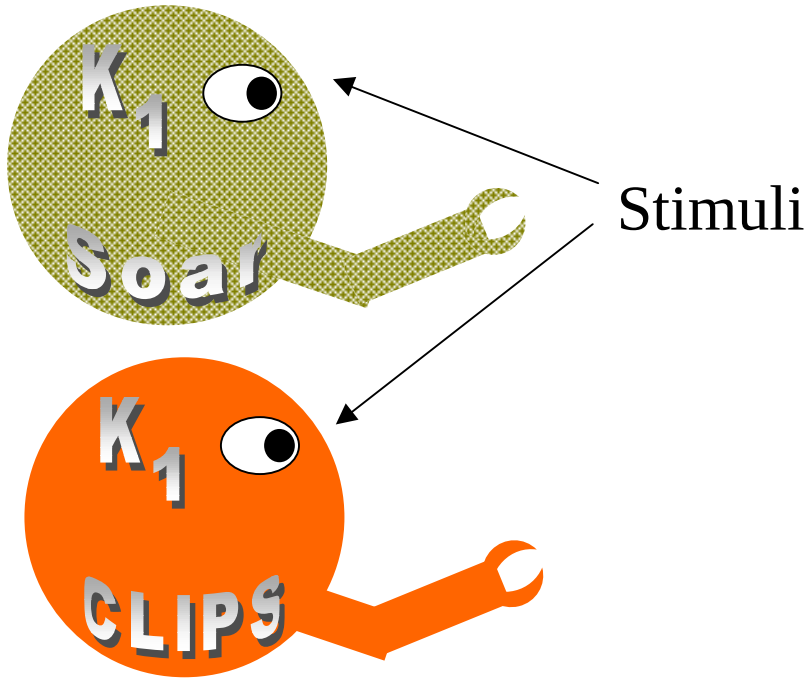

Examining Decision Making Capabilities

- Mutually Exclusive Actions
- Segregation of Control Knowledge
- Two Phase Decision Process
- Three Phase Decision Process
- Goal Constrained Selection
- Utility Based Decision Process
- Analogy Based Decision Process
- ...

Mutex Actions

- Actions are only considered when they when should be invoked

Proposal & Selection

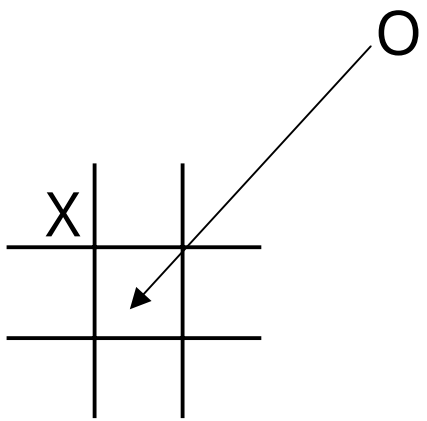

Preconditions:
X in upper left
all others
empty

Effects:
Put O in center

Segregated Control

- Actions may be considered even when they are not invoked, but an architectural mechanism ensures the proper action is always taken

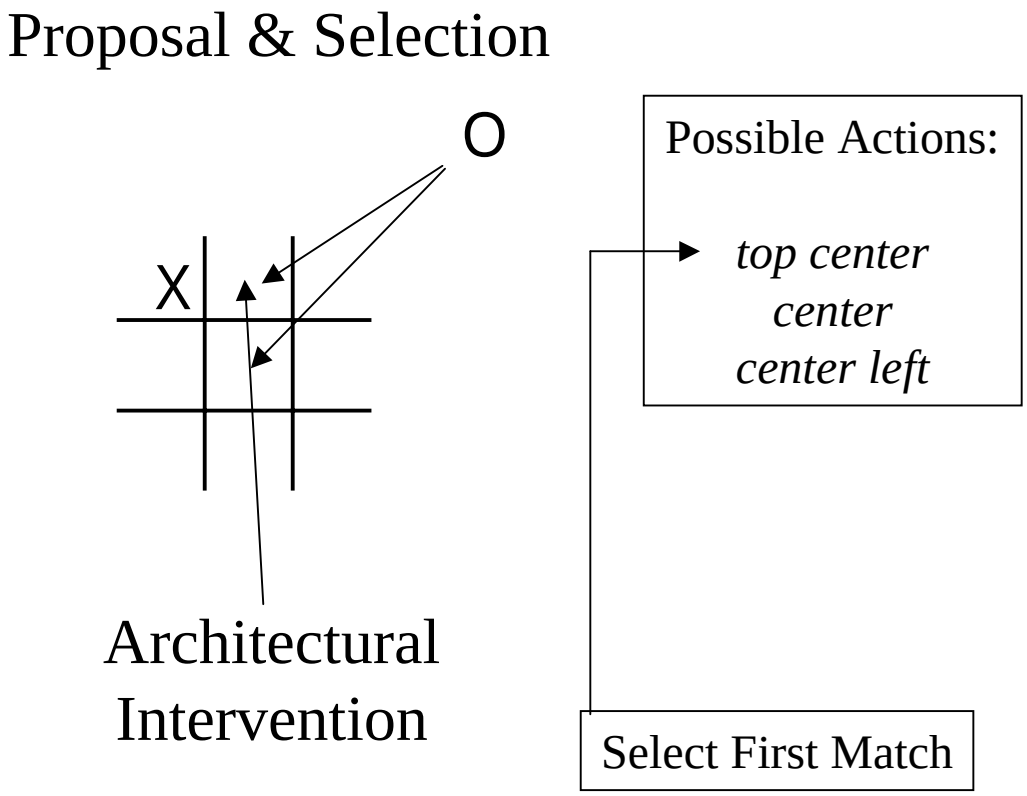

Two Phase Process

- Proposal and Selection take place asynchronously.
- Proposal involves creating symbols represent each action
- Selection pursues one action

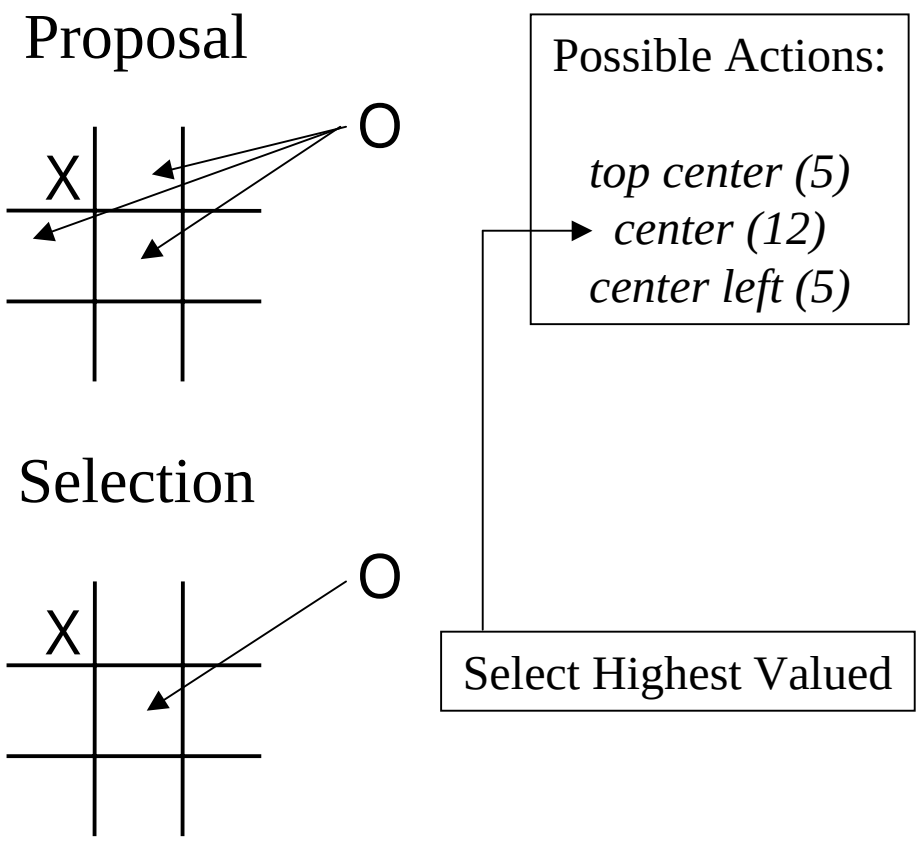

Three Phase Process

- Three distinct phases occur
  - Proposal
  - Preference
  - Selection
- Preferences can be stated in terms of other actions under consideration

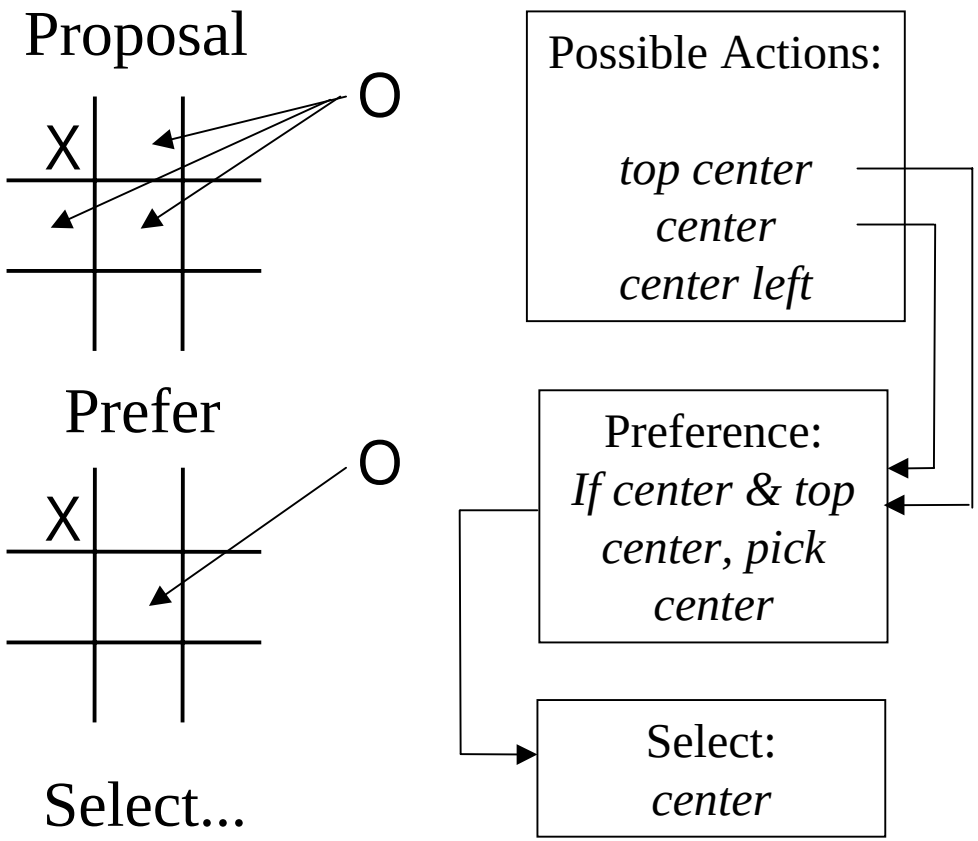

Goal Constrained

- Consider actions even if they cannot be invoked in the current situation
- Currently proposed actions as well as selected goals can be used for preferences
- Build goal list of selected actions

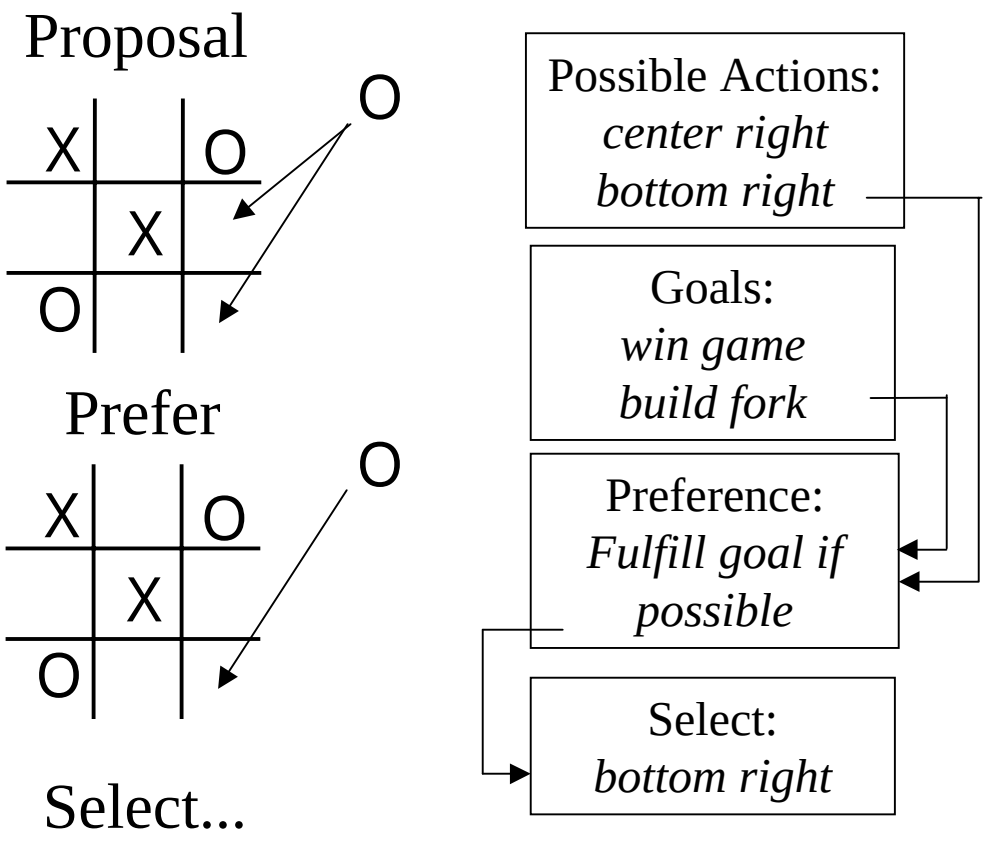

Architectural Properties

CLIPS

- Short term memory -- a list of facts
- Rules are fired serially, matches recalculated after each firing
- Flow control:
  - Rule matching
  - Salience
  - Search strategy

Soar

- Short term memory -- directed graph
- Rules are fired in parallel, matches recalculated after each elaboration cycle
- Flow control:
  - operators
  - preferences
  - 3 phase decision cycle

Runtime Performance in Towers of Hanoi

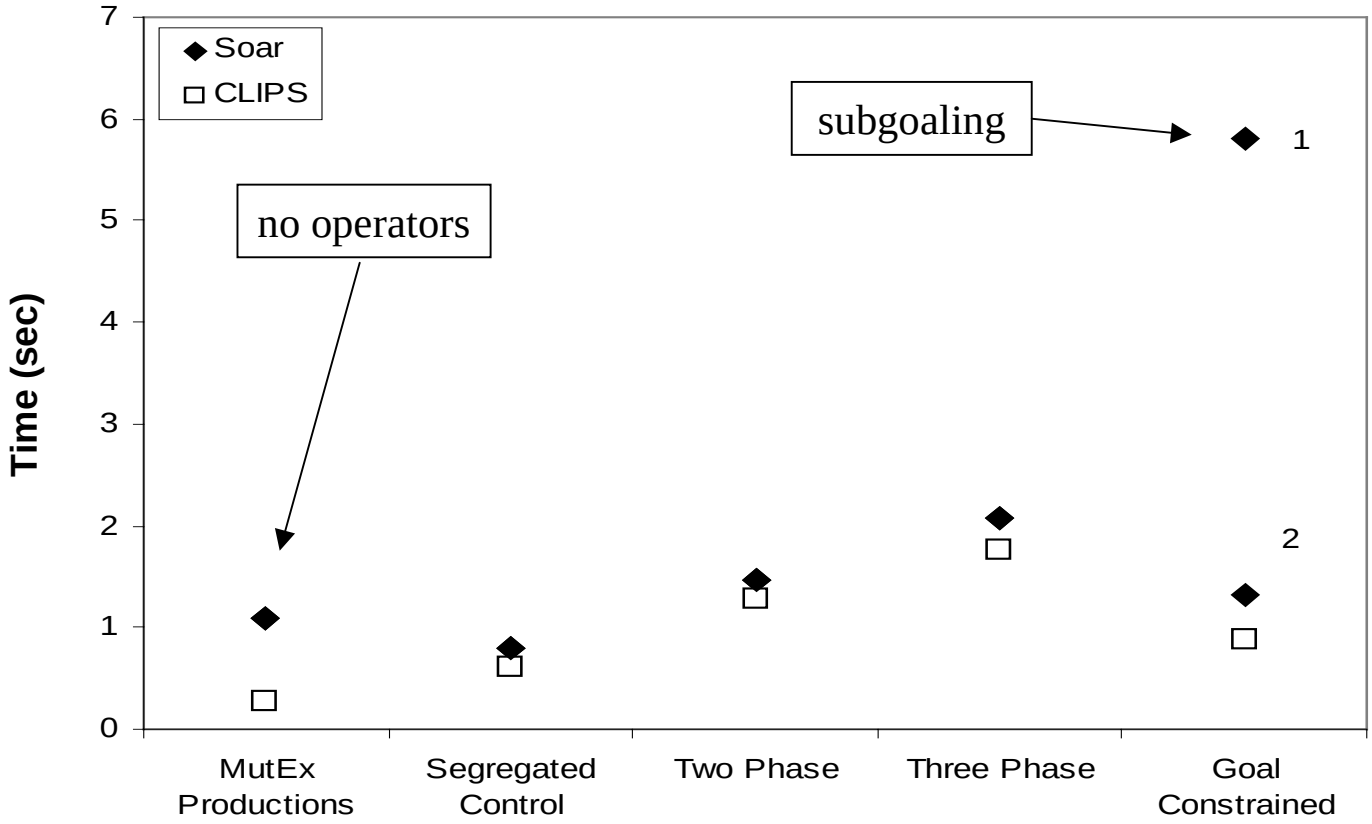

Number of Rules in Towers of Hanoi

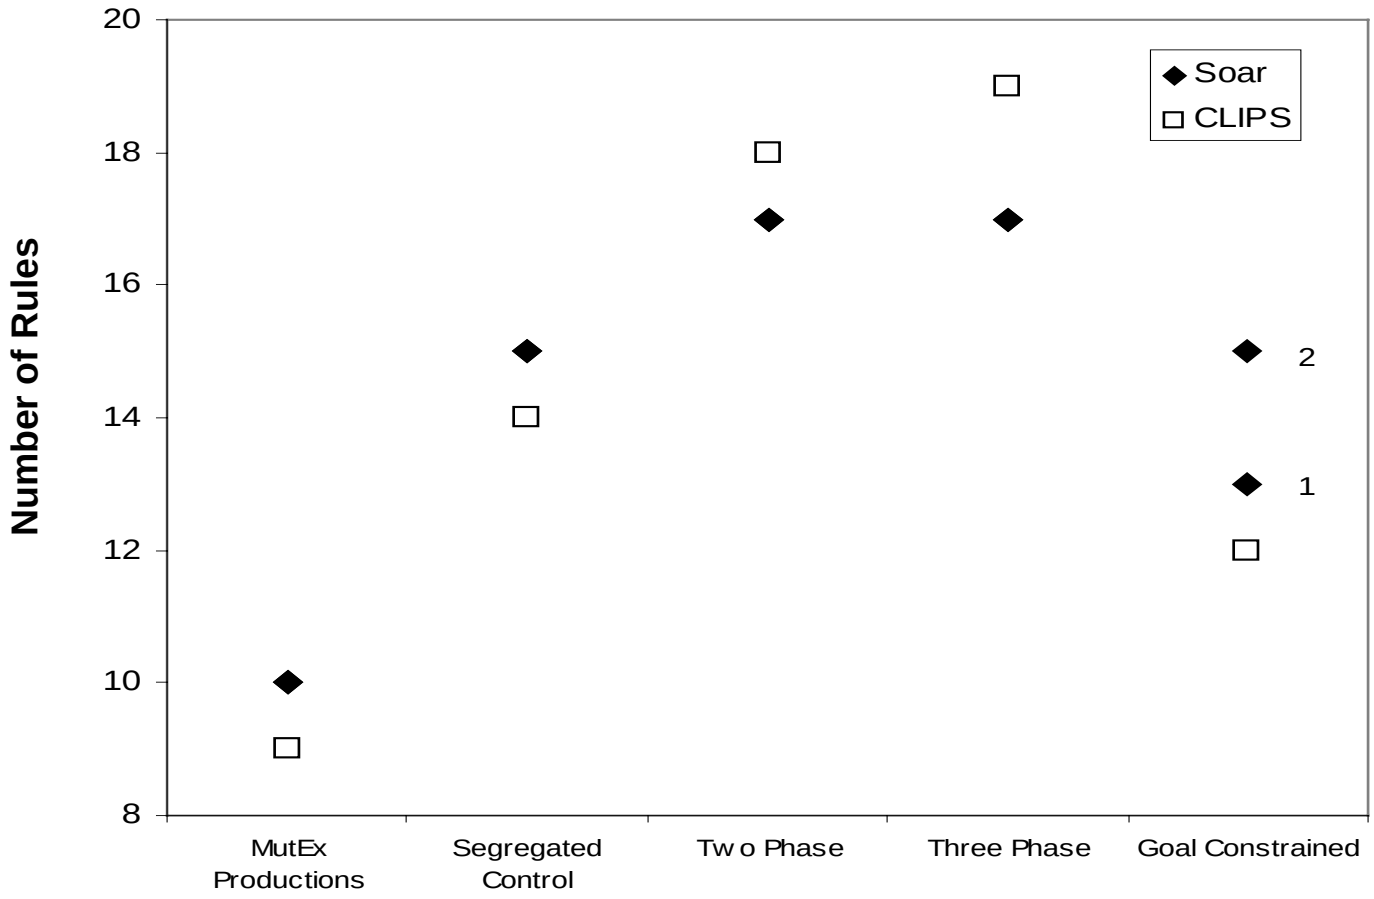

Scaling in ToH

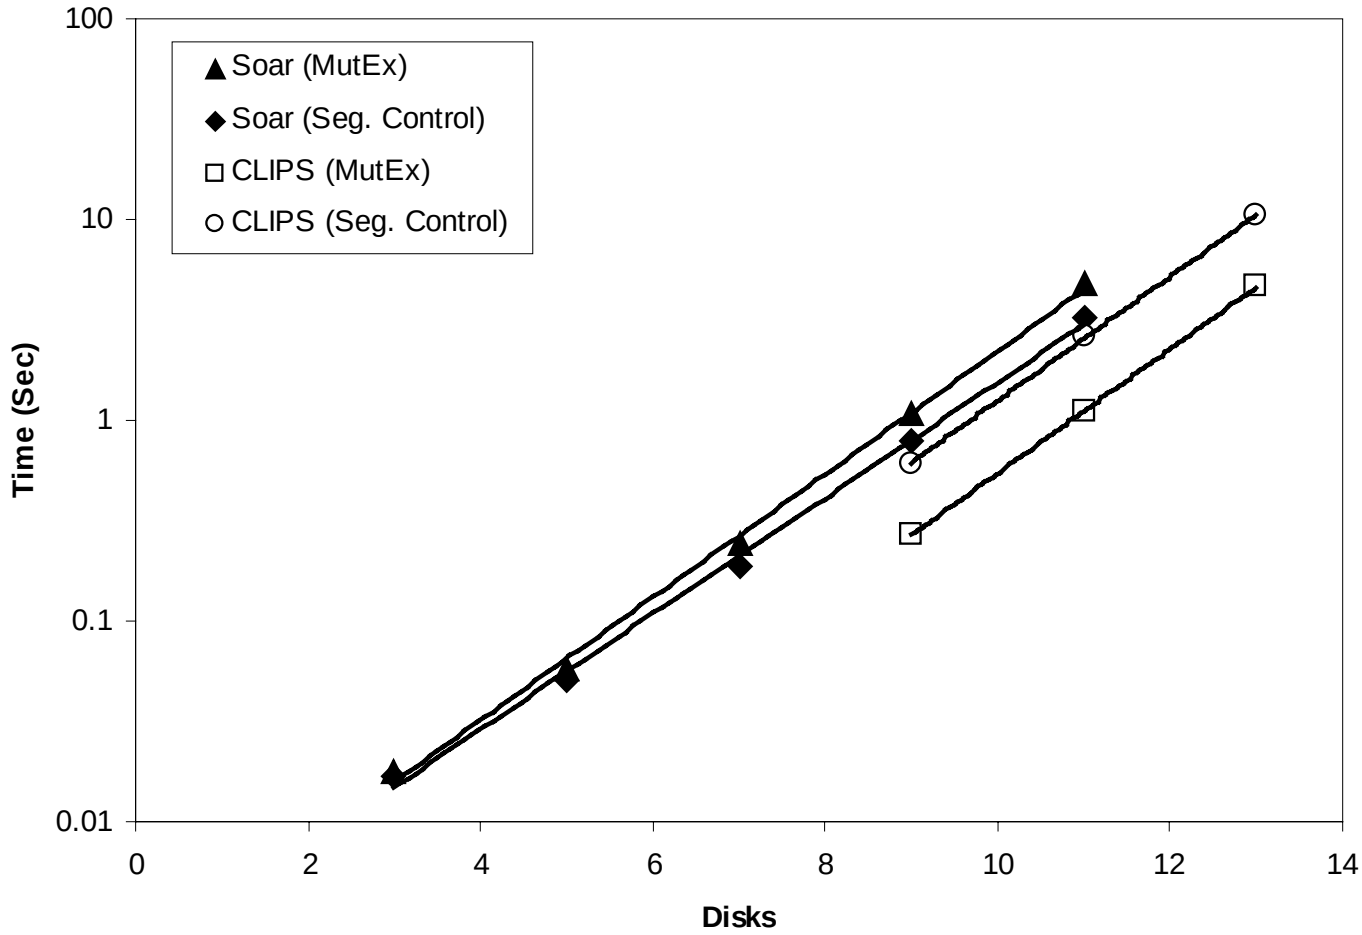

Eaters

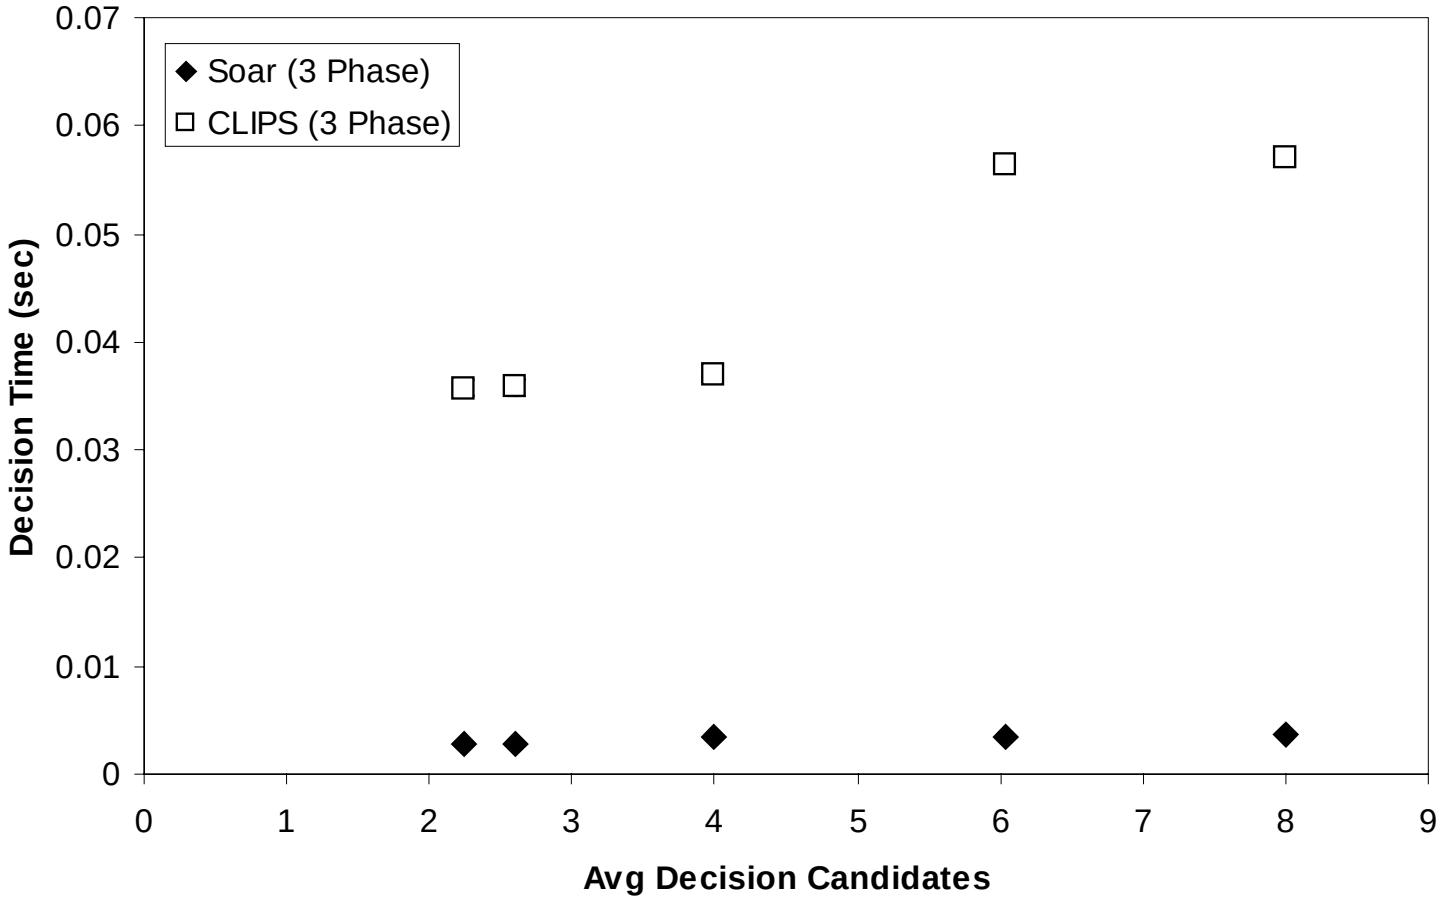
Soar (3 Phase)
CLIPS (3 Phase)
Decision Time (sec)
0.07
0.06
0.05
0.04
0.03
0.02
0.01
0
0
1
2
3
4
5
6
7
8
9
Avg Decision Candidates

Nuggets & Coal

- Nuggets
  - Further evidence that Soar performs well compared to other architectures
  - Yielded leads for further performance improvments
- Coal
  - Difficult to prepare architectures for benchmarking
  - Hard to determine how results gained from this method apply to architectural theories
  - 5 capabilities is just a start

Next Steps & Future Work

- Why does Soar's subgoaling mechanism seem so slow?
- What specific architectural attribute does the nonlinearity in CLIPS's performance trace back to?
- Why does CLIPS perform so much slower than Soar in Eaters, but not in ToH?
- What other tasks can corroborate our hypothesis?
- Increase breadth of tasks
- Examine architectures which have less common elements