

# Benchmarking Soar: Using Soar's Timers

Robert Wray  
wrayre@acm.org

22 May 1999



**Middle Tennessee State University**

## Why Should I Use Soar's Timers?

**Timer:** stop watch

measures time elapsed (usually execution of a procedure)

- Process time (e.g., **ps**) provides only a coarse measure of cost
  - Distinguish between “reasoning” time and other functions (loading, callbacks, input, output, etc.)

**kernel time:** measure of Soar reasoning time only

- Identify costly/inefficient reasoning
  - Which phase requires the most time?
  - How much time is spent processing input?
  - What is the time cost of match, chunking, etc?
- If you don't need to know the time costs of Soar's computations and your application does not demand “maximum” efficiency, you can probably ignore timers.

## How Do I Use the Timers?

Two flags determine if timing data is collected and the level of detail of the statistics.

`soarkernel.h` definitions (`#define`)

- No definitions

Soar will collect some basic timing data (version-dependent)

Default in 7.0.x distributions

- `#define NO_TIMING_STUFF`

Soar will not keep any timing statistics

- `#define DETAILED_TIMING_STATS`

Soar will collect “detailed” (fine-grained) timing statistics

Default in 8.2 distribution

Re-build Soar after defining the level of timing desired

`DETAILED_TIMING_STATS` has no effect if `NO_TIMING_STUFF` is defined.

## stats with NO\_TIMING\_STUFF enabled

Soar 8.2 on newell.mtsu.edu at Thu May 13 09:48:33 1999

91 productions (79 default, 12 user, 0 chunks)

+ 0 justifications

39 decision cycles

170 elaboration cycles

833 production firings

2403 wme changes (1267 additions, 1136 removals)

WM size: 131 current, 174.488 mean, 242 maximum

## What Will the Timers Tell Me?

- Soar6.2
  - Total execution (“run”) time (not process time)
  - DETAILED\_TIMING\_STATS
    - chunking
    - match
    - ownership
- Soar7.0.3/7.0.4
  - Kernel time (no callbacks, input, output)
  - Kernel time reported for individual phases
  - DETAILED\_TIMING\_STATS are also reported by phase
- Soar8.2
  - Soar7.0.3/7.0.4
  - DETAILED\_TIMING\_STATS includes a category for Soar8-specific processing

## Timing Statistics: Soar6.2

Soar 6.2.5 on parrot.eecs.umich.edu at Sat May 15 15:31:29 1999

127 productions (100 default, 24 user, 3 chunks)

+ 0 justifications

Total cpu time: 0.169 seconds

27 decision cycles (6.244 msec/dc)

92 elaboration cycles (3.407 ec's per dc, 1.832 msec/ec)

438 production firings (4.761 pf's per ec, 0.385 msec/pf)

1100 wme changes (577 additions, 523 removals)

WM size: 54 current, 205.741 mean, 308 maximum

## What Will the Timers Tell Me? (Soar7.0.3)

Soar7.0.3 reports **kernel time** as well as total CPU time

- Kernel time includes time for match, WME changes, chunking, ownership calculations, etc.
- Kernel time includes only time spent “run”-ning (not loading, parsing, etc.)
- Kernel time excludes:
  - Input function (e.g., tcl simulator)
  - Output function
  - Callbacks (monitors)

Kernel time is used to report frequencies

```
Kernel CPU Time:      0.290 sec.  
Total CPU Time:      5.870 sec.
```

```
42 decision cycles (6.905 msec/dc)  
137 elaboration cycles (3.262 ec's per dc, 2.117 msec/ec)
```

## What Will the Timers Tell Me? (Soar7.0.3)

Soar7.0.3 reports time per phase as well as total kernel time and total CPU time

- Individual timers are used for each phase

Derived totals

- Sum of the time spent in individual phases
- Example: Sum will always be  $\leq$  Kernel CPU Time
  - Differences due to implementation of timers....
  - Kernel CPU Time used to report frequencies (eg, msec/dc)

Phases:	Input	Pref	W/M	Output	Decision	Derived Totals
Kernel:	0.000	1.040	0.130	0.000	0.040	1.210

Kernel CPU Time: 1.220 sec.



# What Will the Timers Tell Me? (Soar7.0.3)

## DETAILED\_TIMING\_STATS

### Match

Time for adding and deleting WMEs from the RETE

### Ownership

Time for promoting and demoting  
(determine goal to which an identifier is linked)

### Chunking

Time for creating chunks (and justifications)

### Other

Derived total:

$$\text{Other}_{phase} = \text{KernelTime}_{phase} - (\text{Match} + \text{Ownership} + \text{Chunking})_{phase}$$

# stats with DETAILED\_TIMING\_STATS enabled: 7.0.3

Soar 7.0.3. TCL TK on stork.eecs.umich.edu at Sat May 15 12:11:57 1999

192 productions (4 default, 188 user, 0 chunks)  
+ 2 justifications

Phases:	Input	Pref	W/M	Output	Decision	Derived	Totals
Kernel:	0.080	0.130	0.050	0.000	0.010		0.270
===== Detailed Timing Statistics =====							
Match:	0.030	0.000	0.050	0.000	0.000		0.080
Own'ship:	0.030	0.000	0.000	0.000	0.000		0.030
Chunking:	0.000	0.110	0.000	0.000	0.000		0.110
Other:	0.020	0.020	0.000	0.000	0.010		0.050
=====							
Input fn:	5.000						5.000
=====							
Outpt fn:				0.000			0.000
=====							
Callbcks:	0.570	0.000	0.000	0.000	0.000		0.570
=====							
Derived-----							
Totals:	5.650	0.130	0.050	0.000	0.010		5.840

Values from single timers:

Kernel CPU Time: 0.290 sec.  
Total CPU Time: 5.870 sec.

42 decision cycles (6.905 msec/dc)  
137 elaboration cycles (3.262 ec's per dc, 2.117 msec/ec)  
373 production firings (2.723 pf's per ec, 0.777 msec/pf)  
2335 wme changes (1494 additions, 841 removals)  
match time: 0.034 msec/wm change  
WM size: 653 current, 618.804 mean, 704 maximum

# Summary: What Will the Timers Tell Me?

Soar 7.0.3. TCL TK on stork.eecs.umich.edu at Sat May 15 12:11:57 1999

192 productions (4 default, 188 user, 0 chunks)  
+ 2 justifications

Phases:	Input	Pref	W/M	Output	Decision	Derived Totals
Kernel:	0.080	0.130	0.050	0.000	0.010	0.270
===== Detailed Timing Statistics =====						
Match:	0.030	0.000	0.050	0.000	0.000	0.080
Own'ship:	0.030	0.000	0.000	0.000	0.000	0.030
Chunking:	0.000	0.110	0.000	0.000	0.000	0.110
Other:	0.020	0.020	0.000	0.000	0.010	0.050
Input fn:	5.000					5.000
Outpt fn:	0.000					0.000
Callbcks:	0.570	0.000	0.000	0.000	0.000	0.570
Derived Totals:	5.650	0.130	0.050	0.000	0.010	5.840

Derived Kernel Time

Derived from  $\Sigma$  of Phase Timers

Derived Total CPU Time

#define DETAILED TIMING\_STATS

Time spent in input function

Time spent in output function

Time spent in callbacks

Values from single timers:

Kernel CPU Time: 0.290 sec.  
Total CPU Time: 5.870 sec.

Values from individual timers

42 decision cycles (6.905 msec/dc)  
137 elaboration cycles (3.262 ec's per dc, 2.117 msec/ec)  
373 production firings (2.723 pf's per ec, 0.777 msec/pf)  
2335 wme changes (1494 additions, 841 removals)  
match time: 0.034 msec/wm change  
WM size: 653 current, 618.804 mean, 704 maximum

Frequencies are calculated using Kernel CPU Time

## What's New in Soar8.2?

New Phase: Determine Level Phase (DLP)

- Kernel phase timers include this phase
- **stats** reflects additional phase

New category in **DETAILED\_TIMING\_STATS**: “Operand2”

- Includes:

Goal Dependency Set calculations

Determination of the highest active level (DLP)

- Time measured for **Operand2** is separate from other statistics (**Operand2** time is included in **Other**)

+ Category name should be re-labeled

? Separate GDS and determine level timers?

# Soar8.2 stats with NO\_TIMING\_STUFF disabled

Soar 8.2 on newell.mtsu.edu at Thu May 13 09:44:23 1999

91 productions (79 default, 12 user, 0 chunks)  
+ 0 justifications

Phases:	Input	DLP	Pref	W/M	Output	Decision	Derived Totals
Kernel:	0.000	0.000	1.040	0.130	0.000	0.040	1.210
Input fn:	0.000						0.000
Outpt fn:					0.000		0.000
Callbcks:	0.000	0.000	0.000	0.000	0.000	0.010	0.010
Derived	-----+-----						
Totals:	0.000	0.000	1.040	0.130	0.000	0.050	1.220

Values from single timers:

Kernel CPU Time: 1.220 sec.

Total CPU Time: 1.230 sec.

39 decision cycles (31.282 msec/dc)

170 elaboration cycles (4.359 ec's per dc, 7.176 msec/ec)

45 p-elaboration cycles (1.154 pe's per dc, 27.111 msec/pe)

833 production firings (4.900 pf's per ec, 1.465 msec/pf)

2403 wme changes (1267 additions, 1136 removals)

WM size: 131 current, 174.488 mean, 241 maximum

# Soar8.2 stats with DETAILED\_TIMING\_STATS enabled

Soar 8.2 on newell.mtsu.edu at Thu May 13 09:52:37 1999

91 productions (79 default, 12 user, 0 chunks)  
+ 0 justifications

Phases:	Input	DLP	Pref	W/M	Output	Decision	Derived Totals
Kernel:	0.000	0.020	1.520	0.180	0.000	0.130	1.850
===== Detailed Timing Statistics =====							
Match:	0.000	0.000	0.000	0.150	0.000	0.000	0.150
Own'ship:	0.000	0.010	0.000	0.010	0.000	0.000	0.020
Chunking:	0.000	0.000	0.010	0.000	0.000	0.000	0.010
Other:	0.000	0.010	1.510	0.020	0.000	0.130	1.670
Operand2:	0.000	0.010	0.000	0.000	0.000	0.000	0.010
=====							
Input fn:	0.000						0.000
=====							
Outpt fn:					0.000		0.000
=====							
Callbcks:	0.000	0.000	0.000	0.000	0.000	0.000	0.000
=====							
Derived-----							
Totals:	0.000	0.020	1.520	0.180	0.000	0.130	1.850

Values from single timers:

Kernel CPU Time: 1.850 sec.  
Total CPU Time: 1.860 sec.

67 decision cycles (27.612 msec/dc)  
289 elaboration cycles (4.313 ec's per dc, 6.401 msec/ec)  
77 p-elaboration cycles (1.149 pe's per dc, 24.026 msec/pe)  
1410 production firings (4.879 pf's per ec, 1.312 msec/pf)  
4021 wme changes (2076 additions, 1945 removals)  
match time: 0.037 msec/wm change  
WM size: 131 current, 174.426 mean, 241 maximum

## How Are the Timers Implemented?

- Soar implementation provides an interface for using timers:

```
/* Initialize/reset timer to 0 */  
void reset_timer (struct timeval *timer_to_reset);  
  
/* Start a timer */  
void start_timer (struct timeval *start_time);  
  
/* Stop a timer. Store the elapsed time since the call to start_time in  
   accumulated_time */  
void stop_timer(struct timeval *start_time,struct timeval *accumulated_time);  
  
/* Return the number of seconds stored in the timer some_timer */  
double timer_value (struct timeval *some_timer);
```

- Implementation of functions is dependent on operating system

Unix: `rusage/getrusage()`

- Easy to add new timers using existing functions
- Users can (mostly) ignore underlying implementation
  - e.g., Users don't need to know definition of `timeval`

**"run"**

```
start total CPU time timer  
start kernel timer
```

**'do phase'**

```
start phase timer
```

**"input/output function or callback"**

```
stop kernel timer  
stop phase timer  
start I/O or callback timer
```

```
stop I/O or callback timer  
start phase timer  
start kernel timer
```

```
stop phase timer
```

```
stop kernel timer  
stop total CPU time timer
```



# Do the Timers Have Any Limitations?

Uncertainty Principle:

Executing code to measure time of execution

- Individual calls to `start_timer` or `stop_timer` are (mostly) insignificant
- Many calls to timer functions over the course of execution

Empirical Results (Scott Wallace):

- \*  $\approx 1 * 10^{-4}$  seconds/timer function call
- \* Thousands of calls to timers for simple tasks (esp. with `DETAILED_TIMING_STATS`)
- \* Factor of 4 change in time data in some experiments

Recommendations:

- Use timers for development, not deployment
- Avoid using `DETAILED_TIMING_STATS` in applications with real-time demands
- ? Need new timing level: `KERNEL_TIME` (no phase timing)

## Summary

- Timers provide information useful for measuring performance and recognizing inefficiencies
- Soar provides an interface to the timers that abstracts from the actual implementation of the timers (underlying O/S)
- Execution of timer procedures can be costly/may degrade performance

### Possible Action Items

- Distribution default: No compiler directives for timing (basic timing information)
- Replace “Operand2” with a better term in Soar8.2 **stats**
- Are there other categories of calculations we should consider for the **DETAILED\_TIMING\_STATS**?
- Should there be another level of benchmarking?  
**NO\_TIMING\_STUFF**, **KERNEL\_TIME**, **PHASE\_TIME**,  
**DETAILED\_TIMING\_STATS**