


Long-Term Learning with TankSoar




Bill Kennedy

George Mason University

21st Soar Workshop


Sunday, May 6, 2001

Long-Term Learning Fundamental Questions




- Will learning go on forever?
- How much of learned knowledge is ever used?

Previous work related to the fundamental questions




- Will learning go on forever?
 - Bob Doorenbos' 100K chunks in 20 problem spaces: mixed
 - Experiments with physics demo: trending toward stopping
 - Experiments with blocks-world: ended by 150 problems
 - Experiments with SCA: 1 chunk per example, of course

Previous work related to the fundamental questions




- How much of learned knowledge is ever used?
 - Bob Doorenbos' 100,000 chunks in 20 problem spaces: unknown
 - Experiments with physics demo: constant “average” use but ...
 - different chunks used to make up the constant average
 - most recently learned used more
 - earlier chunks continue to be used
 - Experiments with blocks-world: all used, but 2-3 were ~75% of use
 - Experiments with SCA: less than about 10% after 300 examples

Experimenting with TankSoar2.5




- Experimental design:
 - Modify “simple” tank to learn at “T” intersections
 - Extend “score limit” to allow to time out at 4,000 decision cycles (dc)
 - Collect stats every 50 dc
 - Run both alone and with a “simple” tank

Experiments with “simple-learner” alone

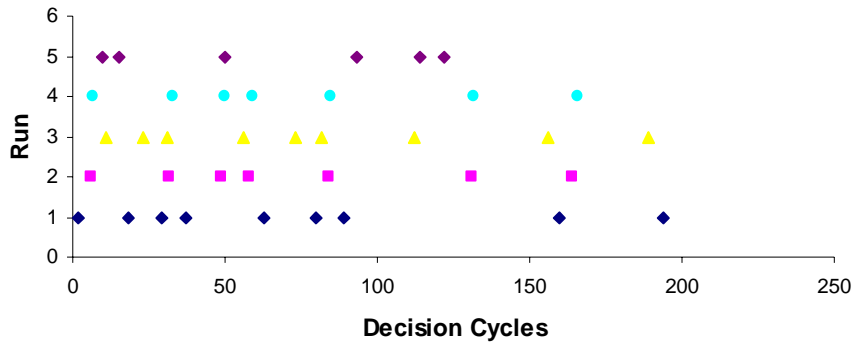


- At least 7 occasions for chunking
- More depending on starting position
- Initial circuit is 140 dc
- After chunking, 132 dc
- Chunking ends & all chunks on circuit are used each cycle

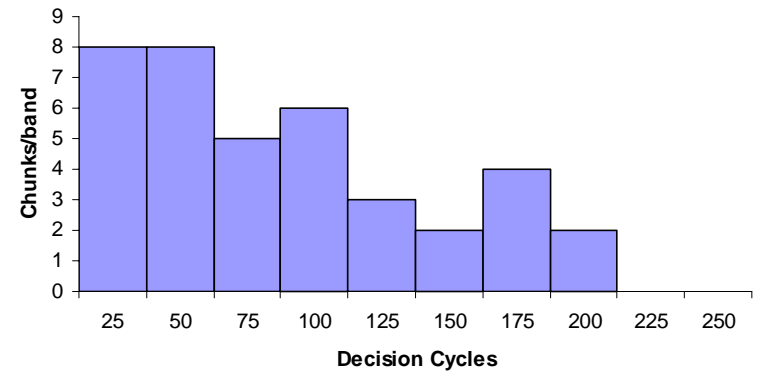
Chunk Learning in “simple-learner” alone



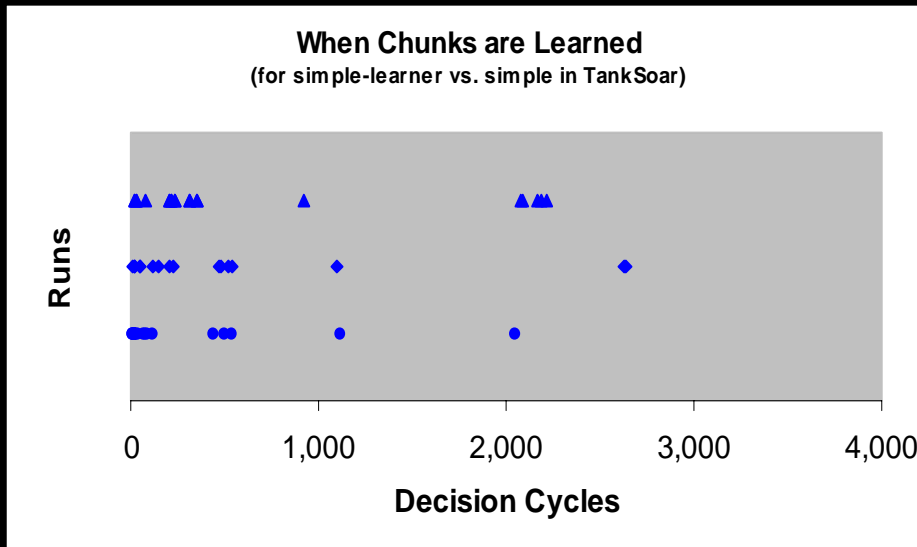

When Chunks are Learned
(for simple-learner alone in TankSoar)



Chunking Rate (in 25 DC bands)
(for simple-learner alone in TankSoar)





Experiments with “simple-learner” with “simple” tank




- Learning forever?
 - Learning continues much longer
 - Learns more chunks

Experiments with “simple-learner” with “simple” tank



Future long-term experimenting with TankSoar



- **Plans:**
 - Make tank's learning more meaningful
 - Use more sophisticated opponents
- **Advice?**
 - Experimental design change?
 - Advice on making learning more meaningful?
 - Other advice?