

SML: A New Interface Into Soar

Douglas Pearson

+ Bob, Jon, Karen, Trevor, Devvan, John

douglas.pearson@threepenny.net



Short-Term Problem

- New debugger
 - Avoid string parsing
 - Use Java for faster tool development
 - Dynamic attachment to existing processes
 - High performance
- gSKI integration
 - Removal of existing command line
- Redux choice to use JESS for rule matcher

Long-Term Problem

- Interfacing to Soar has always been hard
 - Initially limited to implementation language (LISP/C)
 - Extended to Tcl but lead to dependencies on Tcl
 - SGIO supported embedded kernels but only for I/O
 - gSKI added clean interface to kernel but multiple languages, debugging and remote I/O unimplemented.

Key Properties of our Solution

- Supports multiple languages (Java, C++, Tcl currently) while removing Tcl dependency
- Supports uniform interface for I/O (environments) and commands (debuggers)
- Supports embedding kernel within debugger or environment with remote connections between them
- Supports multiple clients (environments, tools, debuggers etc.) connecting to a single kernel
- Supports dynamic connection and disconnection of tools (esp. debuggers) from a running kernel
- Provides a uniform, high-level, data-driven model for the entire Soar interface while achieving high performance
- Moves command line support out of the kernel while providing universal access to it from any client
- Includes a new cleaned up command line interface
- Lots of new capabilities yet in most cases the new interface is substantially faster than 8.5.2
- No production level changes (except 'tcl' -> 'exec')

Connecting to Soar

I/O Commands

`^input-link`
`^output-link`

Run Commands

`run 10 --d`
`step 5`

Debug Commands

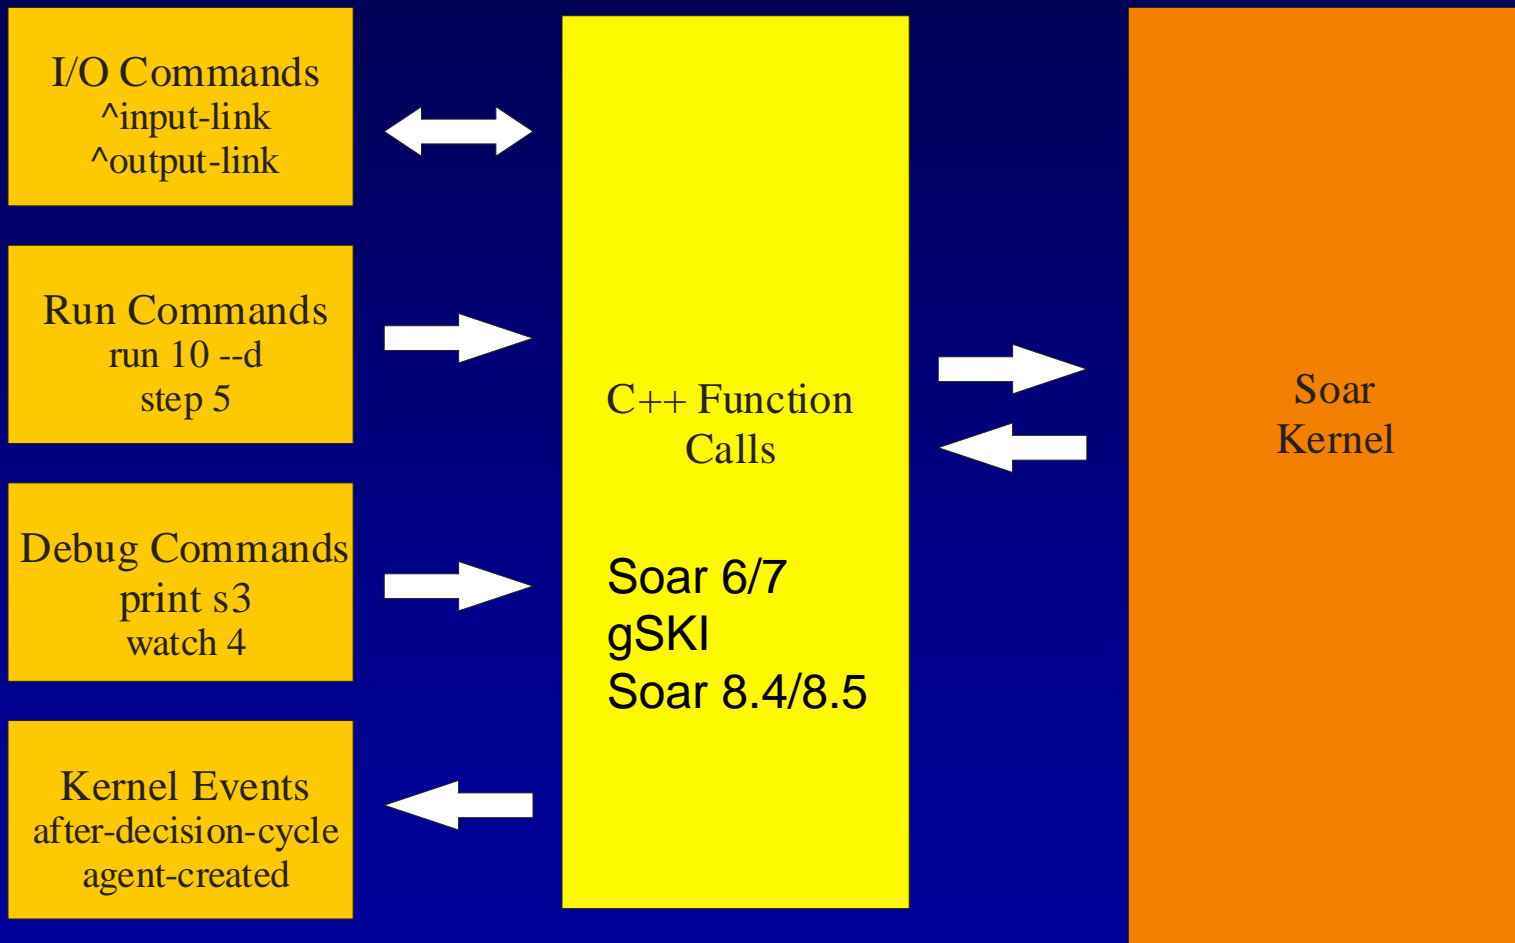
`print s3`
`watch 4`

Kernel Events

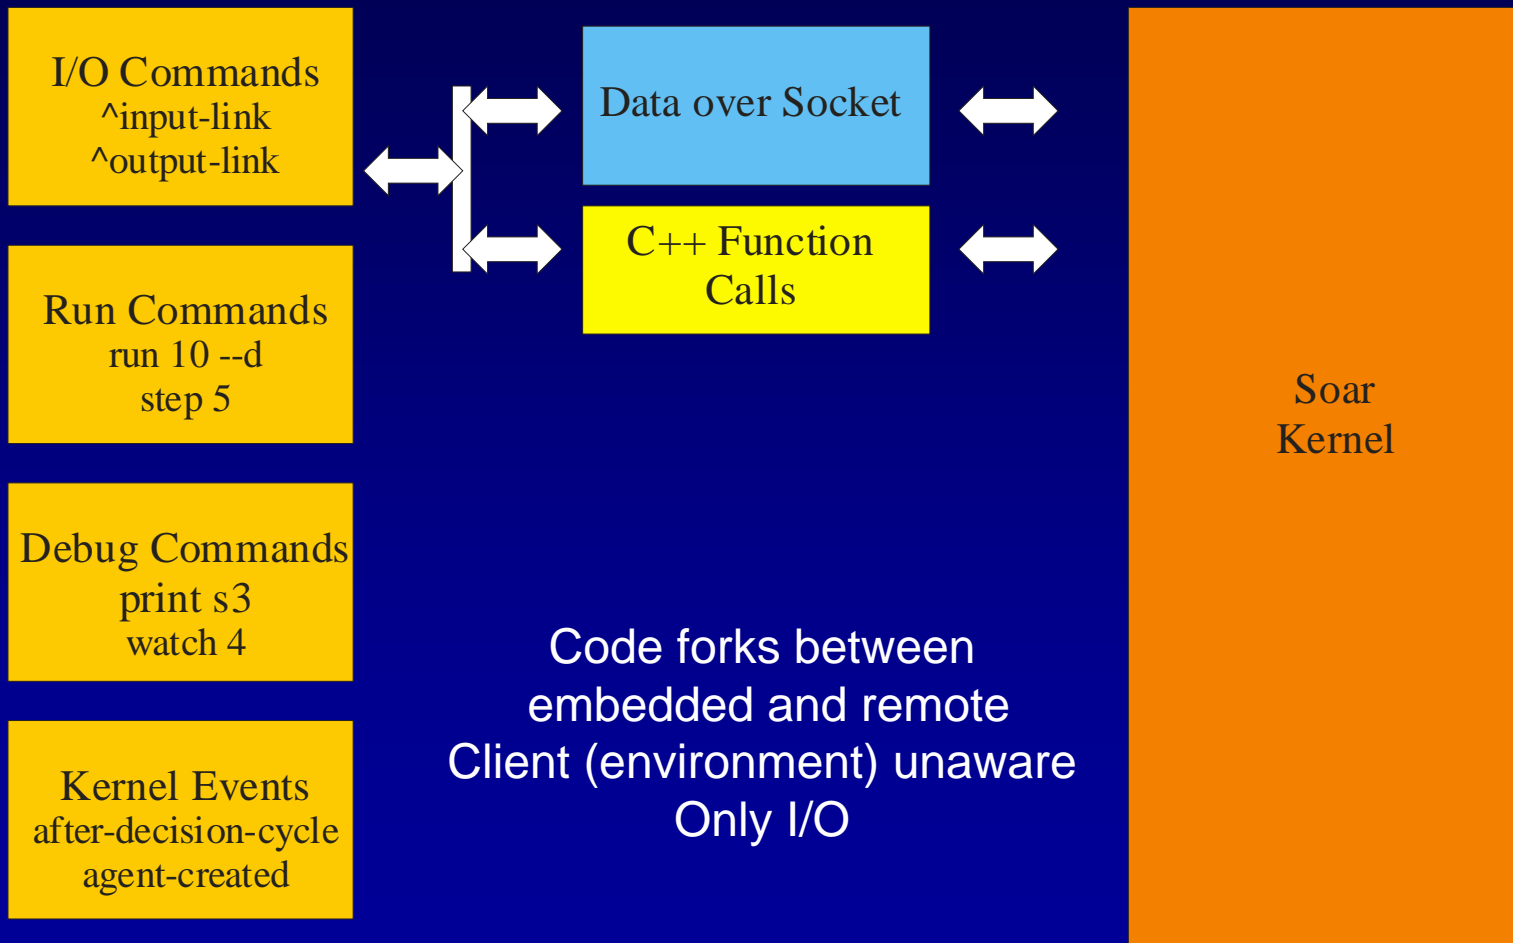
`after-decision-cycle`
`agent-created`

Soar
Kernel

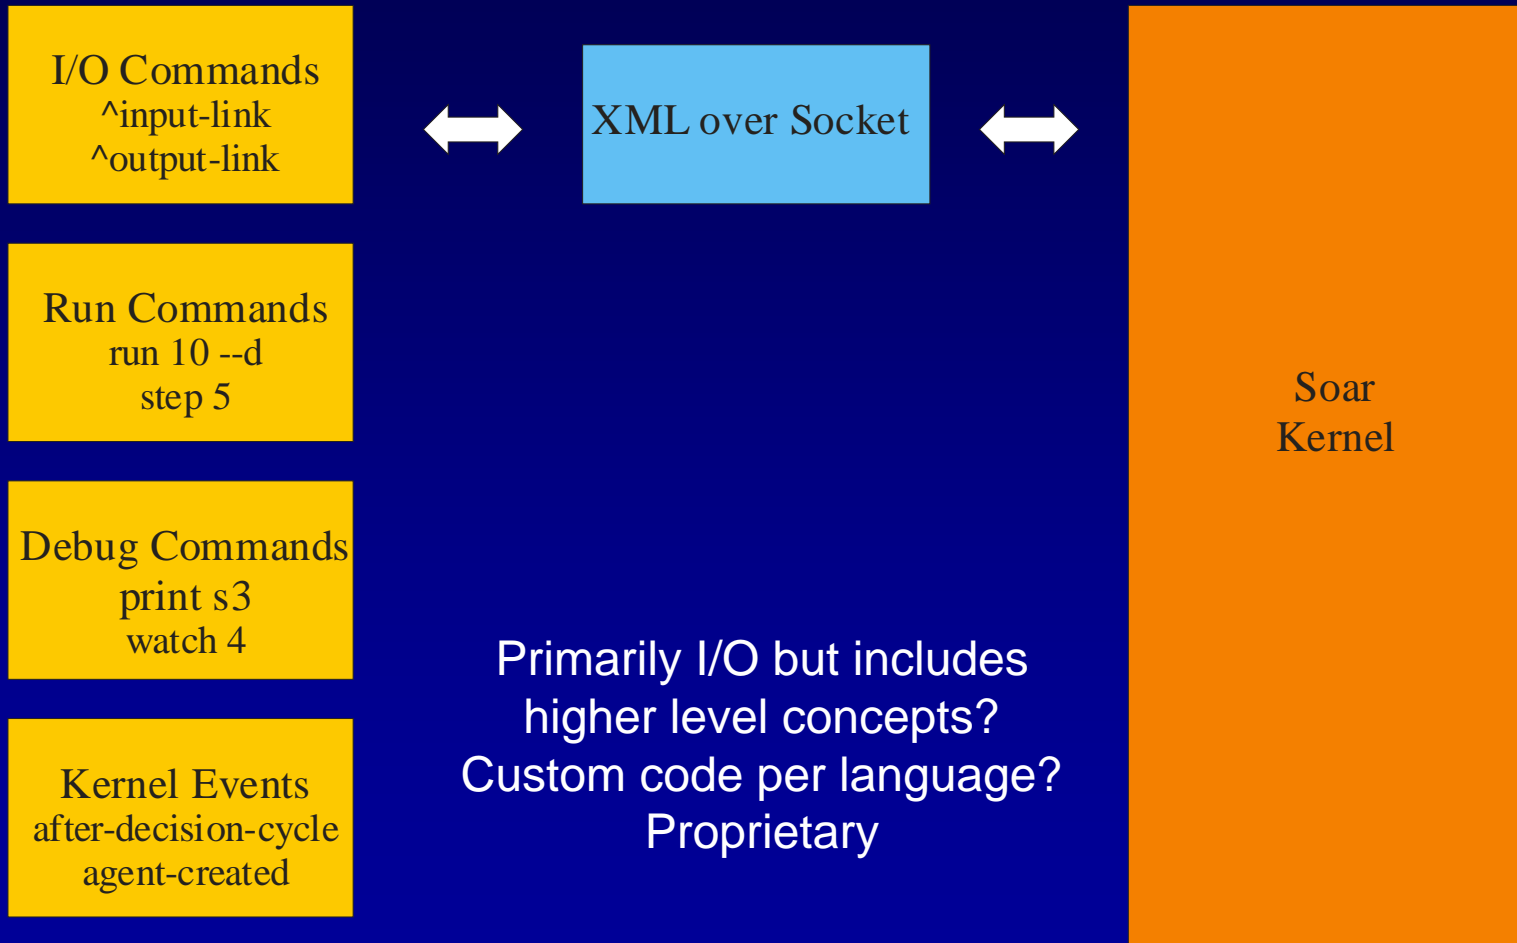
Connecting to Soar Direct Link to Kernel



Connecting to Soar SGIO Style

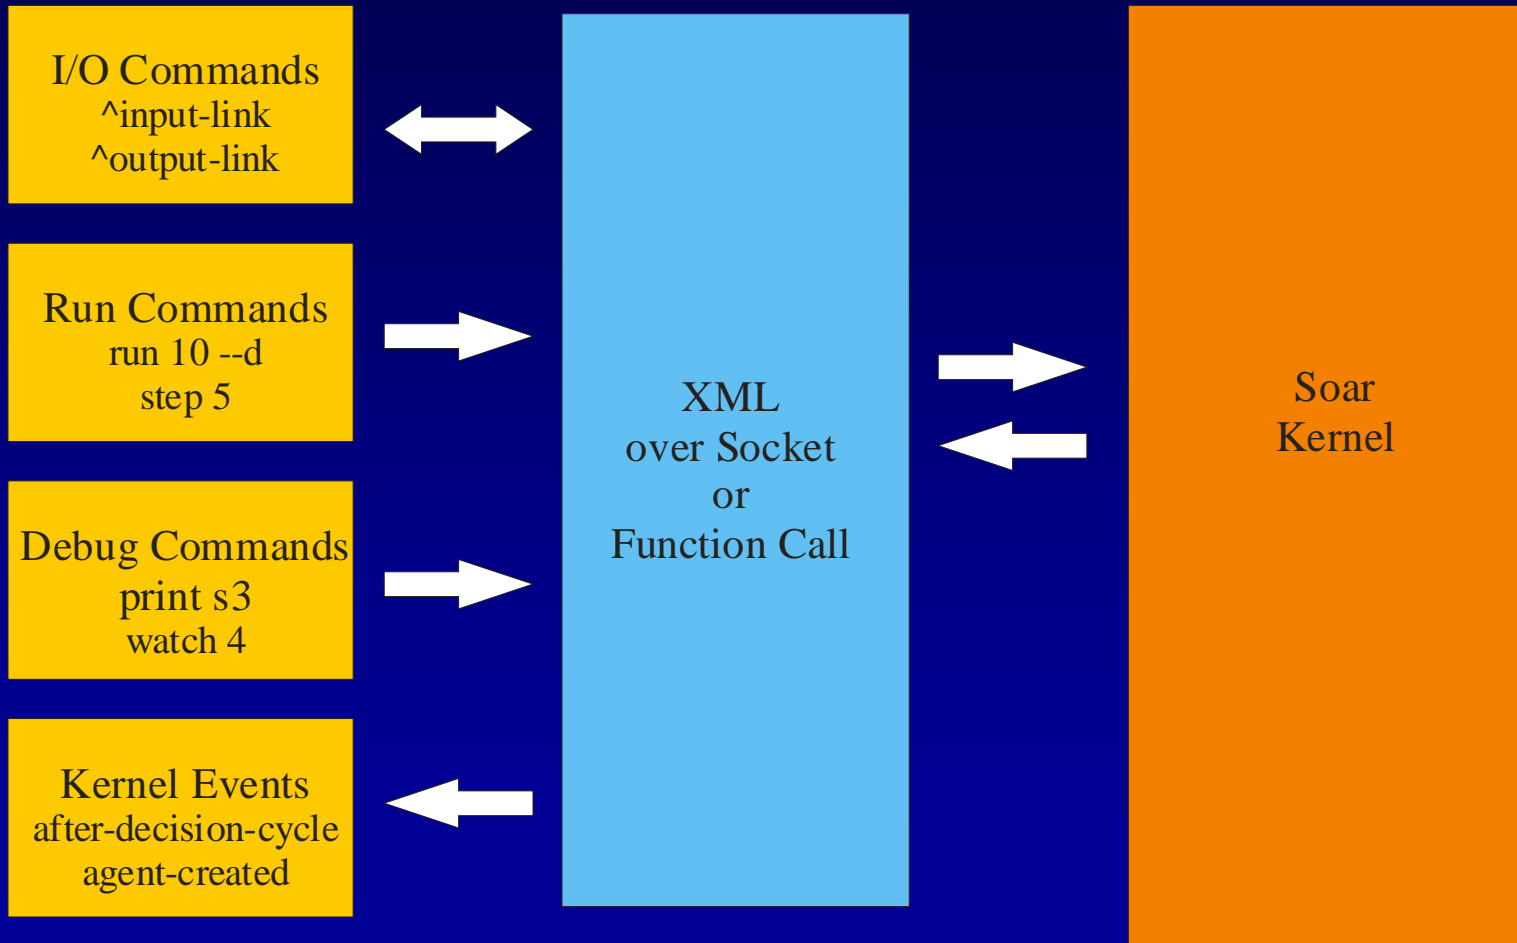


Connecting to Soar ATE Style?



Connecting to Soar

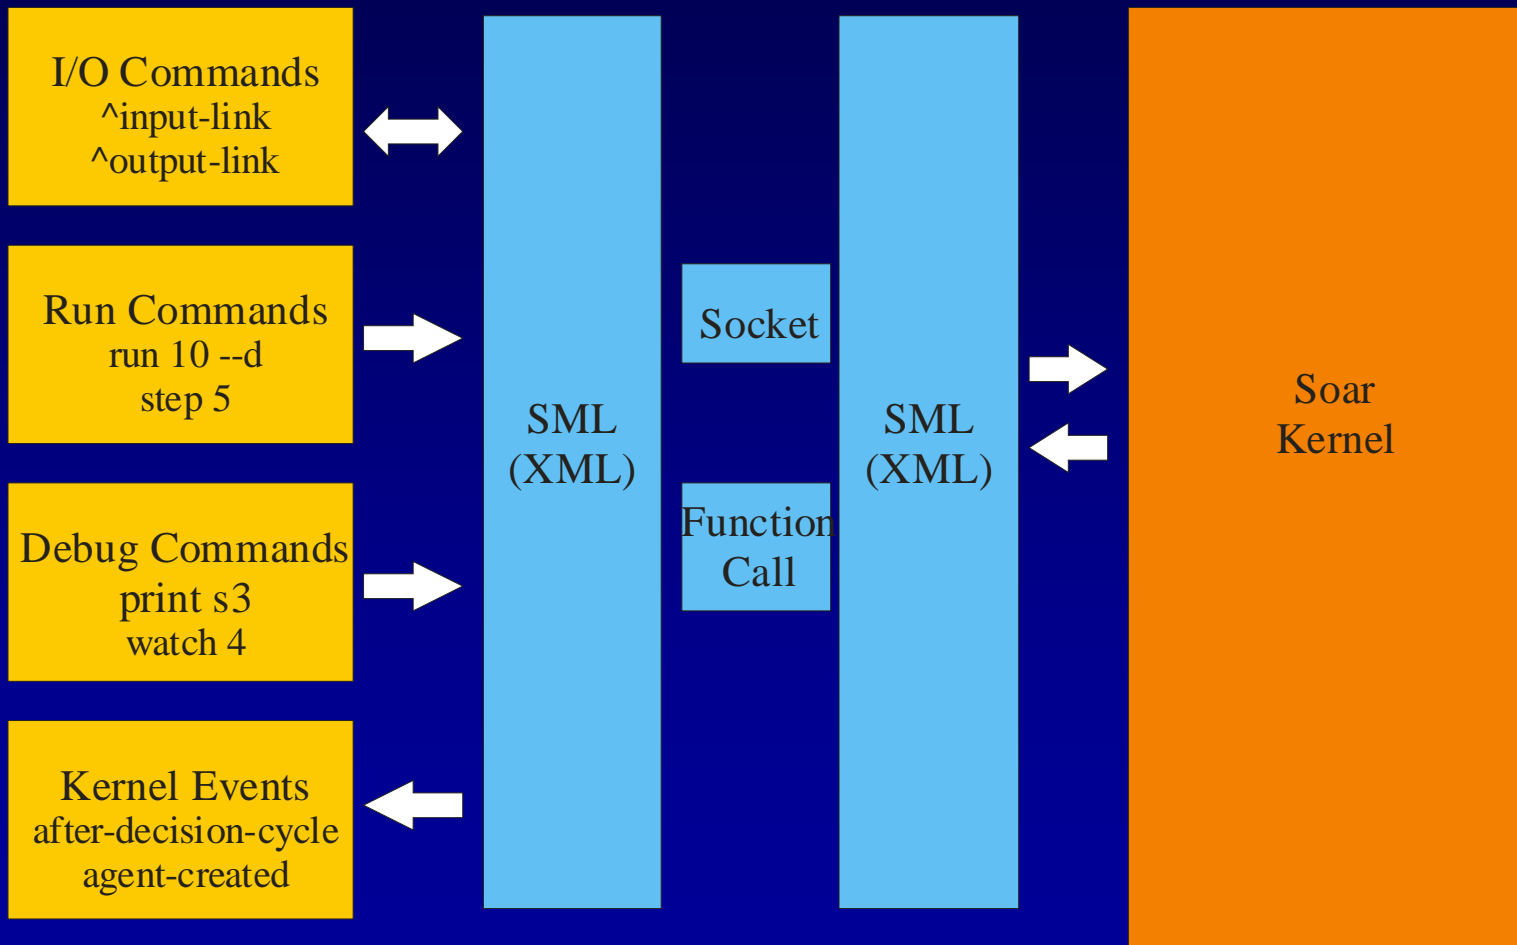
SML Style – Soar Markup Language



Uniform for entire interface

Connecting to Soar

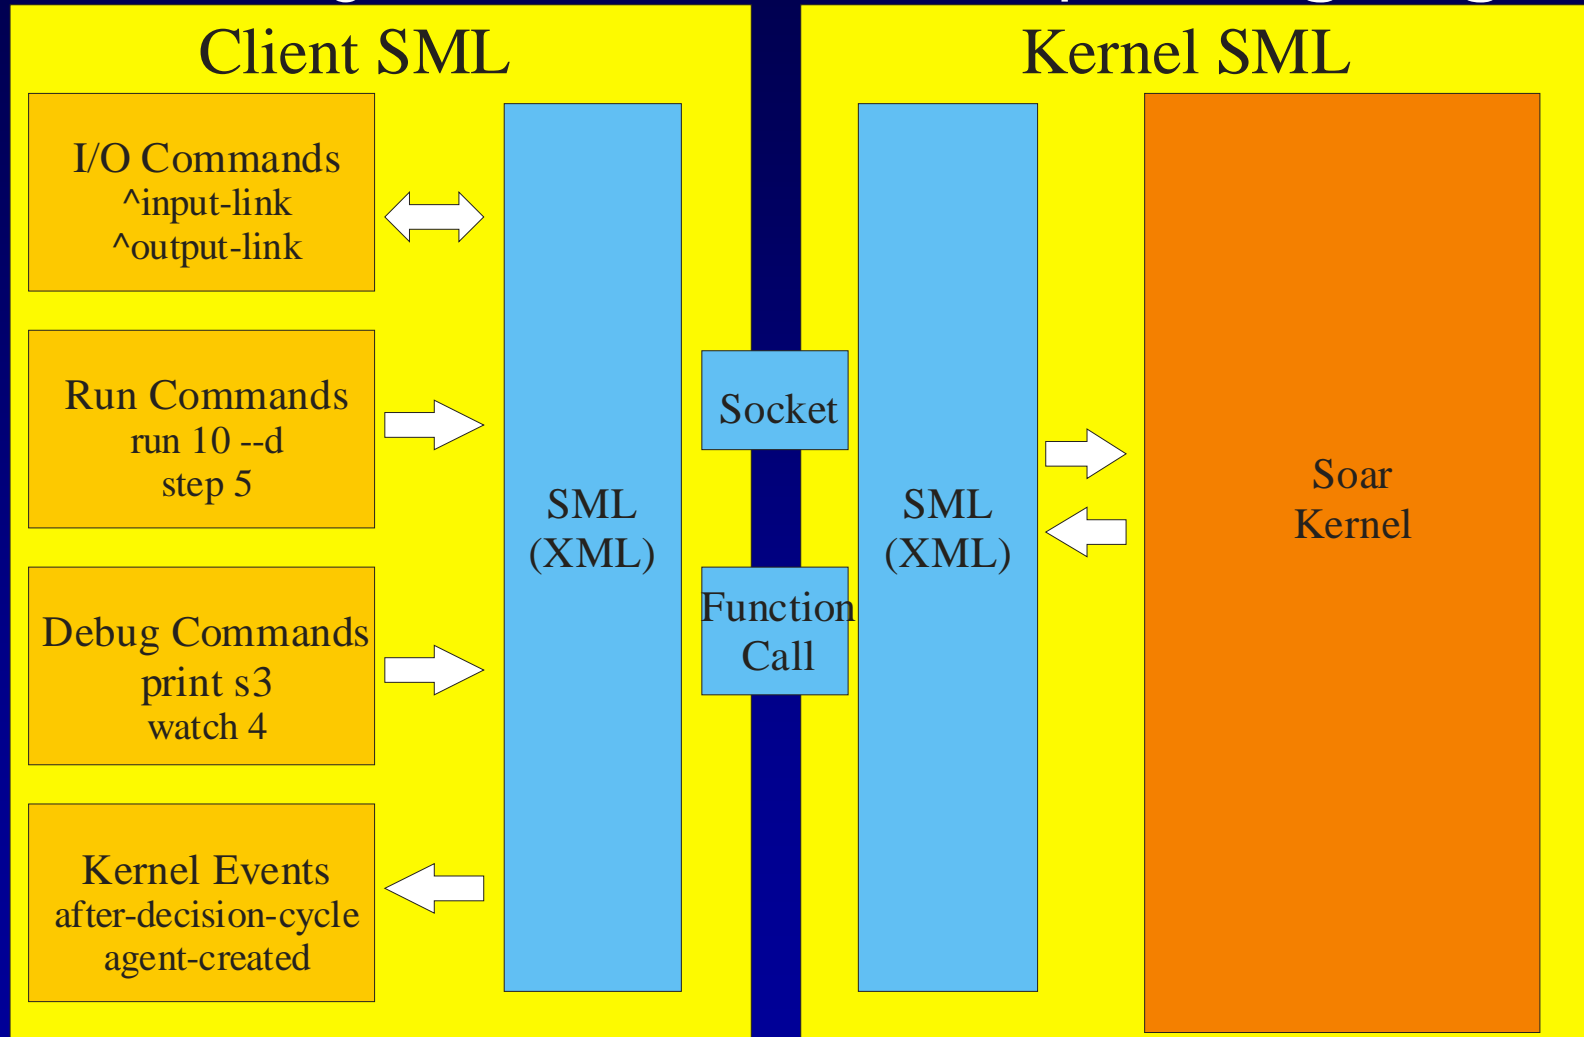
SML Style – Soar Markup Language



Internally use socket or function call
Client, Kernel and (99%) SML code all unaware

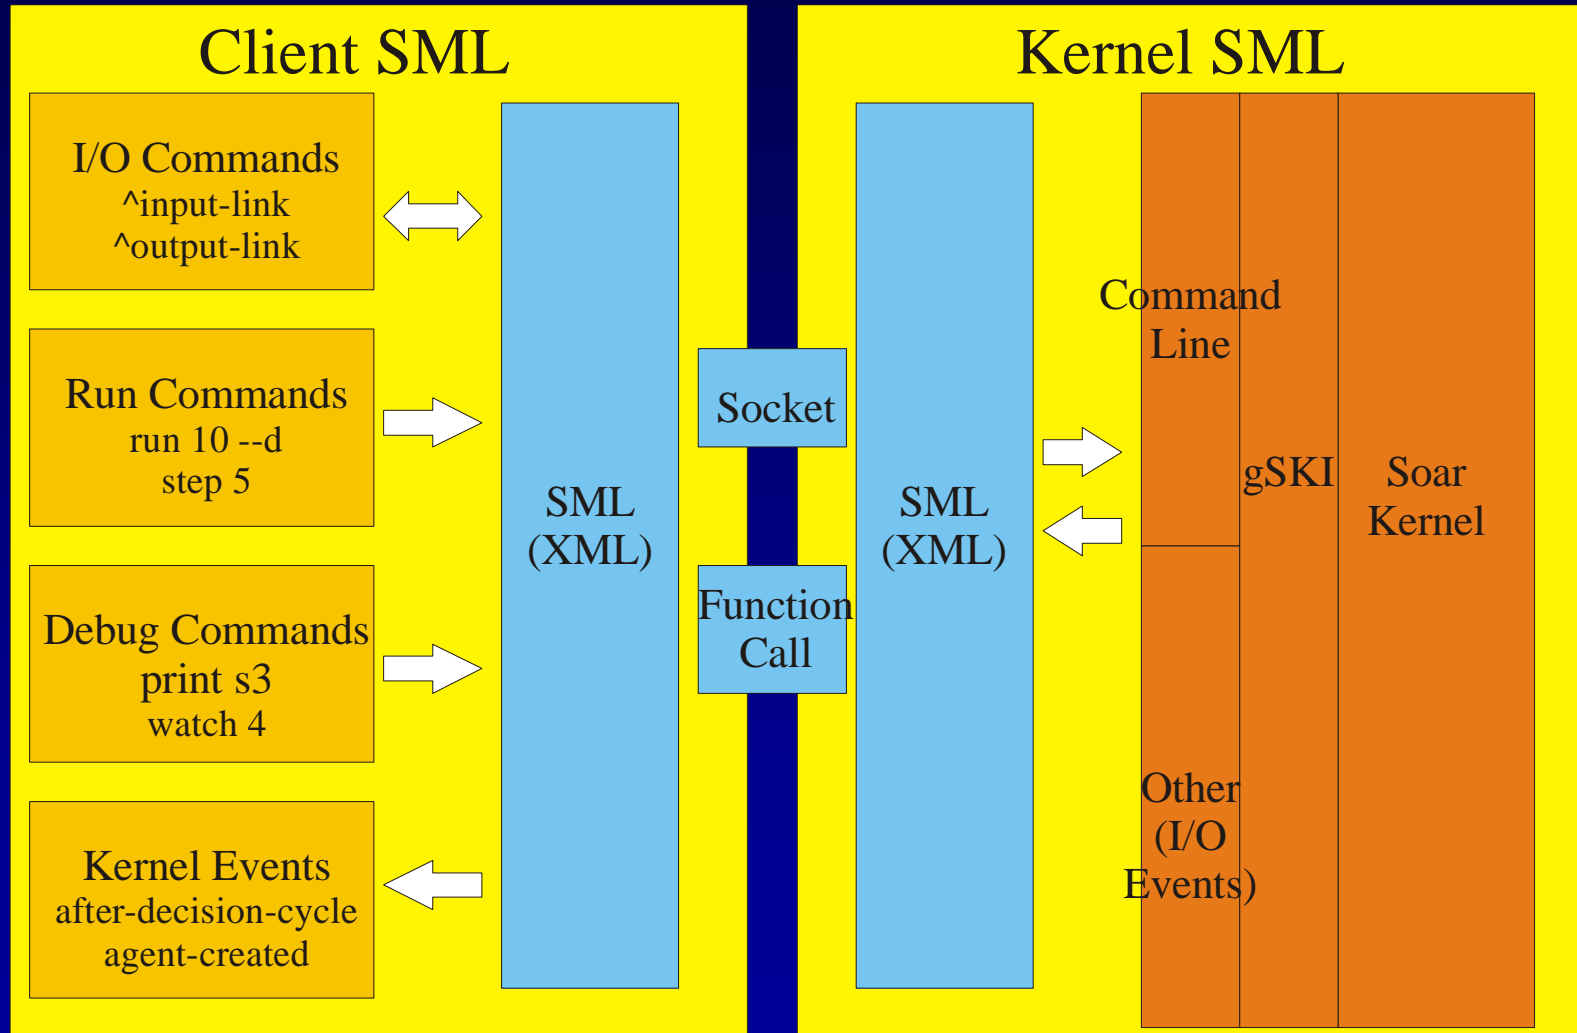
Connecting to Soar

SML Style – Soar Markup Language



Connecting to Soar

SML Style – Soar Markup Language



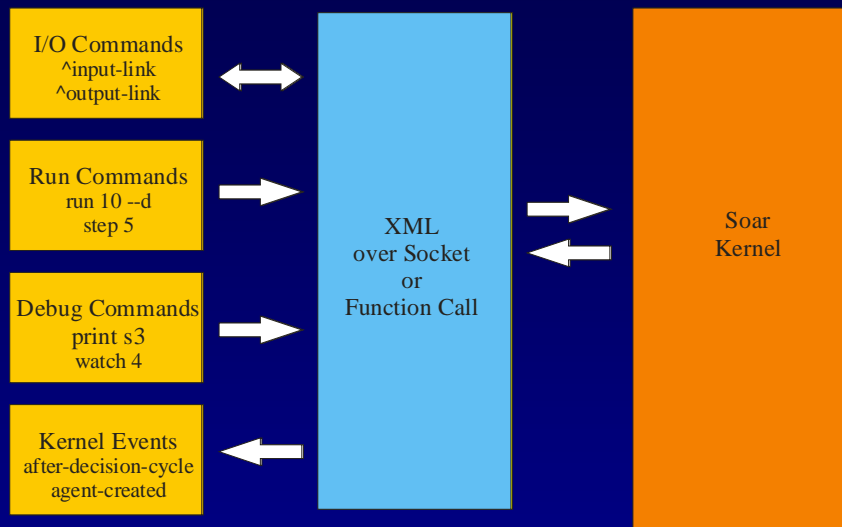
Possibly useful ways to think about SML

- An XML interface into the kernel
 - A data-driven approach rather than function calls
- A remote procedure call protocol for the kernel
 - Although it's not just for remote calls
- A language independent way to access the kernel
- SOAP/Web Service for the kernel
 - Although again not limited to remote access

De-Coupling from the Kernel

- Two central tenets for good software design
 - High cohesion
 - Each class does one thing not several things
 - Loosely coupled
 - Each class knows as little as possible about the classes they interact with
- What is good for classes is good for modules

SML Makes Systems More Loosely Coupled to Soar



8.5/gSKI Interface size (# functions which break client if changed and not recompiled) \approx 1000

SML reduces this to 2 (+ about 20 for ElementXML)

Why does this matter?

Supports a larger community building clients – may not have source available
Allows a single client to support multiple different kernel versions because data-driven

By default builds to 2 libraries with no dependencies
So less coupled to Tcl or other stuff

SML Example Packets

“Print chunk-1”

```
<sml smlversion="1.0" doctype="call" soarVersion="8.6.1" id="657" >  
  <command name="print">  
    <arg param="agent">agent-1</arg>  
    <arg type="string" param="name">chunk-1</arg>  
  </command>  
</sml>
```

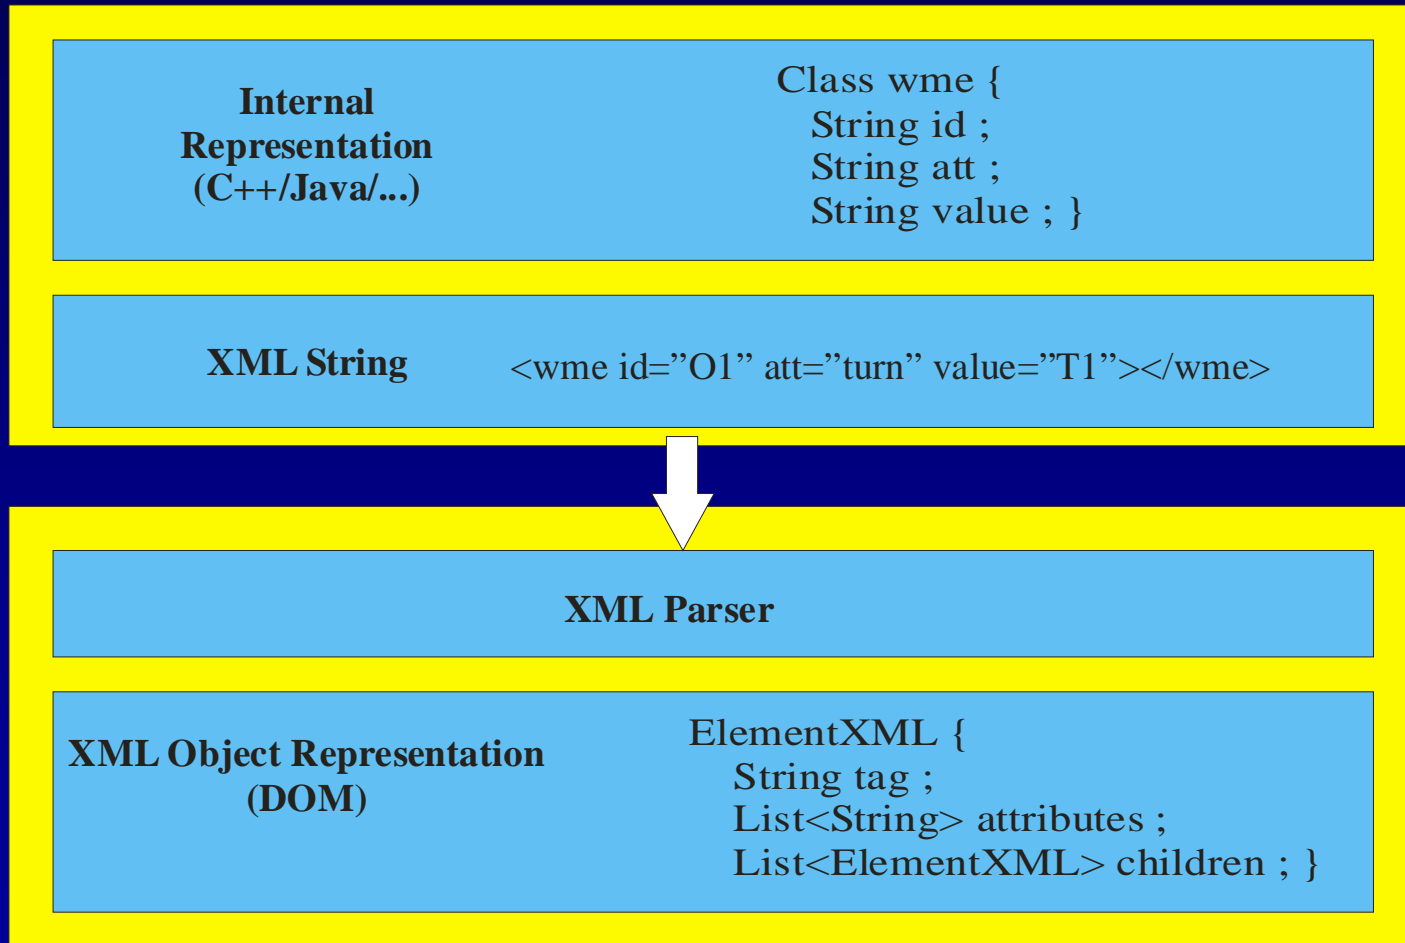
“Output (O1 ^turn T1) (T1 ^heading 045 ^speed 225)”

```
<sml smlversion="1.0" doctype="call" soarVersion="8.6.1" id="1763">  
  <command name="output">  
    <arg param="agent">agent-1</arg>  
    <wme action="add" att="turn" id="O1" tag="7" type="id" value="T1"></wme>  
    <wme action="add" att="heading" id="T1" tag="6" type="int" value="045"></wme>  
    <wme action="add" att="speed" id="T1" tag="12" type="int" value="225"></wme>  
  </command>  
</sml>
```


Maximizing XML performance

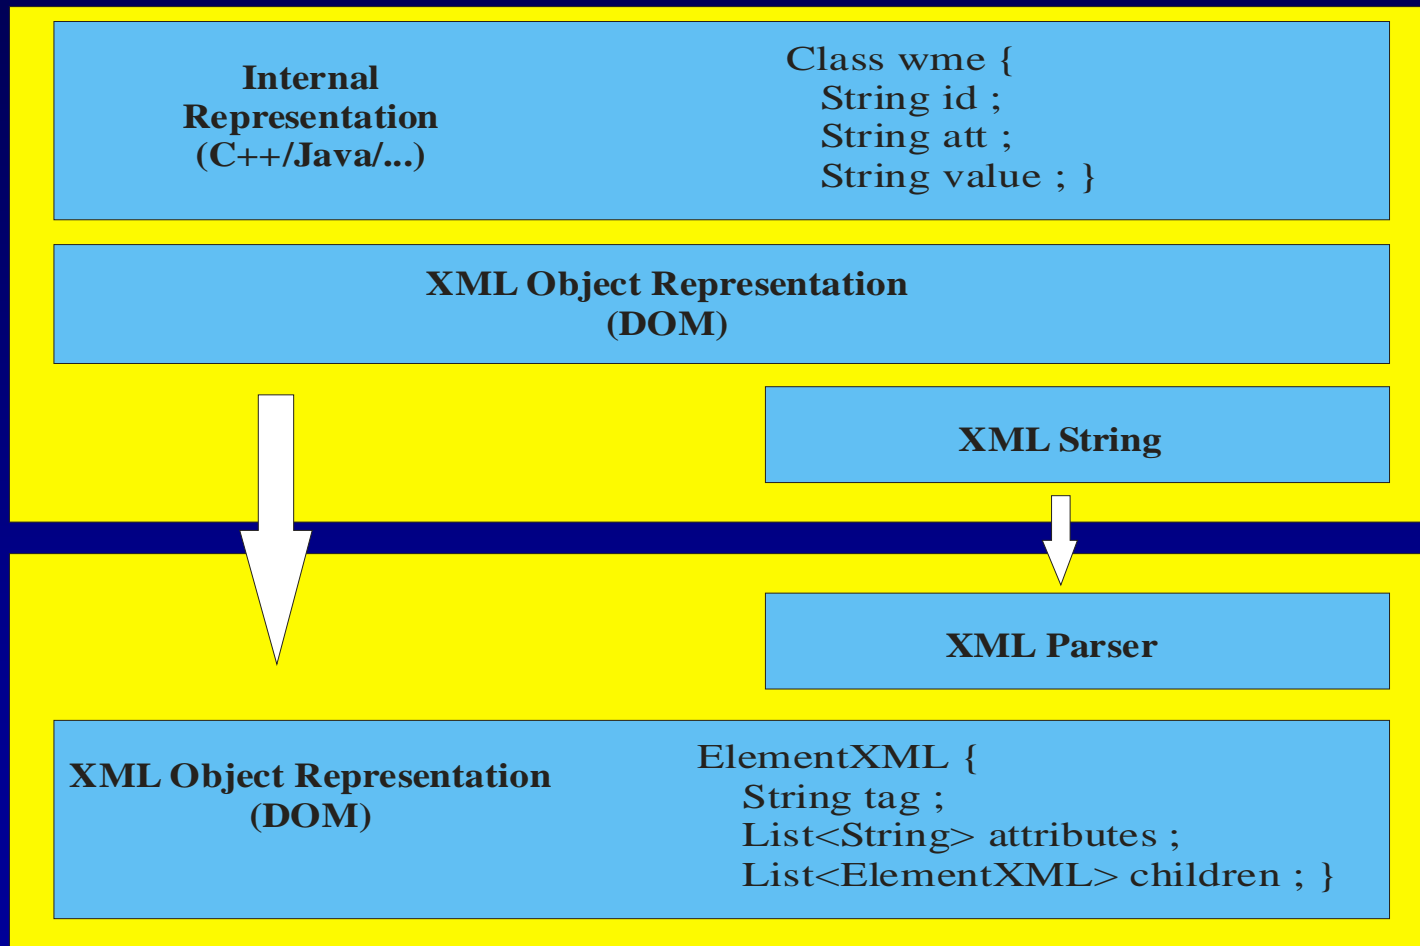
Standard Approach for passing XML

-- Always requires parsing step



Maximizing XML performance

SML Approach – create DOM representation directly
- only convert to XML string if needed



Speed benefit ~50x if always use socket
Speed benefit ~10x if pass string directly

Building an SML Client

- Option A: Just send XML strings to socket
 - Format: 4 byte length + XML string
 - Not recommended, lose embedded speed and need details on XML packets but an option
- Option B: Use client SML functions we provide in
 - C++
 - Java
 - Tcl
 - more...

Creating Kernel and Agents

- Local Kernel
 - Kernel* pKernel = CreateKernelInNewThread()
 - Generally recommended
 - Kernel* pKernel = CreateKernelInCurrentThread()
 - Tcl only supports this form (so far)
 - Call CheckForIncomingCommands() periodically
- Remote Kernel
 - Kernel* pKernel = CreateRemoteConnection(ipAddress, port)
 - Pass null for ipAddress => same machine
- Agents
 - CreateAgent()
 - DestroyAgent()
- Shutdown
 - delete pKernel object
 - disconnects
 - cleans up all memory

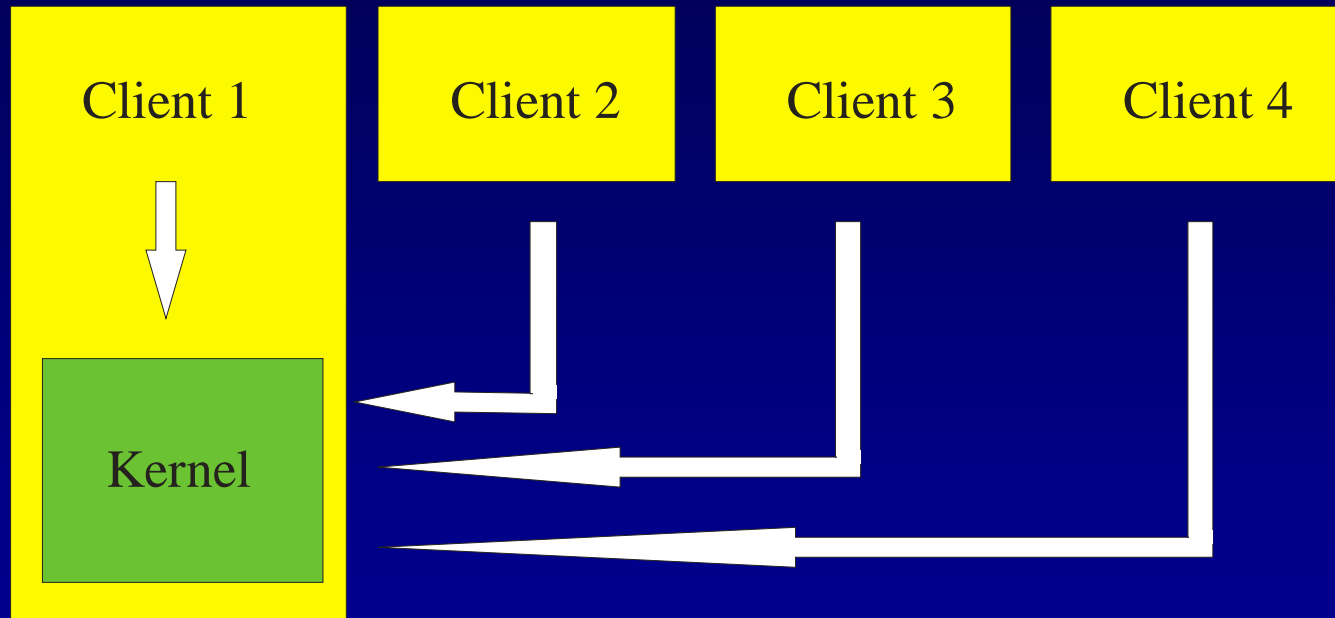
Local vs Remote Kernels

CreateKernel
InNewThread()

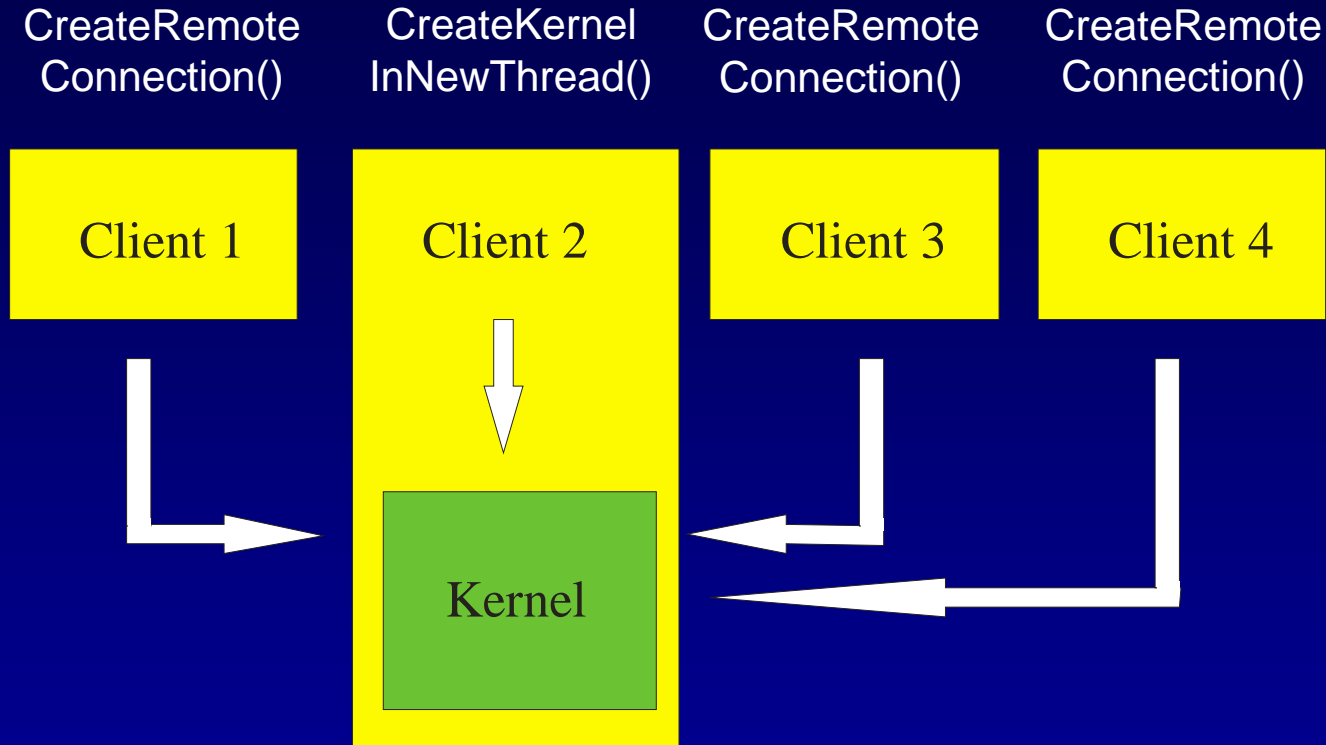
CreateRemote
Connection()

CreateRemote
Connection()

CreateRemote
Connection()



Local vs Remote Kernels



Change one line of code

- Internally different execution paths
 - XML objects + local functions vs XML strings + sockets
- To the client, same interface and capabilities
- No special access for debugger

Input

- Add structures to input link representation
 - Identifier* pInputLink = pAgent->GetInputLink()
 - Identifier* pID = pAgent->CreateIdWME(pInputLink, "plane") ;
 - StringElement* pWME1 = pAgent->CreateStringWME(pID, "type", "Boeing747") ;

Result: (I1 ^input-link I2) (I2 ^plane P1) (P1 ^type Boeing747)

- Send changes to kernel
 - pAgent->Commit() ;
 - Actual wmes added during next input phase
- Update or remove in future cycles
 - pAgent->Update(pWME1, "Cessa") ;
- NOTE: init-soar works and comes for free.
 - Sends over current input state automatically so you can pick up reasoning

Output

- Check for output commands (since last check)
 - Based on “command” notion on output link
 - E.g. (I1 ^output-link I3) (I3 ^move M1) (M1 ^speed 20)
 - `pAgent->GetNumberCommands()` ;
 - `Identifier* pCommand = pAgent->GetCommand(i)` ;
 - `string name = pCommand->GetCommandName()` ;
 - `string speed = pCommand->GetParameter(“Speed”)` ;
 - `pCommand->AddStatusComplete()` ;
 - Same as `pAgent->CreateStringWME(pCommand, “status”, “complete”)` ;
- But not forced to use this model
 - `GetCommand()` returns `Identifier*`. Can access substructure directly.
 - `GetOutputLink()` and walk output link directly.
 - `GetOutputLinkChange()` to walk list of WMEs changed
- Call `ClearOutputLinkChanges()` after reading the output link
 - Allows SML to only report changes to the output link.

Running Soar

- Run
 - Kernel->RunAllAgentsForever()
 - Kernel->RunAllAgents(steps, stepSize)
 - Agent->RunSelfForever()
 - Agent->RunSelf(steps, stepSize)

- Stop
 - Kernel->StopAllAgents()
 - Agent->StopSelf()

Debugging commands

- ExecuteCommandLine(commandLine)
 - ExecuteCommandLine(“watch 3”) ;
 - ExecuteCommandLine(“print o1”) ;
 - ExecuteCommandLine(“excise –all”) ;
- Simple form returns a string
 - “set-library-location”
 - returns: “Current library location: e:\soarmich\soar-library”
- Alternative form returns XML
 - “set-library-location”
 - ```
<sml doctype="response" id="32" smlversion="1.0 soarversion="8.6.1"ack="237">
 <result>
 <arg param="directory" type="string">e:\soarmich\soar-library</arg>
 </result>
</sml>
```
  - GetArgValue(“directory”) [smlNames::kParamDirectory]
- Why not implement all as methods in the client interface?
  - Just saving work
  - Some clients (lots?) will want to embed command line windows
  - XML all documented on Wiki ([http://winter.eecs.umich.edu/soarwiki/Main\\_Page](http://winter.eecs.umich.edu/soarwiki/Main_Page))

# Event Handling

- Register callback handler for an event

```
pAgent->RegisterForRunEvent(smlEVENT_AFTER_DECISION_CYCLE,
 MyRunEventHandler, 0) ;
```

- Called back when event occurs:

```
void MyRunEventHandler(smlRunEventId id, void* pUserData, Agent* pAgent,
 smlPhase phase) ;
```

- Supports all gSKI events (plus a few additions)

- Long list

- Events can be handled over remote connections the same as local ones

# RHS Functions

- Register RHS function

```
pKernel->AddRhsFunction("test-rhs", &MyRhsFunctionHandler, 0);
```

- Referred to via “exec” in production (just like “tcl”)

```
sp {apply*user*exec
 (state <s> ^operator <o> ^io.output-link)
 (<o> ^name move ^space <sp>)
 (<sp> ^row <row> ^col <col>)
-->
 (^test (exec test-rhs | hello | <row> | world |))}
```

- Calls RHS function handler

```
std::string MyRhsFunctionHandler(smlRhsEventId id, void* pUserData, Agent* pAgent, char
 const* pFunctionName, char const* pArgument);
```

- “cmd” allows access to command line as built in command set

```
--> (cmd print <s>)
```

- Handler can be in any language (e.g. register from Java, handler in Java)
- RHS functions can be handled over remote connections the same as local ones
  - If same function registered locally and remote, calls local one

# Threads

- `CreateKernelInNewThread()`
  - Runs kernel in its own thread
  - Means kernel can respond to remote commands (over socket) on its own
  - Otherwise need to call “`CheckForIncomingCommands()`” periodically
- Also `EventThread` helps client remain responsive
  - Register for an event (e.g. `AFTER_DECISION_CYCLE`)
  - Go to sleep => entire system locks up waiting for this client
  - Solution is a thread which receives these events and pushes them through client callback to get response
- Both just make it easier to build a client
  - Neither thread is required – `CreateClientInCurrentThread()`
  - For maximum performance turn them off and handle these issues yourself

# SWIG and Language Support

- SWIG

- Automatic generation of Java and Tcl implementations
- Fast to create and maintain
- Custom code added to handle callbacks
- Code looks the same in each language

## C++

```
sml::Kernel* pKernel = Kernel::CreateRemoteConnection(true, null, Kernel::GetDefaultPort());
```

```
if (pKernel->HadError())
{
 cout << pKernel->GetLastErrorDescription() << endl ;
 return false ;
}
```

## Java

```
sml.Kernel kernel = Kernel.CreateRemoteConnection(true, null, Kernel.GetDefaultPort());
```

```
if (kernel.HadError())
{
 System.out.println(kernel.GetLastErrorDescription());
 return false ;
}
```

# Common Environment Control Loop

Asynchronous environments

Act as soon as output is sent from an agent

Synchronous environments

Fixed amount of Soar processing before update world

while (!stopped)

- Run(1) (or Run-til-output)
- Collect-output
- Update-state-of-world
- Send-input

Assumes environment triggers run.

Assumes run(1) completes a decision

Assumes user doesn't issue "stop-soar" (needs to pause/stop environment)

# Preferred Environment Control Structure

1. Register for “update world” event (e.g. AFTER\_ALL\_OUTPUT\_PHASES)
2. Handler takes form:
  - Collect-output
  - Update-state-of-world
  - Send-input
3. Run controls
  - Kernel::RunAllAgentsForever()
  - Kernel::RunAllAgents(1)

As agents run, fire event as reach end of output phases.

Allows run to come from debugger or environment or other tools.

Allows for arbitrary interruption of run commands (e.g. breakpoints/stop-soar).



# Debugger Performance Comparison

| Towers of Hanoi   | TSI (8.5.2) | Java Debugger (Text) | Java Debugger (Tree)                    |
|-------------------|-------------|----------------------|-----------------------------------------|
| Watch 1 (Run 100) | 1.25 secs   | 0.73 secs            | 0.75 secs                               |
| Watch 5 (Run 100) | 59.68 secs  | 2.14 secs            | 1.51 secs<br>0.58 secs (full filtering) |

- Within process faster than 8.5.2 (for the debugger)
- Embedding Soar in environment so remote => local
  - Not measured yet, but huge speed up

# Documentation

- PDFs included in Release
  - SML Quick Start Guide
  - Moving from SGIO to SML
  - Soar XML Interface Specification
  - Intro to the Soar Debugger
  - How to guide – adding an event to gSKI and SML
- Online on the Wiki
  - [http://winter.eecs.umich.edu/soarwiki/Main\\_Page](http://winter.eecs.umich.edu/soarwiki/Main_Page)
- Questions, bugs, building something cool, need help?
  - [soar-sml-list@umich.edu](mailto:soar-sml-list@umich.edu)

# Nuggets

- Supports multiple languages (Java, C++, Tcl currently) while removing Tcl dependency
- Supports uniform interface for I/O (environments) and commands (debuggers)
- Supports embedding kernel within debugger or environment with remote connections between them
- Supports multiple clients (environments, tools, debuggers etc.) connecting to a single kernel
- Supports dynamic connection and disconnection of tools (esp. debuggers) from a running kernel
- Provides a uniform, high-level, data-driven model for the entire Soar interface while achieving high performance
- Moves command line support out of the kernel while providing universal access to it from any client
- Includes a new cleaned up command line interface
- Lots of new capabilities yet in most cases the new interface is substantially faster than 8.5.2

# Nuggets

- Building a command line interface in a client is pretty simple
  - Allows embedding console windows into environment (e.g. games)
- Readily supports other solutions
  - Other debuggers (e.g. TSI-8.6)
  - Other editors (edit-production is generic)
- 8.6.1 will be out by end of June
  - Linux and Mac support
  - Soar7 mode
  - Tcl Eaters
  - Java Missionaries and Cannibals
  - New environment event model
  - Kernel XML generation much faster
  - New debugger features (fast tree view, filtering etc.)
  - Lots of other bug fixes
- 8.6.1r5 Used for tutorial at the workshop w/o problems
  - Tcl Eaters with Java Debugger with C++ Kernel

# Coal

- Not all commands generate fully structured XML
  - Esp. print, but should do so shortly
  - Still XML but not as rich as we'd like
- Threading models not completely resolved yet
  - Esp. Tcl
- Lots of new code
  - There will be some bugs lurking
  - Not all combinations fully tested yet
  - Interfaces will probably change a bit as everything settles