



Player Modeling in IDA

Dr. Brian Magerko
University of Michigan



Interactive Drama





The Boundary Problem

- Author defines a *story space*
 - The space of intentionally dramatic stories possible
 - Consonance between player actions and authored content
- Player actions may lead to a world state outside the boundary of the story space

The dramatic experience may stall or even halt...






Contributions

1. Incorporation of semi-autonomous characters
2. *A complete* story director with..
 - a. Player knowledge modeling
 - b. Player prediction
 - c. Both reactive and preemptive direction for addressing boundary problems
 - d. Two-tiered selection of content
3. A story representation that supports:
 - a. Reactive & preemptive direction
 - b. Pacing
4. Evaluation via archetypes

IDA: An Interactive Drama Architecture



Haunt 2: Built in Unreal Tournament

percepts

user actions

direction

story content



Human player




AI Director



Author



Director



- Omniscient & omnipotent intelligent agent
- Executes “direction”
 - Enacts story content
 - Mediates between player actions & plot
 - Influences player behavior to stay within story space
 - Based on current actions
 - Based on predicted actions



Knowledge Maintenance

- Records when world changes significantly
- Observes story world omnisciently
- Creates new WMEs to reflect changes in the state of the world
- Hypothesizes basic situational awareness of entities in the story world
- Applies ontology of everyday world knowledge





Hypothesized Knowledge

- Player is in lobby, invisible
- Sally is in the lobby
- Innkeeper is in the lounge
- Corpse is in the library
- Thermos is in the lobby

The player knows:


- Exists(Sally)
- Exists(Thermos)
- In(Sally , Lobby)
- In(Thermos, Lobby)
- Thermos is for drinking





Plot Monitoring

- Occurs when plot content matches with current knowledge base
 - State of the world
 - Mental states
- Plot:
 - Content: plot points
 - Preconditions
 - Actions
 - Structure:
 - Ordering constraints
 - Timing constraints





Plot selection

- What if there is more than one candidate for performance?
- Two-tiered process
 - Player-motivated
 - Director-selected
- Selected plot point marked



Reactive Direction

- Immediate response to boundary problems
- Relies on more *effective* director strategies
- Boundary problems signaled by an *occurrence of timing violations*




Preemptive Direction

- Preemptive response to predicted boundary problems
- Relies on more *subtle* director strategies
- Queries predictive model of player behavior

More believable than relying only on reactive measures or directing continuously

Example

- User has been moving to new rooms...
 - Predicted goal: *Explore*
- Predicted timing constraint violation






Single Modeling Run



world state


state copy



result

{Success, 0.42}

model





Player Prediction

- Success
 - Precondition is fulfilled before a timing constraint is violated
- Failure
 - Timing constraint is violated
- Result determines:
 - *if* the director preemptively directs
 - *how* it directs preemptively



Probabilistic Sampling

- Player model
 - Returns a tuple, $M \rightarrow (R, P)$
 - R: success / failure
 - P: probability of particular sequence of actions chosen
 - Runs created iteratively until an author-defined limit ρ is reached
- Director computes confidence in the user fulfilling plot content (C_m)
 - Function of likelihood of each run & its result
 - $C_m > \alpha \rightarrow$ success
 - $C_m < \alpha \rightarrow$ insignificant success *or* failure



Connecting Modeling to Direction

- Direction chosen by best *score*
- Scores dynamically assigned as function of modeling result, C_m
- Each director action rated by a scoring function:
$$Score_{action} = (Sub * S_{action} + Eff * E_{action}) / 2$$
 - S_{action} / E_{action} : authored rating for a particular director action
 - Sub / Eff: weights assigned at run-time
- Example: *transport-player* {S = 0.05, E=.9}



Quantitative Evaluation

- Play as archetypes
 - Explorer
 - Chaser
 - General
- Experimental groups
 - Modeling on
 - No modeling
 - Model still run
 - Preemptive direction “turned off”
- Measures
 - Boundary problems
 - Frequency of direction
 - Average subtlety of direction



Results

- No difference for avg subtlety and # of directions
 - Director should direct the same #
 - Small variability in ratings or problem in authorship



Results (cont.)

- Boundary problems
 - modeling < no_modeling ($p < 0.01$)
 - Within-groups:
 - General: no effect
 - Chaser: no effect
 - Explorer: $p < 0.01$
- Why?
 - Lack of general coverage of predictive model
 - Explorer easily “most accurate” hand-encoded model
 - Points to the need for adaptive modeling



Future Work

- Authoring tools
- Categorization of director strategies
- Player modeling
 - Unified model
 - Knowledge model
 - Behavioral model
 - Skill model
 - Adaptive models
- Adaptive story for education
 - Story as engagement
 - Experience tailored for dramatic and learning value



Nuggets

- *A complete* story director with..
 - Player knowledge modeling
 - Player prediction
 - Both reactive and preemptive direction
 - Two-tiered selection of content
- A story representation that supports
 - Reactive & preemptive direction
 - Pacing
- Evaluation via archetypes
- Successful defense (!)



Coal

- *Haunt 2* did not wind up being a complete or robust experience
- Starting over with a new domain
- No rigorous user study done in evaluation
- Possible bias in evaluation
 - Parameter values
 - Archetype & individual sensitivity to direction