

# Hierarchical Reinforcement Learning and Soar

Shelley Nason

Soar Workshop 2006

# Soar-RL: Soar with an architectural RL mechanism

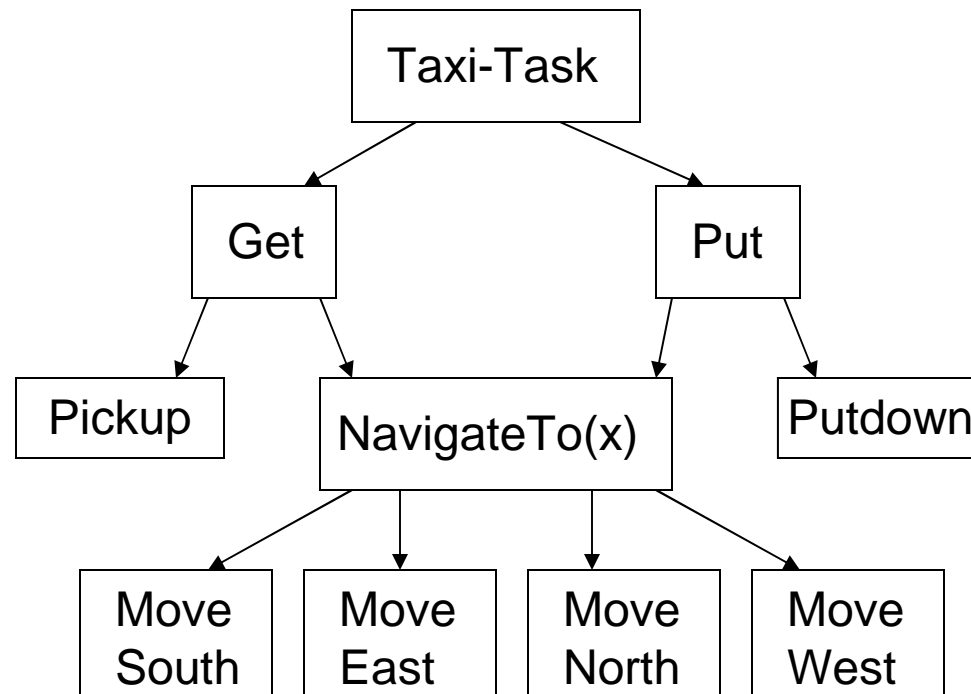
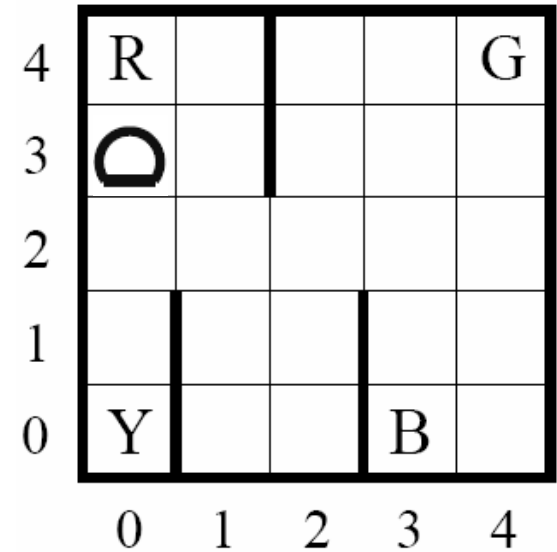
- Reinforcement learning:  
Learning how to act so as to maximize the expected cumulative value of a (numeric) reward signal
- In Soar terminology, RL learns operator comparison knowledge
- Uses:
  - Learning without model of operator effects
  - Learning in nondeterministic domains
  - Learning in non-goal-based tasks
- Limitation of old work:  
RL incompatible with Soar impasses

# RL and Soar impasses

- Learning impasses (i.e., tie impasse)
  - A place for RL? Interesting question.
  - Not what this talk is about
- Michigan-style goal-stack (op no-change)
  - Hierarchical task decomposition
    - Hierarchical RL
  - Hierarchical RL: RL with temporally-extended actions,
    - Actions of variable timespan
    - Actions whose expected reward, timespan, and next state depend on subtask policy
  - Impassed operator → temporally-extended action

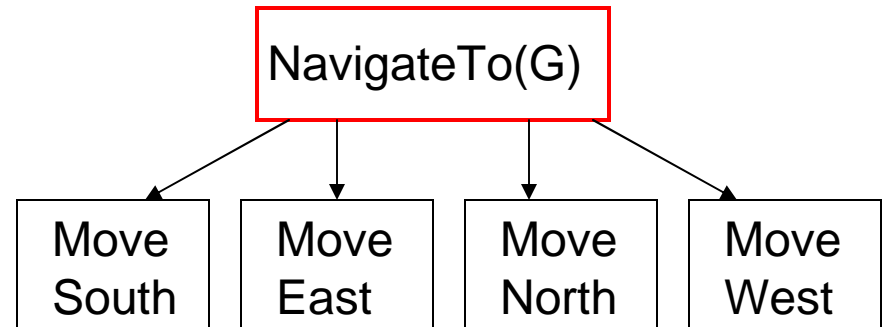
# Example Hierarchical Task: Taxi Domain

- A higher-level task: Navigate to **G**
- A lower-level task (primitive action): Move East



# Hierarchical RL:

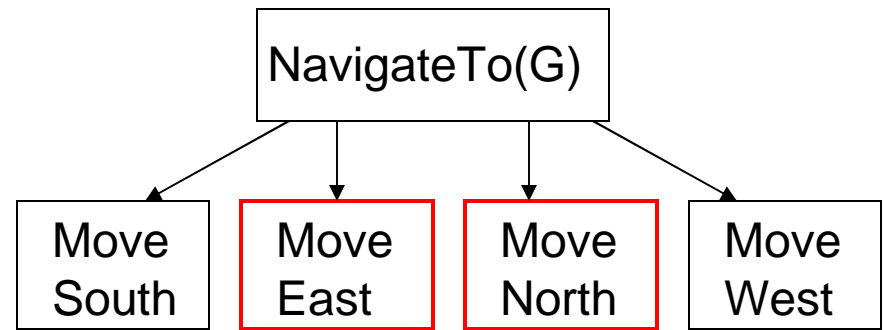
## Higher-level operator



- Higher-level operator (i.e., `NavigateTo(G)`)
  - Treated as temporally extended action.
  - Value includes:
    - Rewards received while operator selected
    - Value of next state, discounted by length of time operator was selected

$$\begin{array}{l}
 \mathbf{S1} \quad r_1 \quad + \lambda r_2 \quad + \lambda^2 r_3 \quad + \lambda^3 r_4 + \lambda^4 \text{Pred(Putdown)} \\
 \quad \text{NavTo(G)} \quad \text{NavTo(G)} \quad \text{NavTo(G)} \quad \text{NavTo(G)} \quad \text{Putdown} \\
 \hline
 \mathbf{S2} \quad \text{MoveEast} \quad \text{MoveNorth}
 \end{array}$$

# Hierarchical RL: Lower-level operator



- Lower-level operators (i.e., MoveEast)
  - Learning in subtask - try to divorce subtask value function from context
    - Subtask value function does not predict beyond end of subtask
    - Task-specific reward for subtask end state

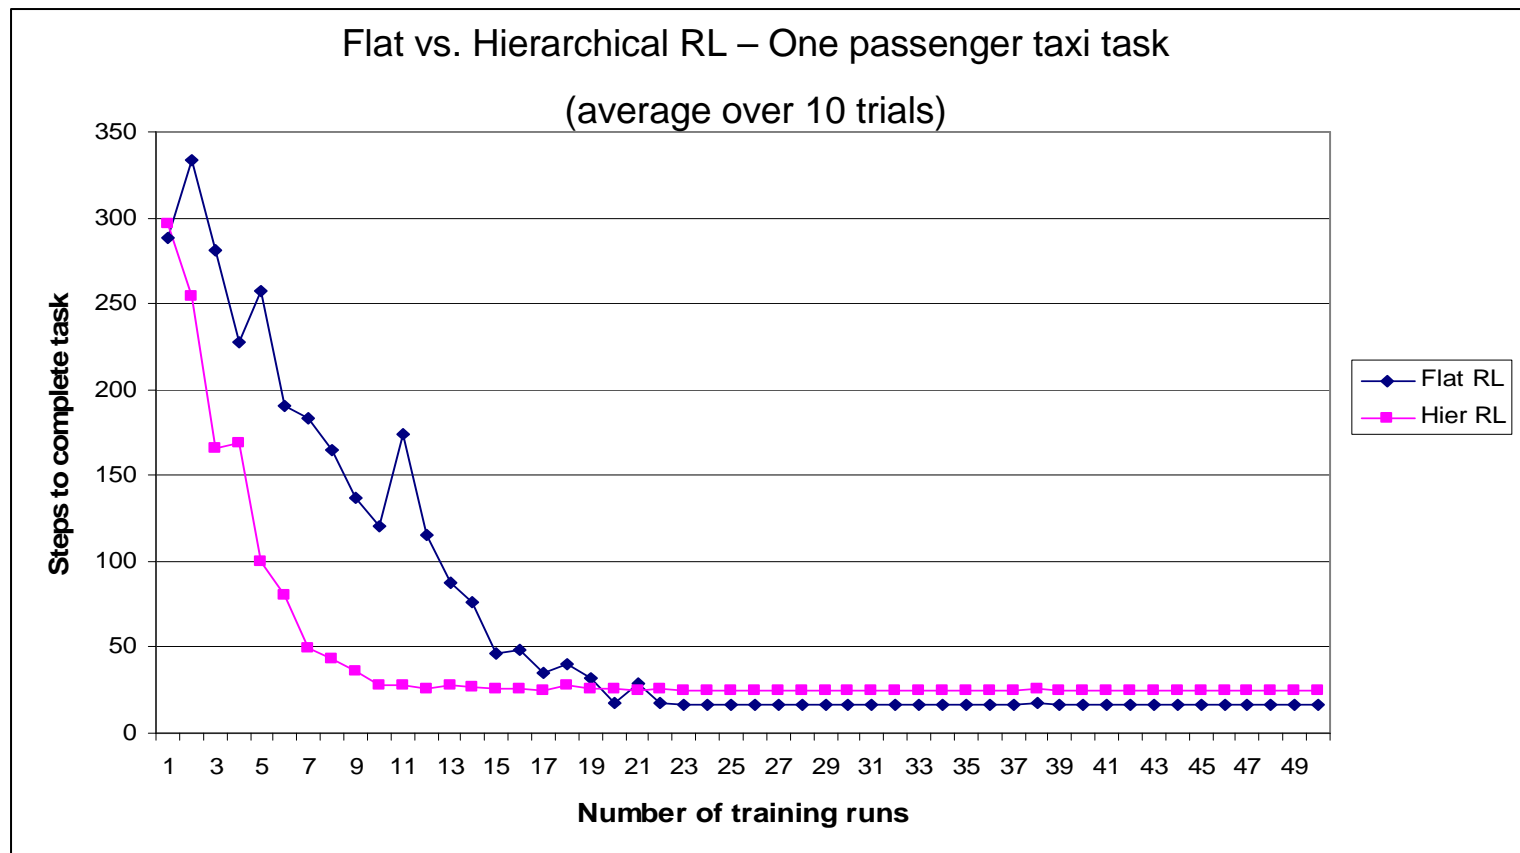
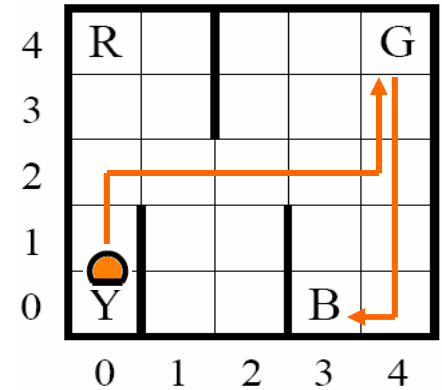
**S1** NavTo(G) NavTo(G) NavTo(G) NavTo(G) Putdown

---

**S2** MoveEast MoveNorth  
 $r_1 + \lambda * \text{Pred}(\text{MoveNorth})^{r_2}$

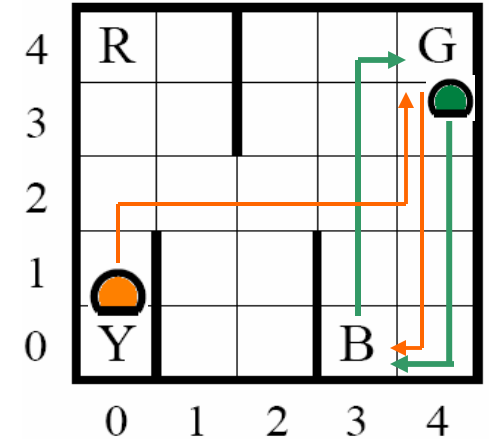
# What is gained?

- Faster learning in single task
  - More immediate feedback since rewards received by end of subtask

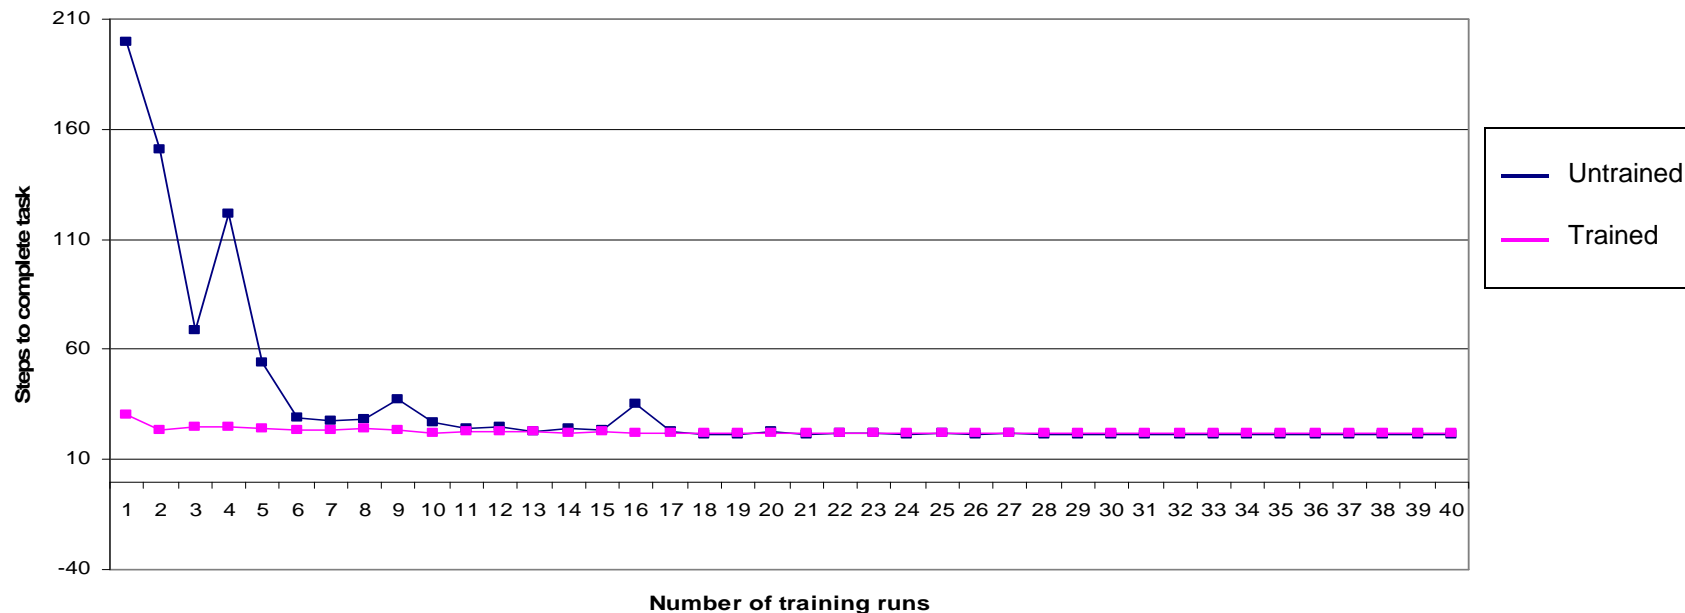


# What is gained?

- Transfer of subtask policy between higher-level tasks
  - Made possible by generalizing away higher-level context in subtask value functions



Untrained vs. Trained RL – Transfer one-passenger taxi task  
(average over 10 trials)





# Learning to choose among high-level operators

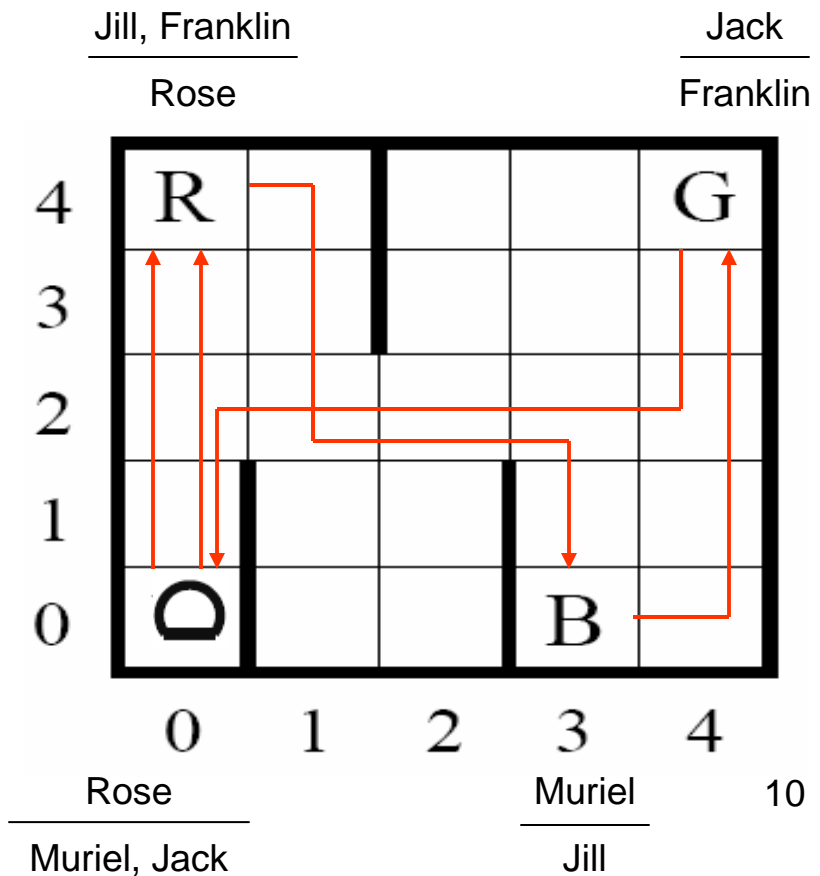
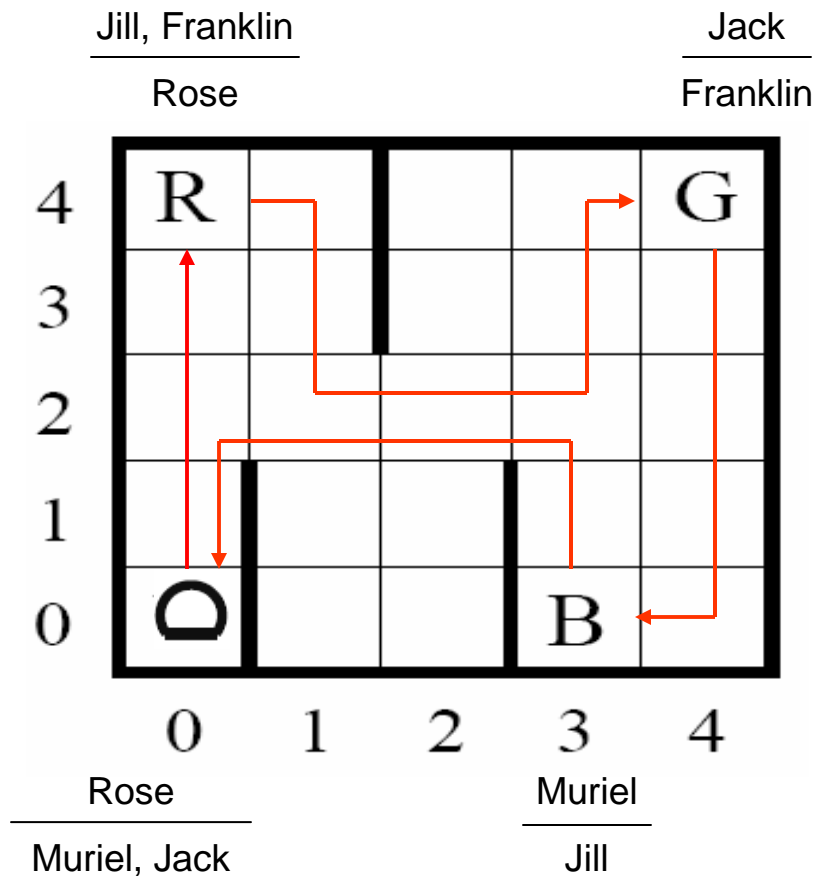
- The task:
  - 5 passengers with various locations, destinations
  - Max of 2 taxi occupants
  - Learn the shortest path
- Decision-making:
  - Using learned policies for navigate
  - Select among  $\text{Get}(\text{Pass}_i)$  &  $\text{Put}(\text{Pass}_i)$  tasks
- Unhappy result – Convergence to non-optimal policies

# 5-passenger taxi task results

- Optimal policy – 61 steps
- In 5 trials converged to policies of length 75, 87, 69, 86, and 78

Optimal

Best Achieved



# 5-passenger taxi task –

## What went wrong? Insufficient exploration

- The Good – Estimated values for the converged-upon policy accurate
- The Bad – Exploration:  
Experimental exploration strategy in which exploration probability decreases with visits to state
  - The Good – Automatically handles exploration → exploitation shift
    - Useful when different subtasks trained to varying degrees
    - Useful when different levels of hierarchy trained to varying degrees
  - The Bad – Lacks sufficient subtlety
    - Exploration in earlier states must continue longer than in later states
    - Exploration at higher levels of hierarchy must continue longer than at lower levels
- Exploration must improve before attempts to learn simultaneously at multiple levels of the hierarchy – which is possible in theory

# Nuggets & Coal

- Nuggets
  - Using Soar's operator no-change impasses for hierarchical task decomposition provides a nice structure for implementing established hierarchical RL ideas
- Coal
  - Relies on impasses for task decomposition
  - Exploration strategies insufficient
  - Soar-RL rewards difficult to use – particularly making rewards appear right before a subgoal retracts