



Transfer Learning Experiments with Soar & the UCT

Nicholas Gorski
with John Laird and Taylor Lafrinere

Outline

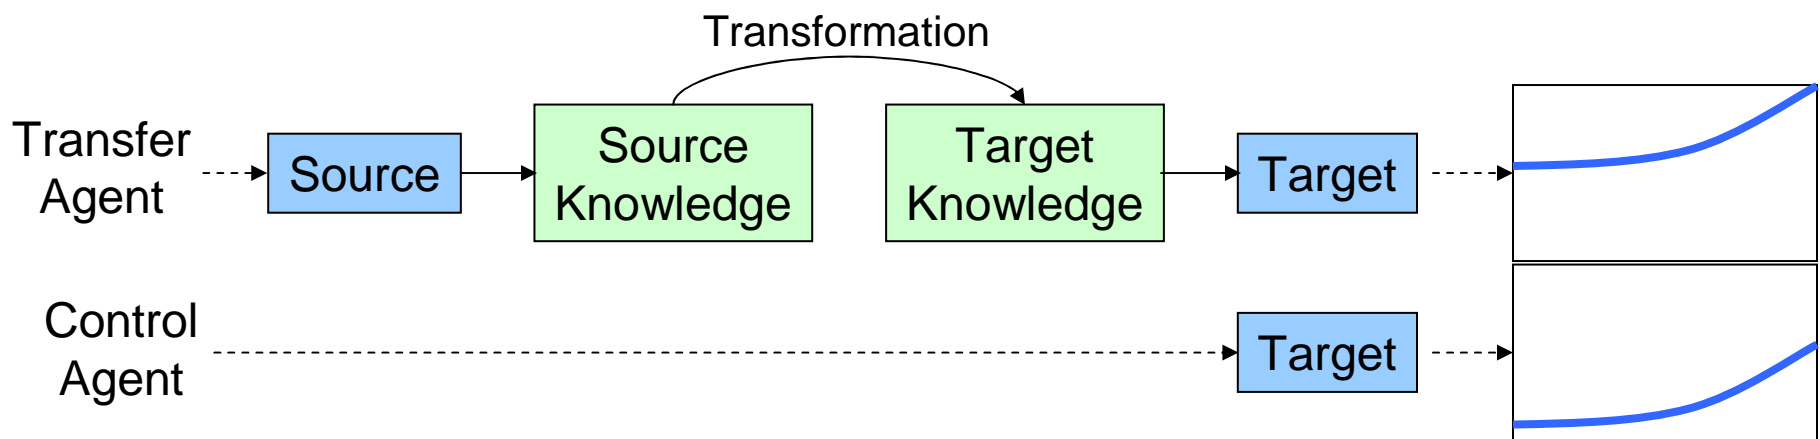
- Transfer Learning Background
- Urban Combat Testbed (UCT)
- Results
- Spatial Reasoning in the UCT
- Movies!

Project Background

- 3 year DARPA Transfer Learning (TL) initiative
- Soar grouped with ICARUS & Companion
 - Y1: Urban Combat Testbed (completed F06)
 - Y2: GGP (ongoing, evaluation in F07)
- Last workshop, reported some initial experiments

Transfer Learning (TL)

- Similarities to multi-task learning, inductive learning, and “learning to learn”
- Transfer Learning:
 - performs in source problem
 - applies learned knowledge to a target problem via transformation
 - performs in target problem, applying previously learned knowledge




Urban Combat Testbed

- Software suite consisting of
 - First Person Shooter video game engine
 - Scenarios designed to test for specific types of transfer
- Complex domain
 - Large and continuous
 - Noisy actions
 - Many different objects & obstacles
 - Doors, windows, barriers, pits, water, electrical barriers, etc.




TL Scenarios in the UCT



- Agent must navigate from start to goal
- To reach goal, it must climb ladder/drainpipe
- Generalization: drainpipe can be climbed because ladder was climbable

Y1 Results

Comparing Agent and Human Performance



Jump-start Discussion

■ Y1 results

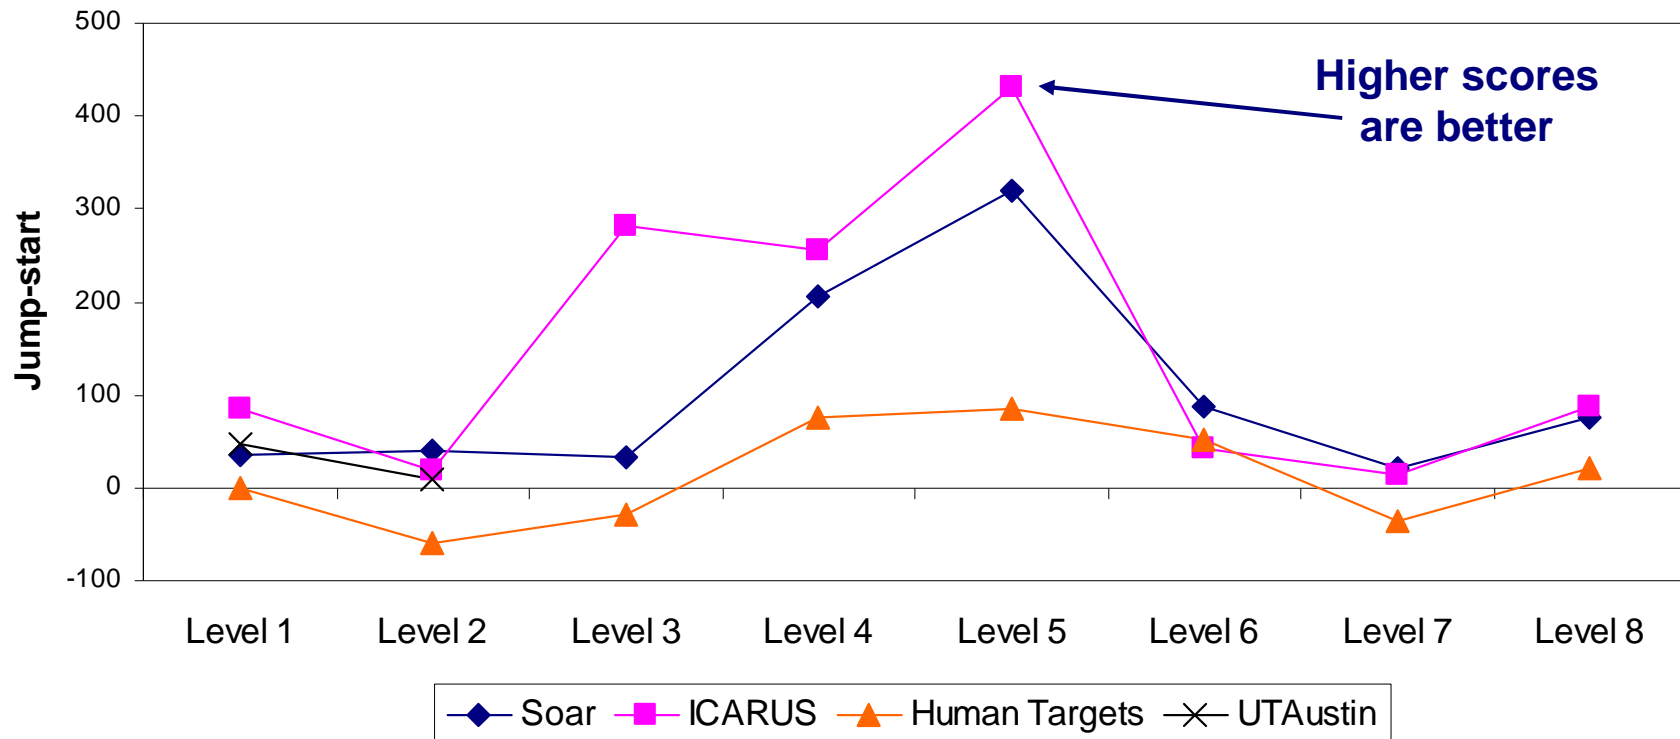
- Single methodology used by all teams
- Used for go/no-go decision for Y2

■ Used “Jump-start”

- Magnitude of initial differences in performance (not normalized)
- Rewards poorly performing agents

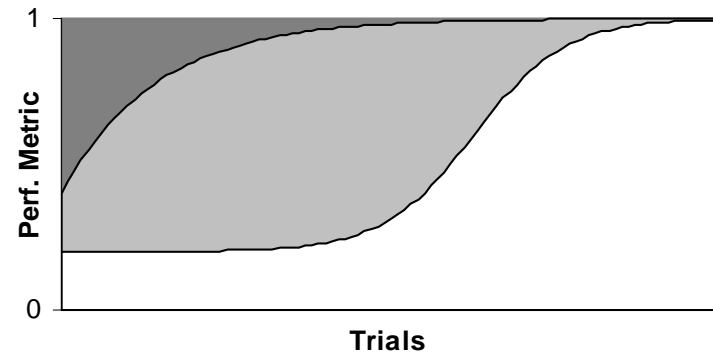
Y1 Results

Comparing Agent and Human Performance



Calibrated Transfer Ratio (CTR)


$$1 - \frac{\int_{t_0}^{t^*} OPT(t) - AB(t) dt}{\int_{t_0}^{t^*} OPT(t) - B(t) dt}$$




- Interpretation: “the amount of available improvement achieved”
- Disadvantage: requires knowledge of optimality
- Advantage: more meaningful

UM Evaluation: CTR vs Jump-start

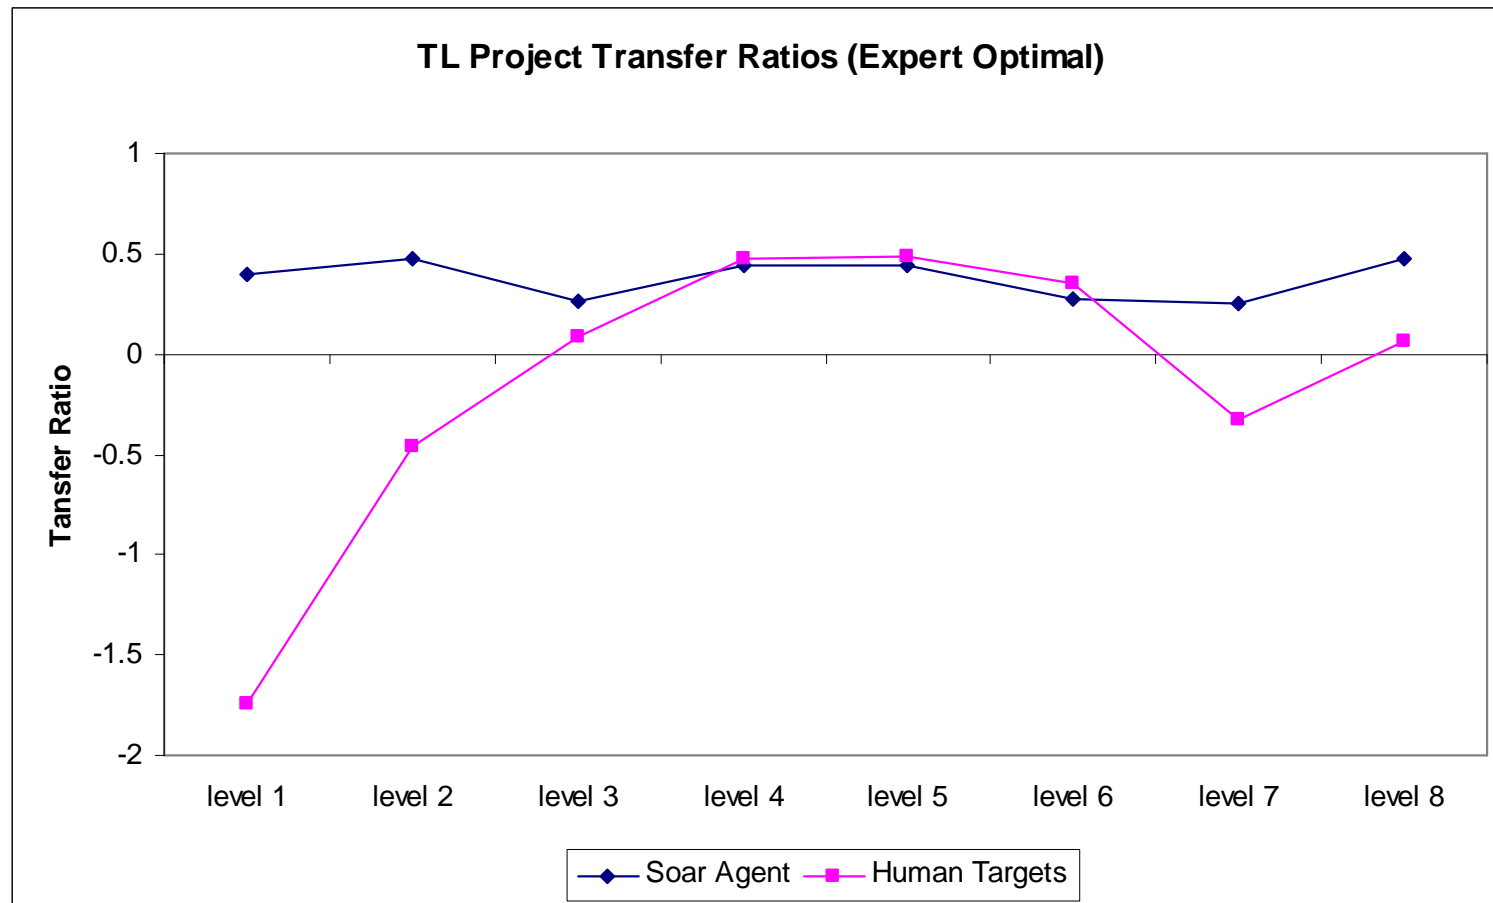
UCT Transfer Ratios of Soar Agent



UCT Jump-starts of Soar Agent



UM Evaluation: Soar vs Human



Outline

- Transfer Learning Background
- Urban Combat Testbed (UCT)
- Results
- **Spatial Reasoning in the UCT**
- Movies!

Navigation in UCT

- Agent perceives 3D space as set of positive and negative convex polyhedrons
 - Mapped to 2D convex polygons by SML middleware
 - “Gateways” are intersections of free space regions
- UCTBot navigates from region to region by:
 - Moving to a gateway
 - Moving through a gateway
 - Suboptimal navigation
 - Doesn't cut close to corners
 - Uses partitioning even when moving in wide open terrain


Obstacle Detection & Avoidance

- UCTBot is “blind” to obstacles and some gateways
- Detection: relies on velocity
- Avoidance:
 - Some obstacles can be surmounted
 - Test all available actions
 - Robust for most obstacle types
 - For blocking obstacles, find paths around them
 - (Mostly) robust
- Learning: which obstacle/gateway is blocking?


(Partially) Motivated SRS

- Better navigation improves performance
 - In untrained agents
 - In transfer agents
 - To honestly evaluate TL, must optimize both transfer & control cases
- Possible applications:
 - Route finding, obstacle avoidance
 - But would also allow for multi-agent tasks

Learning to climb



Searching indoors



Using weapons



Gold Nuggets & Lumps of Coal

- Met Y1 goals
- Laid groundwork for more motivating TL experiments in Y2
- Motivated SRS
- Developed CTR
- UCT cut from Y2
- Scenarios and domain lacked motivating transfer
- Didn't get to use SRS
- CTR not adopted for internal evaluations