

ADAPT: A Cognitive Architecture for Robots

D. Paul Benjamin

Pace University Robotics Laboratory

Deryle Lonsdale

Brigham Young University Department of Linguistics
and English Language

Damian Lyons

Fordham University Robotics Laboratory

ADAPT

Adaptive Dynamics and Active Perception for Thought

The goal of the ADAPT project is to create a robot that can model its environment accurately in real time, and use that model to perform tasks and interact with people using natural language.

We are not interested in robot programming for particular tasks but in investigating embodying cognition.

The emphasis in ADAPT is on solving problems by reformulation (reperception).

The structure of the ADAPT architecture is based on linguistics.

We are building ADAPT by implementing the RS (Robot Schemas) language in the Soar cognitive architecture.

ADAPT uses a sophisticated, multimedia internal world model. Comprehension is modeled as search to reconstruct the environment within this world model.

Overview

Overall structure and philosophy

RS/Soar

Coherence Theory

Virtual World

Visualization and Semantics

Lessons Learned the Hard Way:

Everything is sensory-motor.

Perception is an active, goal-directed process.

Robotics requires a high degree of true concurrency.

Analyzing interactions is just too hard.

Active Perception

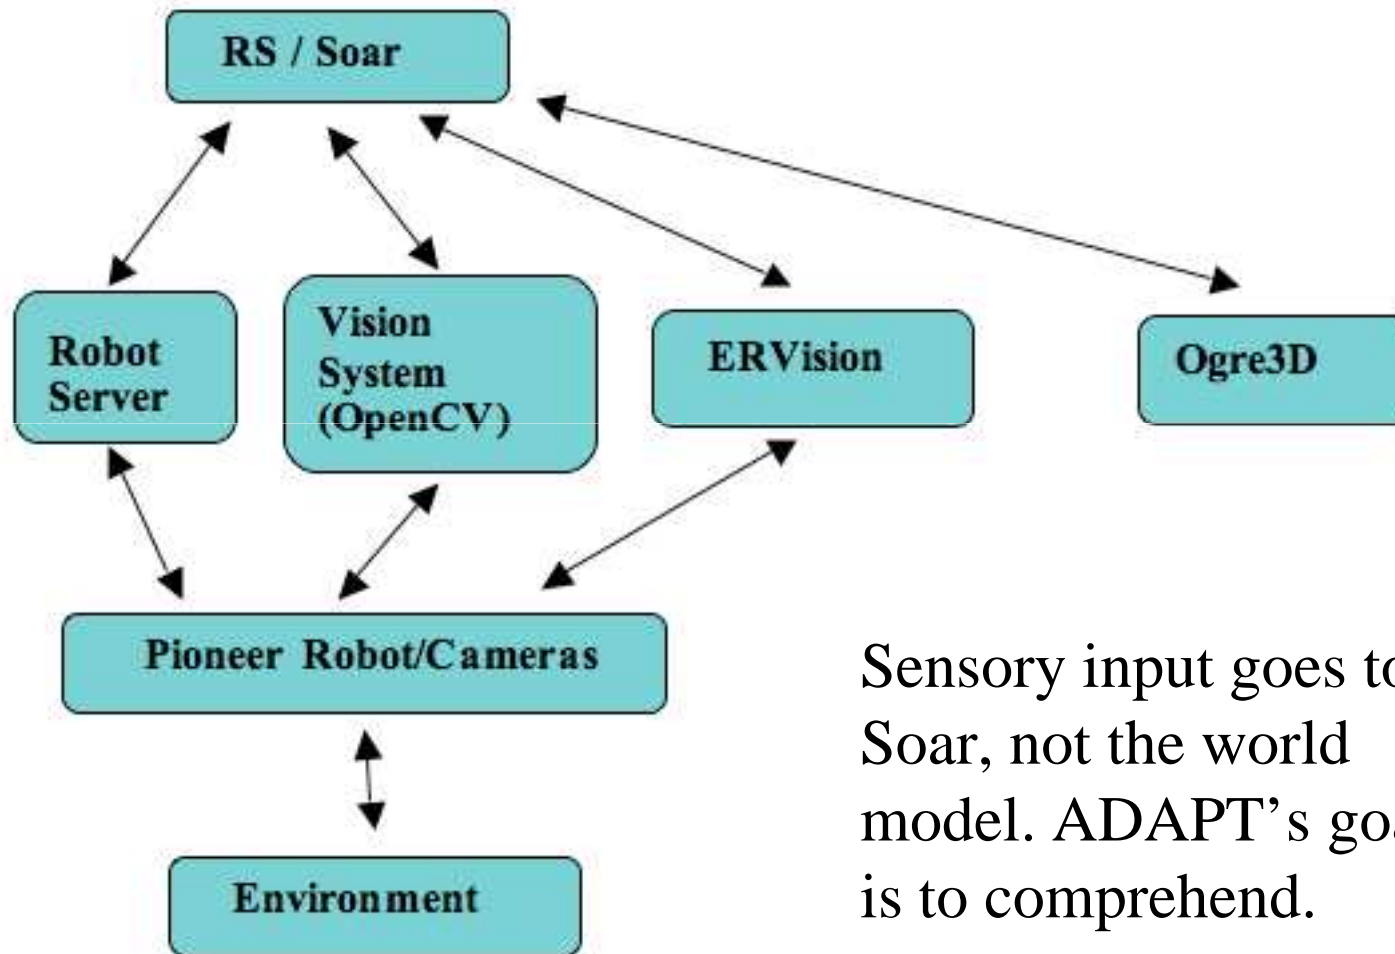
Active perception is top-down and goal-directed, so that perception becomes a problem solving process. This is in contrast to the way perception is usually approached in AI and cognitive science. Cognitive robotics means more than just using a cognitive architecture on a robot; it means treating robotics as a cognitive domain, which includes treating perception as problem solving.

Perception is the hard problem in robotics. This distinguishes robotics from tasks like chess.

We view perception as the representation problem. Perception is the construction and modification of problem-solving representations.

The representation problem is at least hyperexponential ($O^{2^{2^n}}$).

ADAPT's Structure



Sensory input goes to Soar, not the world model. ADAPT's goal is to comprehend.

ADAPT integrates distributed, concurrent control with cognitive structures

The RS (Robot Schemas) language is the basis of the robotics capabilities of ADAPT. RS is precise and mature.

RS is a CSP-type programming language for robotics, that controls a hierarchy of concurrently executing schemas.

$$\text{Joint}_i(s)() = [\text{Jpos}_i()(\mathbf{x}), \text{Jset}_i(s, \mathbf{x})(\mathbf{u}), \text{Jmot}_i(\mathbf{u})()]^{c0}$$

$$c0: \quad (\text{Jpos}_i, \mathbf{x}) \quad (\text{Jset}, \mathbf{x}) \quad (\text{Jset}, \mathbf{u}) \quad (\text{Jmot}_i, \mathbf{u})$$

$\text{Jpos}_i()(\mathbf{x})$ continuously reports the position of joint i on port \mathbf{x}

$\text{Jmot}_i(\mathbf{u})()$ accepts a signal on port \mathbf{u} and applies it to the actuator of joint i

$\text{Jset}_i(s, \mathbf{x})(\mathbf{u})$ accepts a setpoint on port s and iteratively inputs a joint position on port \mathbf{x} and outputs a motor signal on port \mathbf{u} to drive the joint position to the setpoint

A Sensory-motor Schema Hierarchy

$$\text{Touch}_i = [\text{Tact}_i()(\nu), \text{Gmove}_i(\nu)(y), \text{Joint}_i(y)()]^{c1}$$

$$c1: \quad (\text{Tact}_i, \nu) \quad (\text{Gmove}, \nu) \quad (\text{Gmove}, y) \quad (\text{Joint}_i, y)$$

Tact_i reports on tactile contact on the i -th join on its port ν .

Gmove increments the setpoint of the joint actuator as long as it gets a no-contact signal on its port ν .

Touch_i implements a guarded move of the i -th link.

RS has a Formal Semantics

$P = (Q, L, X, \delta, \beta, \tau)$ where

Q is the set of states

L is the set of ports

$X = (X_i / i \in L)$ is the event alphabet for each port

$XL = \{ (i, X_i) / i \in L \}$ i.e., a disjoint union of L and X

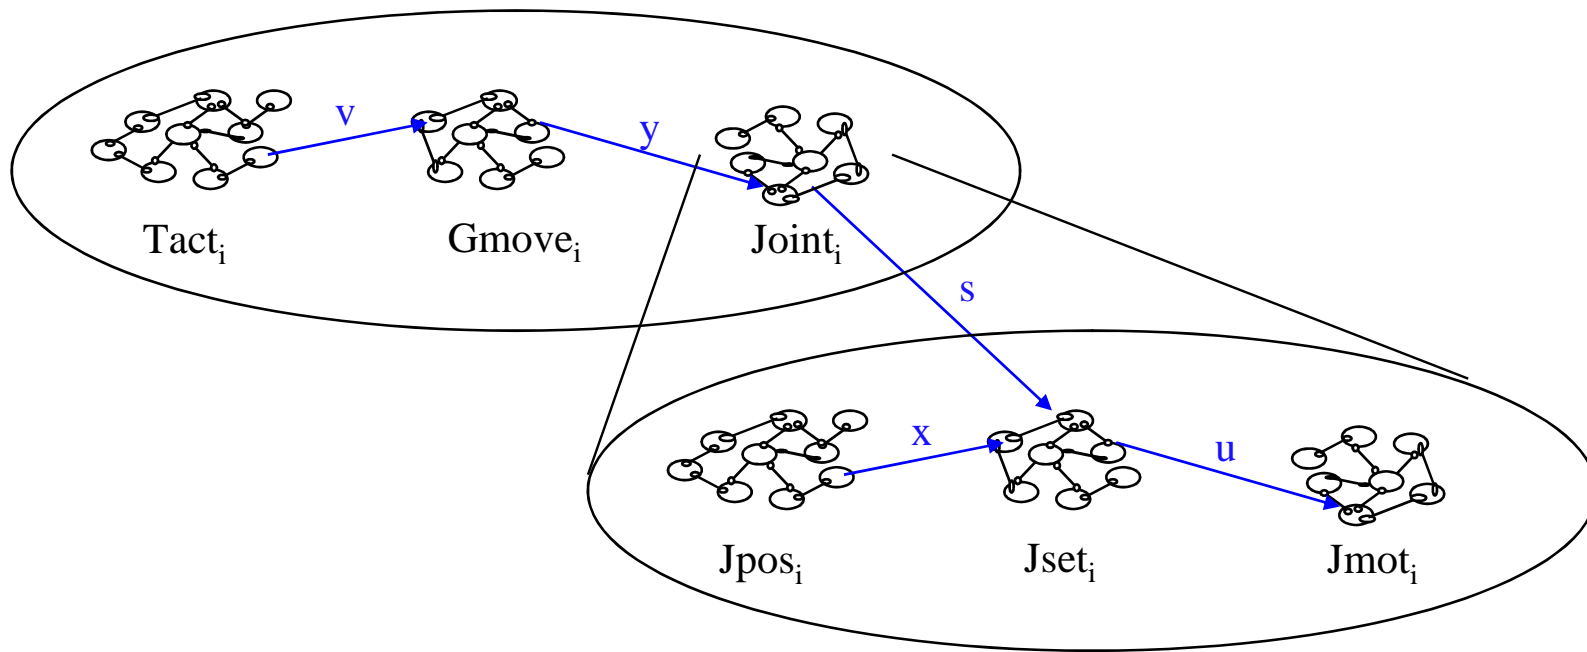
$\delta: Q \times XL \rightarrow 2^Q$ is the transition function

$\beta = (\beta_i / i \in L)$ $\beta_i: Q \rightarrow X_i$ is the output map for port i

$\tau \in 2^Q$ is the set of start states

RS has a Formal Semantics

The behavior of every RS schema is defined using port automata. This provides precision to the semantics and also a constructive means of reasoning about the behavior and meaning of schemas.



Process Composition in RS


1. Sequential Composition: $T = P;Q$. The process T behaves like the process P until that terminates, and then behaves like the process Q (regardless of P 's termination status).
2. Concurrent Composition: $T = (P \mid Q)^c$. The process T behaves like P and Q running in parallel and with the input ports of one connected to the output ports of the other as indicated by the port-to-port connection map c . This can also be written as $T = \left(\prod_{i \in I} P_i \right)^c$ for a set of processes indexed by I .
3. Conditional Composition: $T = P\langle v \rangle : Q_v$. The process T behaves like the process P until that terminates. If P aborts, then T aborts. If P terminates normally, then the value v calculated by P is used to initialize the process Q , and T then behaves like Q_v .
4. Disabling Composition: $T = P\#Q$. The process T behaves like the concurrent composition of P and Q until either terminates, then the other is aborted and T terminates. At most one process can stop; the remainder are aborted.
5. Synchronous Recurrent Composition: $T = P\langle v \rangle :: Q_v$. This is a recursively defined as follows:
 $P :: Q = P : (Q; P :: Q)$.
6. Asynchronous Recurrent Composition: $T = P\langle v \rangle :: Q_v$. This is recursively defined as follows:
 $P :: Q = P : (Q \mid (P :: Q))$.

Operator Precedence: The operator precedence from loosest to tightest is as follows: Concurrent; Disabling; Sequential; Conditional; Synchronous Recurrent; Asynchronous Recurrent.

Implementation of RS in Soar

Schemas, facts, and hypotheses are nodes in a graph. Links implement the composition operations, as well as other relations, including deductive and evidential inference.

Automata that implement a schema are built as needed.



Advantages of RS

Formal Semantics

Complete Representation of Distributed Control

Maturity

Invariants to monitor tasks/environment

Disadvantages of RS

No synthesis method

No cognitive plausibility

No learning

Goal: to use chunking to learn RS schemas

RS has been successfully used in factory automation.

QuickTime™ and a
Microsoft Video 1 decompressor
are needed to see this picture.

However, this was a hand coded planner using Allen's
interval logic.

Coherence Theory

A Method of abductive inference for evidential reasoning with a claim to psychological plausibility

Represents evidence, facts and hypotheses in a constraint network. Each of these can be incomplete and/or inaccurate.

Searches for a set of accepted (believed) hypotheses that satisfies the maximal weighted sum of constraints

Used by Johnson et al.: (unfinished)

Thagard: scientific discovery, jury decisions, problem solving

Ranney & Thagard: students solving physics problems

Mead & Miller: perception of social relationships

Implementation of Coherence Theory

Rather than use a connectionist approach, ADAPT implements coherence theory in Soar.

Evidence, facts and hypotheses are connected by constraints of varying types (deductive or evidential).

A model checker proves some of the hypotheses to be true/false. True hypotheses are *Accepted*; false ones *Rejected*. Other hypotheses are randomly assigned initially.

Operators compute the amount of constraint satisfaction to be gained by changing each hypothesis to the other set. The hypothesis with the highest gain is switched. Repeat until local maximum reached.

Fast: typically under 250 ms.

Experience in Cybersecurity Domain

DARPA-funded project with BBN Technologies and Adventium Labs.

Goal: To defend a network of > 30 hosts from insider threats, catching at least 50% of attacks with no more than 10% false positives, and response in less than 250 ms.

Computer network is simulated in JESS.

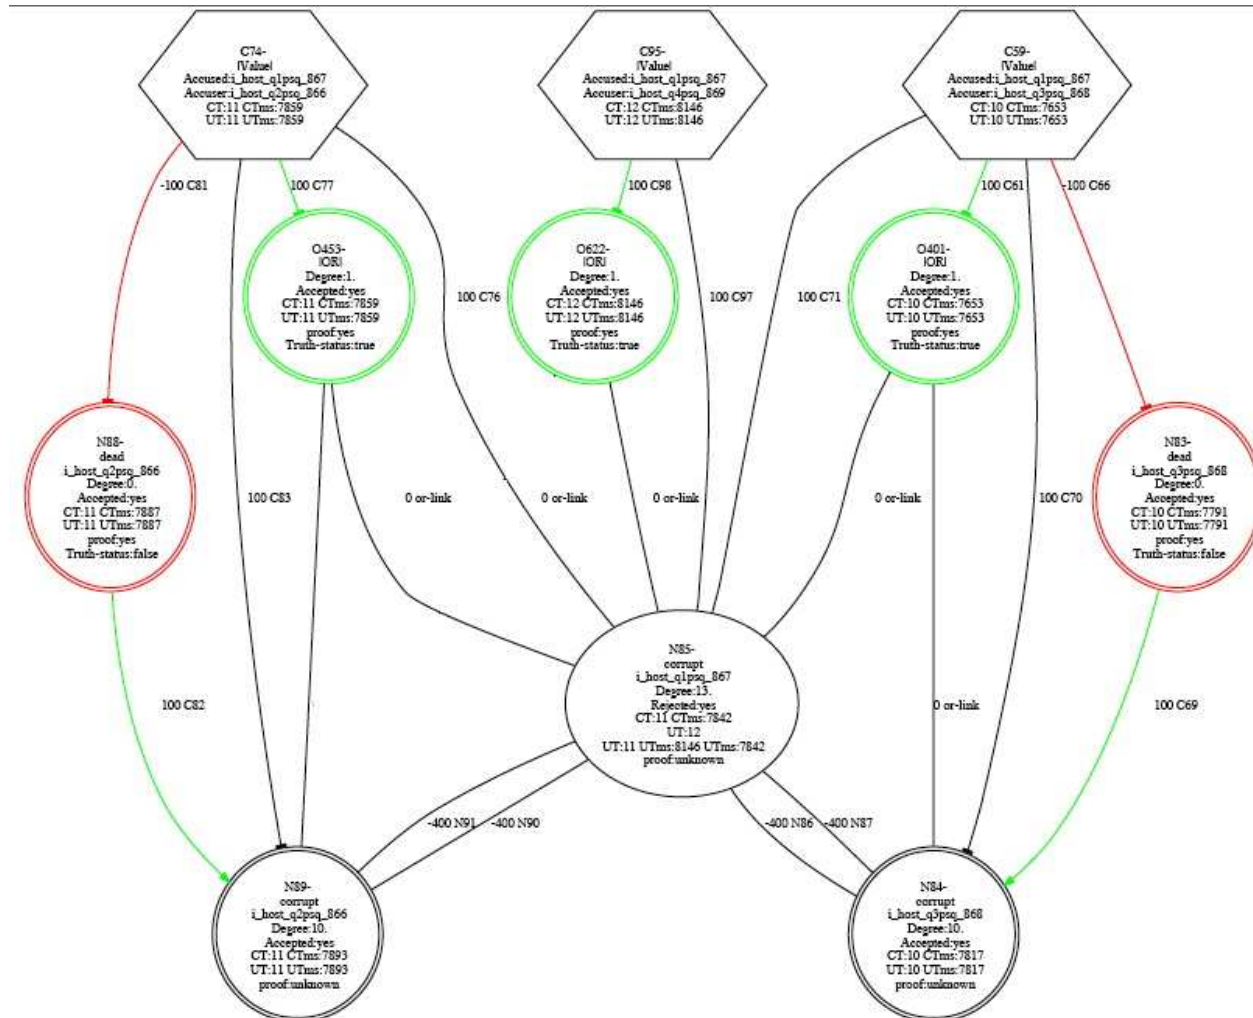
Successfully handles all scenarios from 2005 tests, as well as tests from a problem generator.

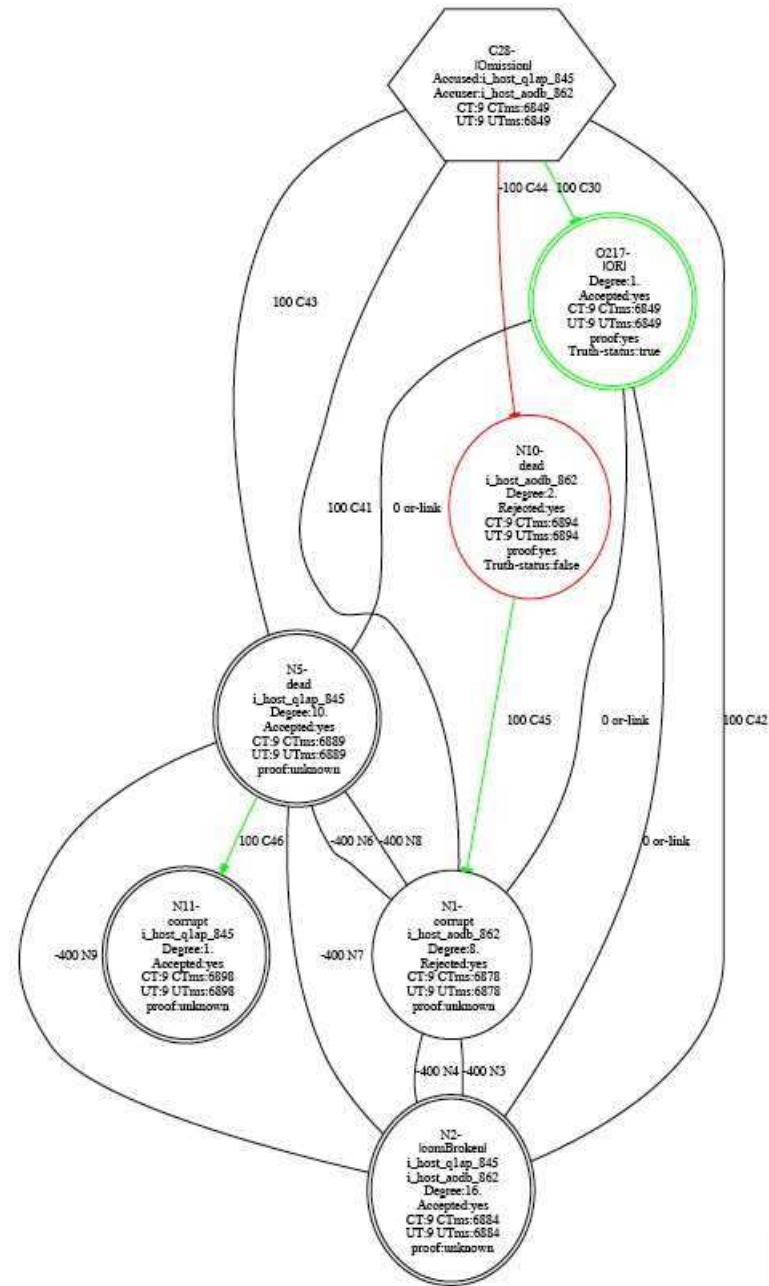
Fast and accurate. Hypothesis networks often > 100 nodes, sometimes > 500 nodes.

Red Team test to occur in late May, 2008.

Cybersecurity Example Hypothesis Network

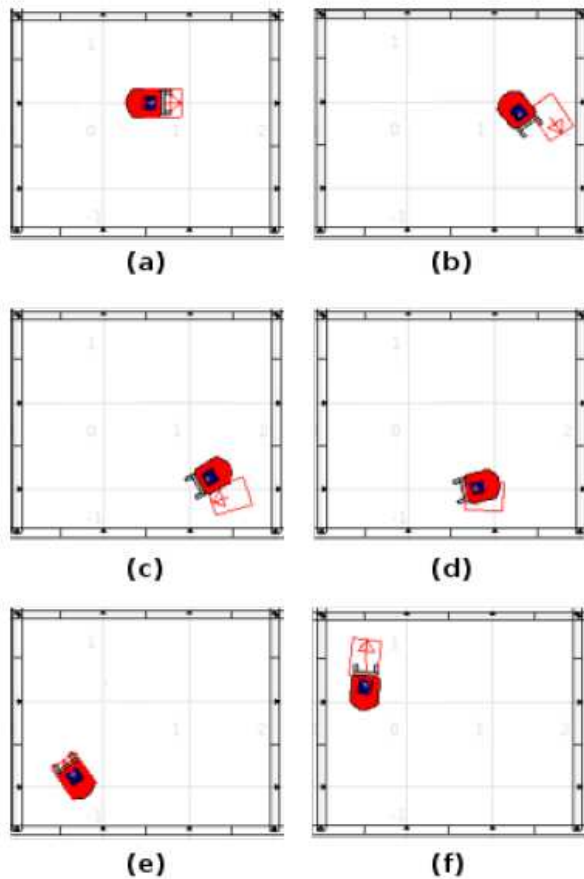
Run #11 from November, 2005. The hypothesis network has two components:





SoarSLAM Example

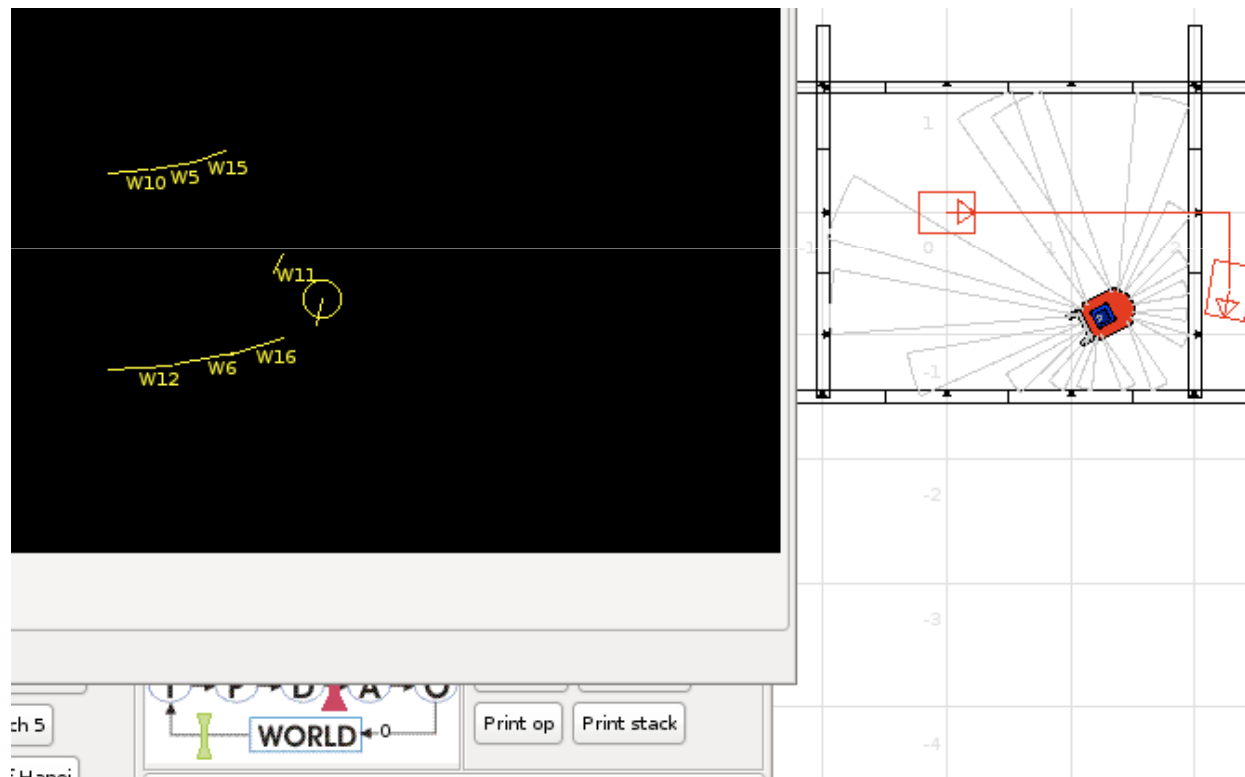
Sonar SLAM written in Soar. Runs on the Pioneer robots and simulator. Successfully maps floors of office buildings.



Dead reckoning errors accumulate. The empty red rectangle shows where the robot thinks it is.

SoarSLAM Example

Sonar SLAM creates local maps consisting of wall segments. Hypotheses are the robot's position relative to local maps.



SoarSLAM Example

SoarSLAM switches local maps and resets the odometry.

The screenshot displays the Soar Debugger interface for a remote SoarRobot. The main window is divided into several panels:

- Log Panel (Left):** Shows a sequence of actions and observations. Key events include:
 - 2032: 0: 02025 (turn-right)
 - 2033: ==>s: s14019 (operator tie)
 - 2034: 0: 02026 (remove-action)
 - 2035: 0: 02027 (move-forward)
 - 2036: ==>s: s14074 (state no-change)
 - 2037: 0: 02031 (|ChangeRobotPosition2|)
 - 2038: 0: 02033 (|ChangeRobotPosition2|)
 - 2039: ==>s: s14130 (state no-change)
- Match Rule Panel (Top Right):** Displays a match rule for "matches slam*monitor*wall*data". The rule is:

```
1 (state <s> ^superstate nil)
>>> {<s> ^|wallData| <w*1>}
      {<w*1> ^|Wall| <uw>}
      {<uw> ^|Time| <t>}
      {<uw> ^|Length| <l>}
      {<uw> ^|R| <r>}
      {<uw> ^|Th| <th>}
      {<uw> ^|Ref| <ref>}
      {<ref> ^|Th| <rth>}
      {<uw> ^|End| <end>}
      {<ref> ^|V| <ry>}
      {<end> ^|V| <ey>}
      {<ref> ^|X| <rx>}
```
- Stack Panel (Middle Right):** Shows the current stack with the entry "p --stack" and a value "s1".
- Map Panel (Right):** A 2D grid-based map showing the robot's current position (a red square) and its field of view (a red cone). The map includes walls and a path.
- Control Panel (Bottom):** Contains buttons for "Step", "Run", "Run 1 -p", "Stop", "Matches", "Print <s>", "Print <ts>", "Clear", "Watch 1", "Watch 3", "Watch 5", "Init-soar", "Source", "cd", "Excise all", and "Towers of Hanoi". It also features a state transition diagram with nodes I, P, D, A, O and a "WORLD" node.

Lesson and Question

This simple example shows that robotics algorithms can be realized in a cognitively plausible manner within a cognitive architecture.

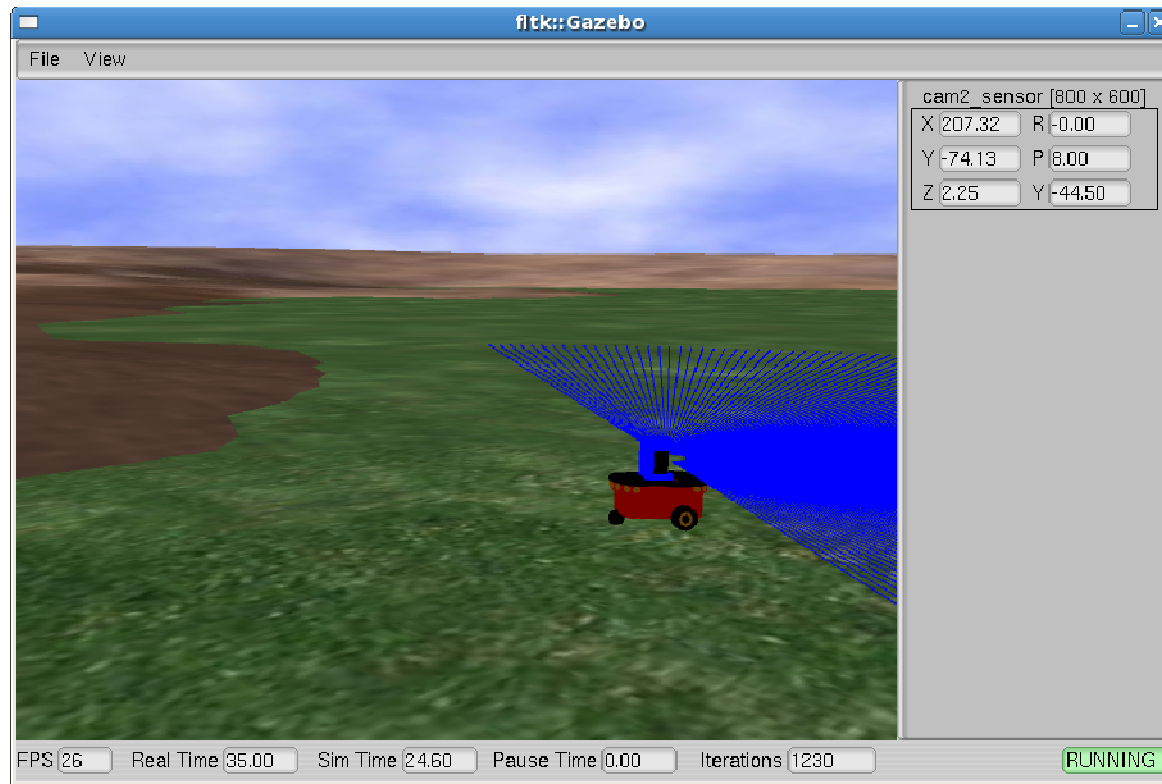
Is this cognitive science?

The basic loop of ADAPT is:

- 1 - check Soar's output link to see if there are any commands, which may be either motion commands for the robot or modeling commands for the World Model,
- 2 - blend the motion commands that are to be sent to the robot,
- 3 - send all robot commands both to the robot and to the virtual robot in the World Model,
- 4 - send all other commands to the World Model,
- 5 – periodically (every tenth of a second) fetch data from the robot to be put into Soar's working memory,
- 6 - periodically fetch data from the Vision System, compare it to visual data from the World Model, and put any significant differences into Soar's working memory.

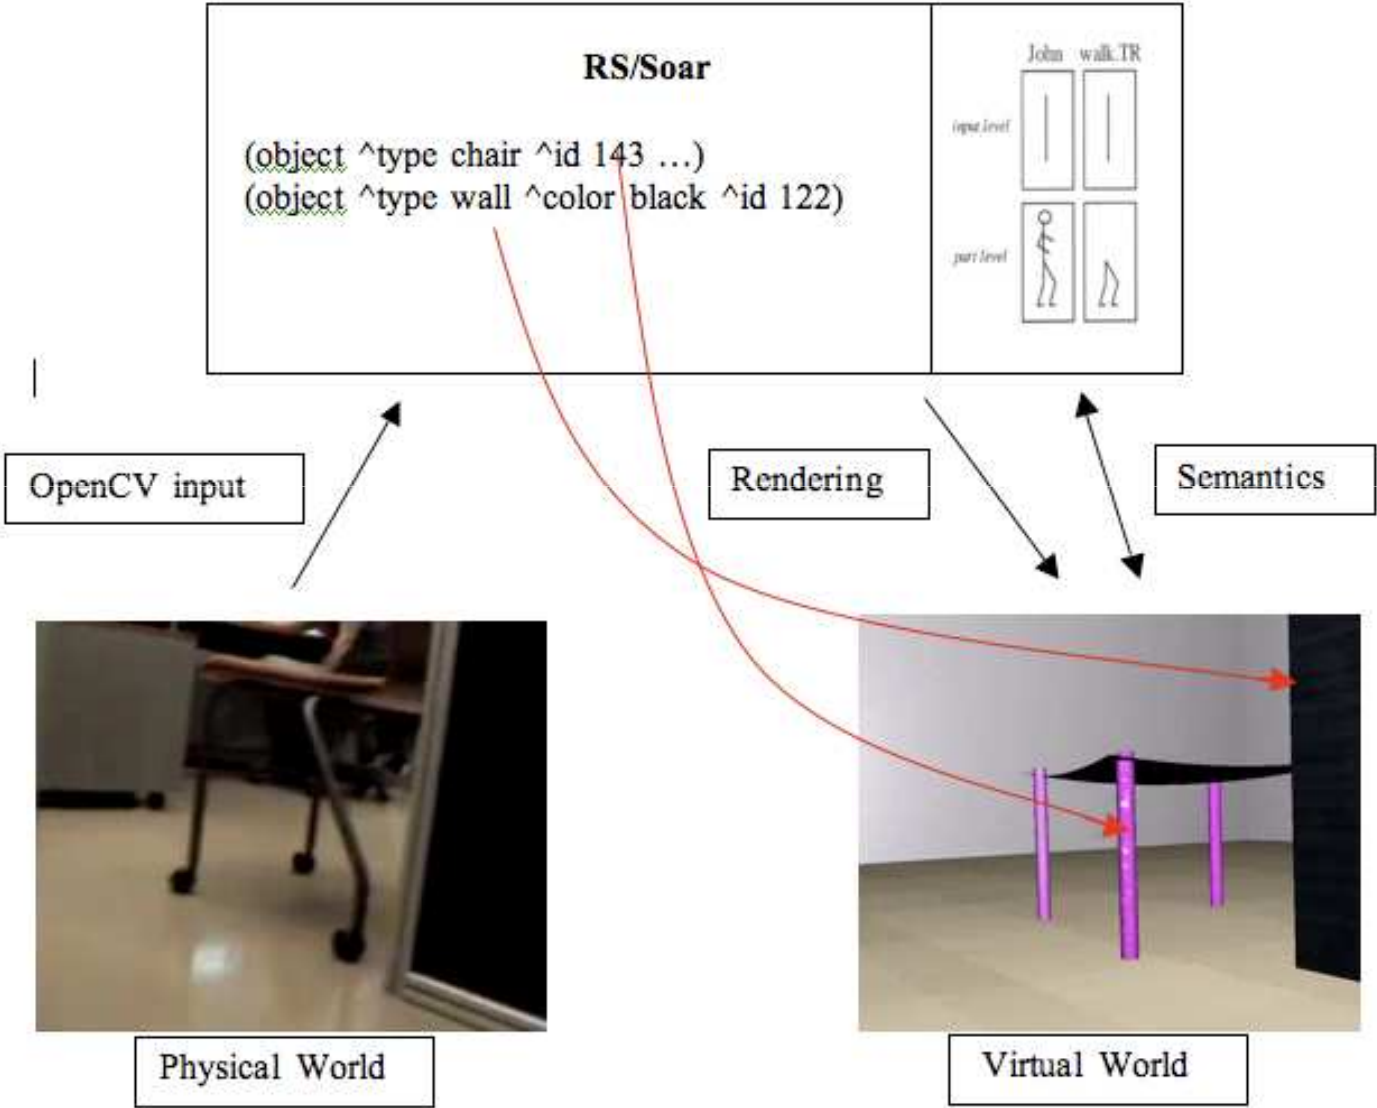
ADAPT's Mental Model

Ogre3D Video game platform: physics and graphics



Goal: to create a working copy of the environment

ADAPT's Mental Model



Videos showing the Modeling Process

QuickTime™ and a
H.263 decompressor
are needed to see this picture.

QuickTime™ and a
YUV420 codec decompressor
are needed to see this picture.

Virtual World

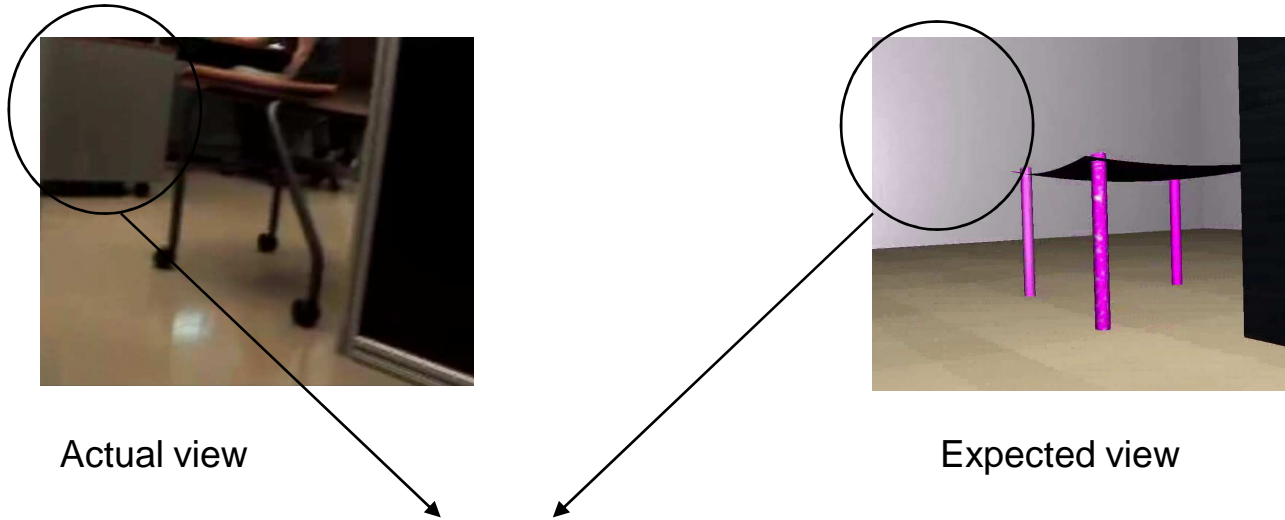
Reconstruction currently uses a hand-built library of objects and schemas.

Ogre is used for:

- motion planning in dynamic environments
- predictive vision
- language comprehension.

ADAPT searches to reconstruct the environment. This approach is used for recall, instead of data chunking.

Predictive Vision

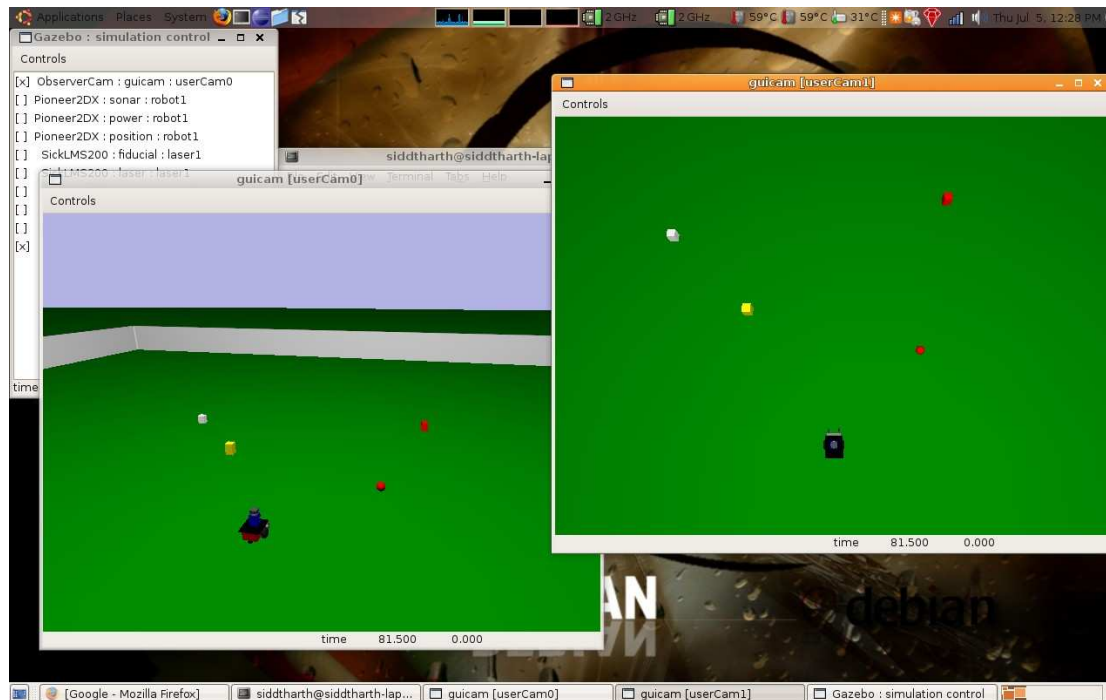


Sum-of-squares difference exceeds threshold.
Soar operator proposed to focus on difference.

Expensive vision operations are goal-dependent,
greatly reducing their frequency of use.

Visualization for Semantics

Comprehension requires visualization.



Observer camera

Neighborhood camera

Visual Context



The Neighborhood camera moves with the robot and defines the visual context.

The circle defines “near”; it is movable
But of a fixed size.

If the robot is told to move near the small red block, it will plan motions to take it inside the black circle.


The Visual Context is Associated with the Task



If the robot is then told to pick up the small red block, the new task changes the context.

The new context is smaller, causing the camera to zoom in.

Changing the Visual Context

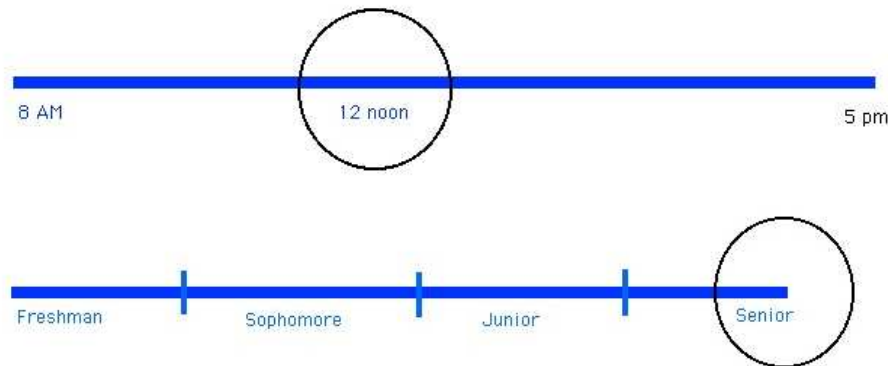


In the new context, the robot is no longer near the small red block, because although “near” is visually the same neighborhood, it denotes a much smaller region.

If the robot is told to pick up the small red block, it must move inside the black circle, because picking up something requires being closer than moving near it.

Tasks Determine Contexts

Visualization uses the virtual camera to define a context within which terms have a single meaning. This is intended to fit NL-Soar's comprehension mechanism.

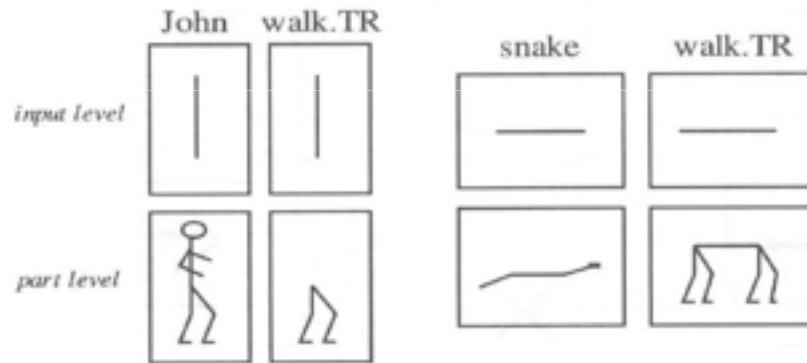


ADAPT must search among contexts instead of among meanings for terms.

This approach could conceivably be used for terms that are not inherently physically grounded.

Cognitive Semantics

The idea of using visualization for linguistic semantics is not new. Langacker's Cognitive Grammar described a way to do this.



Holmqvist partially implemented this grammar, but didn't finish.

Our approach improves the methodology because the virtual world is 3D and dynamic.

Cognitive Plausibility and Soar Theory

Duncker: “Problem solving consists of a sequence of phases; each phase is a reformulation of the problem (Newell).”

Polk & Newell: “We propose that the central processes in deductive reasoning are linguistic (encoding, reencoding, and generation) rather than reasoning-specific skills.”

“... for deduction tasks for which the necessary information is provided verbally, the heart of deduction involves repeatedly reencoding the problem, ...”

Status

Individual capabilities have been demonstrated: **NUGGETS!**

RS in Soar for basic navigation

Coherence theory in navigation, cybersecurity

Modeling vision data in Ogre in real time

Use of visual contexts for semantics in navigation

Predictive vision attends to and models changes

But: ADAPT is still in three big pieces **COAL**

Everything is hand built

RS/Soar, vision, NL-Soar: what are the constraints?

A big problem is the virtual world software

Thanks!

References

ADAPT:

Benjamin, D. Paul, Damian Lyons and Deryle Lonsdale, "Embodying a Cognitive Model in a Mobile Robot", Proceedings of the SPIE Conference on Intelligent Robots and Computer Vision, Boston, October, 2006.

Benjamin, D. Paul, Damian Lyons and Thomas Achtemichuk, "Obstacle Avoidance using Predictive Vision based on a Dynamic 3D World Model", Proceedings of the SPIE Conference on Intelligent Robots and Computer Vision, Boston, October, 2006.

Benjamin, D. Paul, Damian Lyons and Deryle Lonsdale, "Designing a Robot Cognitive Architecture with Concurrency and Active Perception", Proceedings of the AAAI Fall Symposium on the Intersection of Cognitive Science and Robotics, Washington, D.C., October, 2004.

RS:

Lyons, D.M. and Hendriks, A., "Exploiting Patterns of Interaction to Select Reactions", Special Issue on Computational Theories of Interaction, Artificial Intelligence **73**, 1995, pp.117-148.

Lyons, D.M. and Arbib, M.A., "A Formal Model of Computation for Sensory-based Robotics", IEEE Transactions on Robotics and Automation **5**(3), Jun. 1989.

Lyons, D., and Arkin, R.C., "Towards Performance Guarantees for Emergent Behavior", (Submitted) IEEE Int. Conf. on Robotics and Automation, New Orleans LA, April 2004.

References

Port Automata:

Martha Steenstrup, Michael A. Arbib, Ernest G. Manes, *Port Automata and the Algebra of Concurrent Processes*. JCSS 27(1): 29-50, 1983.

Cognitive Semantics:

Holmqvist, K., “Implementing Cognitive Semantics”, Lund: Department of Cognitive Science, 1993.

Holmqvist, K., “Conceptual Engineering”, in *Cognitive Semantics: Meaning and Cognition*, Allwood and Gardenfors (Eds.), John Benjamins, pp.153-171, 1999.

Langacker, R., “Foundations of Cognitive Grammar, Vol. I”, Stanford University Press, 1987.

Langacker, R., “Foundations of Cognitive Grammar, Vol. II”, Stanford University Press, 1991.

Langacker, R., “Concept, Image, and Symbol”, Berlin, New York, Mouton de Gruyter, 1991.

References

Reformulation:

Benjamin, D. Paul, “On the Emergence of Intelligent Global Behaviors from Simple Local Actions”, *Journal of Systems Science*, special issue: Emergent Properties of Complex Systems, Vol. 31, No. 7, 2000, 861-872.

Benjamin, D. Paul, “Formulating Patterns in Problem Solving”, *Annals of Mathematics and AI*, **10**, pp.1-23, 1994.

Benjamin, D. Paul, “Reformulating Path Planning Problems by Task-preserving Abstraction”, *Journal of Robotics and Autonomous Systems*, **9**, pp.1-9, 1992.