

# The SVS Spatial/Visual Reasoning System

Samuel Wintermute  
University of Michigan

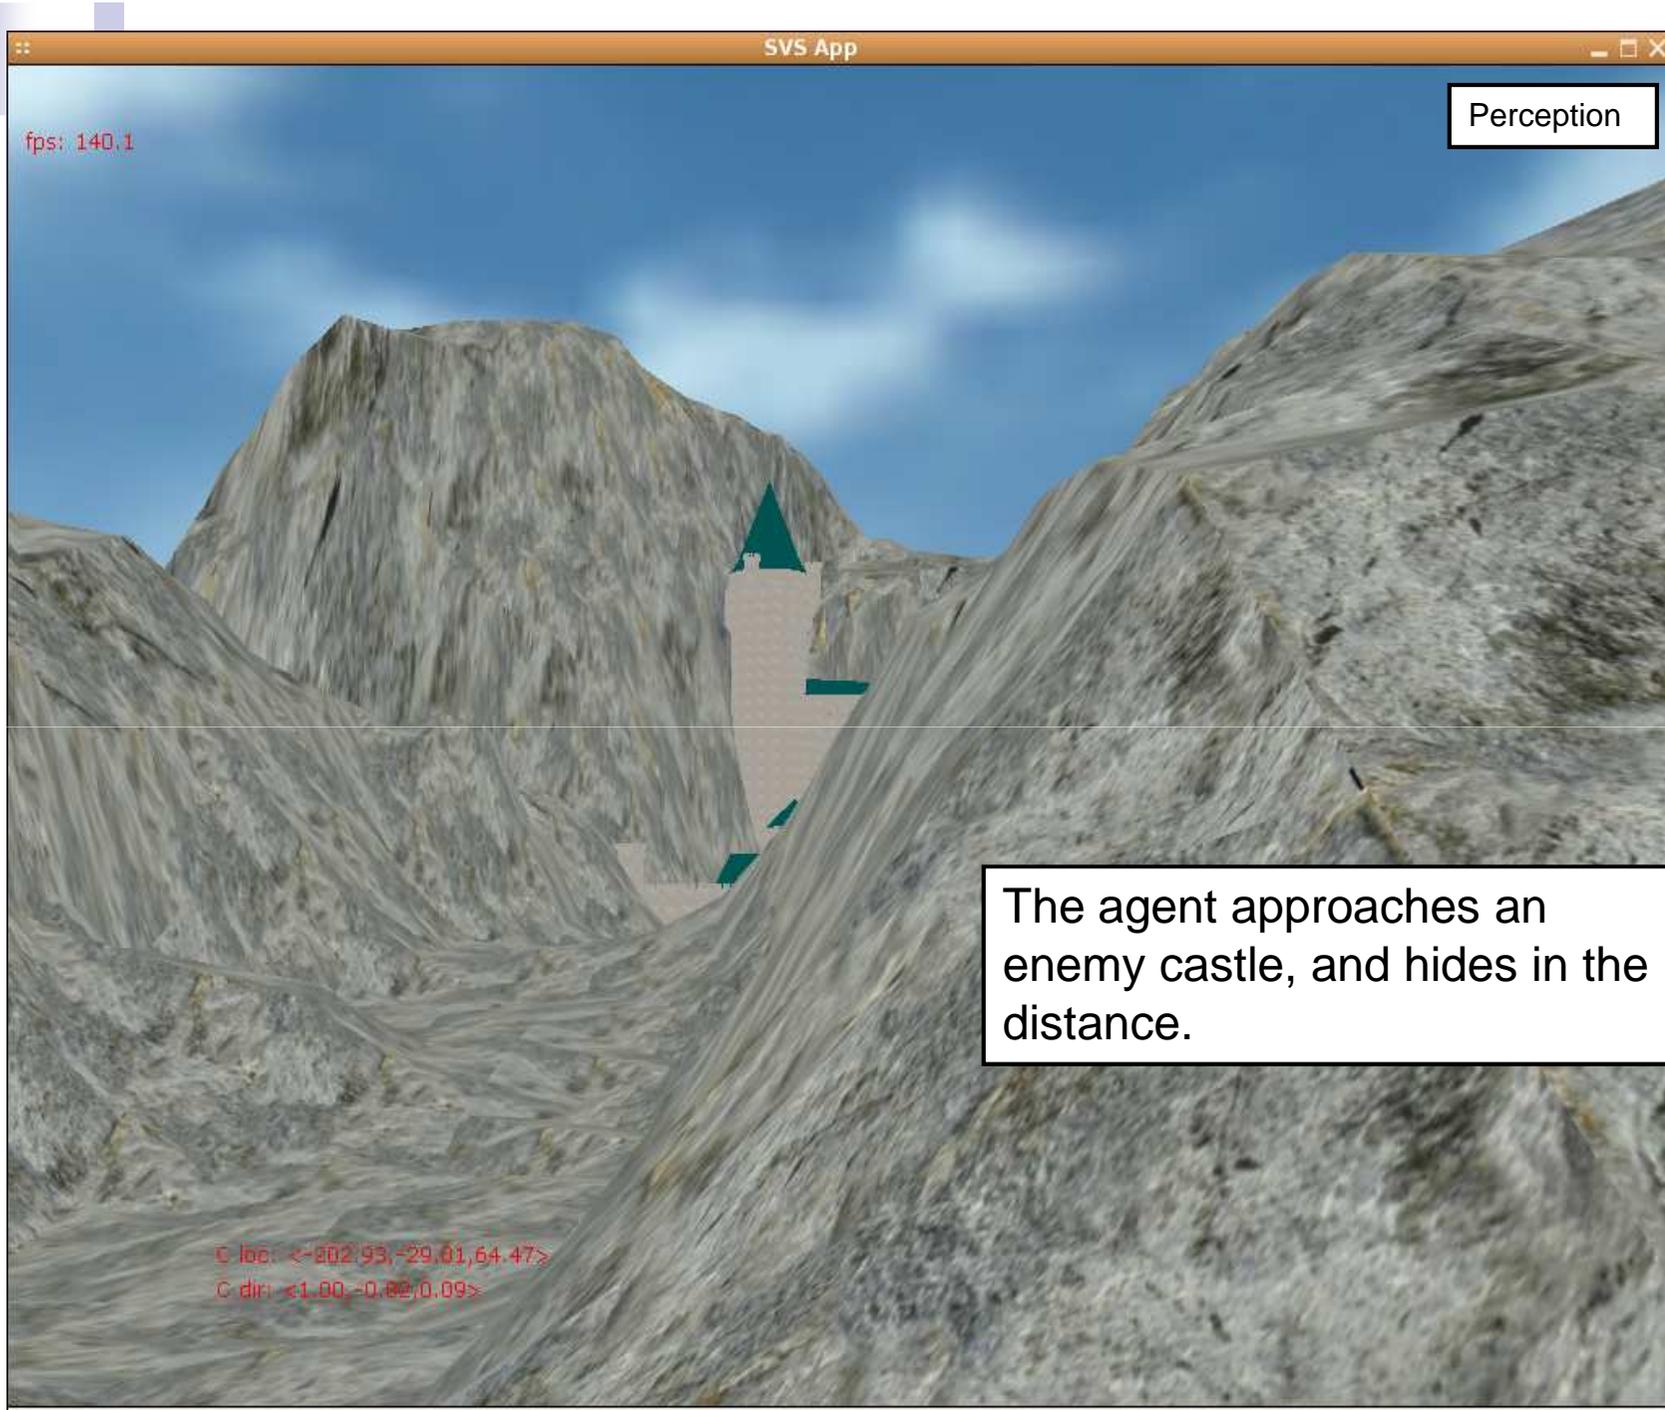
Soar Workshop 28, Ann Arbor, MI

# Outline

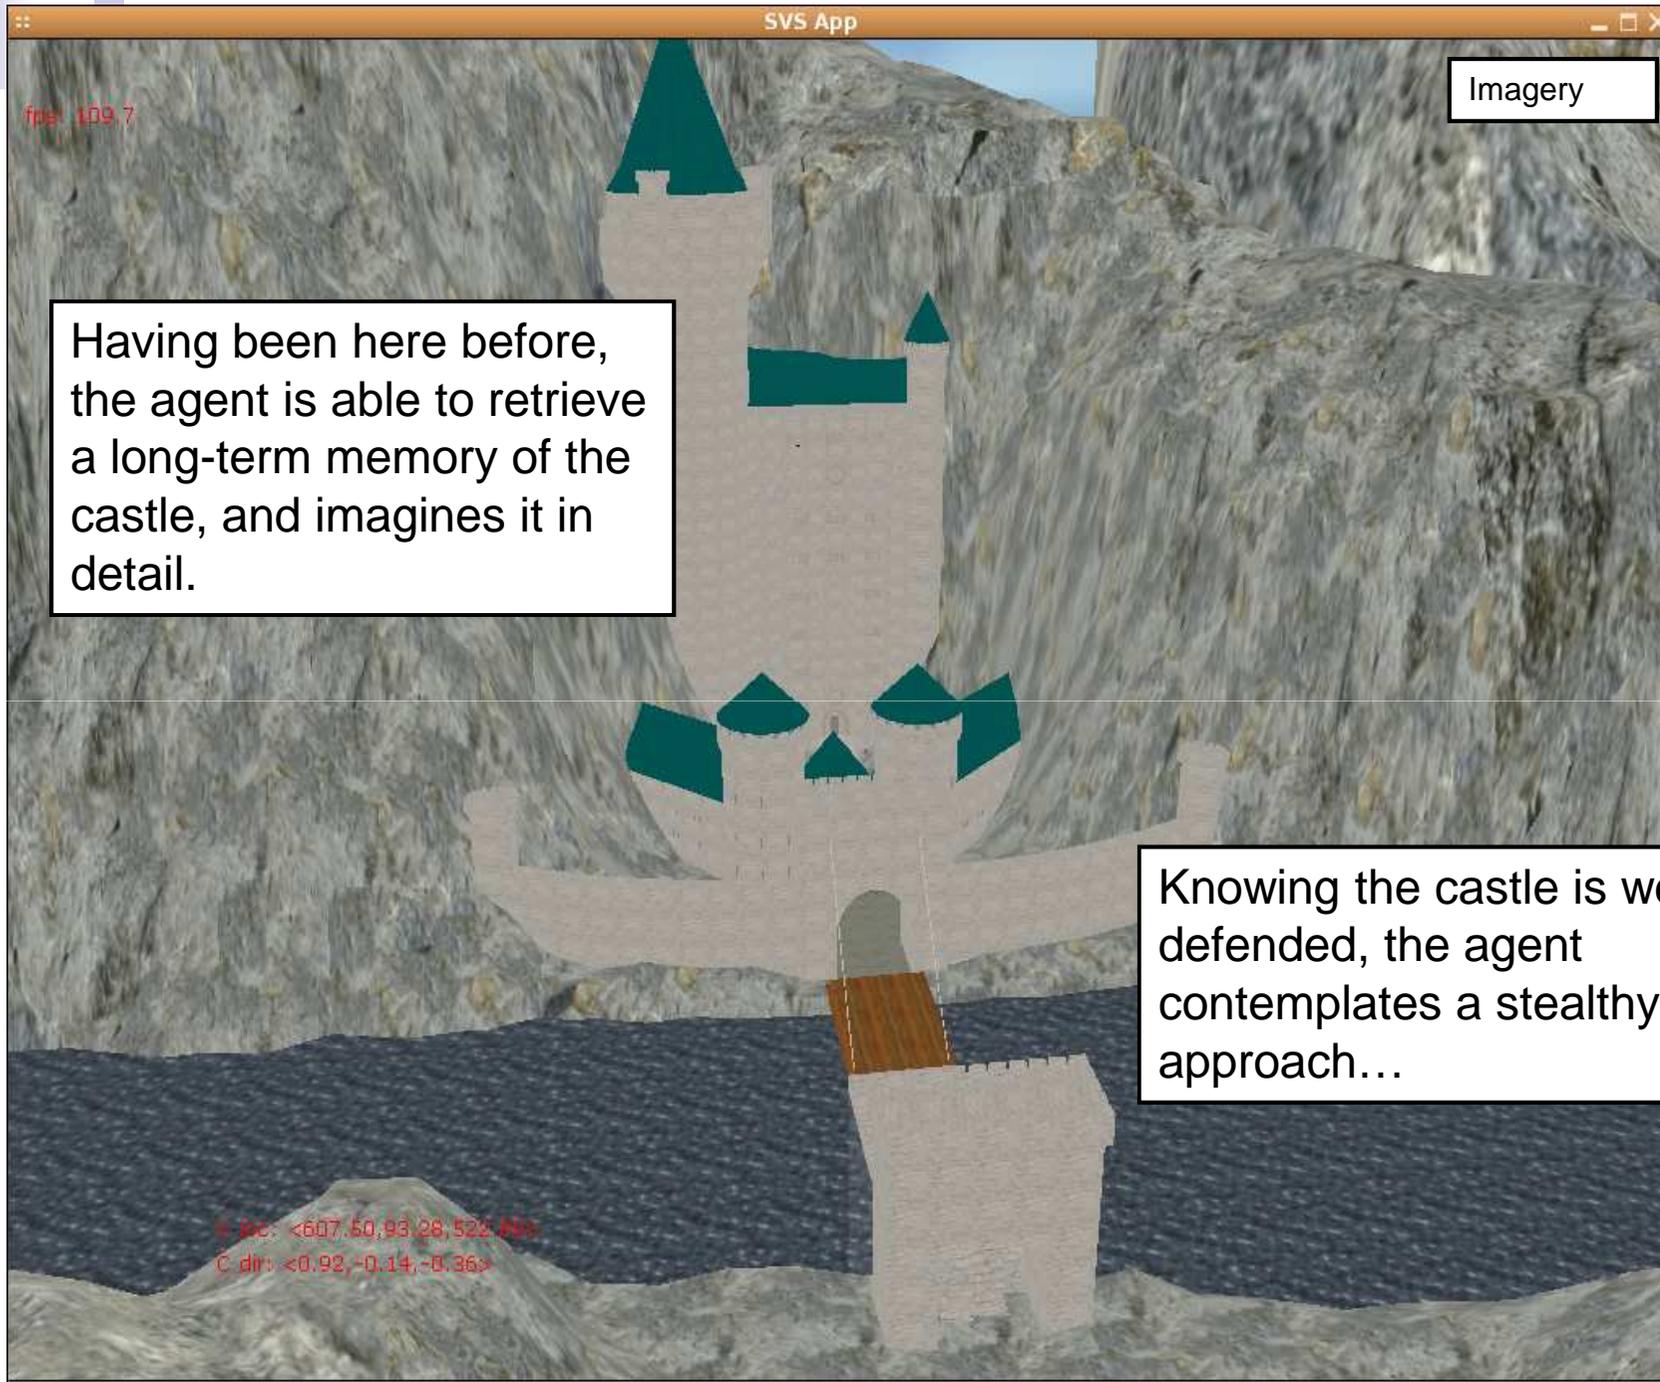
- Motivation
- Design
- SVS Scene Graphs
- Spatial Generation and Extraction
- Visual Generation and Extraction
- Example Problem
- Nuggets and Coal

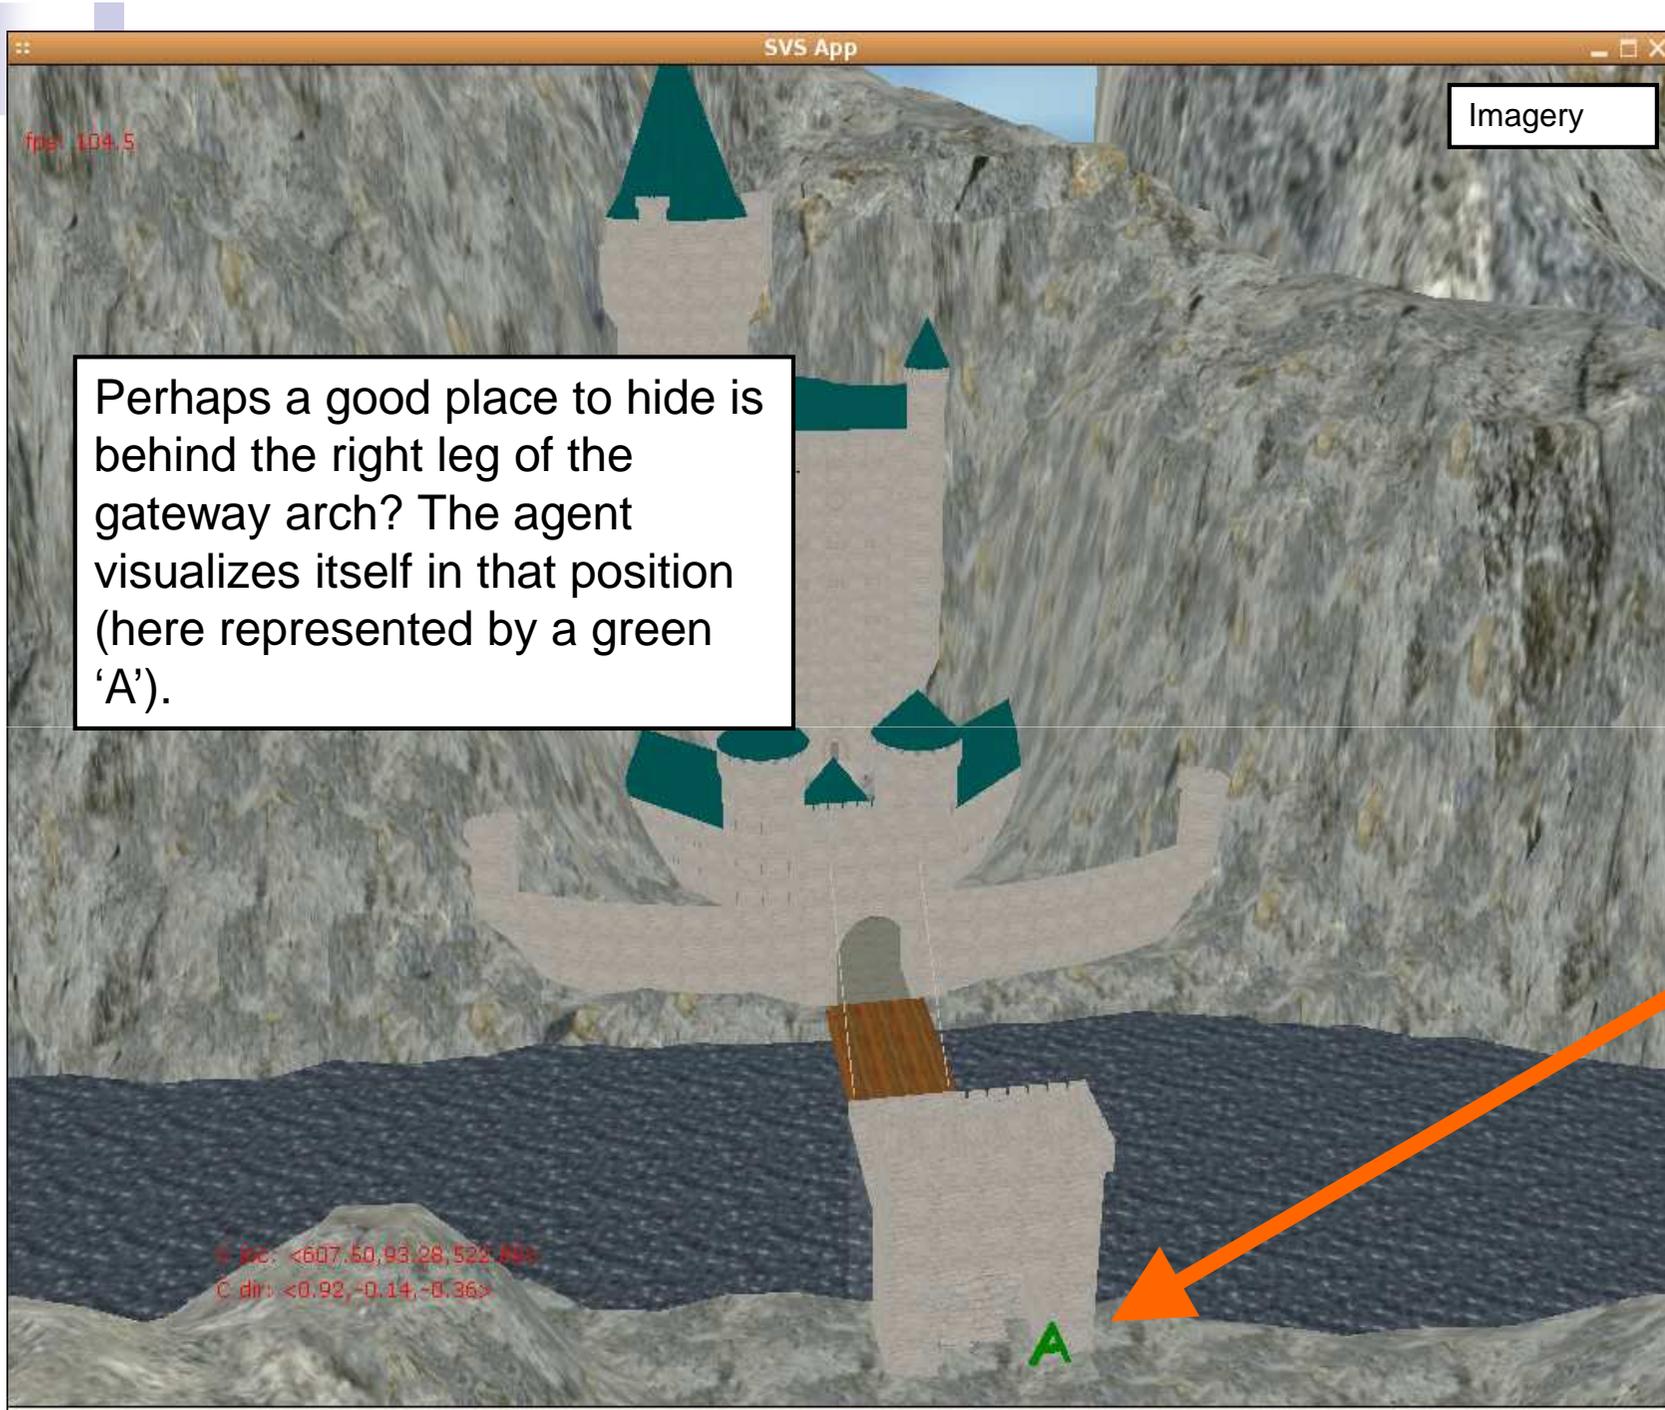
# Motivation

- Much work has been done in the last few years to enhance Soar with lower-level spatial and visual reasoning
- Two systems, SVI and SRS, have been developed
- SVS is the unification of SVI and SRS, with additional features

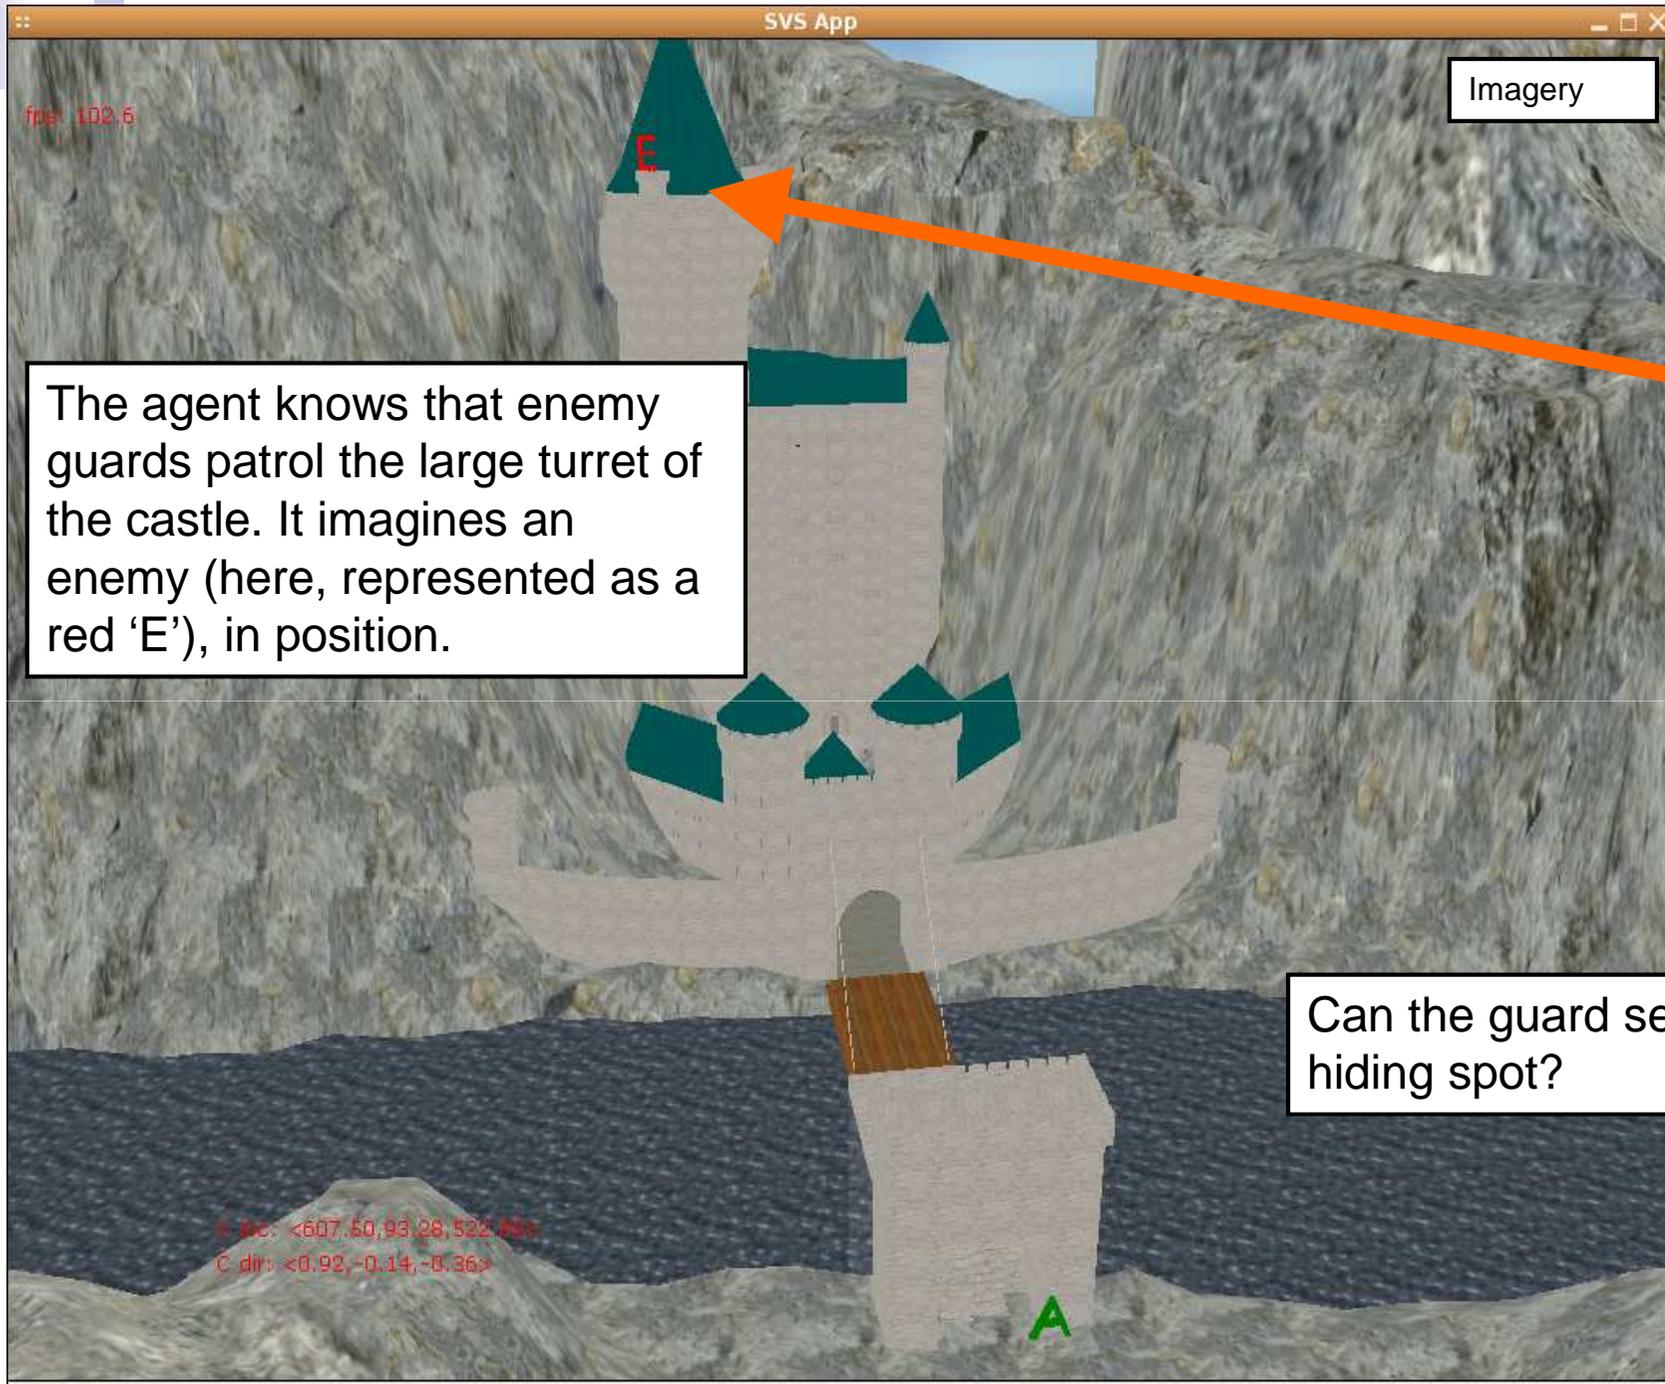


The agent approaches an enemy castle, and hides in the distance.



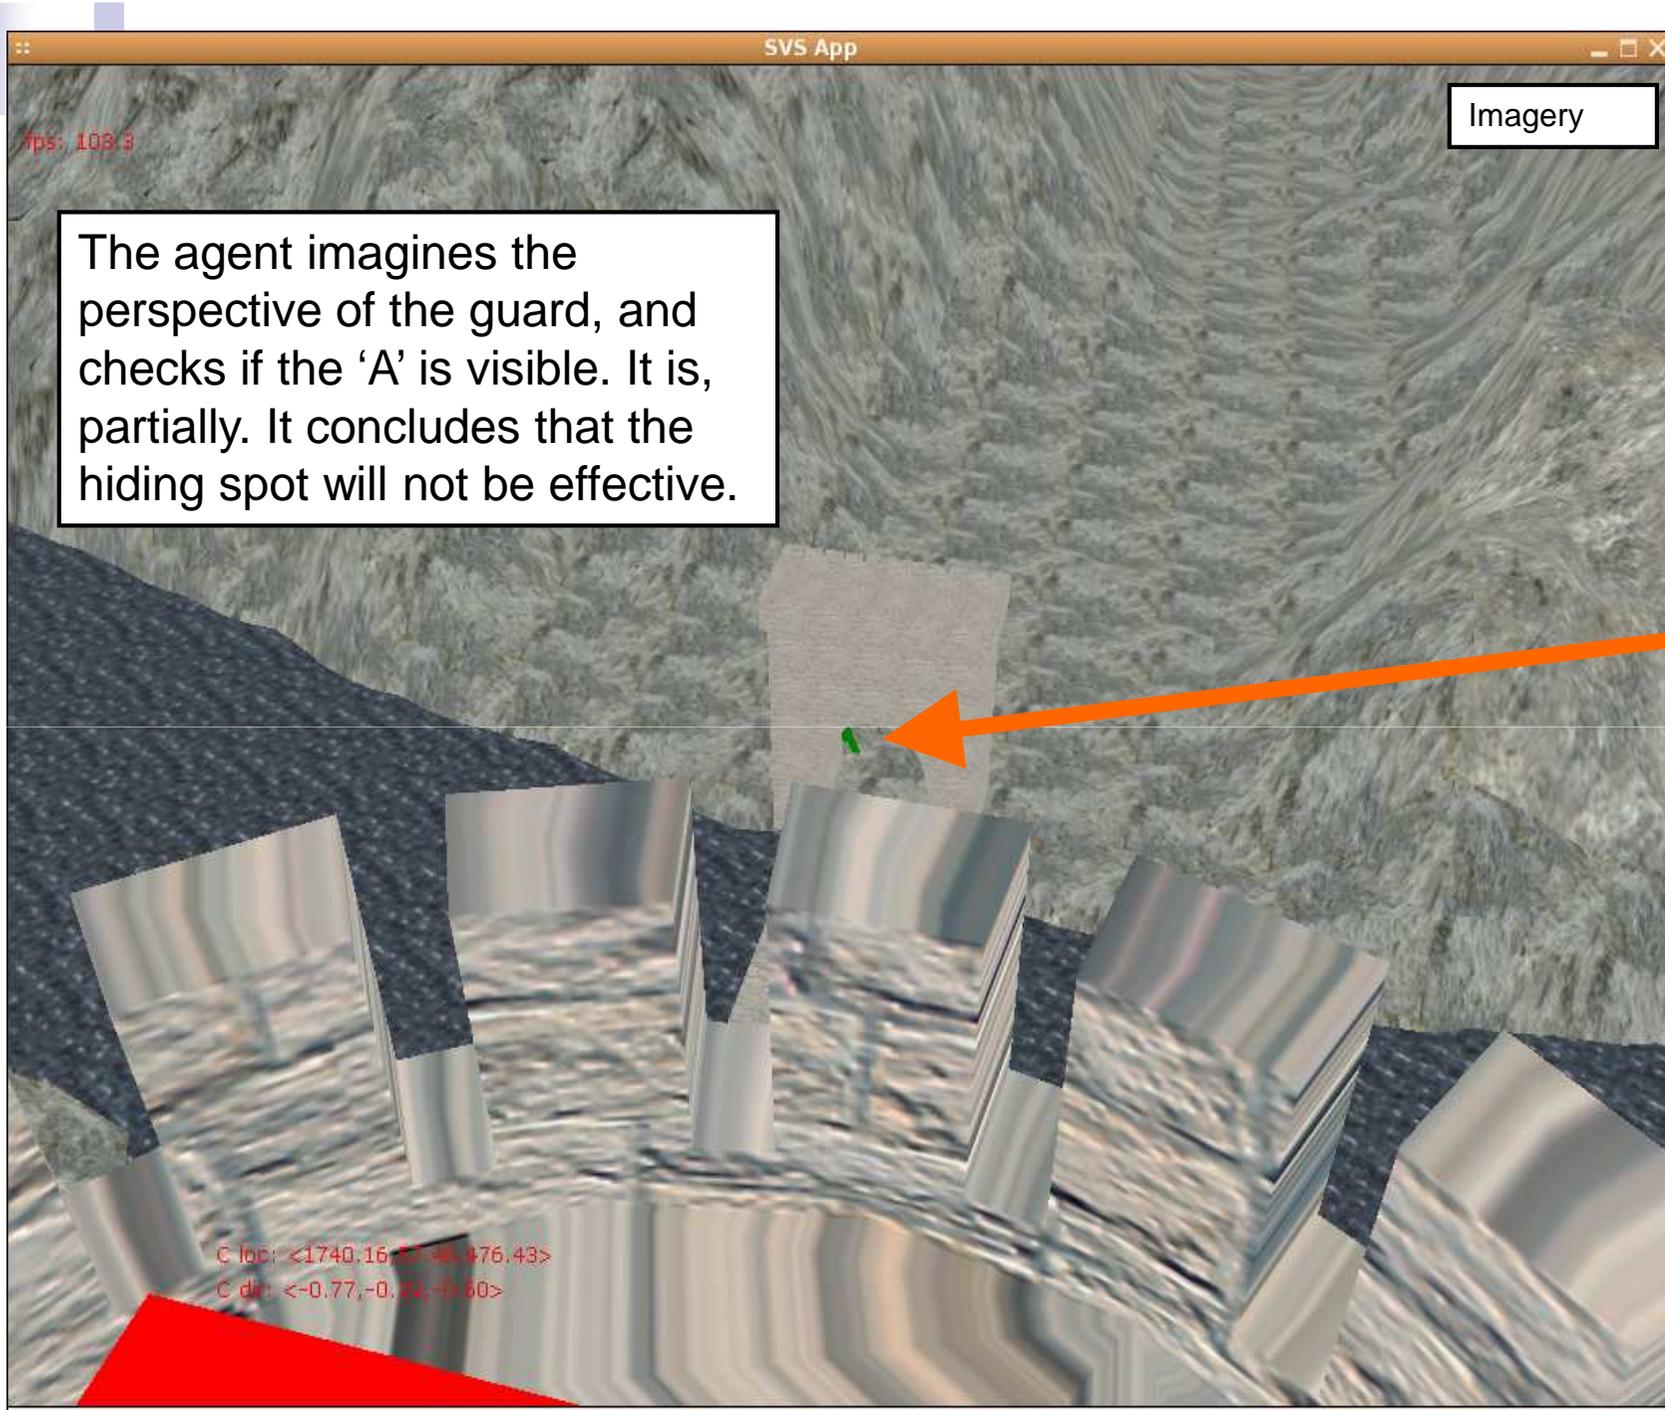


Perhaps a good place to hide is behind the right leg of the gateway arch? The agent visualizes itself in that position (here represented by a green 'A').

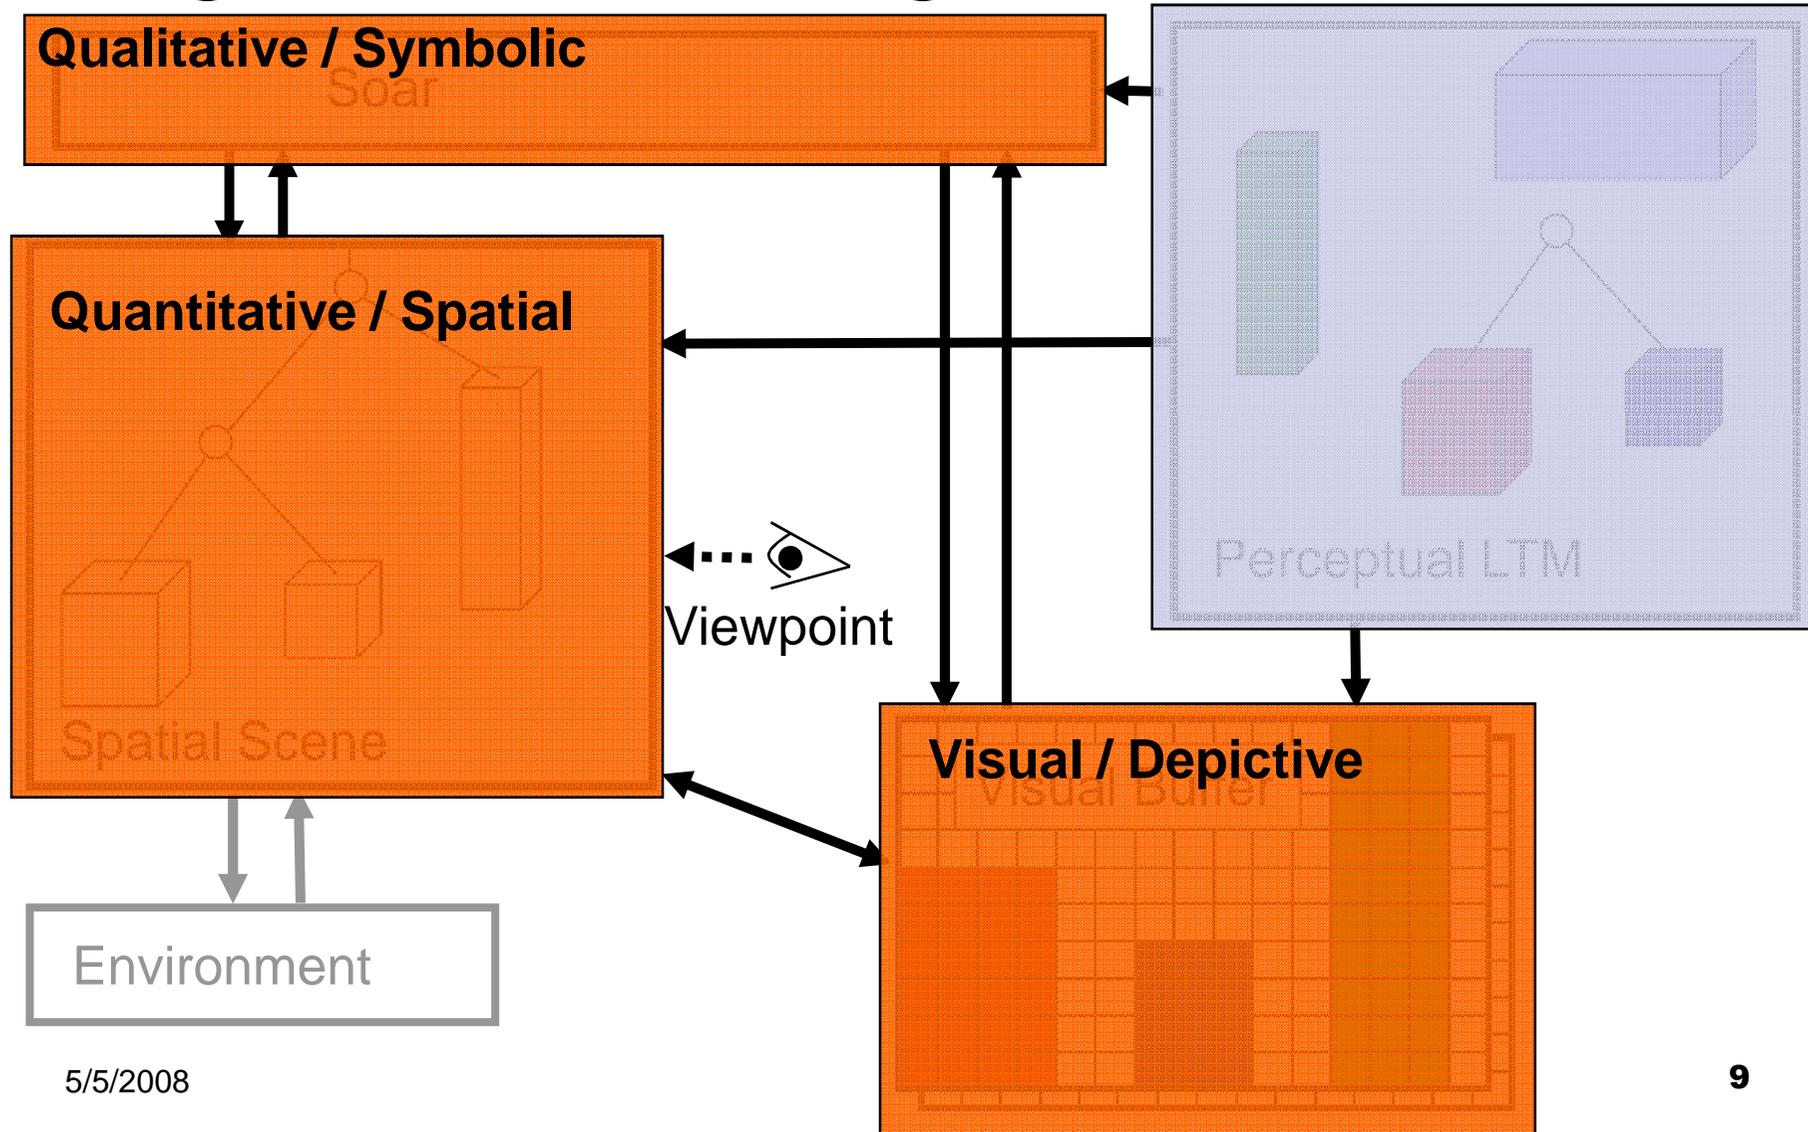


The agent knows that enemy guards patrol the large turret of the castle. It imagines an enemy (here, represented as a red 'E'), in position.

Can the guard see the hiding spot?



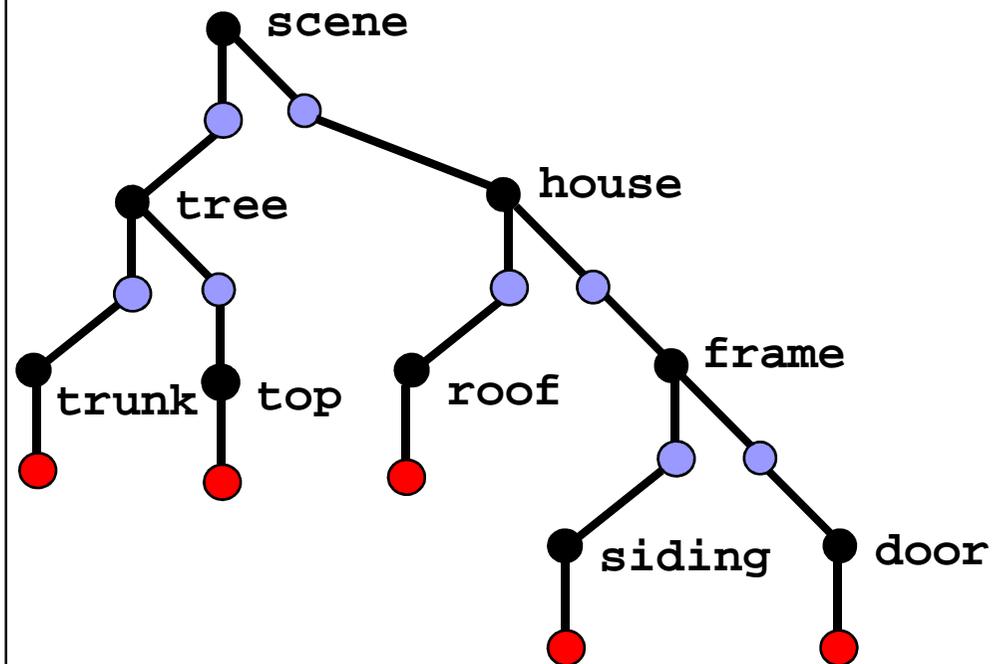
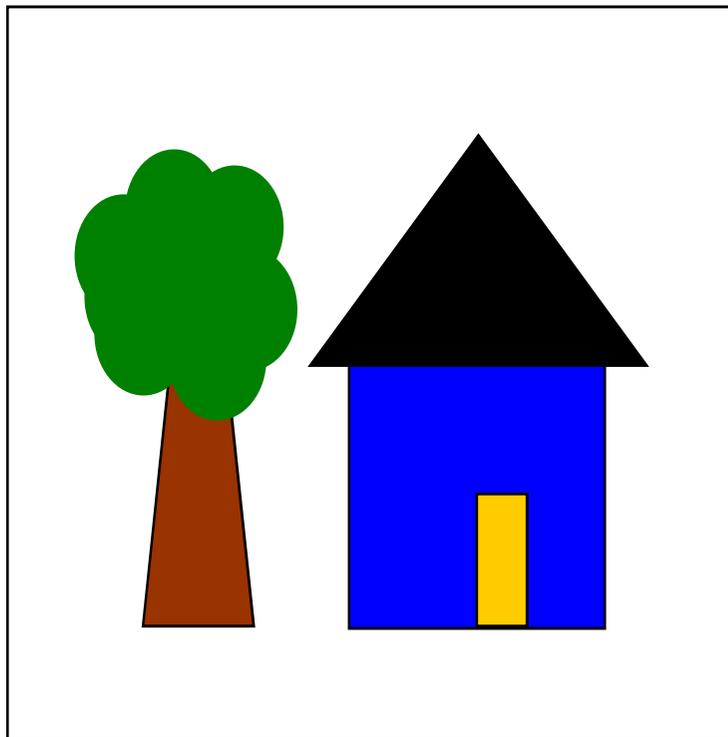
# High-Level Design



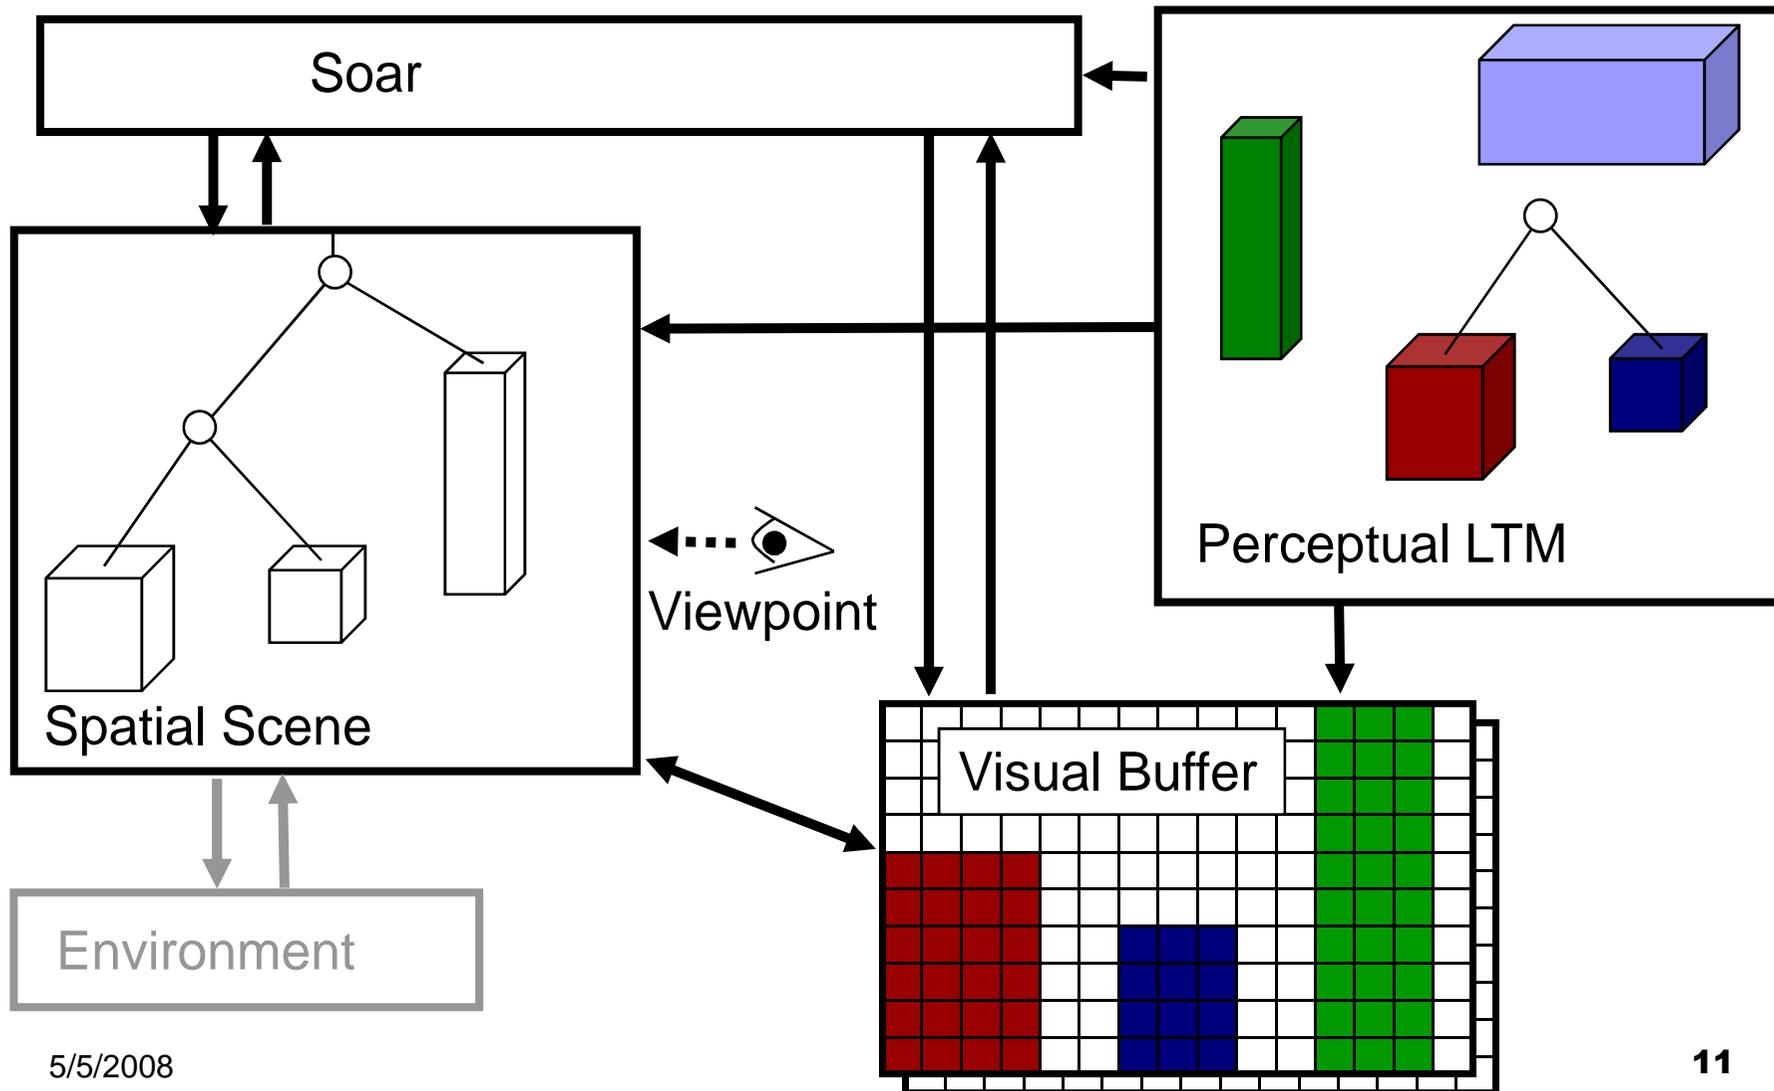
# SVS Scene Graphs

- object
- transformation
- texture

- SVS now makes scene graphs (in Spatial STM and LTM) accessible to Soar
- Soar can expand and contract nodes

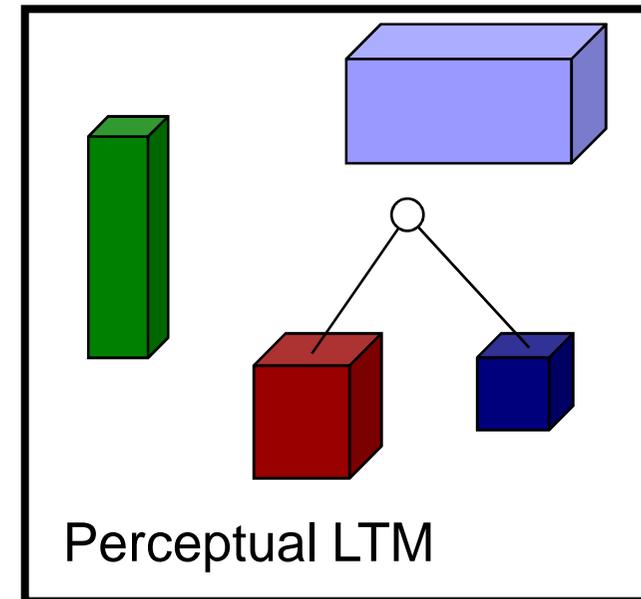


# High-Level Design

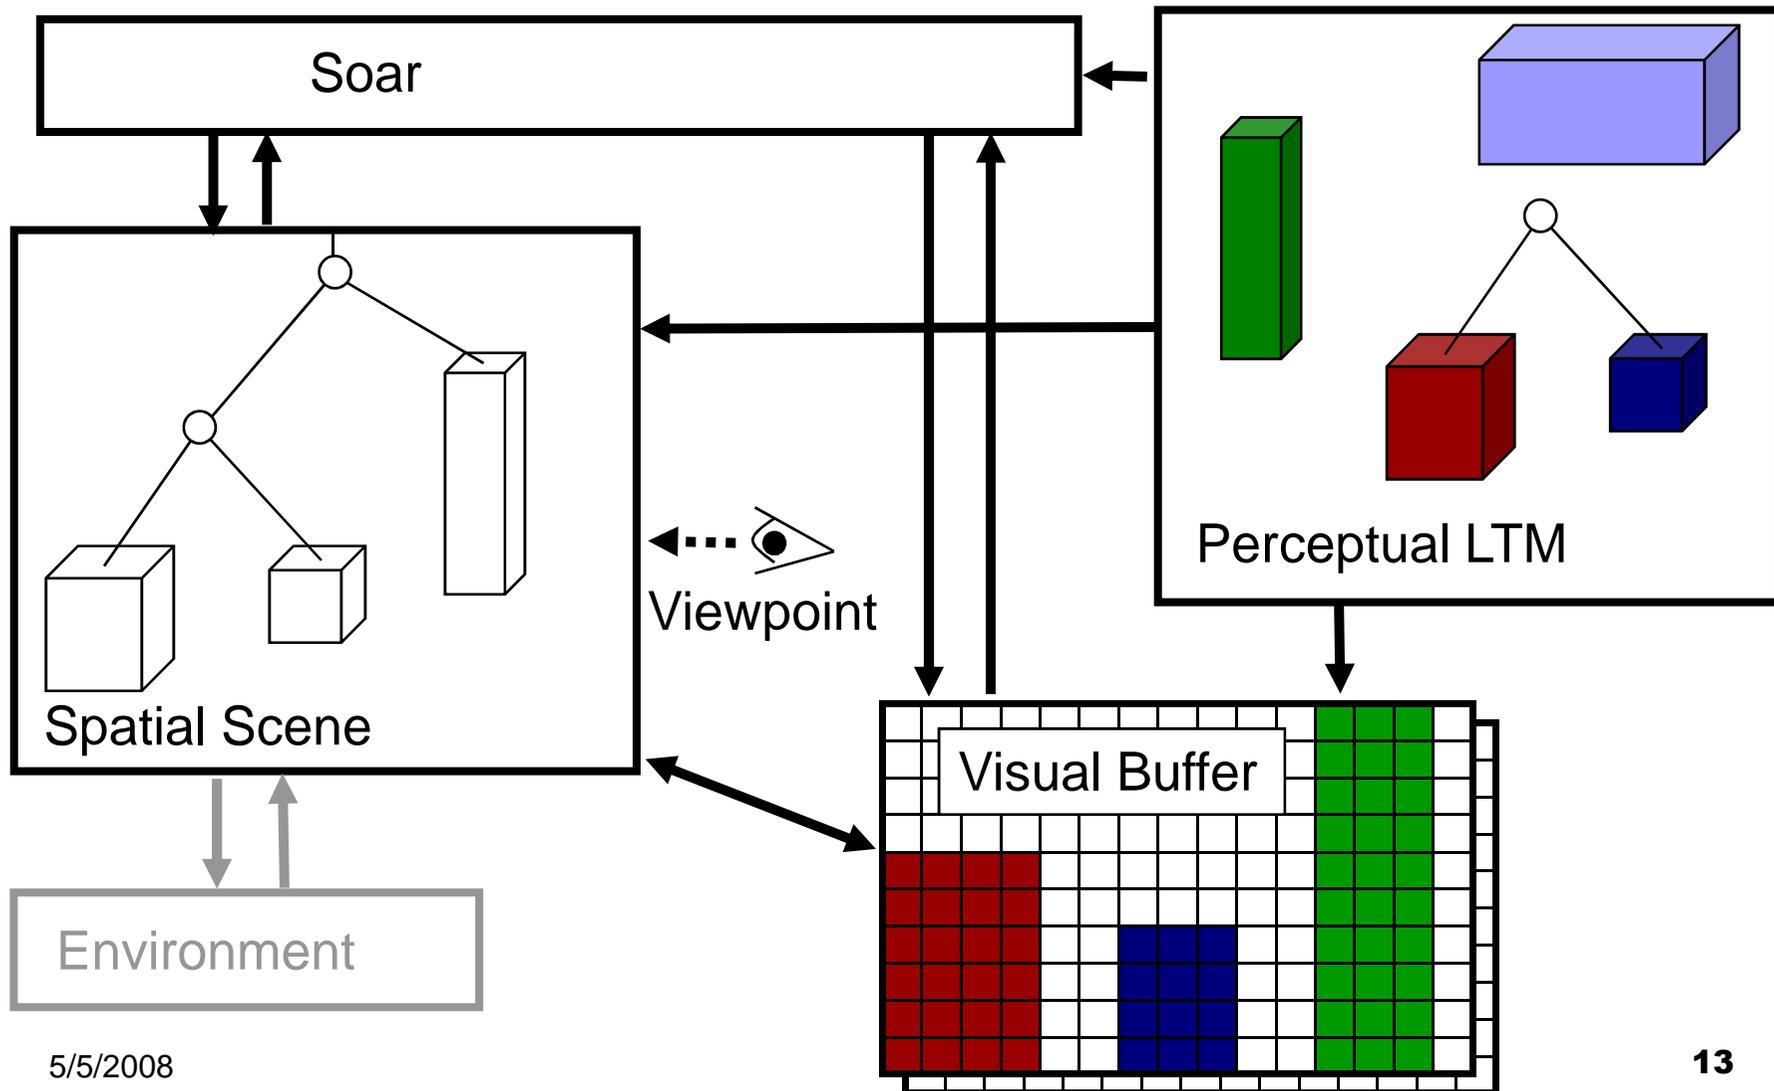


# Perceptual LTM

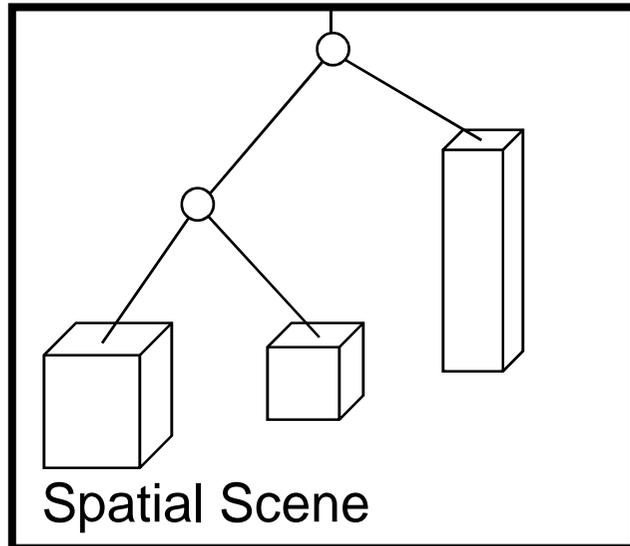
- Stores prototypical objects (shapes), textures, and transformations
- Hierarchically organized in scene graphs



# High-Level Design

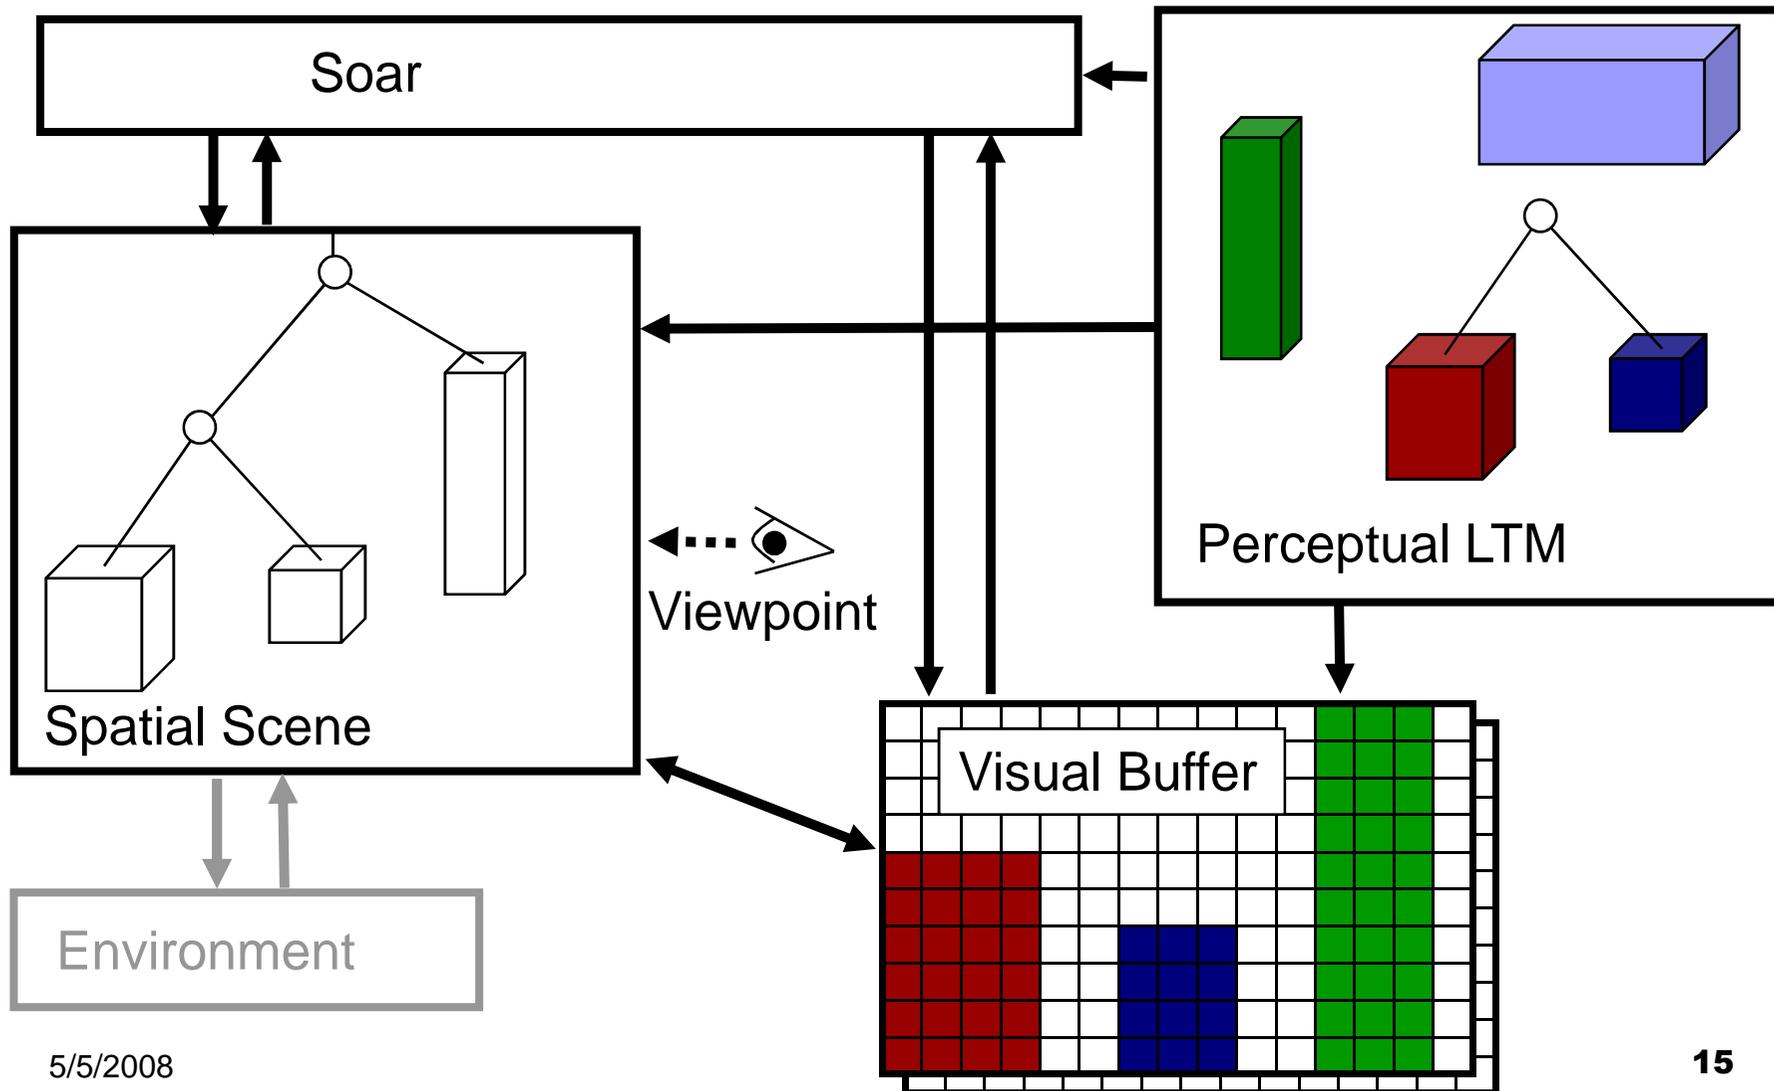


# Spatial Scene



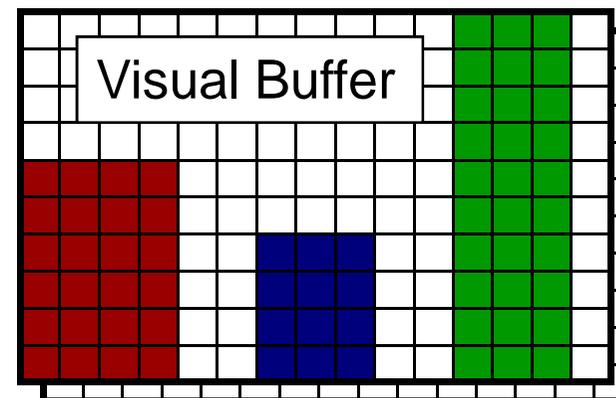
- 3D spatial short-term memory
- Scene graph
  - But grounded in coordinates

# High-Level Design

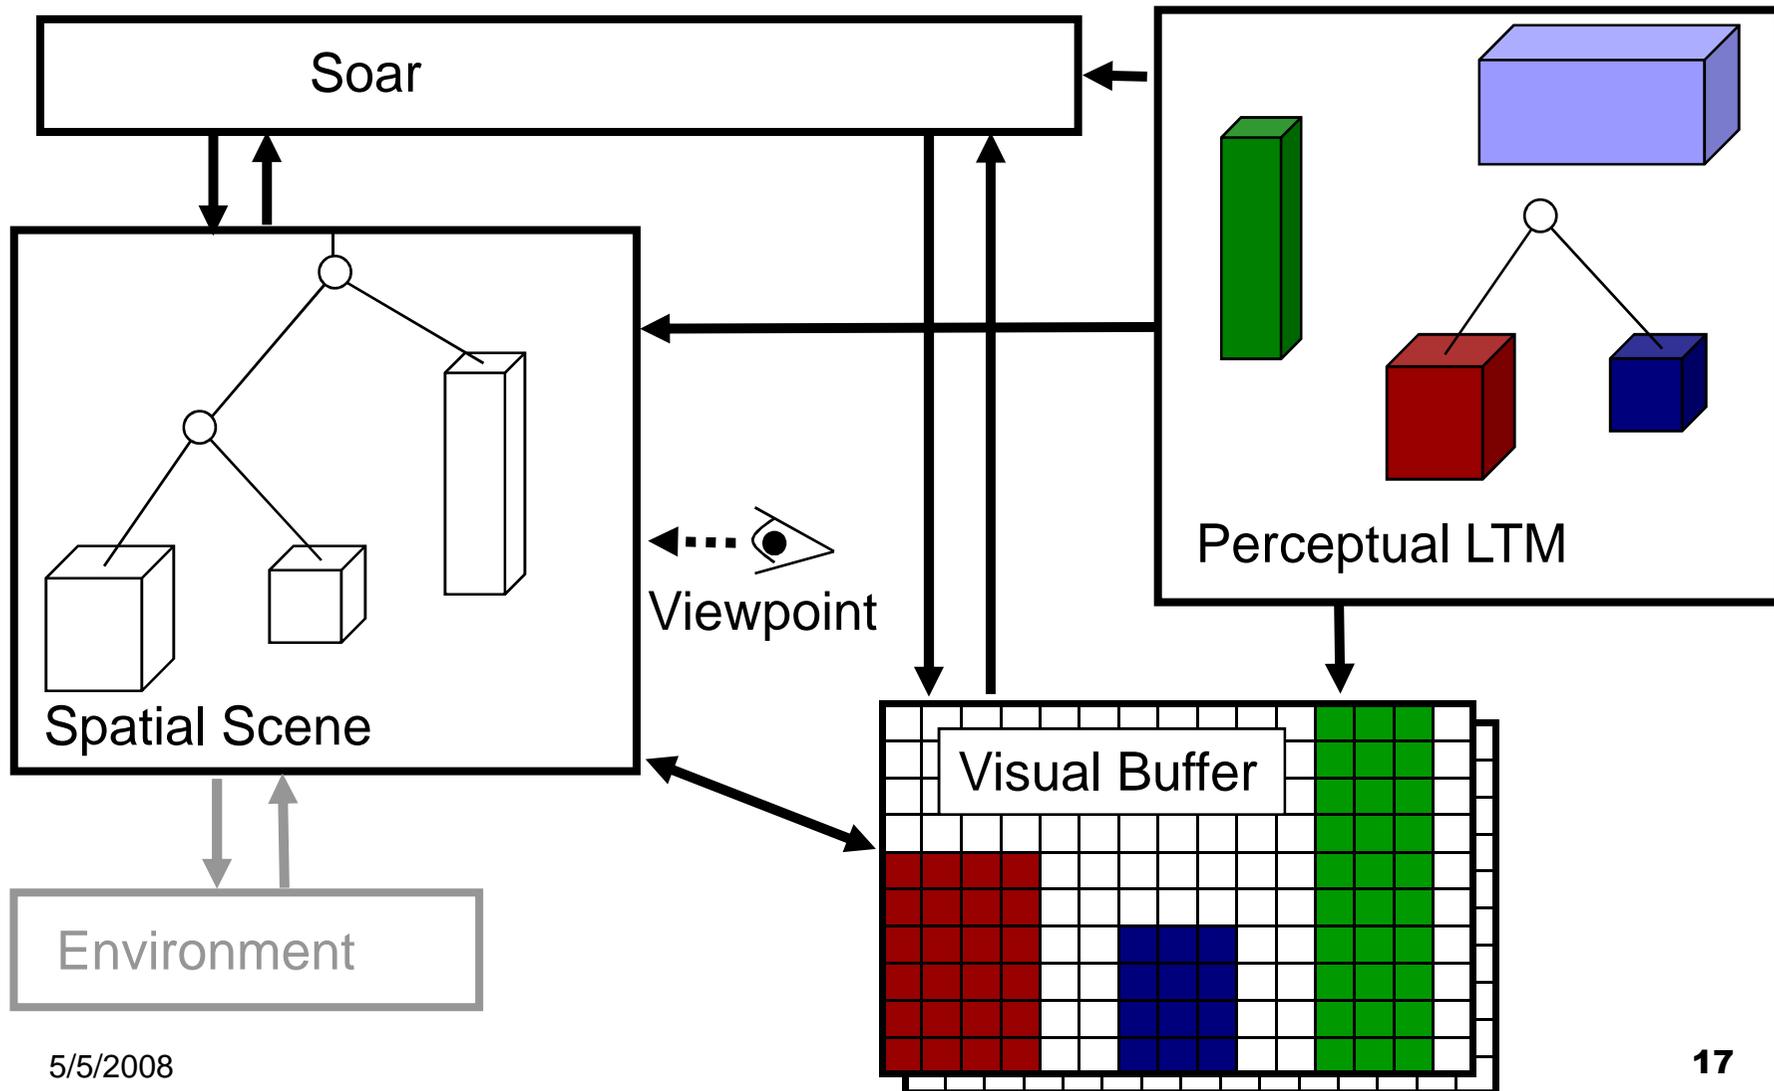


# Visual Buffer

- Short-term memory
- Contains a set of 2D, depictive representations
- Derived from the scene, at a given viewpoint
  - Theoretically, the opposite is true



# High-Level Design



# Soar Interface

- Interfaces are (mostly) qualitative: all quantitative information stays within SVS
- SVS provides contents of all memories to Soar
  - Scene graphs are expandable/collapsible
- Soar can add new information to SVS
  - Generate new objects in the scene, and new depictions in the visual buffer
- Soar can retrieve information from SVS
  - Extract predicates from the scene and visual buffer through queries

# Outline

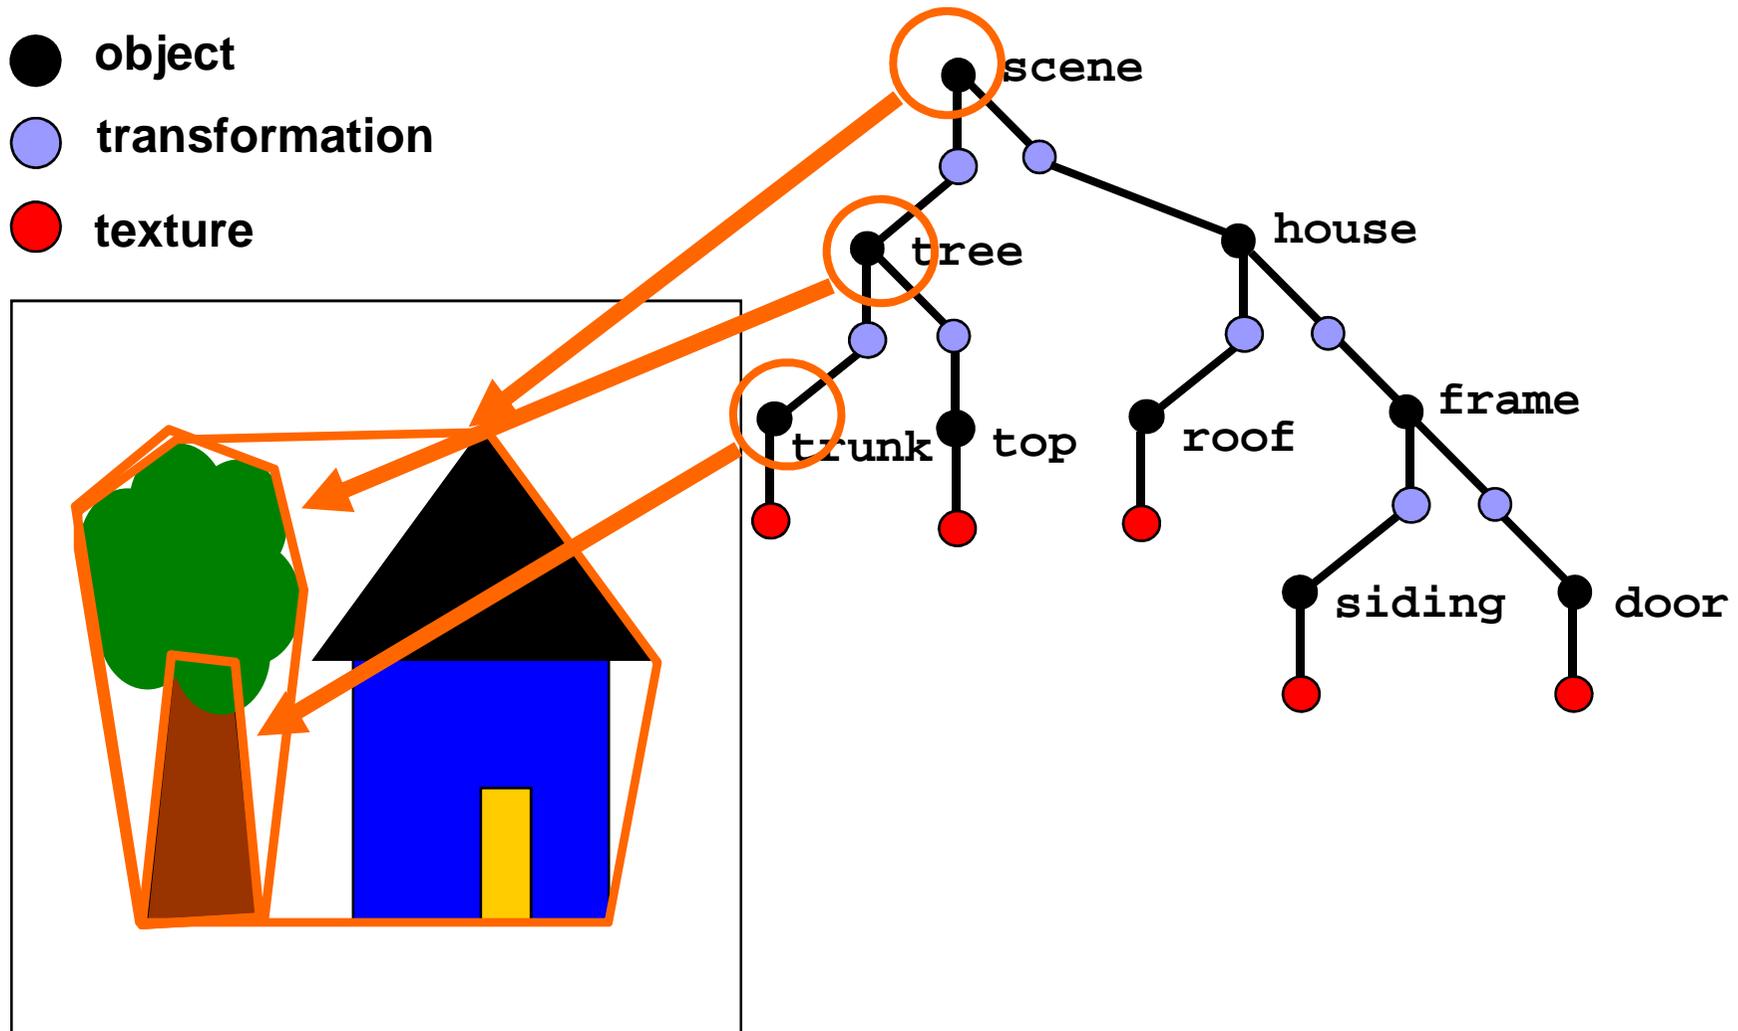
- Motivation
- Design
- **SVS Scene Graphs**
- Spatial Generation and Extraction
- Visual Generation and Extraction
- Example Problem
- Nuggets and Coal

# Scene Graph Example: Objects

● object

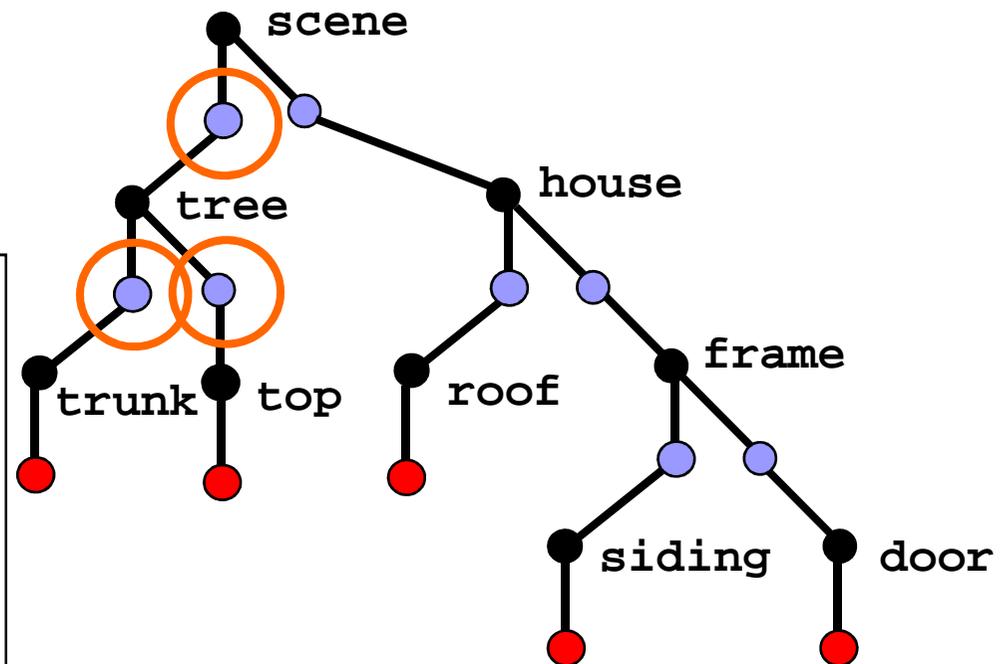
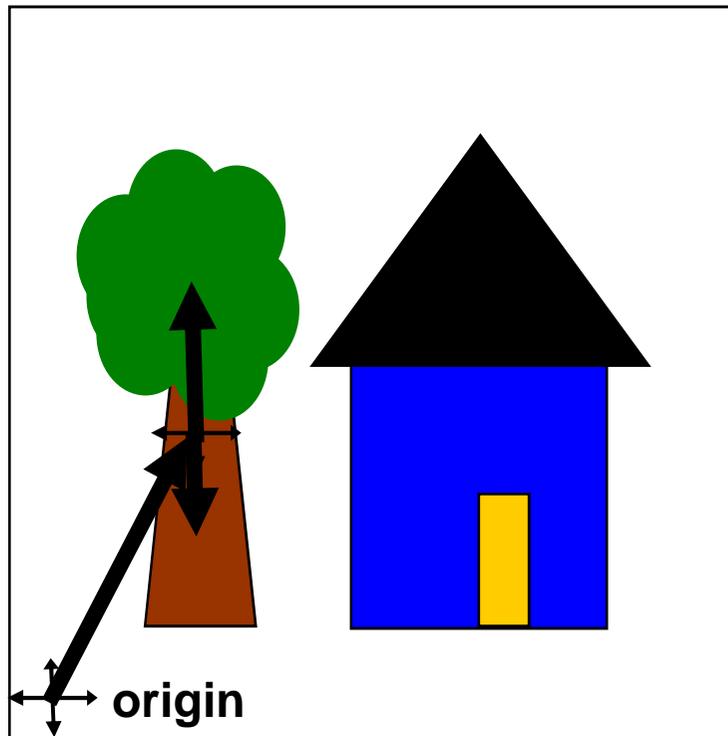
● transformation

● texture



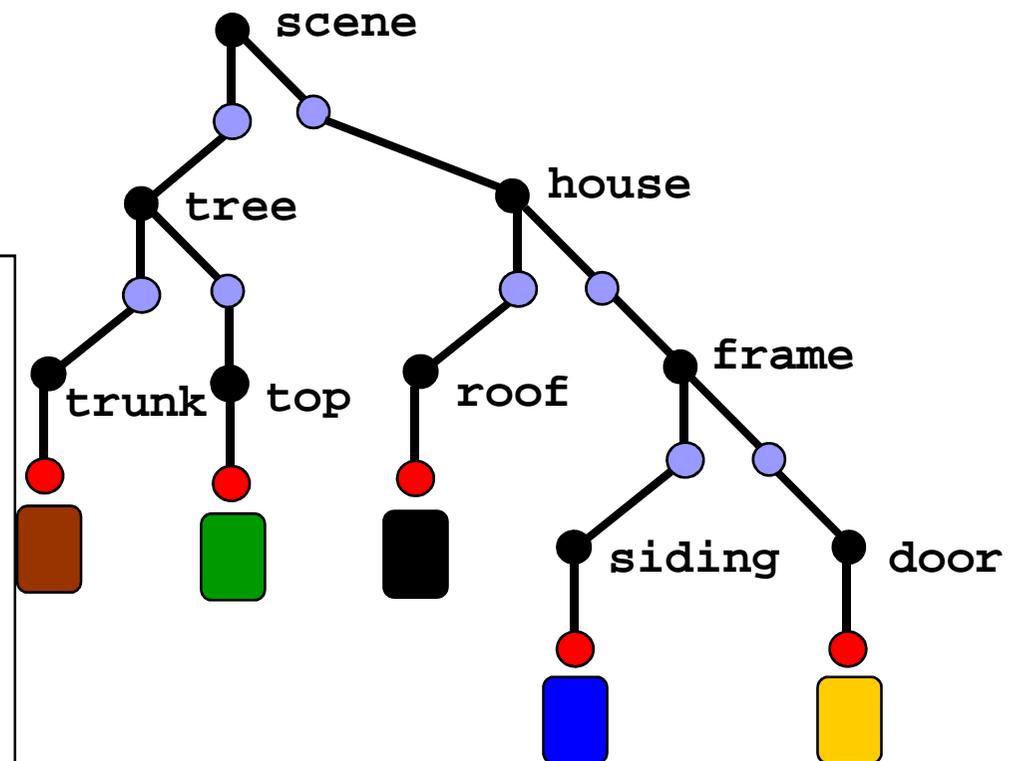
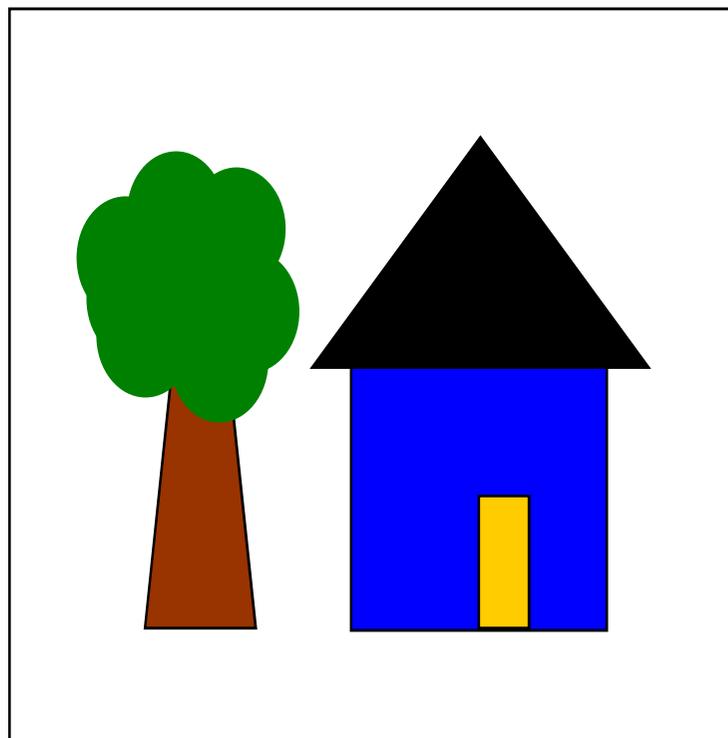
# Scene Graph Example: Transformations

- object
- transformation
- texture



# Scene Graph Example: Textures

- object
- transformation
- texture



# Outline

- Motivation
- Design
- SVS Scene Graphs
- **Spatial Generation and Extraction**
- Visual Generation and Extraction
- Example Problem
- Nuggets and Coal

# Spatial Generation

- SVS supports imagery where Soar directly provides a scene graph
  - Scene graphs are composable, any valid combination of nodes yields a scene
  - These operations are called “graph generations”
- SVS also supports imagery where objects and/or transformations are qualitatively described
  - Scene graph is implicit in the description
  - These operations are called “projections”
- Motion models allow SVS to represent movement (my other talk)
- All images (spatial and visual) exist as long as the command is on the output-link
  - Images can be manipulated by changing the command in place

# Accessing Long-Term Memory

- Items in LTM have a ^class-id
  - e.g. car, car37, car-transform37, car-texture37
- Items in the scene have a ^class-id and an ^instance-id
  - e.g. ^class-id car, ^instance-id car:i23
- A retrieval from LTM happens when a ^class-id appears in a generate command
  - Objects are copied all the way down to the leaves
  - Graph generations can access LTM nodes identically to nodes in the scene

# Spatial Graph Generation Example

```
^generate-spatial
```

```
  ^transformation
```

```
    ^parent scene
```

```
    ^source house-trans
```

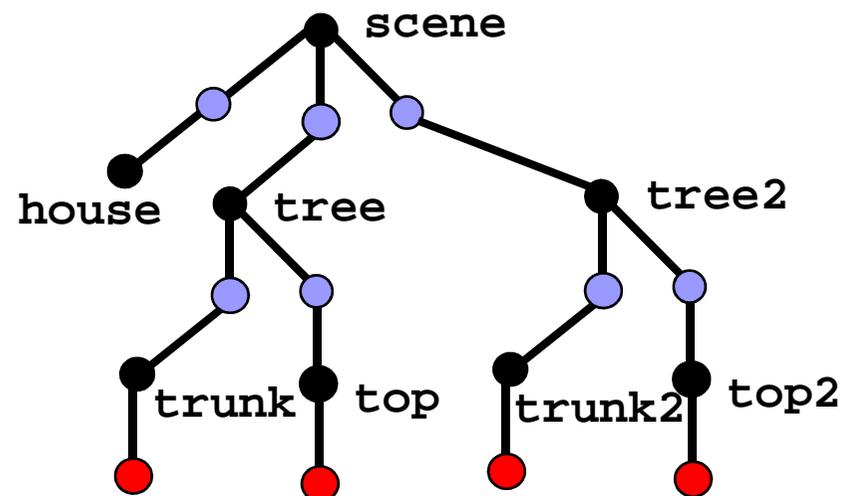
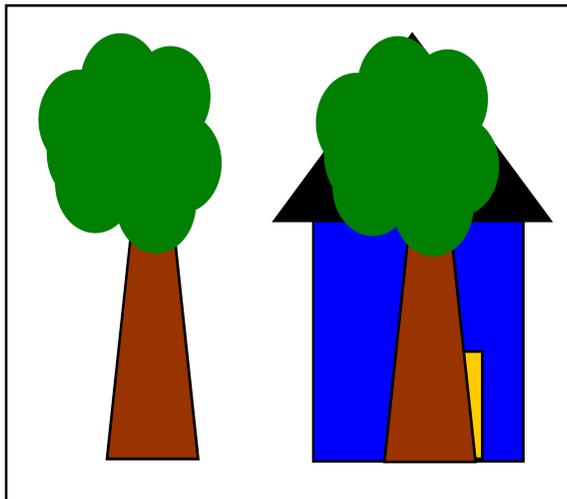
```
    ^new-id tree-trans2
```

```
    ^child-new-id tree2
```

```
^object
```

```
  ^source tree
```

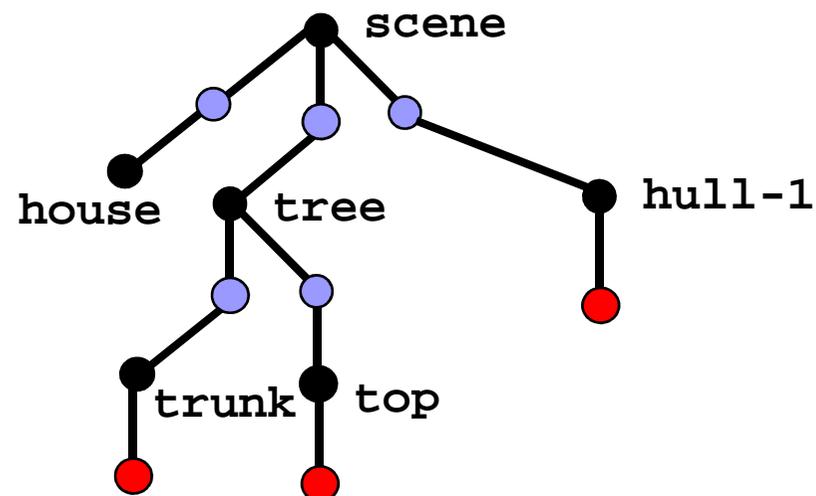
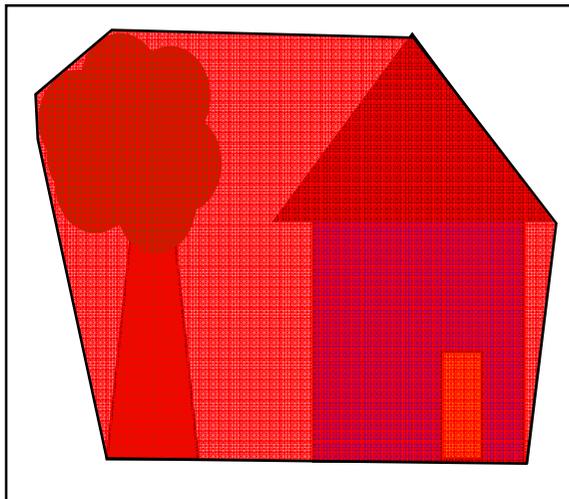
```
  ^new-id tree2
```





# Spatial Projection Example

```
^generate-spatial  
  ^projection  
    ^hull  
      ^instance-id tree  
      ^instance-id house  
    ^transform-parent scene  
    ^texture-source.class-id red
```



# Spatial Extraction Example

`^extract-spatial`

`^relationship adjacent`

`^primary-object frame`

`^reference-object ?`

`^value true`

`^extract-spatial-result`

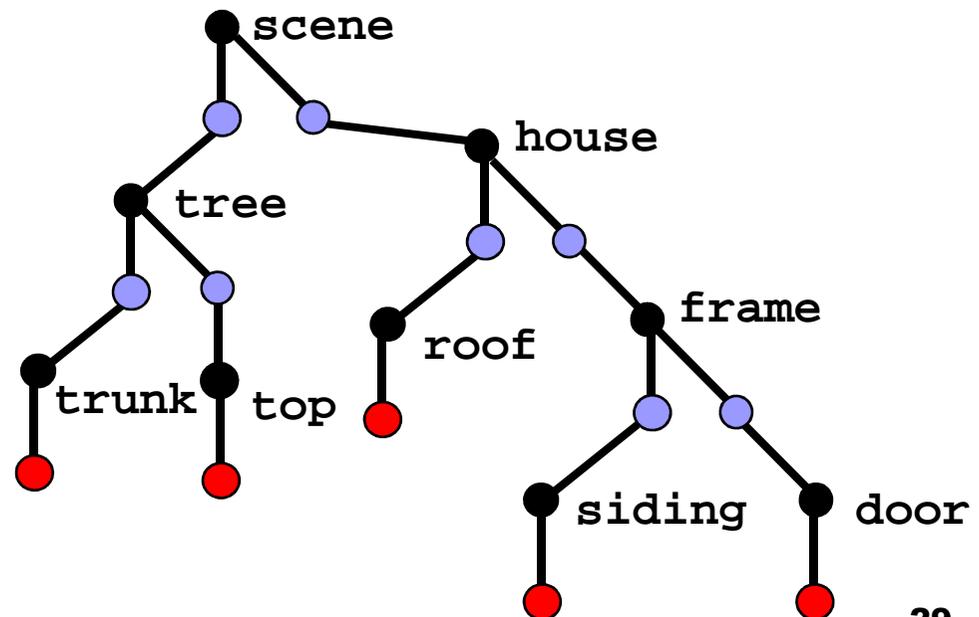
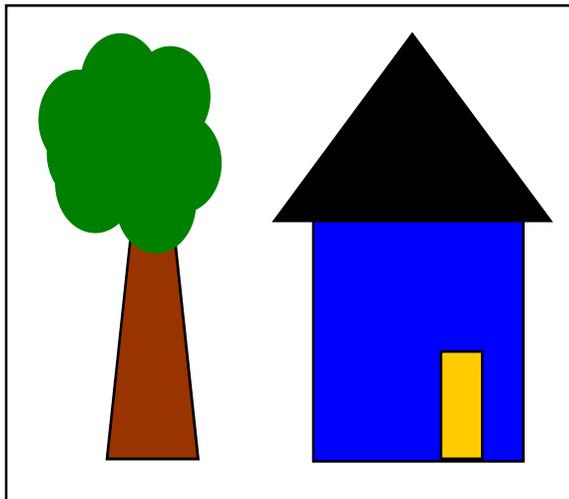
`^retrieved *5`

`^relationship adjacent`

`^primary-object frame`

`^reference-object {door, siding, roof, house, scene}`

`^value true`



# Outline

- Motivation
- Design
- SVS Scene Graphs
- Spatial Generation and Extraction
- **Visual Generation and Extraction**
- Example Problem
- Nuggets and Coal

# Visual Generation

- Visual generation commands add depictions to the visual buffer
- The scene generator renders all or part of the spatial scene to the VB, from a given camera position
- Other visual generators make new depictions without accessing the spatial scene
  - Depictive manipulation (Scott Lathrop's talk) is a subset of this

# Visual Scene Generation

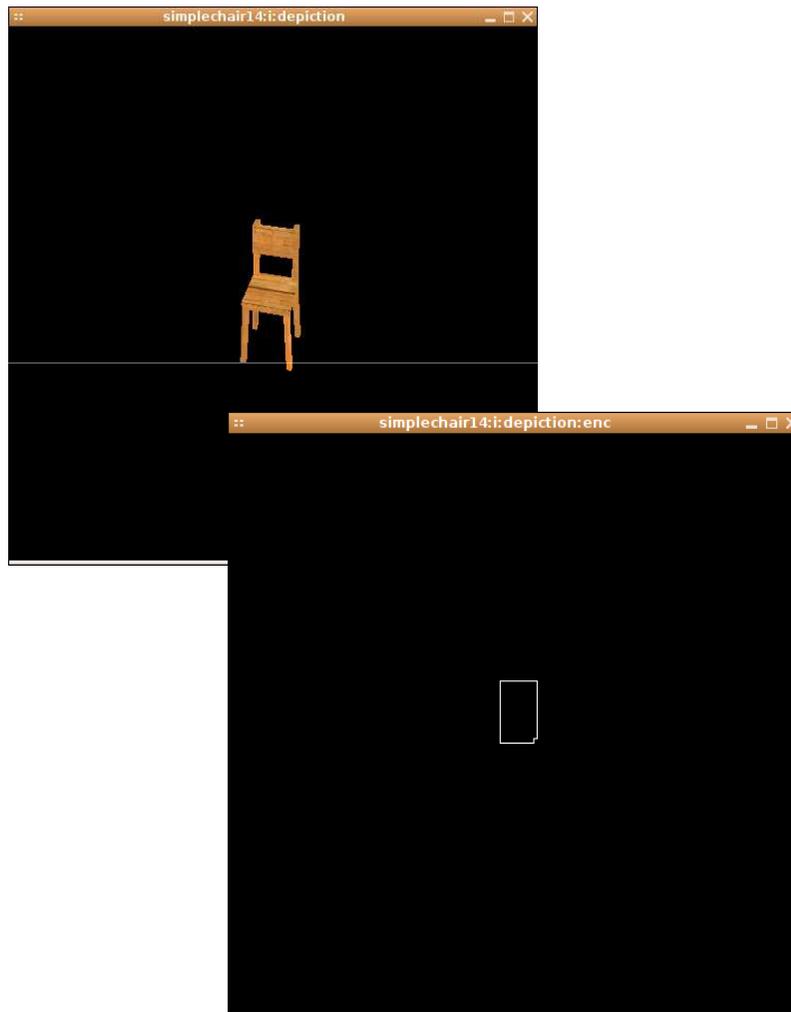


from the perspective of chandelier01 looking towards table04, generate the scene

from the perspective of chandelier01 looking towards table04, generate table04

from the perspective of chandelier01 looking towards table04, generate chair14, ignoring table04

# Visual Generation and Extraction



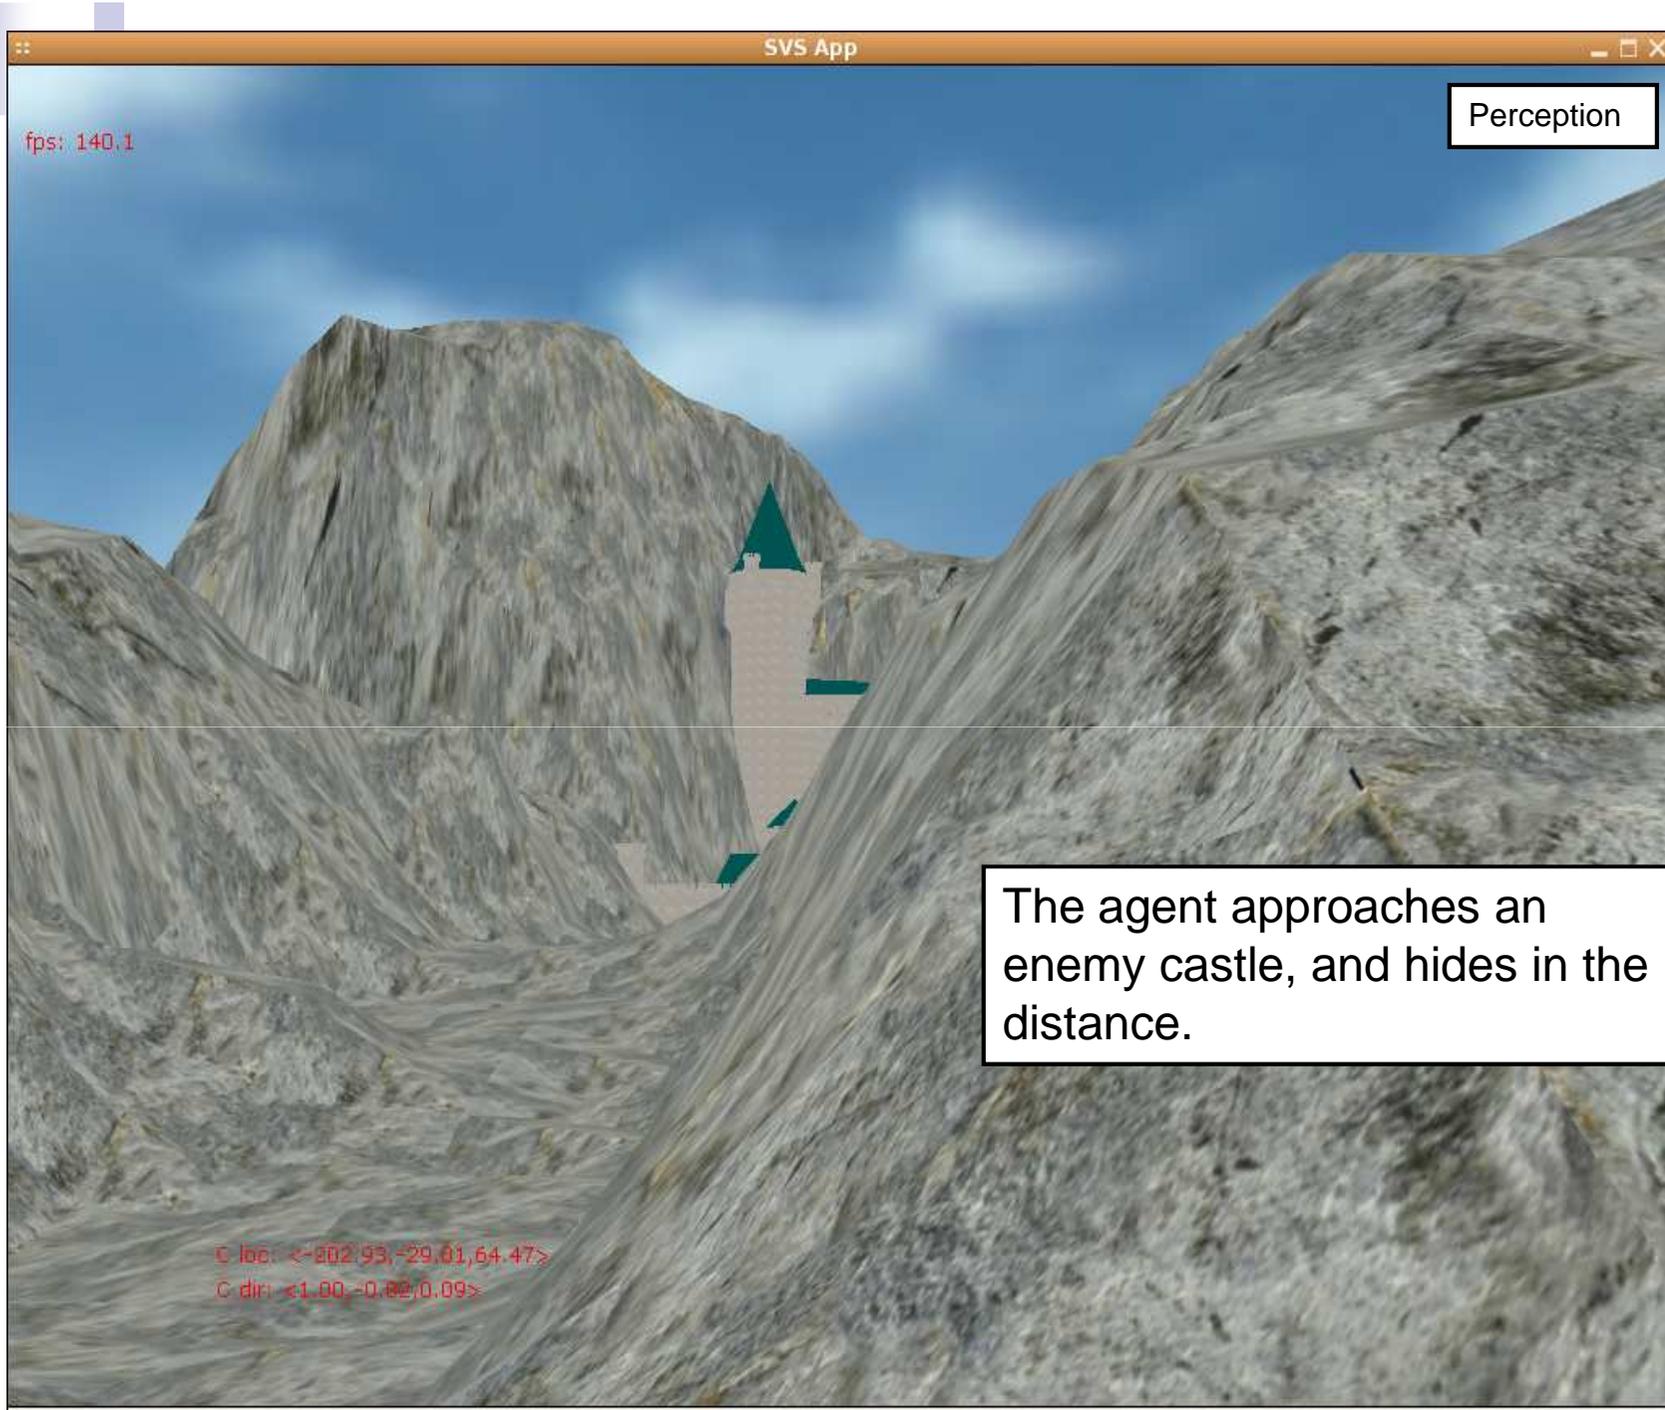
generate a depiction  
outlining the enclosed  
space in chair14:depiction

query: are any pixels in  
chair14:depiction:enc filled  
in?

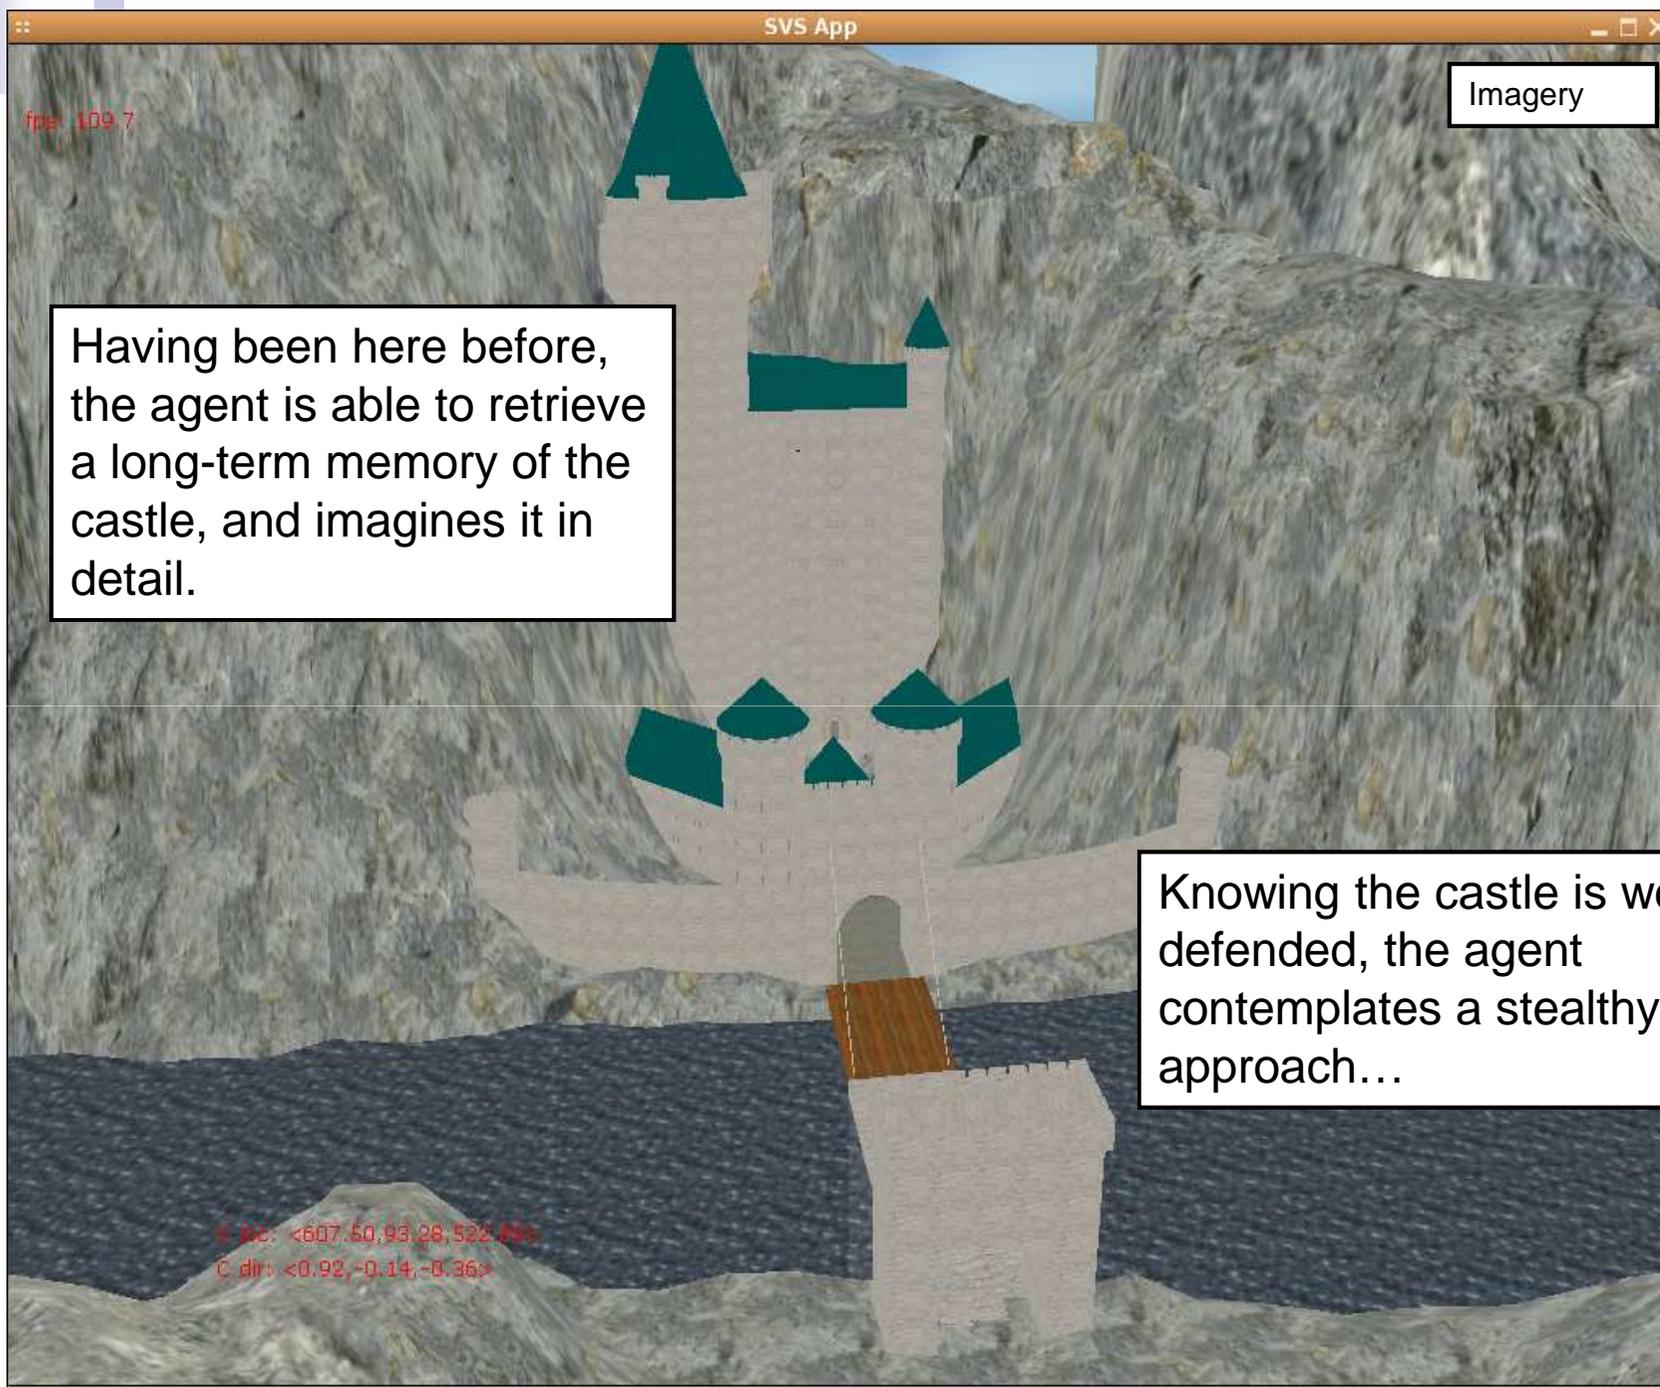
result: yes

# Outline

- Motivation
- Design
- SVS Scene Graphs
- Spatial Generation and Extraction
- Visual Generation and Extraction
- **Example Problem**
- Conclusion, Nuggets and Coal

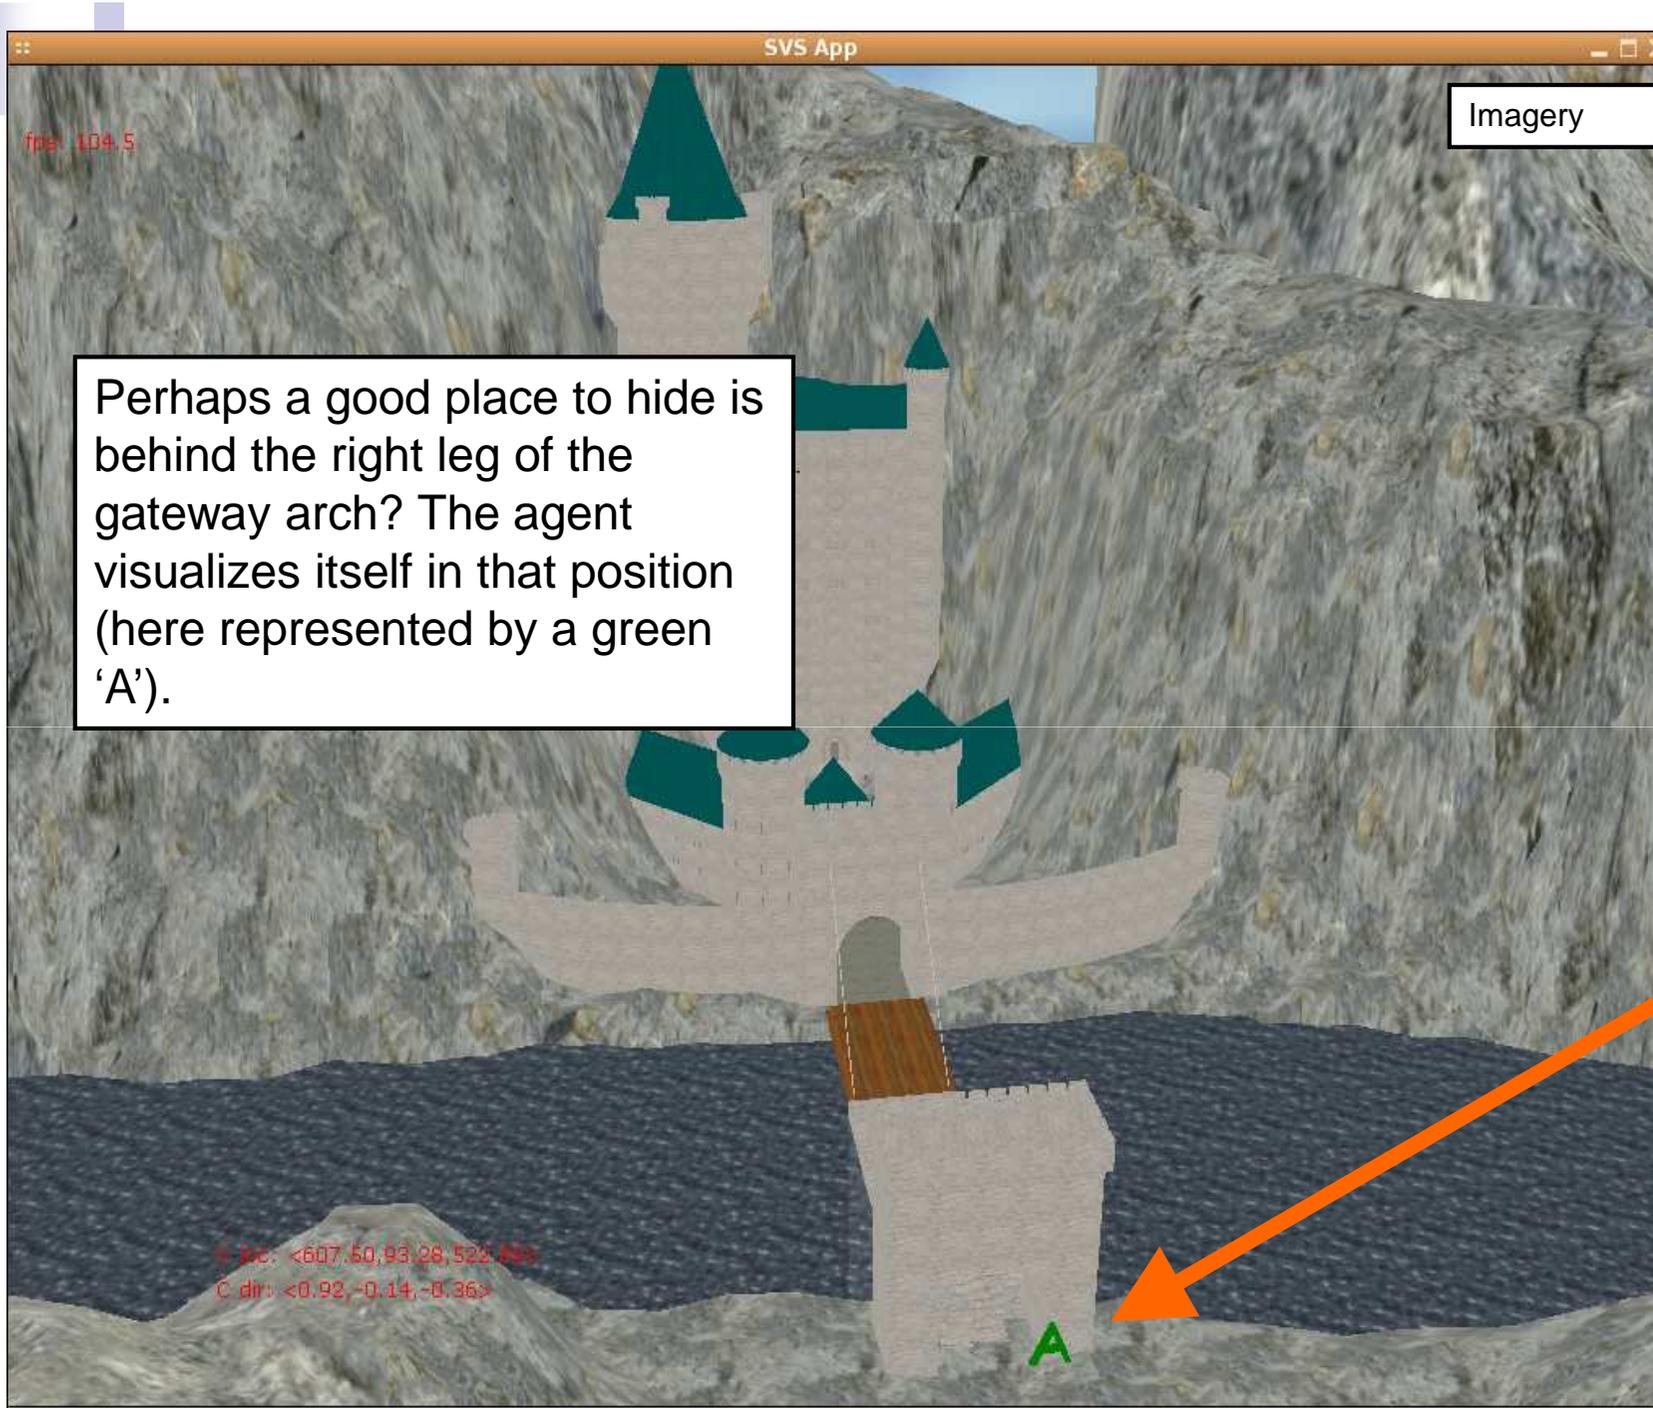


The agent approaches an enemy castle, and hides in the distance.

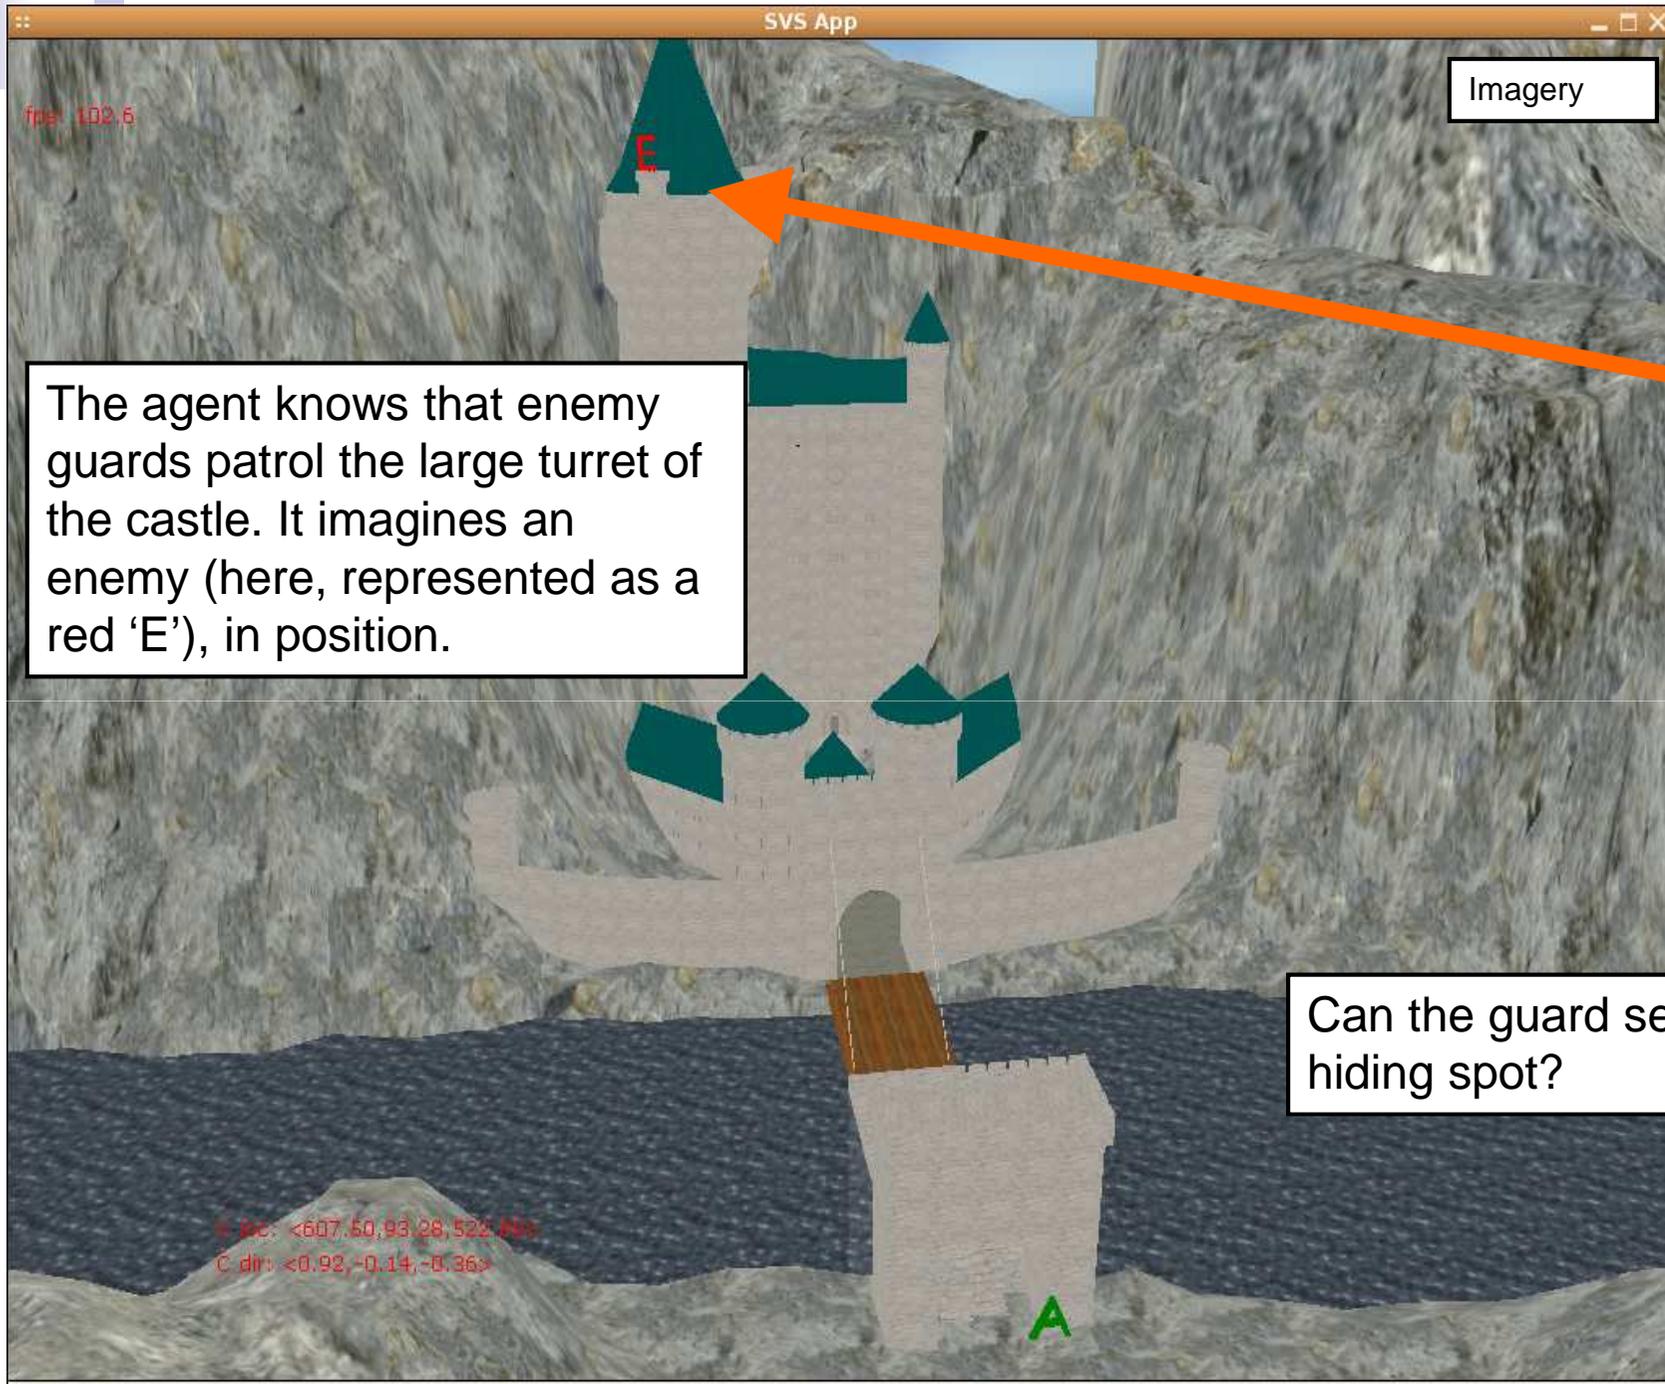


Having been here before, the agent is able to retrieve a long-term memory of the castle, and imagines it in detail.

Knowing the castle is well-defended, the agent contemplates a stealthy approach...

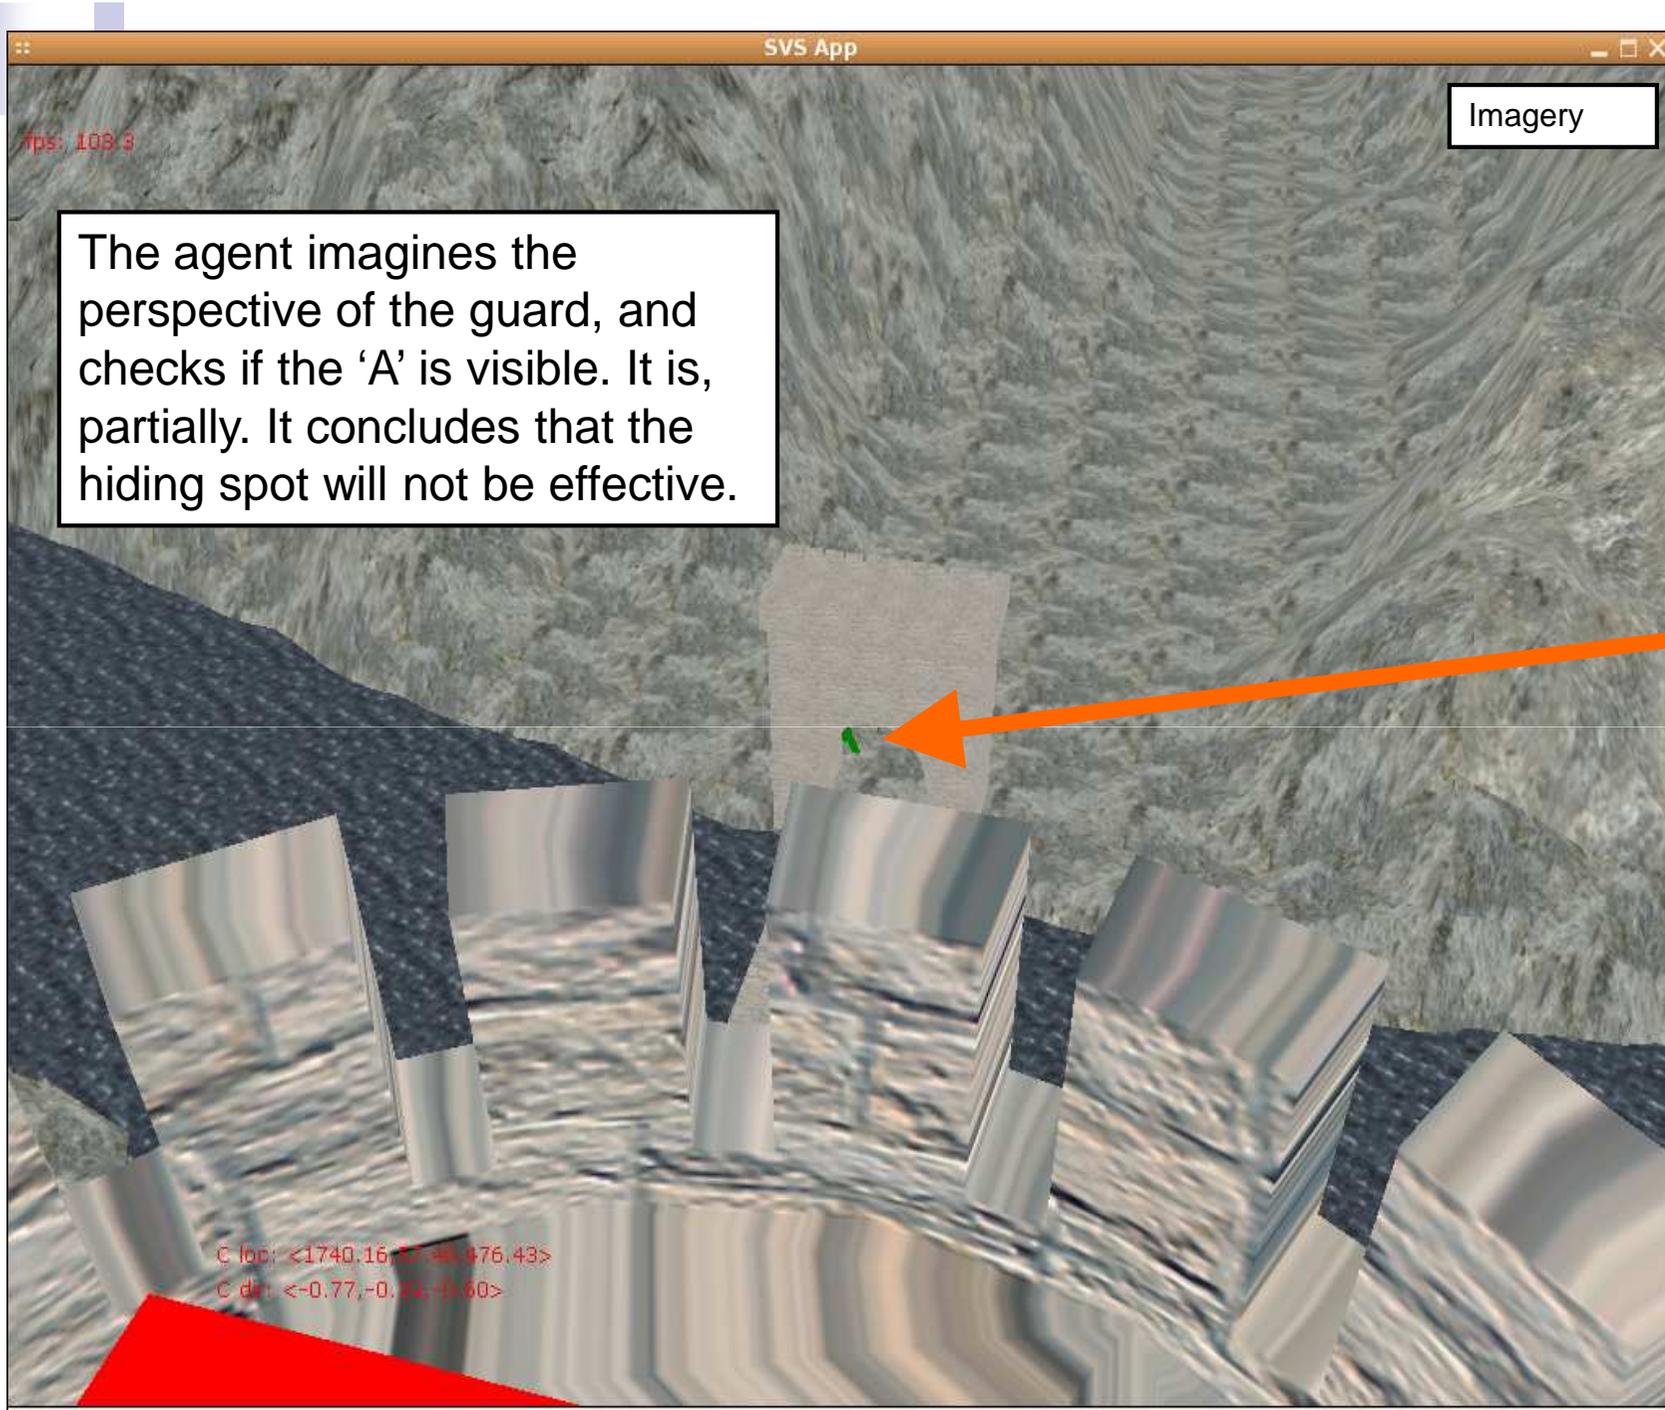


Perhaps a good place to hide is behind the right leg of the gateway arch? The agent visualizes itself in that position (here represented by a green 'A').



The agent knows that enemy guards patrol the large turret of the castle. It imagines an enemy (here, represented as a red 'E'), in position.

Can the guard see the hiding spot?



The agent imagines the perspective of the guard, and checks if the 'A' is visible. It is, partially. It concludes that the hiding spot will not be effective.

# Nuggets and Coal

## ■ Nuggets

- Architecture seems to be becoming simpler, not more complex
- Soar now has direct access to the scene graph
- Episodic memory integration is close to being possible

## ■ Coal

- Requires the environment (or low-level perceptual system) to provide a scene graph
- Integration with environments still requires lots of custom code
- Scene graph representation in SVS is dictated by graphical modelers and their tools
- Implementation isn't done yet