

**Introducing
Constrained Heuristic Search
to the
Soar Cognitive Architecture**
(29th Soar Workshop, University of Michigan)

**Sean A. Bittle
Mark S. Fox**

June 25th, 2009



university of toronto
mechanical & industrial engineering



Agenda

- Introduction: The problem, objectives of research
- Review:
 - Soar Cognitive Architecture
 - Constraint Satisfaction Problems (CSP)
 - Hyper-Heuristics
 - Constrained Heuristic Search (CHS)
- Design of CHS-Soar
 - Learning via subgoaling and chunks (Soar 8.6.3)
 - Experiments and Results
 - Reinforcement learning (Soar 9.0)
- Future Work and Issues

Introduction

General problem solving and domain independent learning are central goals of AI research on cognitive architectures.

Problem:

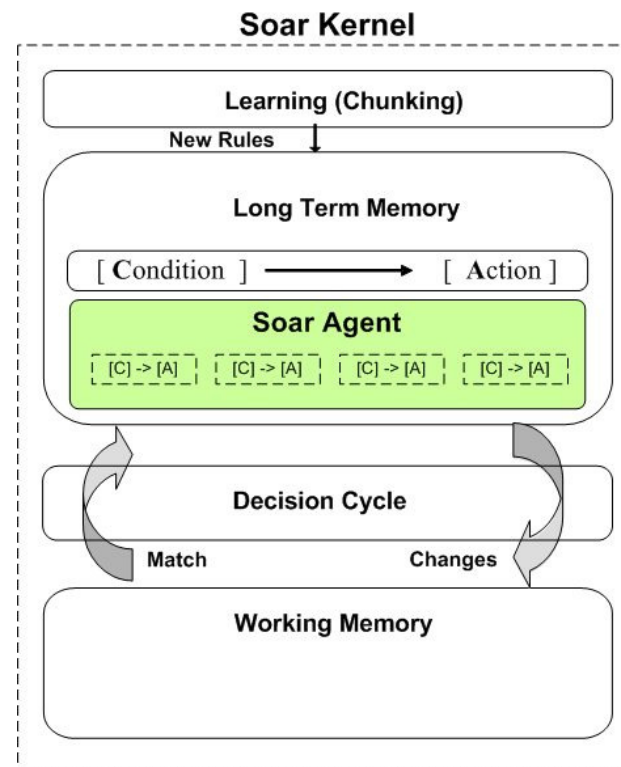
- However, there are few examples of domain independent learning in cognitive architectures

Objective:

- Demonstrate Soar can learn and apply domain independent knowledge


Soar Cognitive Architecture

- Current work based on Soar version 8.6.3, and Soar 9.0



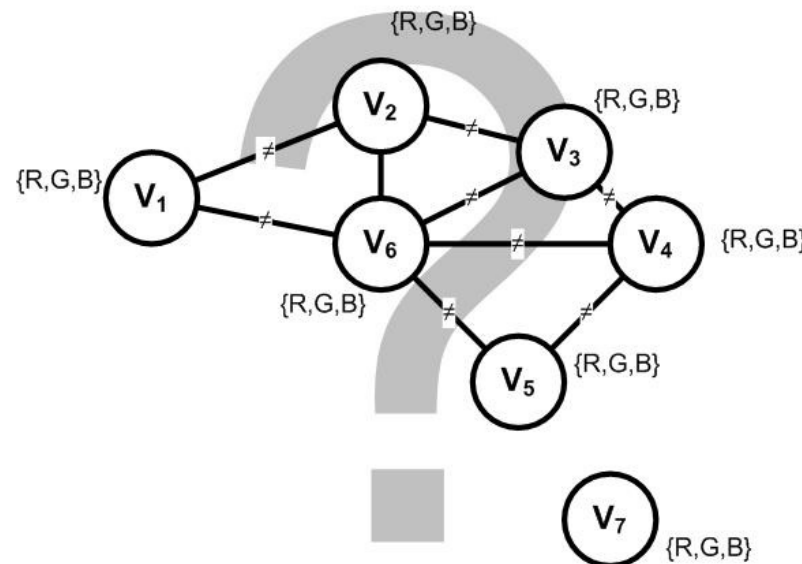
Constraint Satisfaction Problems (CSP)

- Constraint satisfaction is a sub-domain of constraint programming dealing with problems defined over a finite domain
- More formally, CSP consists of a finite set of:
 - Variables (X_1, X_2, \dots, X_n)
 - Constraints (C_1, C_2, \dots, C_n)
 - Each variable has a finite domain D_i of possible values
- Useful to represent CSP as a binary constraint graph



Constraint Satisfaction Problems (CSP)

- Backtrack search is the general approach used to solve a CSP
- General-purpose methods can provide ways to improve backtrack search efficiency:
 - Can we detect inevitable failure early? → Propagation
 - Which variable should be assigned next? → Variable Ordering
 - In what order should its values be tried? → Value Ordering



Use heuristics to guide variable and value ordering

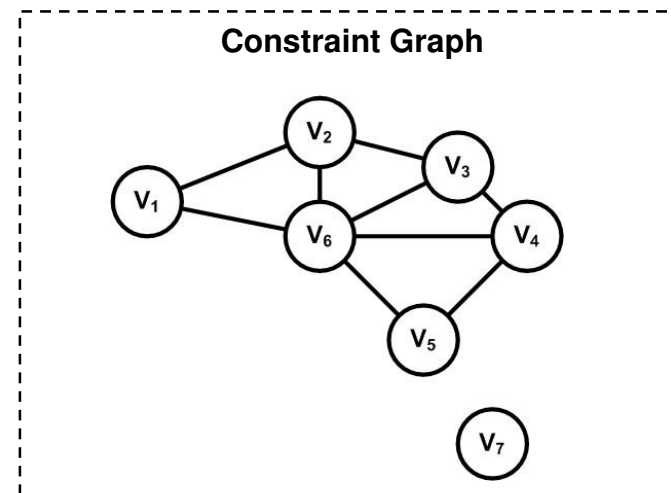
Hyper-Heuristics

Problem with variable and value ordering heuristics is “effective generality”

- Hyper-Heuristics are *“Heuristics to Choose Heuristics”*
- A hyper-heuristic is a high-level heuristic which uses some type of learning mechanism in order to choose (switch) between various low-level heuristics
- Most popular learning approach based on using a Genetic Algorithm (GA)

Constrained Heuristic Search (CHS)

- Developed by Fox, Sadeh, Bayken, 1989
- CHS is a problem solving approach that combination of constraint satisfaction and heuristic search where the definition of the problem space is refined to include:
 - *Problem Topology*
 - *Problem Textures*
 - *Problem Objective*




Constrained Heuristic Search (CHS)

What are Texture Measures?

- A texture measurement is a technique for distilling information embedded in the constraint graph into a form that heuristics can use
- A texture measurement is not a heuristic itself, but can be considered the constituent parts of a heuristic

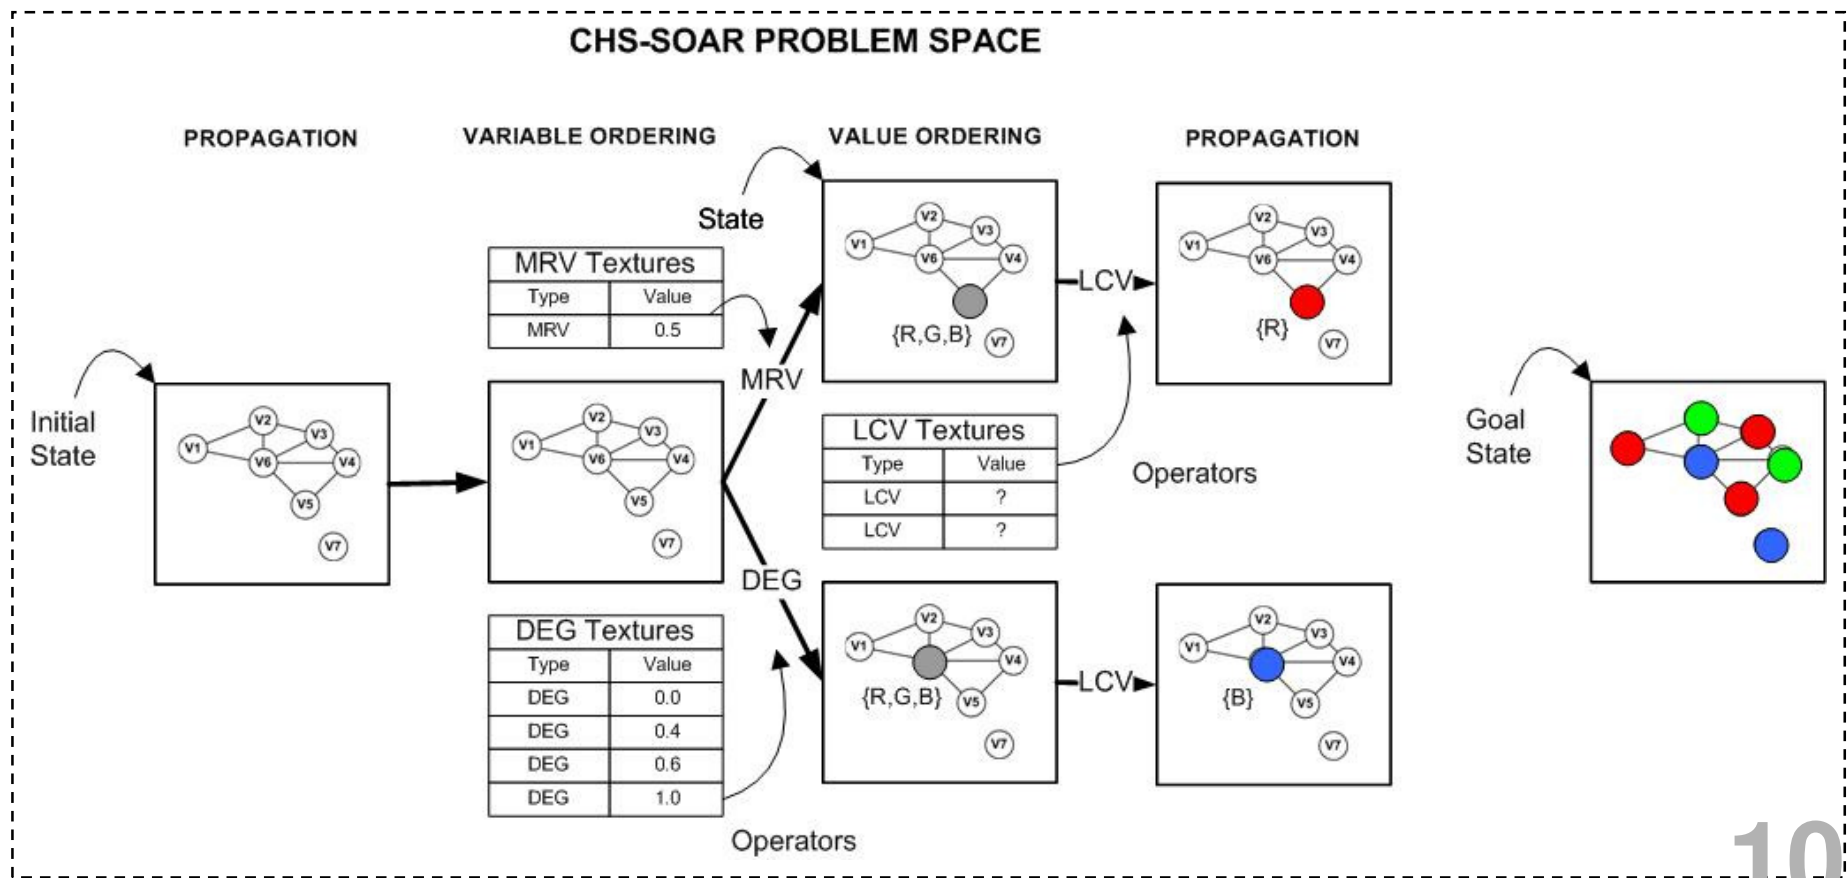
Ordering	Name	Texture	Heuristic
Variable	Minimum Remaining Values (MRV)	D_i , number of remaining values in domain of variable.	Select the variable with the smallest D_i , value e.g. pick the variable with the <u>fewest</u> legal values.



Design of CHS-Soar

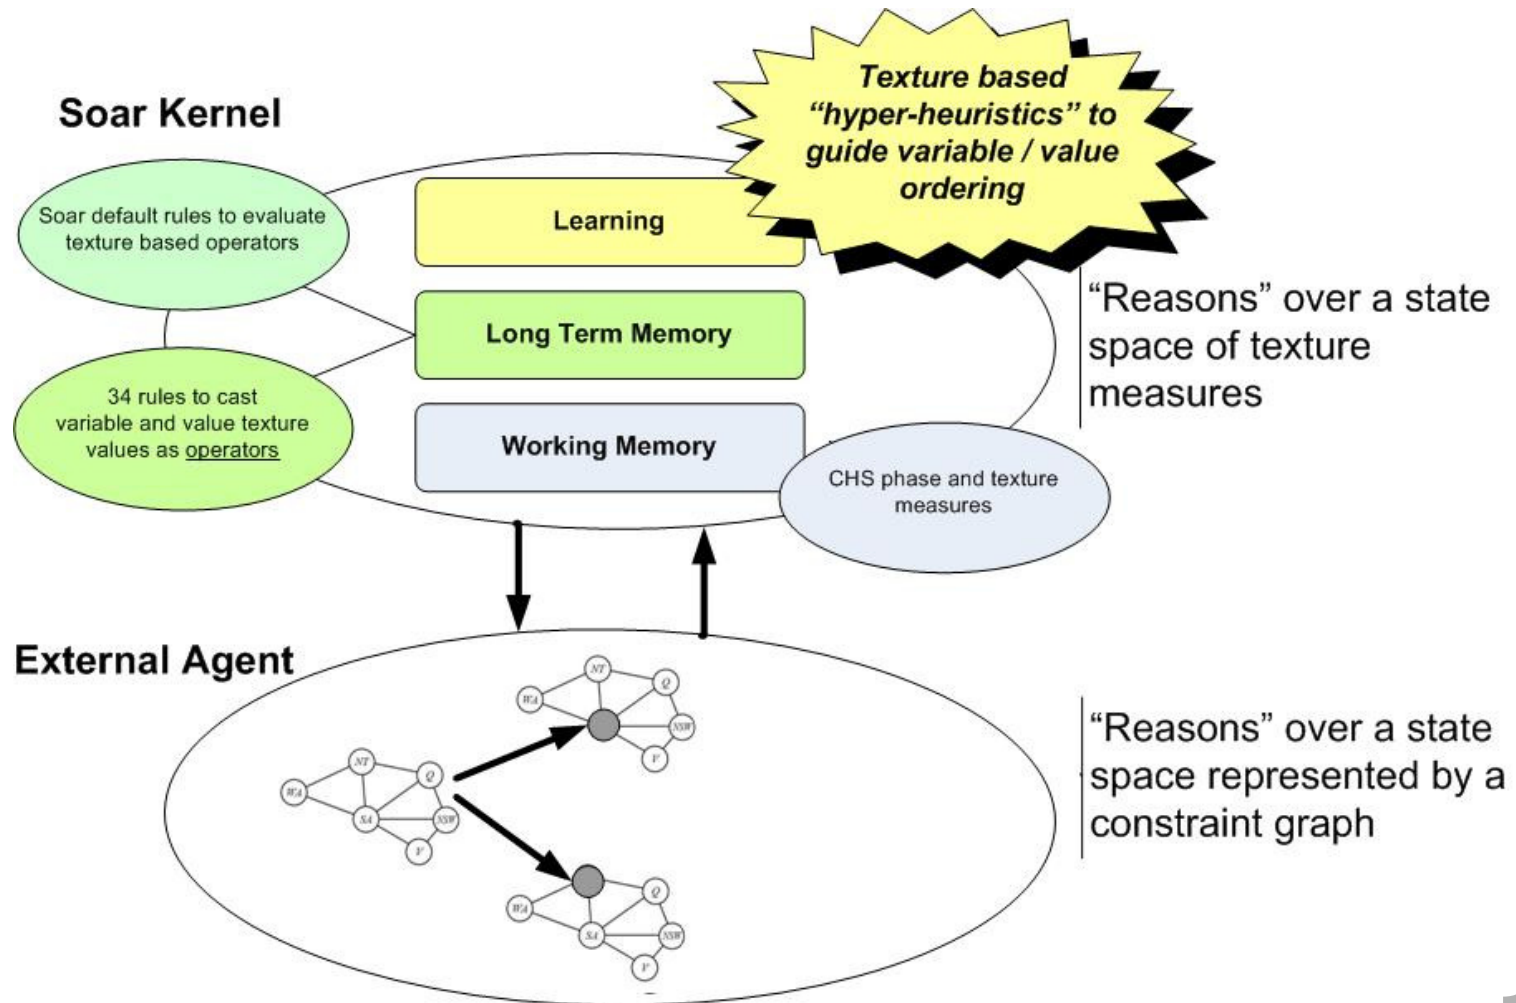
“How Does CHS-Soar Solve Problems?”

CHS-Soar problem solving is formulated by applying operators to states within a problem space in order to achieve a goal



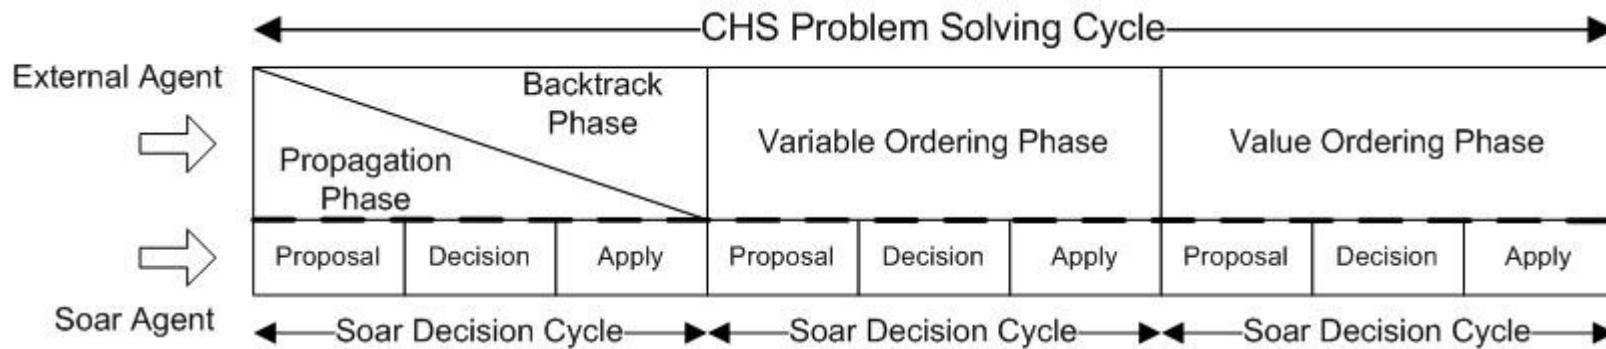

Design of CHS-Soar

“What are we trying to Learn?”



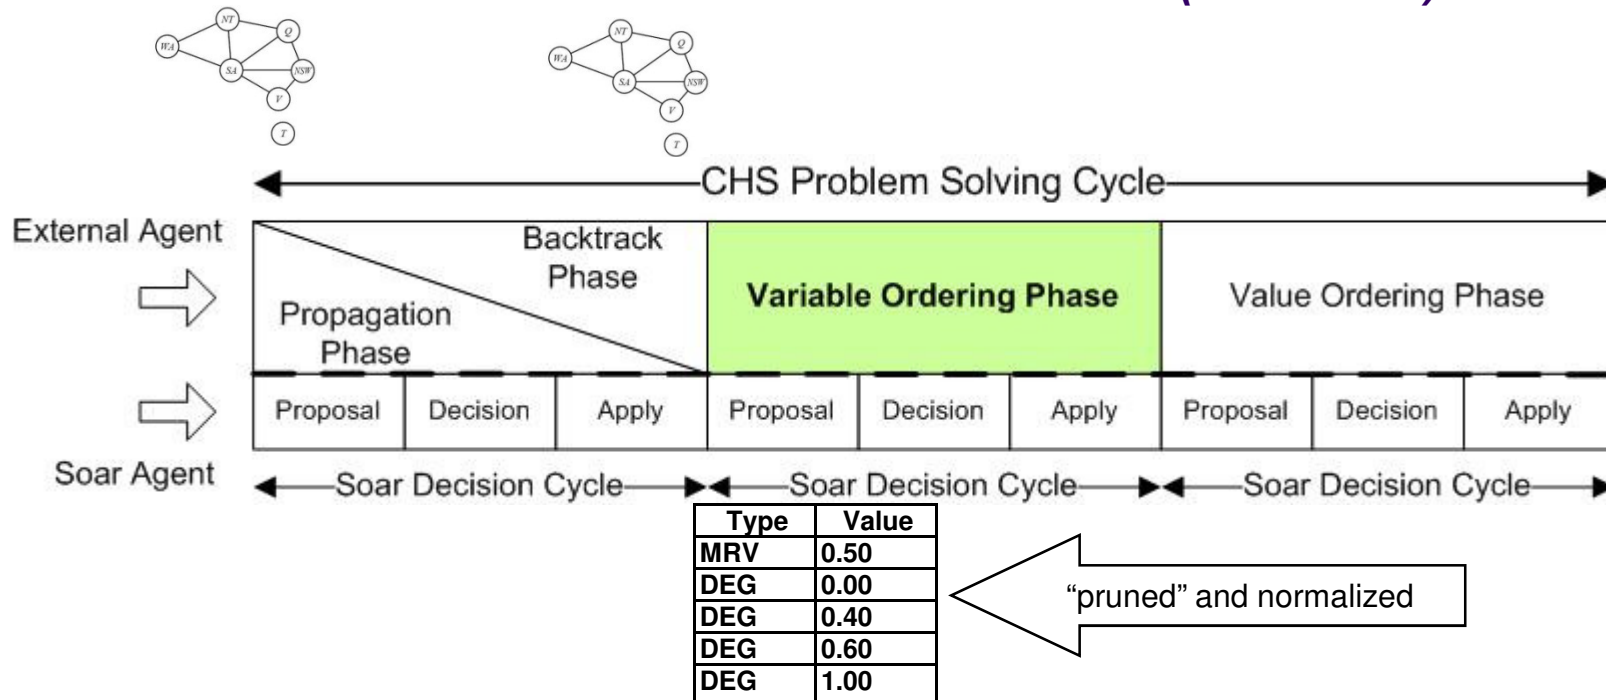
Design of CHS-Soar

“How does CHS-Soar Solve Problems (and Learn)?”



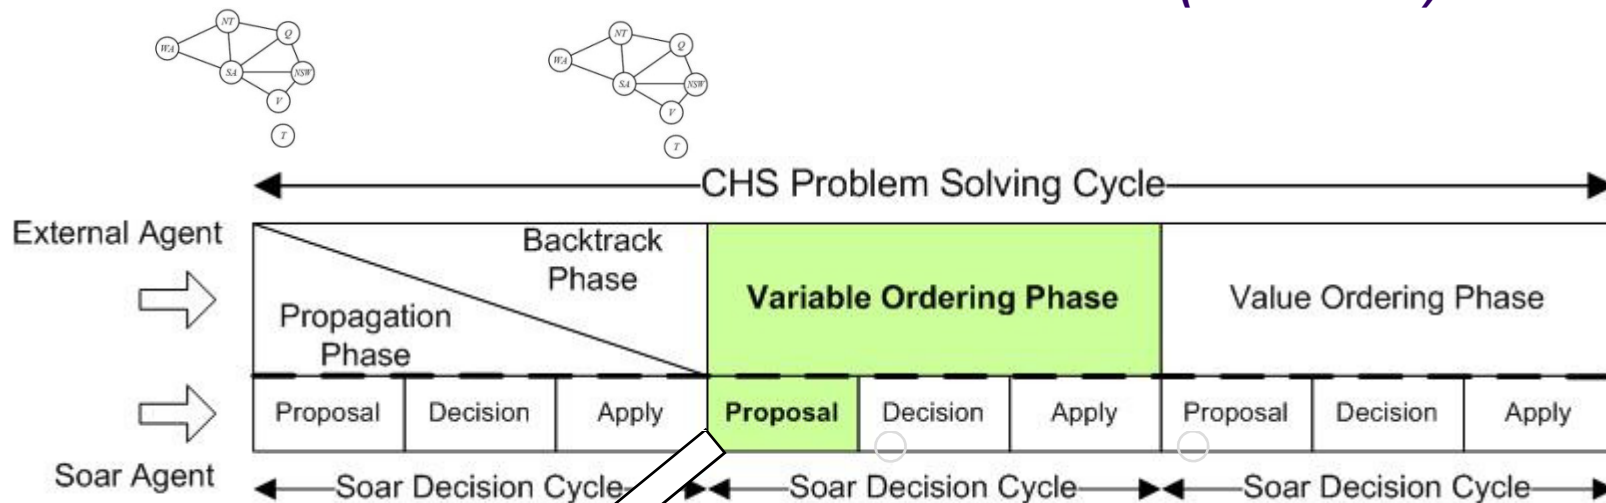
Design of CHS-Soar

“How does CHS-Soar Solve Problems (and Learn)?”



Design of CHS-Soar


“How does CHS-Soar Solve Problems (and Learn)?”



Type	Value
MRV	0.50
DEG	0.00
DEG	0.40
DEG	0.60
DEG	1.00

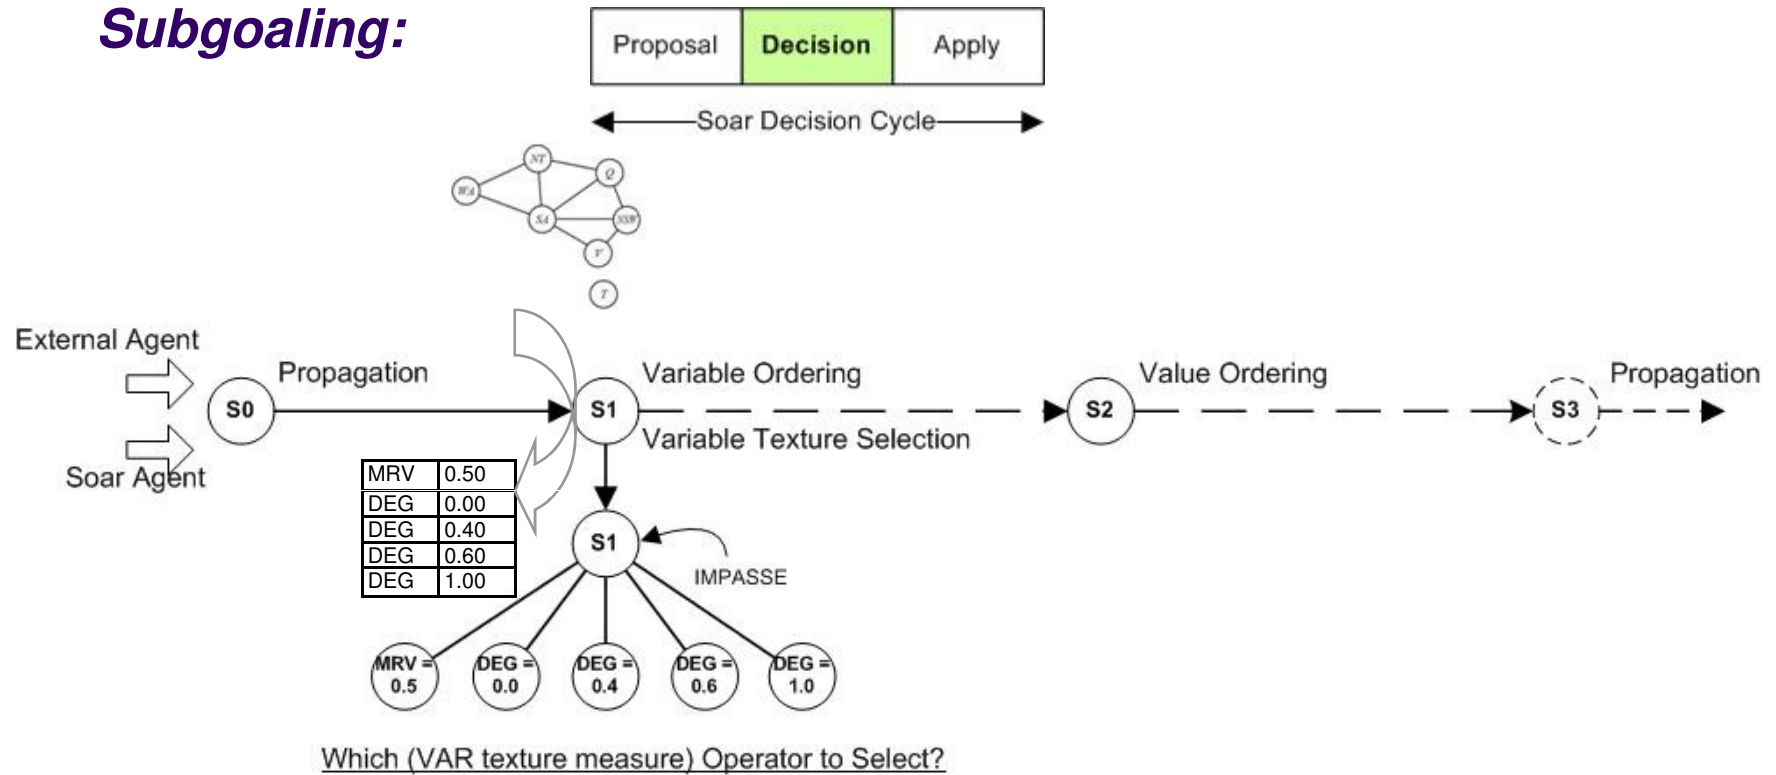
```

Variable Texture Propose Rule
sp {propose*SelectVariableTexture
  (state <s> ^name CHS-Soar
    ^problem-space <p>
    ^phase SelectVariableTexture
    ^top-state.vartextype <type>
    ^<type> <value> { <value> <= 1.0 } )
  -->
  (<s> ^operator <op> + )
  (<op> ^name SelectVariableTexture
    ^type <type>
    ^value <value> ) }
  
```



Design of CHS-Soar

Subgoaling:

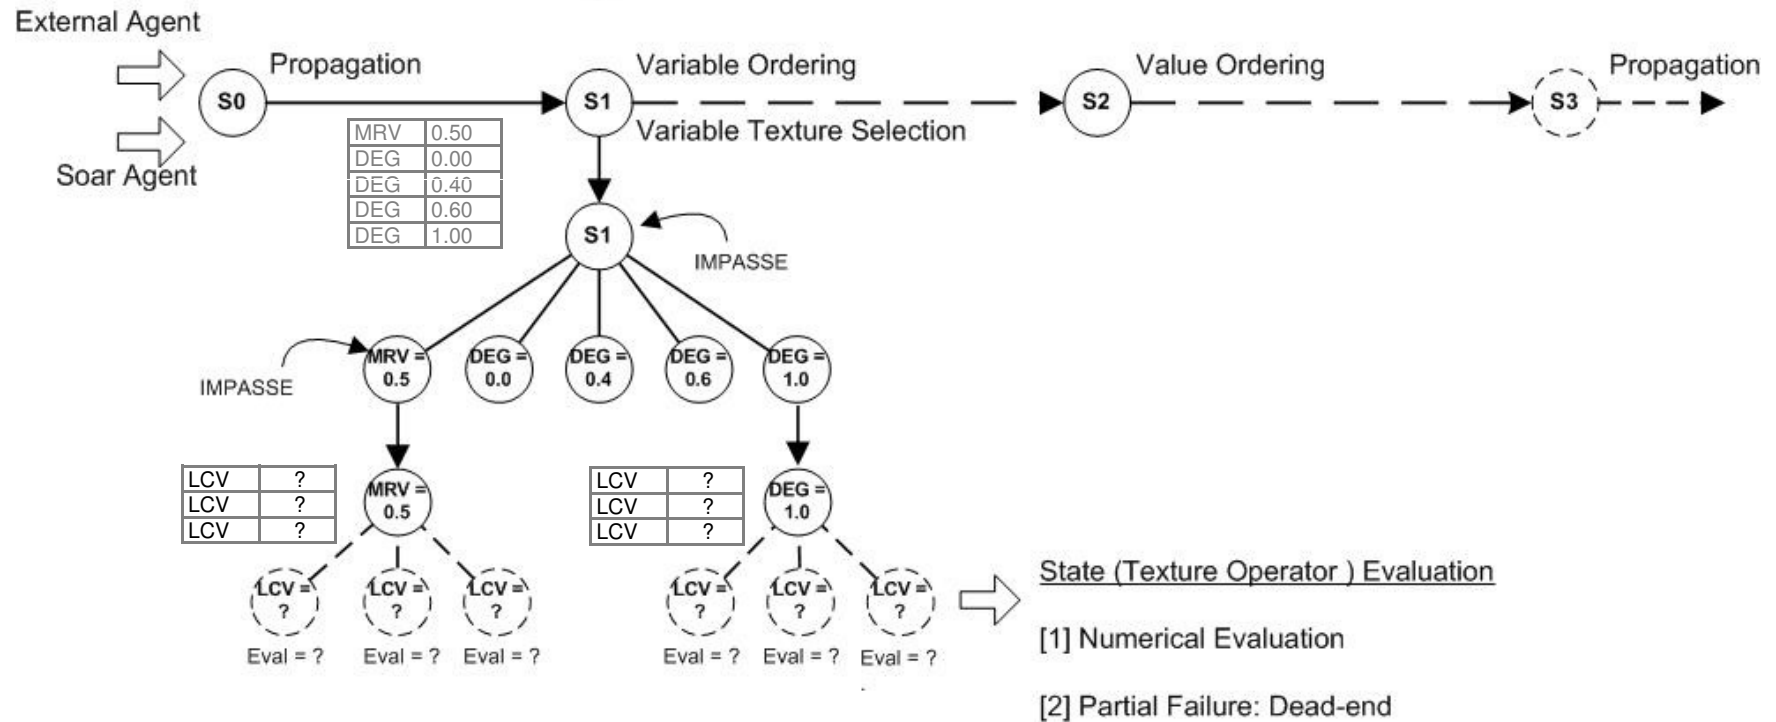



Design of CHS-Soar

Subgoaling:

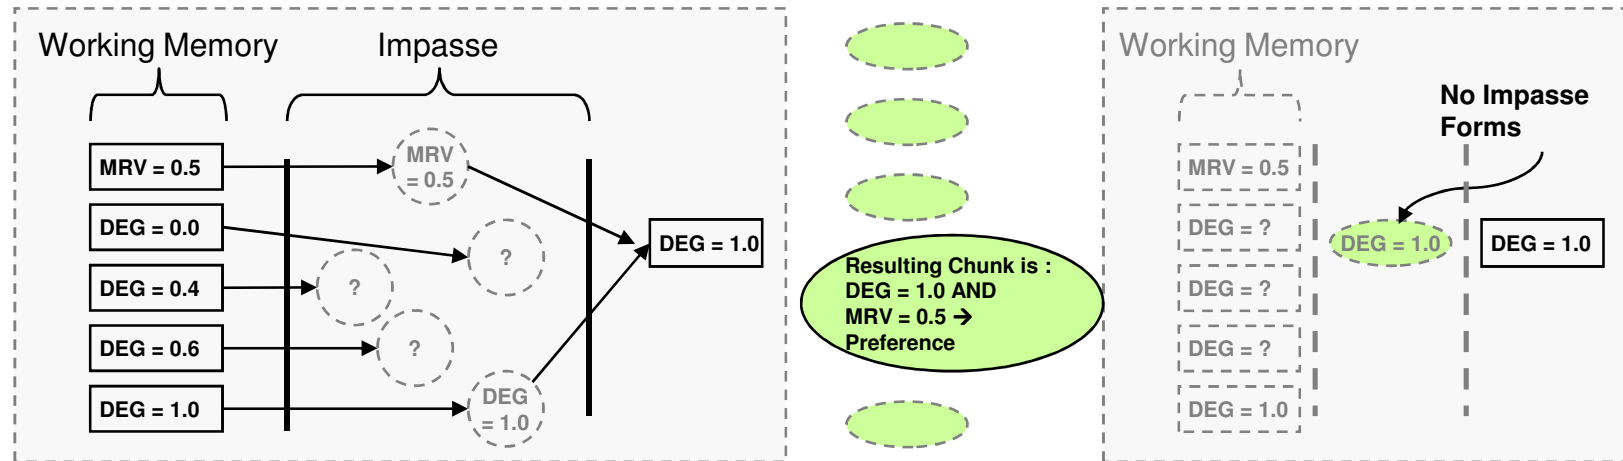


← Soar Decision Cycle →



Design of CHS-Soar

Subgoaling-Chunking:



Standard Soar Chunk (Water Jugs)

```
sp {chunk-54*d150*tie*2
:chunk
(state <s1> ^name water-jug ^operator <o1> +
^problem-space <p1>
^desired <d1> ^jug <j1> ^jug <j2>)
(<o1> ^name fill ^jug <j1>)
(<p1> ^name water-jug)
(<j1> ^contents 0 ^volume 3)
(<j2> ^contents 0 ^volume 5)
(<d1> ^jug <j3>)
(<j3> ^contents 1 ^volume 3)
-->
(<s1> ^operator <o1> >) }
```

CHS-Soar Binary Chunk (decoupled from problem type)

```
sp {chunk-514*d513*tie*4
:chunk
(state <s1> ^phase |SelectVariableTexture|
^top-state <s1> ^name |CHS-Soar| ^desired <d1>
^operator <o1> + ^operator <o2> + ^problem-space <p1>)
(<d1> ^better higher)
(<o1> ^value 1. ^name |SelectVariableTexture| ^type |DEG|)
(<o2> ^value 0.14 ^name |SelectVariableTexture| ^type |MRV|)
(<p1> ^name |CHS-Soar|)
-->
(<s1> ^operator <o2> < <o1>) }
```



Experiments and Results

Experiments conducted to investigate:

1. Intra (e.g. within) problem type learning and problem solving
2. Inter (e.g. across) problem type learning and problem solving

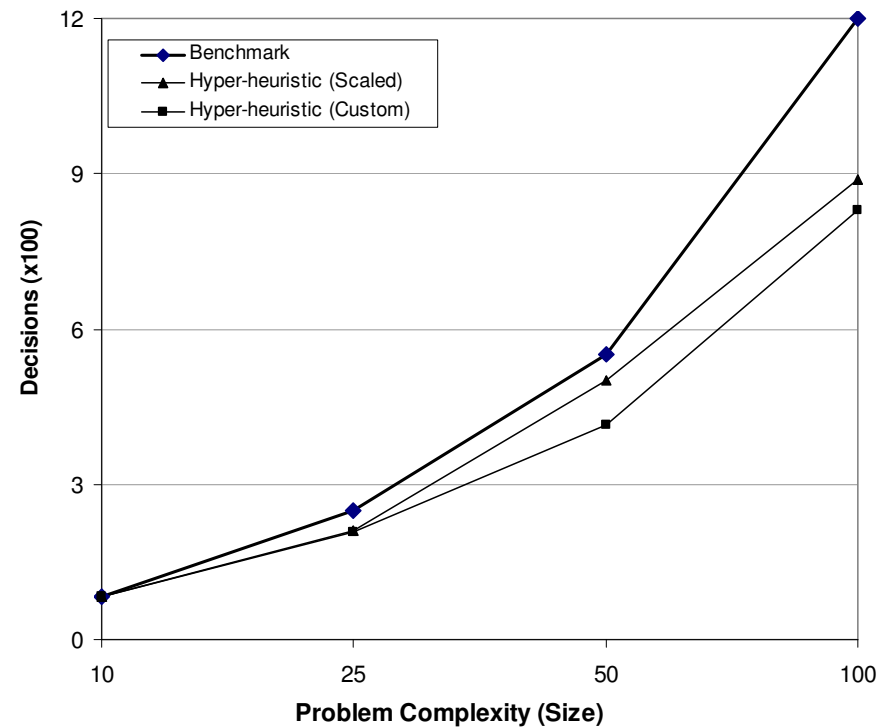
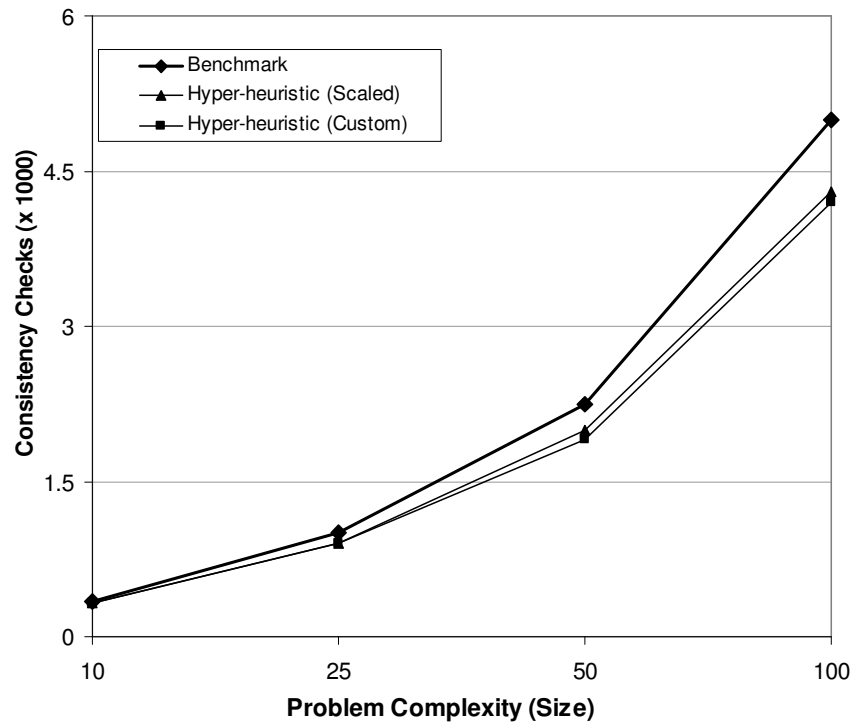
Problem types considered to date:

- Towers of Hanoi, Water Jugs
- Job Shop Scheduling (JSS)
- Map Coloring
- Radio Frequency Assignment Problem (RFAP)
- N-Queens
- Random CSP's
- Vehicle Routing Problem (VRP)

Experiment 1:

Intra-Problem Solving and Learning

Map Coloring Problem


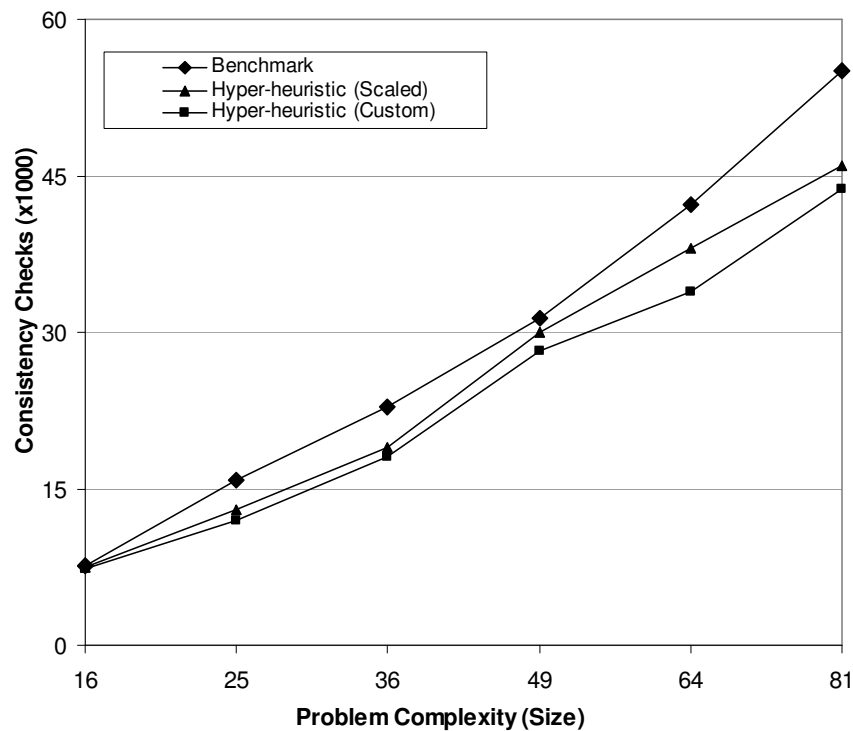


Learned non-max/min texture based rules delivery superior problem-solving performance over traditional heuristics

Experiment 1:

Intra-Problem Solving and Learning

Job Shop Scheduling Problem


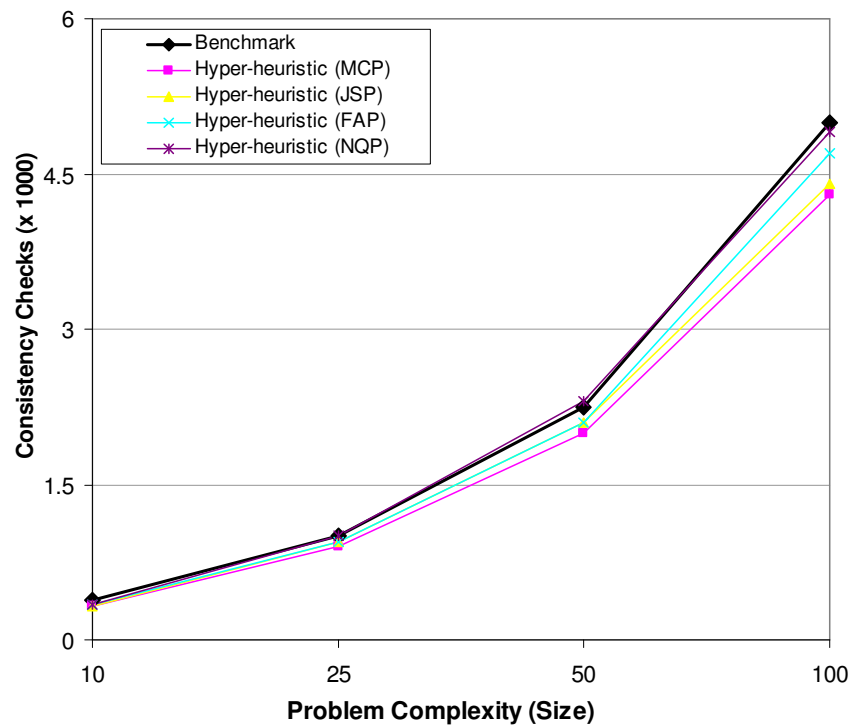


Learned non max/min texture based rules can scale to deliver superior problem-solving performance over traditional heuristics

Experiment 2:

Inter-Problem Solving and Learning

Map Coloring Problem

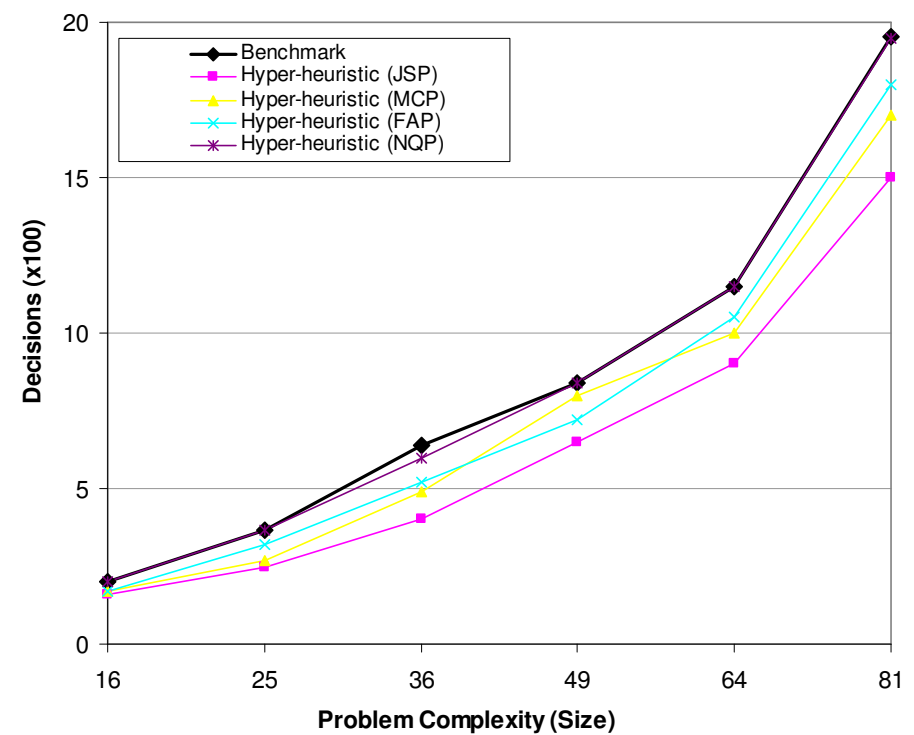



Learned rules while solving one problem type can be successfully be applied in solving different problem types and deliver superior problem-solving performance

Experiment 2:

Inter-Problem Solving and Learning

Job Shop Scheduling Problem



Learned rules while solving one problem type can be successfully be applied in solving different problem types that scale and deliver superior problem-solving performance

Design of CHS-Soar-RL

Issues with Subgoal-ing-Chunking (Soar 8.6.3):

- Chunk preferences are fixed (drawback)
- Chunking - subgoal-ing, allows us to “look-ahead” (benefit)

Issues with Reinforcement Learning (RL, Soar 9.0):

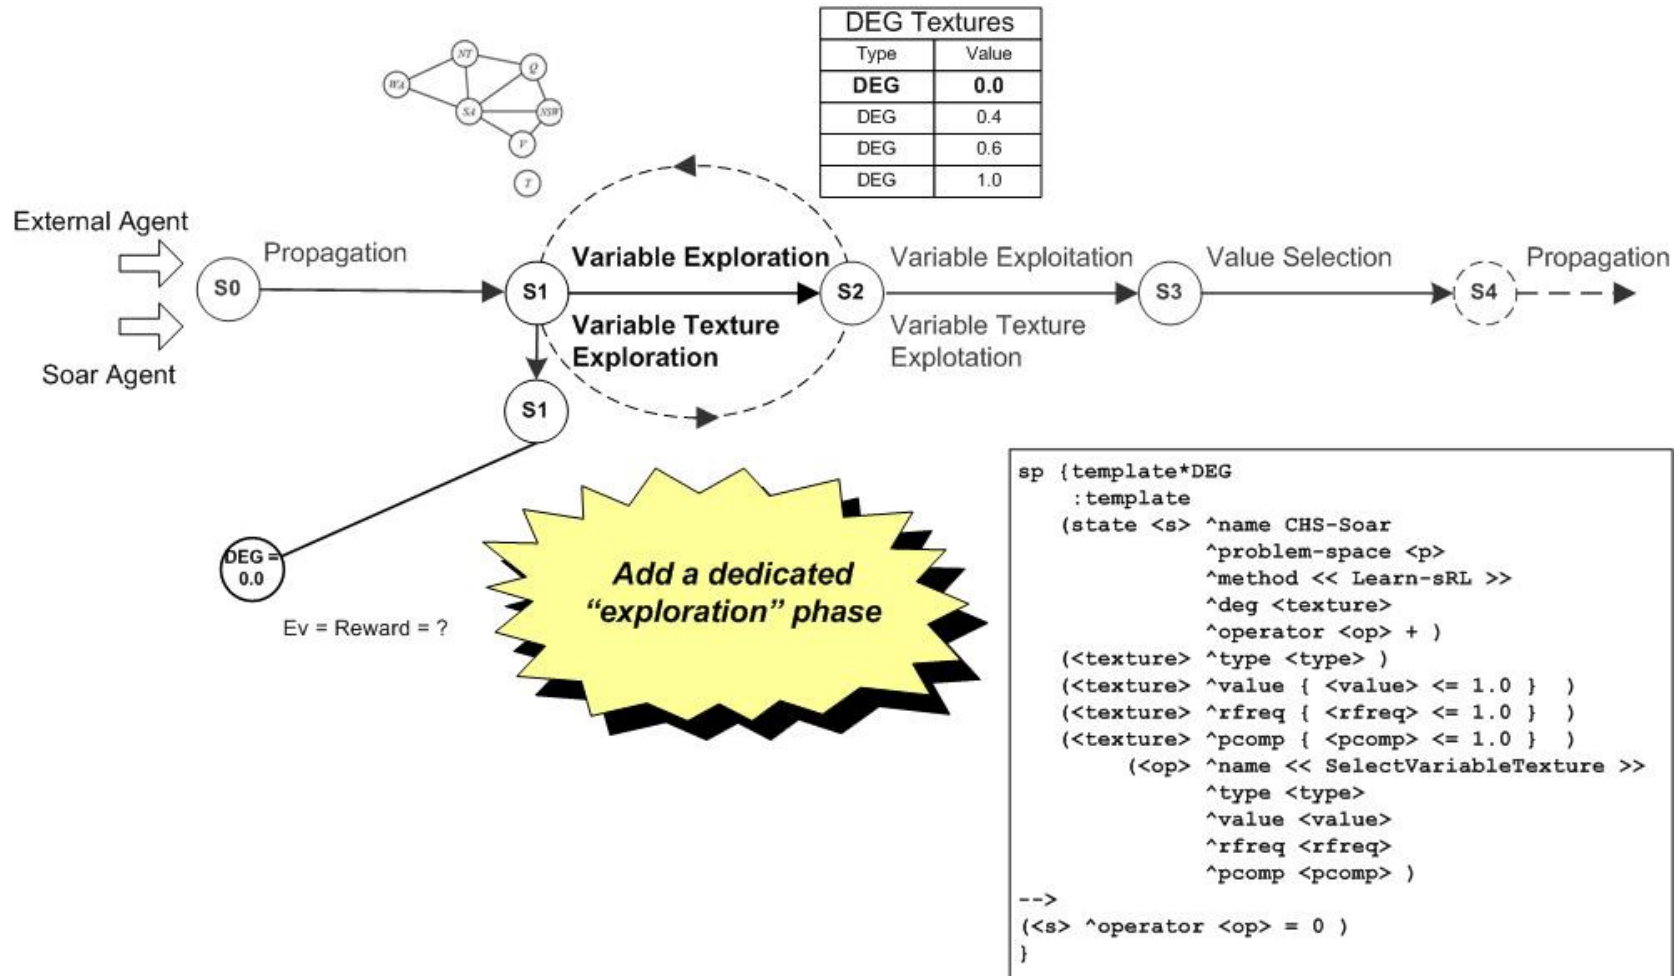
- RL rules can change numerical preferences (benefit)
- Does not allow us to subgoal in order to “look-ahead” (drawback)

Design goal of CHS-Soar-RL is to combine the benefits of both

- Allow us to “look-ahead”
- Use RL which allow num. preferences to change

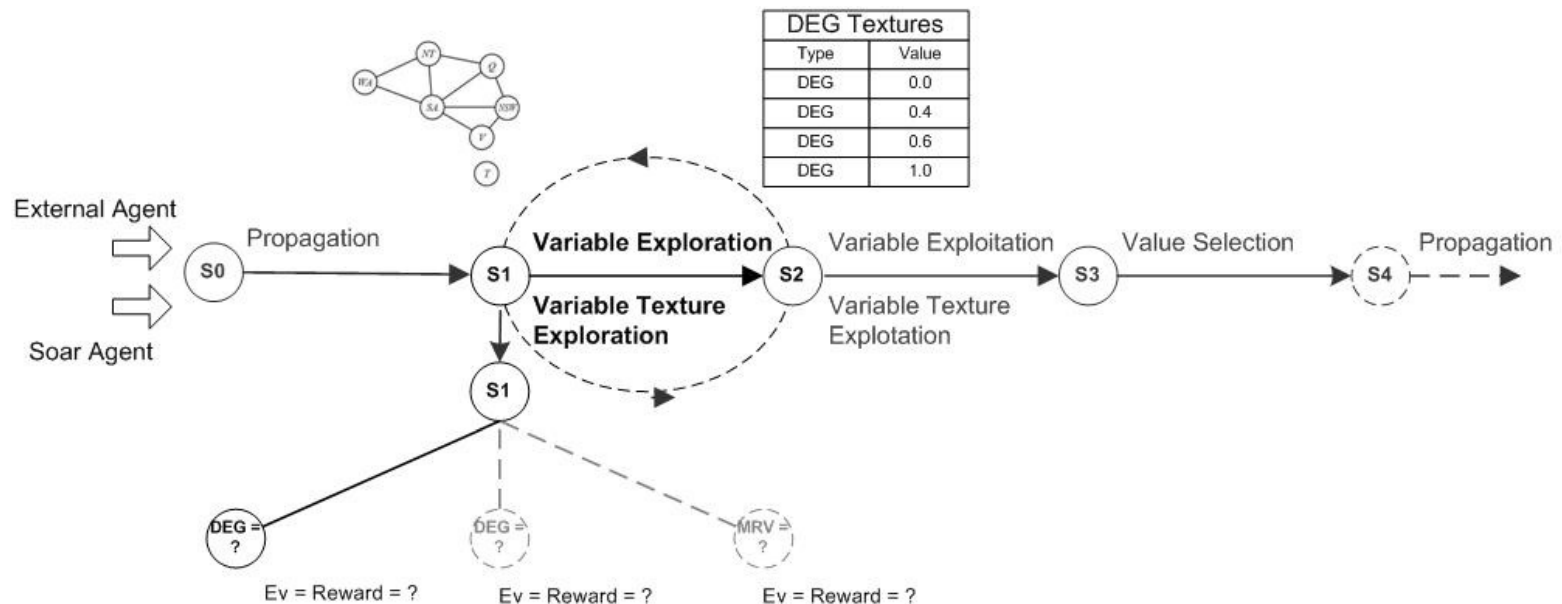
Design of CHS-Soar-RL

How can we “look-ahead” with RL?



Design of CHS-Soar-RL

How can we “look-ahead” with RL?



```

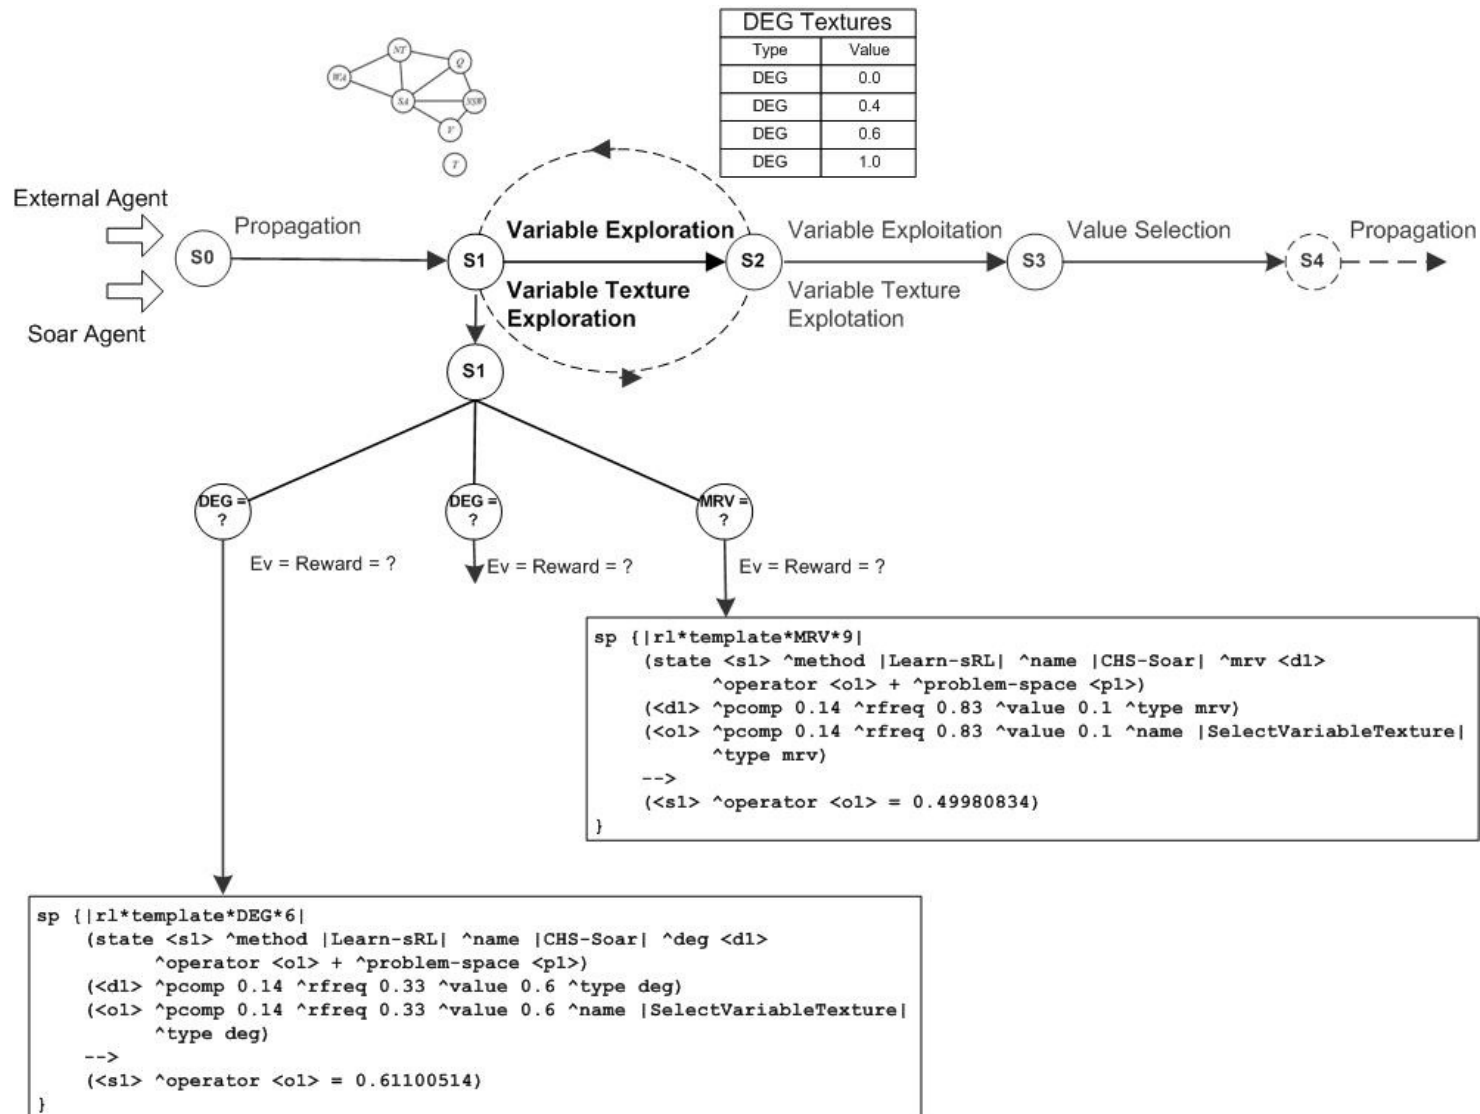
sp {propose*ExploreVariableTexture_DEG
  (state <s> ^name CHS-Soar
    ^problem-space <p>
    ^phase ExploreVariableTexture
    ^method << Learn-sRL >>
    ^deg <texture>)
  (<texture> ^type <type> )
  (<texture> ^value { <value> <= 1.0 } )
  (<texture> ^rfreq { <rfreq> <= 1.0 } )
  (<texture> ^pcomp { <pcomp> <= 1.0 } )
  -->
  (cmd rl --set learning on)
  (<s> ^operator <op> + )
  (<op> ^name SelectVariableTexture
    ^type <type>
    ^value <value>
    ^rfreq <rfreq>
    ^pcomp <pcomp> )
}
  
```

```

sp {apply*ExploreVariableTexture_DEG
  (state <s> ^name CHS-Soar
    ^problem-space <p>
    ^operator <op>
    ^phase { <phase> = ExploreVariableTexture }
    ^method << Learn-sRL >>
    ^reward-link <r>)
  (<op> ^name SelectVariableTexture
    ^type { <type> = deg }
    ^value <value>
    ^rfreq <rfreq>
    ^pcomp <pcomp> )
  -->
  (<s> ^name CHS-Soar)
  (<r> ^reward.value (float (exec explorevariabletexture <s> || <type> || <value> )))
  (<s> ^phase GetVariableTexture_DEG)
  (<s> ^phase <phase> -)
}
  
```

Design of CHS-Soar-RL

How can we “look-ahead” with RL?





Nuggets

- Demonstrated integration of rule and constraint based reasoning
- Demonstrated CHS-Soar ability to reason about a small group of well known variable and value texture measures leading to improved solutions over traditional unary heuristics
- Demonstrated the ability to learn hyper-heuristics while solving one problem type can be successfully be applied in solving different problem types and deliver superior problem-solving performance over traditional combinations of unary heuristics
- Soar's rule based encoding dramatically expands the expressiveness of the hyper-heuristic by encoding the constituent textures of each heuristic-not simply the low level heuristics



Coal

- Limited only to CSP problems (and challenge of CSP representation)
- Effort to calculate textures can outweigh benefits
- Many “intermediate” texture measures evaluations provide no insight
- Textures are “proxies” for actual variables and value leads to random selections
- Scalability or results to more realistic CSP problems?
- Ability to export results for other CP solvers (i.e. ILOG) to use



Questions

*Introducing Constrained Heuristic Search
to the
Soar Cognitive Architecture*