

Graphical Models for Cognitive Architecture

Resolving the Diversity Dilemma

Paul S. Rosenbloom

Department of Computer Science & Institute for Creative Technologies
University of Southern California


June 26, 2009

The Diversity Dilemma


- ▶ Should an architecture's mechanisms be uniform or diverse?
- ▶ *Uniformity*: Minimal mechanisms combining in general ways
 - ▶ Appeals to simplicity and elegance
 - ▶ The “physicist’s approach”
 - ▶ *The Challenge*: Achieving full range of required functionality/coverage
- ▶ *Diversity*: Large variety of specialized mechanisms
 - ▶ Appeals to functionality and optimization
 - ▶ The “biologist’s approach”
 - ▶ The Challenge: Achieving integrability, extensibility and maintainability
- ▶ Want best of both worlds, but a choice seems inevitable
- ▶ 2
- ▶ Functionality tends to win, leading to the predominance of

Example: Soar

- ▶ Through version 8 was a uniform architecture
- ▶ Version 9 has become highly diverse



Soar 3-8



Soar 9

Proposal for Resolving the Dilemma

- ▶ Dig beneath architecture for uniformity at *implementation level* that supports diversity/functionality in architecture (and above)
 - ▶ Implementation level is normally just Lisp, C, Java, etc.
 - ▶ Impacts efficiency and robustness but usually not part of theory unless based on neural networks
- ▶ Base implementation level on *graphical models* for a uniform approach to symbol, probability and signal processing
 - ▶ Related to neural networks but broader
- ▶ Reconceive architectures via new implementation level
 - ▶ Reimplement, enhance and hybridize existing architectures
 - ▶ Develop new architectures
- ▶ ⁴ *Improve elegance, functionality, extensibility, integrability and*


Graphical Models

- ▶ Efficient computation with multivariate functions
 - ▶ By decomposition over partial independencies
 - ▶ For constraints, probabilities, speech, etc.

- ▶ Come in a variety of related flavors

- ▶ **Bayesian networks:** Directed, variable nodes

- ▶ E.g., $p(u, w, x, y, z) = p(u)p(w)p(x|u, w)p(y|x)p(z|x)$




- ▶ **Markov networks:** Und., variable nodes & clique potentials

- ▶ Basis for *Markov logic* and *Alchemy*

- ▶ **Factor graphs:** Und., variable & factor nodes

- ▶ E.g., $f(u, w, x, y, z) = f_1(u, w, x)f_2(x, y, z)f_3(z)$



- ▶ Compute marginals via variants of

- ▶ Sum-product (message passing)

- ▶ 5 Monte Carlo (sampling)

Potential for the Implementation Level

- ▶ State-of-the-art algorithms for *symbol*, *probability* and *signal* processing all derivable from the *sum-product algorithm*
 - ▶ Belief propagation in Bayesian networks
 - ▶ Forward-backward in hidden Markov models
 - ▶ Kalman filters, Viterbi algorithm, FFT, turbo decoding
 - ▶ Arc-consistency in constraint diagrams
- ▶ Potential to go beyond existing architectures to yield an effective and uniform basis for:
 - ▶ Fusing symbolic and probabilistic reasoning (mixed)
 - ▶ Unifying cognition with perception and motor control (hybrid)
 - ▶ Bridging from symbolic to neural processing
- ▶ Raises hope of a uniform implementation level that integrates broad functionality at the architecture level

Scope of Sum-Product Algorithm

		<i>Message/Variable Domain</i>	
		Discrete	Continuous
<i>Message/Variable Range</i>	Boolean	Symbols	
	Numeric	Probability (Distribution)	Signal & Probability (Density)

- ▶ *Mixed* models combine Boolean and numeric ranges
- ▶ *Hybrid* models combine discrete and continuous domains
- ▶ *Hybrid mixed* models combine all possibilities
- ▶ *Dynamic hybrid mixed* models add a temporal dimension

Research Strategy

▶ Goals


- ▶ Evaluate extent to which graphical models can provide a uniform implementation layer for existing architectures
- ▶ Develop novel, more functional architectures
 - ▶ Enhancing and/or hybridizing existing architectures
 - ▶ Starting from scratch leveraging strengths of graphical models

▶ Initial approach

- ▶ Reimplement and enhance the Soar architecture
 - ▶ One of the longest standing and most broadly applied architectures
 - ▶ Exists in both uniform (Soar ≤ 8) and diverse (Soar 9) forms
- ▶ Start from the bottom up, implementing uniform version while looking for opportunities to more uniformly incorporate Soar 9's diversity plus critical capabilities beyond all versions of Soar

Progress to Date

- ▶ *Elaboration cycle* implementation via factor graphs
 - ▶ **Production match**
 - ▶ Production firing
- ▶ *Decision cycle* implementation via Alchemy (Markov logic)
 - ▶ **Elaboration phase**
 - ▶ Decision procedure
- ▶ With both also went beyond existing capability
 - ▶ *Lower complexity bound* for production match
 - ▶ Most recently, also began extension of WM beyond symbols
 - ▶ *Mixed* elaboration phase with simple *semantic memory* and *trellises*
- ▶ Still preliminary, partial implementations
 - ▶ Sufficient to demonstrate initial feasibility
 - ▶ Insufficient for full evaluation of impact on uniformity and functionality



Simple Mapping of Production Match onto Factor Graphs

P1: Inherit Color

C1: ($\langle v_0 \rangle \wedge \text{type} \langle v_1 \rangle$)

C2: ($\langle v_1 \rangle \wedge \text{color} \langle v_2 \rangle$)


-->

A1: ($\langle v_0 \rangle \wedge \text{color} \langle v_2 \rangle$)

Model as a Boolean function:

$$P_1(v_0, v_1, v_2) =$$

$$C_1(v_0, v_1) C_2(v_1, v_2) A_1(v_0, v_2)$$



WM is 3D Boolean array (obj x att x val)

1 when triple in WM

0 otherwise

Messages are Boolean vectors

1 when variable value possible

0 when variable value ruled out

Constant tests hidden in factors

WM is embedded in factors

Confuses binding combinations

May not check if rule completely



matches

Factor Graph Results

- ▶ Four issues have been resolved, yielding a new match algorithm
 - ▶ Tracks variable binding combinations only as needed
 - ▶ Complexity bound is exponential in *treewidth* rather than conditions
 - ▶ Avoids some duplicate instantiations on a cycle
 - ▶ Combines discrimination (α) and join (β) activities in uniform graph
- ▶ Solutions to binding confusion and rule matching increase number of rule variables processed at variable nodes
 - ▶ Yields exponential growth in message size and processing cost
 - ▶ Need to leverage tendency towards uniform values in WM and messages to reduce space and time costs

Hierarchical Memories and Messages

- ▶ N dimensional variant of *quad/octrees* (*exptrees*)
 - ▶ If entire space has one value, assign it to region
 - ▶ Otherwise, partition space into 2^N regions at next level, and recur
- ▶ WM & messages are *piecewise constant* functions
- ▶ Recently extended to *piecewise linear* functions
 - ▶ E.g., in 3D: $f(\langle x,y,z \rangle, r) = A_r + B_{r,1}x + B_{r,2}y + B_{r,3}z$
 - ▶ Natural compact representation for probabilities, signals, images, etc.
 - ▶ Also handles symbols by setting the Bs to 0
 - ▶ *Implemented mem. but not yet all of sum-product*



Example Match Times

P1: Inherit Color

C1: (<v0> ^type <v1>)

C2: (<v1> ^color <v2>)

-->

A1: (<v0> ^color <v2>)

With solutions to all four problems, rule graph comprises 8 factor nodes and 8 variable nodes.

WM is 16^3 in size, with 4 wmes

Sum of Products Redistribute P over S

Arrays

Exceeded heap space

1.7 sec.

Hierarchies

132 sec.

.25 sec.

~7

Unoptimized Lisp

~500

Implementing Soar's Elaboration Phase via Alchemy (Markov logic)

- ▶ Markov logic = First order logic + Markov networks
 - ▶ Compiles weighted FOL into a ground Markov logic network
 - ▶ Node for each ground predicate
 - ▶ Weight for each ground clause (clique potentials)
 - Along with links among all nodes in ground clause
 - ▶ Goals for implementation
 - ▶ Explore a *mixed* elaboration phase (rules & probabilities)
 - ▶ Explore semantic (fact) memory and *trellises*
 - ▶ Enable bidirectional message flow across rules
 - ▶ Normal elaboration cycle only propagates information forward
 - ▶ But need bidirectional settling for correct probabilities and *trellises*
-

Encoding

- ▶ Convert productions into logical implications
 - ▶ Define types for objects and values of triples
 - ▶ colors={Red, Blue, Green} and objects = {A, B, C, D, E, F}
 - ▶ Define predicates for attributes
 - ▶ Color(objects, colors) and Type(objects, objects)
 - ▶ Specify implications/clauses for rules
 - ▶ (Type(v0, v1) ^ Color(v1, v2)) => Color(v0, v2).
 - ▶ Add weights to clauses as appropriate
 - ▶ P1: Inherit Color
 - ▶ C1: (<v0> ^type <v1>)
 - ▶ C2: (<v1> ^color <v2>)
- ▶ Initialize evidence (db file) with WM
 - ▶ Color(C, Red), Color(D, Blue), Type(A, C), Type(B, D)
 - ▶ C3: (<v0> ^color <v2>)
- ▶ Semantic memory: weighted ground predicates: 10 Color(F, Green)
- ▶ Trellis: define via a pair of implications (*accept & reject* prefs.)
 - ▶ Size(step, size) => Size(step+1, size*2).
 - ▶ (Size(step, size1) ^ size1!=size2) => !Size(step, size2).

Alchemy Results

- ▶ Mapping basically works (modulo trellis strangeness)
 - ▶ Mixed representation with simple semantic memory and trellises
 - ▶ Match occurs via graph compilation not message propagation
 - ▶ As Alchemy compiles first-order clauses to ground network
 - ▶ All symbolic reasoning in compilation and probabilistic in propagation?
 - ▶ Falls short of uniform processing in the graph itself
 - ▶ Implies a three phase decision cycle
 1. Compile/match to generate a ground/instantiated network
 2. Perform probabilistic inference in the ground network
 3. Decide
 - ▶ Exptrees yield variants of Alchemy's *laziness* and *lifting*
 - ▶ Deal with *default values* and *groups of elements* processed in
-

Locality Implications

- ▶ Alchemy, and systems like it, get stuck in local minima
 - ▶ Generally considered a problem, but is it really?
- ▶ If Alchemy maps onto Soar's decision cycle then it only needs to perform K-Search
 - ▶ Conceptualize K-Search functionally as yielding local minima?
 - ▶ If so, then finding global minima, in general, requires PS-Search
- ▶ Implication would be that Alchemy should just yield local minima, but it also needs PS-Search on top of it
 - ▶ The same might then be said for all one-level, logical and/or probabilistic inference systems

Locality Implications (cont.)

- ▶ Taking this a step further, we can hypothesize functionally that:
 - ▶ Elaboration Cycle (10 ms): Local propagation of information
 - ▶ Decision Cycle (100 ms): Global propagation but only local minima
 - ▶ Problem Space Search (≥ 1 sec): Global minima (via sequence of local minima)
- ▶ But this implies that the elaboration cycle can't do global propagation of information
 - ▶ Explicit global: Creating unique identifiers
 - ▶ Implicit global: Non-monotonic (negated conditions, operator applications)
 - ▶ Accessing all of working memory?
- ▶ Could Soar function if global propagation were limited to the decision cycle?

-
- ▶ ¹⁸ I may need to answer this for a graphical implementation of Soar

Minerals

Gold

- ▶ New approach to cognitive architecture
 - ▶ Via a uniform graphical implementation level
 - ▶ Uncertain symbolic processing
 - ▶ Signal processing in inner loop
 - ▶ Potential bridge to neural
 - ▶ May resolve diversity dilemma
 - ▶ Improving elegance, scope, integrability and maintainability
- ▶ Early results on elaboration cycle/phase are encouraging
 - ▶ New match algorithm with improved complexity bound
 - ▶ Mixed elaboration phase with semantic memory and trellises

Coal

- ▶ Far from complete architecture
 - ▶ Combine two experiments
 - ▶ Add decisions, impasses, chunking
 - ▶ Incorporate Soar 9 extensions
- ▶ Locality may be Achilles heel
 - ▶ Or mapping from mechanism to implementation may be so complex as to lose benefits of uniformity in implementation
- ▶ May be too slow for actual use
- ▶ A common implementation level need not guarantee clean integrability