

Augmenting Soar with Non-Symbolic Processing
via the IO-Link

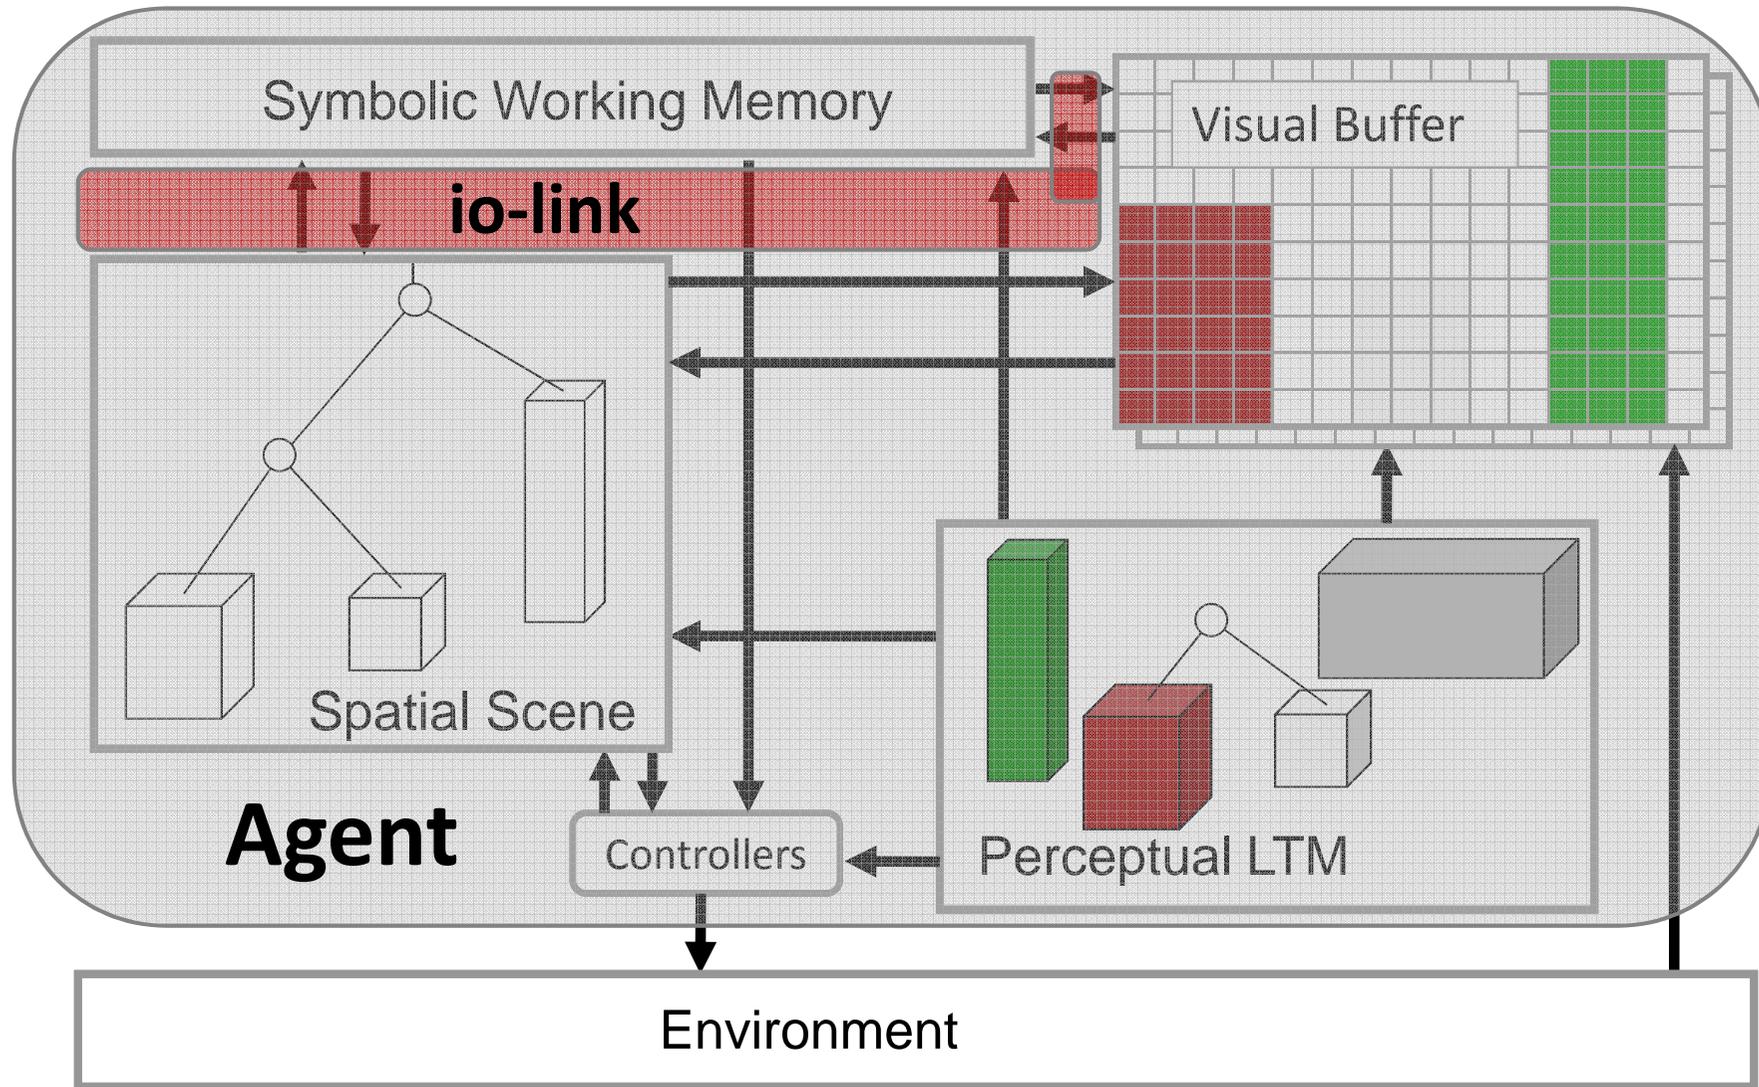
29th Soar Workshop

Samuel Wintermute
University of Michigan

Soar and SVS

- ▶ SVS adds spatial and visual processing to Soar
- ▶ This involves new non-symbolic memories
 - ▶ Memories cannot be “retrieved” into working memory
- ▶ Conceptual problem:
 - ▶ How should these memories be integrated with symbolic working memory?
- ▶ Engineering problem:
 - ▶ What should the actual software look like?

Soar/SVS Architecture

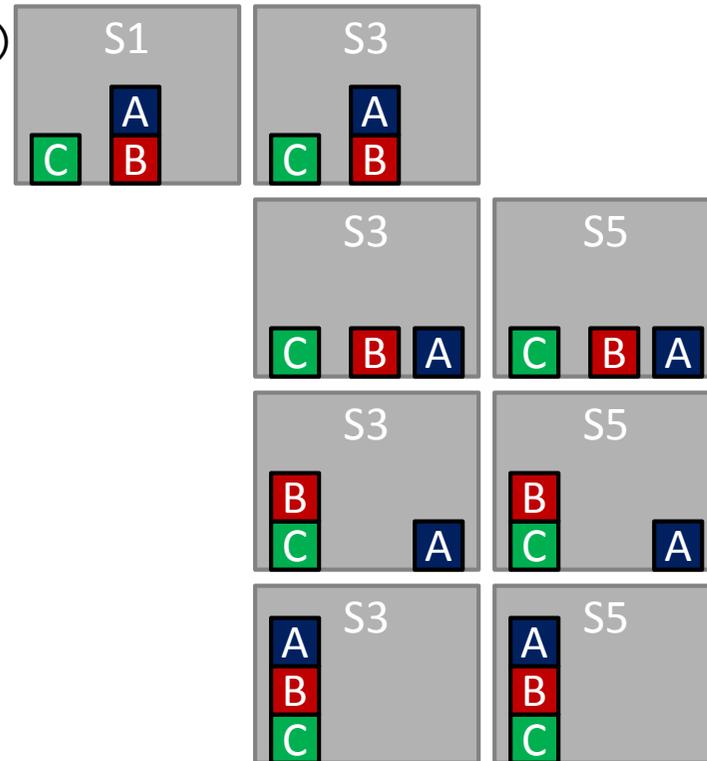


Goal of Integration

```
1: 0: 01 (initialize-blocks-world-look-ahead)
2: ==>S: S2 (operator tie)
3:   0: 08 (evaluate-operator)
4:   ==>S: S3 (operator no-change)
5:     0: C1 (move-block)
6:     ==>S: S4 (operator tie)
7:       0: 022 (evaluate-operator)
8:       ==>S: S5 (operator no-change)
9:         0: C2 (move-block)
10:         0: 028 (move-block)
11:         0: 019 (move-block)
12:         0: 018 (move-block)
13: 0: 05 (move-block)
14: ==>S: S6 (operator tie)
15:   0: 041 (evaluate-operator)
16:   ==>S: S7 (operator no-change)
17:     0: C3 (move-block)
18:     0: 046 (move-block)
19: 0: 038 (move-block)
20: 0: 035 (move-block)
blocks-world achieved
```

Goal of Integration

```
1: 0: 01 (initialize-blocks-world-look-ahead)
2: ==>S: S2 (operator tie)
3:   0: 08 (evaluate-operator)
4:   ==>S: S3 (operator no-change)
5:     0: C1 (move-block)
6:     ==>S: S4 (operator tie)
7:       0: 022 (evaluate-operator)
8:       ==>S: S5 (operator no-change)
9:         0: C2 (move-block)
10:        0: 028 (move-block)
11:        0: 019 (move-block)
12:        0: 018 (move-block)
```

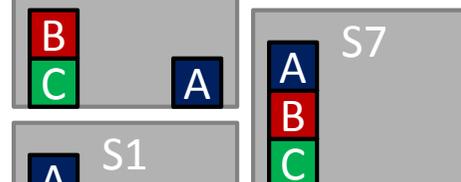
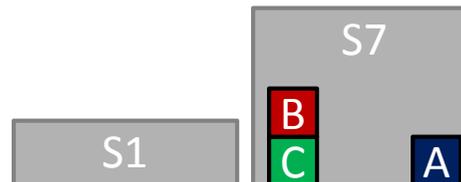
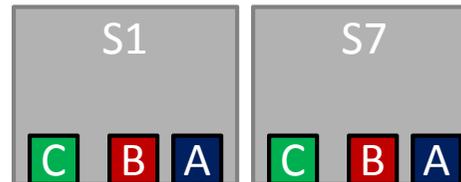
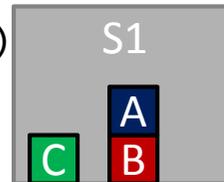


Goal of Integration

```

1: 0: 01 (initialize-blocks-world-look-ahead)
2: ==>S: S2 (operator tie)
3:   0: 08 (evaluate-operator)
4:   ==>S: S3 (operator no-change)
5:     0: C1 (move-block)
6:     ==>S: S4 (operator tie)
7:       0: 022 (evaluate-operator)
8:       ==>S: S5 (operator no-change)
9:         0: C2 (move-block)
10:        0: 028 (move-block)
11:        0: 019 (move-block)
12:        0: 018 (move-block)
13: 0: 05 (move-block)
14: ==>S: S6 (operator tie)
15:   0: 041 (evaluate-operator)
16:   ==>S: S7 (operator no-change)
17:     0: C3 (move-block)
18:     0: 046 (move-block)
19: 0: 038 (move-block)
20: 0: 035 (move-block)
blocks-world achieved


```



Imagery Rules


```
sp {apply*move-a-to-table
    (state <s> ^block-a <a>
      ^operator.name move-a-to-table)
    (<a> ^on <on>)
-->
    (<a> ^on <on> -
      ^on table)
}
```

Imagery Rules

```
sp {apply*move-a-to-table
  (state <s> 
    ^operator.name move-a-to-table)
-->
  (<s>  )
}
```

- ▶ In what scenes should the rule match?
- ▶ How should the scene be modified by the rule?
- ▶ Define and name these properties
 - ▶ Create a qualitative symbolic interface

Imagery Rules



- ▶ Imagery objects aren't in WM, but are tightly related to it
- ▶ How can this relationship be implemented?

Option: Interfacing via Action Operators

```
2: ==>S: S2 (operator tie)
3:   0: 08 (evaluate-operator)
4:   ==>S: S3 (operator no-change)
      0: XX (add-imagery-structure) * N
      0: XX (extract-property) * N
5:   0: C1 (move-block)
      0: XX (remove-imagery-structure)
      0: XX (add-imagery-structure)
      0: XX (extract-property) * N
      0: XX (remove-imagery-structure) * N
6: 0: 09 (move-block)
```

- ▶ Treat imagery as an external environment modified via actions
- ▶ Problems
 - ▶ Slow
 - ▶ Requires lots of knowledge
 - ▶ No truth maintenance
 - ▶ When are images added and removed?
 - ▶ When are properties valid?
- ▶ Not a tight integration with working memory

Option: Use EpMem/SMem Interface

- ▶ Separate interfaces at every state
 - ▶ Some truth maintenance benefit
- ▶ Parallel access, not necessarily operator-based

- ▶ Episodes and semantic structures are symbolic
 - ▶ Storage and retrieval are interesting, actually using the data is just regular Soar processing

- ▶ Imagery objects are not symbolic
 - ▶ “Storage” (WM → imagery) is different
 - ▶ Imagery structures are temporary, not long-term
 - ▶ “Retrieval” (imagery → WM) is different
 - ▶ Imagery objects can not be copied to WM, properties of objects are retrieved instead
 - ▶ Truth maintenance is a problem

- ▶ More dynamic interface is needed

Working Memory Integration

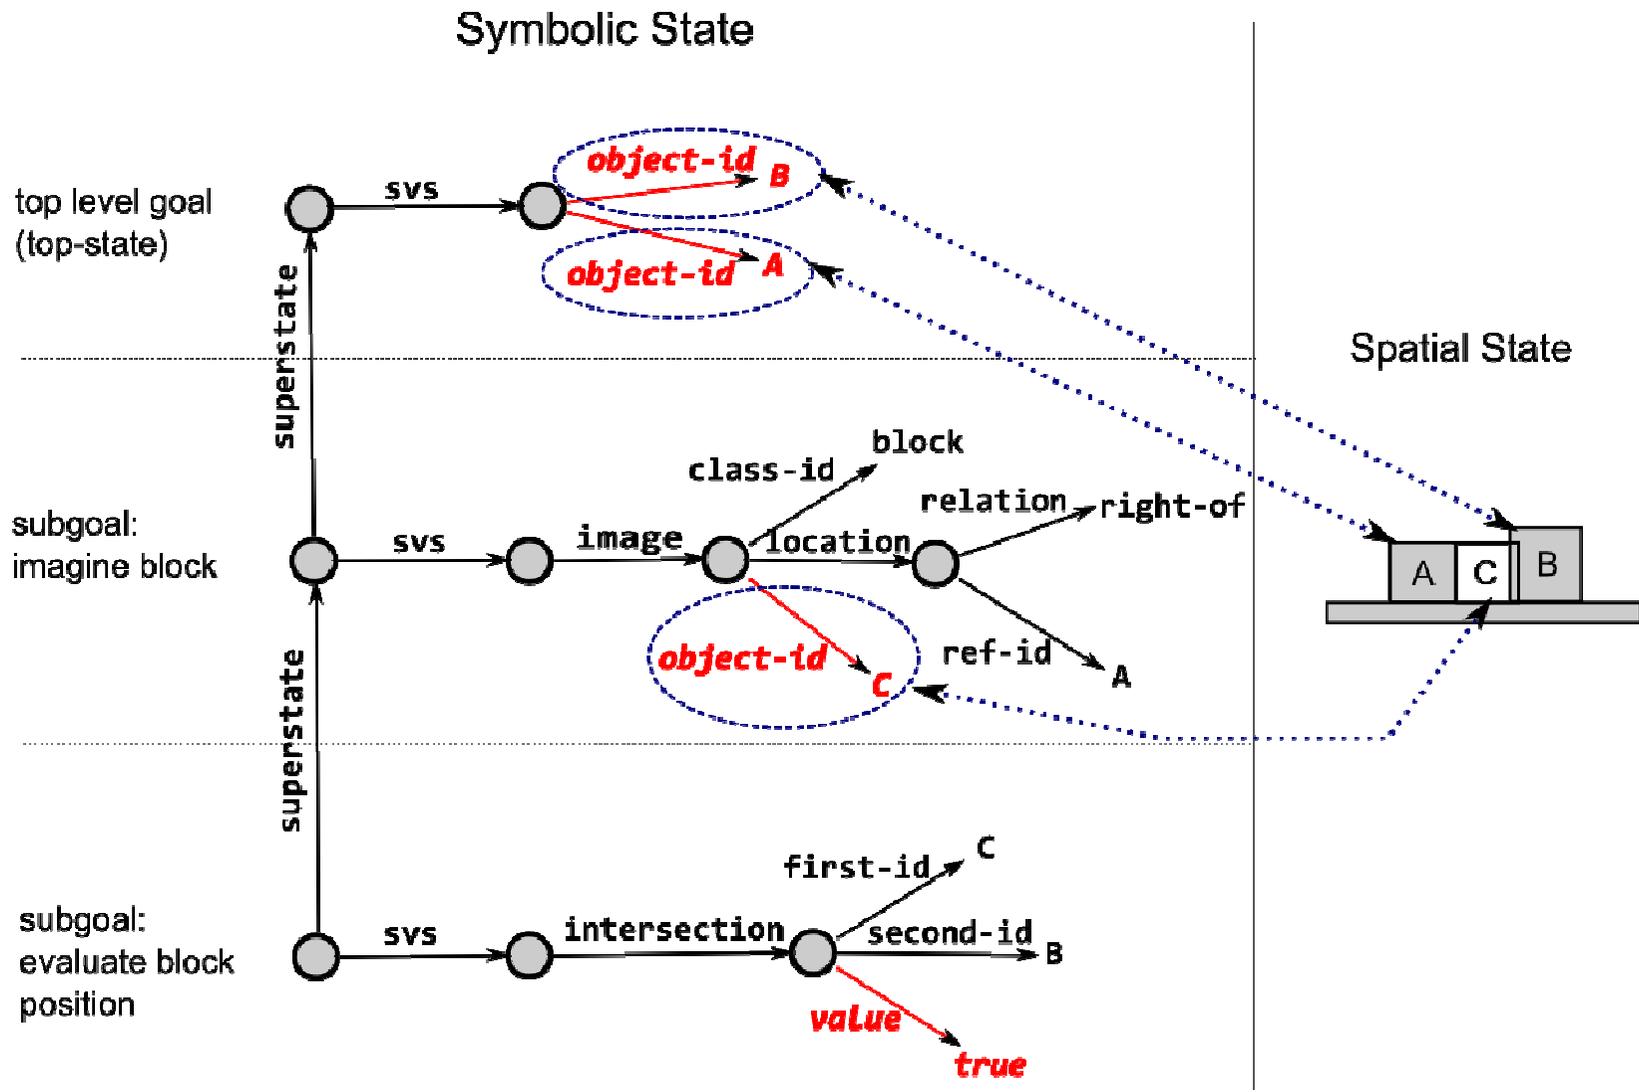
- ▶ State-local **svs** structure connects to imagery system
- ▶ Most imagery operations are equivalent to a hidden set of elaboration productions

```
svs {create*image
  (state <s> ^svs <svs>)
  (<svs> ^image <im>)
  (<im> ^property foo
    ^property bar)
-->
  (<svs> ^object-id image22)
}
```

```
svs {evaluate*intersect-query
  (state <s> ^svs <svs>)
  (<svs> ^intersection <i>)
  (<i> ^first-object image22
    ^second-object image37)
-->
  (<i> ^value true)
}
```

- ▶ Imagery structures persist with WM structures, and update when WM changes
- ▶ WM structures queried from imagery persist with query structure, and update when WM or imagery state changes

Working Memory Integration



SVS Implementation

- ▶ All communication happens at top-state over the io-link
- ▶ Why?
 - ▶ Project integration is simple
 - ▶ Software evolved from originally using an operator-action interface
 - ▶ It works well
- ▶ Remainder of talk will cover implementation

Commands on SML Side

- ▶ Problem: If commands aren't associated with operators, many can be present simultaneously
- ▶ Solution: Allow this, but carefully consider monotonicity

- ▶ Most SVS commands are monotonic
 - ▶ Can be processed in (pseudo) parallel
 - ▶ Even “persistent” structures are monotonic: these are like elaborations of O-supported WM structures
 - ▶ These commands must be *reversible*, and retractions must be monitored
 - ▶ Intra-command interactions are still possible
 - ▶ Current SVS solution is to process commands in one carefully-ordered wave
 - ▶ Must guarantee working memory is consistent with commands by input phase

- ▶ Some commands are non-monotonic
 - ▶ Actual actions passed to the external world
 - ▶ Internal commands with global effects

Soar Implementation Basics

- ▶ Default productions are defined to
 - ▶ build **svs** structures on each state
 - ▶ copy commands to output-link from **svs** structures
 - ▶ associate commands with responses from input-link
 - ▶ fill in decisions requiring multiple i/o phases
 - ▶ keep output-link consistent with subgoal **svs** structures

Decision Cycle Integration

- ▶ Problem: imagery structures can be i-supported, requiring multiple waves to make a decision, but i/o happens between decisions
- ▶ Careful ordering on SML side handles some of this
- ▶ Some cases require interleaved imagery and rule-matching, and can't be done in one decision
 - ▶ Partial solution: propose filler operator if any commands are present without responses

Subgoal Integration

- ▶ Problem: o-supported subgoal imagery must be removed when subgoal goes away
- ▶ Problem: i-supported subgoal imagery modifies **output-link**, usually creating o-supported results
- ▶ Solution: default cleanup production
- ▶ If a command is ever present on **output-link** without an equivalent command on an **svs** WME (on any state), it is removed
 - ▶ Production must be o-supported, but does not need its own operator

Subgoal Integration: Interaction Between States

- ▶ Problem: There is only one instance of each memory in SVS, but multiple subgoals may be present in Soar
- ▶ Solution: monotonic commands prevent most problems
 - ▶ Commands in substates cannot interfere with results of commands in superstates if all are monotonic
 - ▶ Non-monotonic commands must be issued at top-state
- ▶ Superstate processing can still access imagery objects created in substates
 - ▶ Could be fixed by notifying SVS of which state commands belong to

Conclusion

- ▶ Nuggets:
 - ▶ Rich, efficient interaction with non-symbolic memories over the io-link is possible without Soar kernel modification
 - ▶ Resulting interface is useful and intuitive
 - ▶ Task knowledge can be concisely represented
- ▶ Coal:
 - ▶ Multiple waves of i/o operation are impossible during the decision cycle, resulting in extra decisions
 - ▶ Lack of direct connections to subgoals can cause minor problems with chunking and GDS

Efficiently Processing Changes on SML Side

- ▶ Problem: If commands exist for many cycles and have deep structure, what if symbolic processing modifies them?
- ▶ Solution: Detect changes in SML WME structures
- ▶ Every new WME added via SML has a unique timetag, timetags always increase
- ▶ If commands are trees, hashing can be done $O(\# \text{ of WMEs})$
 - ▶ Parse through WME tree, find highest timetag
 - ▶ Simultaneously count how many WMEs are in the tree
 - ▶ Hash is {WME count , highest timetag}
 - ▶ Adding a WME will result in a new count and new timetag
 - ▶ Modifying a WME will result in a new timetag
 - ▶ Deleting a WME will result in a lower count
 - ▶ Hash can detect changes quickly, but can't detect *what* changed
 - ▶ Currently no general-purpose solution to this..

Binding Output and Input

- ▶ Problem: **status complete** is insufficient feedback for commands
- ▶ Solution: **request-ids**
 - ▶ In Soar, use **make-constant-symbol** to add an id whenever a command appears
 - ▶ Externally keep track of id, return result with it on **input-link**
 - ▶ In Soar, link the result structure to the command structure
 - ▶ **request-id** need not concern agent developers

Binding Input and Output: Example

1: (svs.command <c>
 <c> ^parameter one
 ^parameter two)

3: (<c> ^request-id constant23)

5: (output-link.command <c-copy>
 <c-copy> ^parameter one
 ^parameter two
 ^request-id constant23)

7: (io.input-link.response <r>
 <r> ^request-id constant23
 ^property a
 ^property b)

9: (<c> ^response <r>)

2: (svs.command <c>)

-->

(<c> ^request-id (make-constant-symbol))

4: (svs.command <c>)

(<c> ^request-id <id>)

-->

(output-link ^command (deep-copy<c>))

6: process command, create response

8: (svs.command <c>)

(input-link.response <r>)

(<c> ^request-id <id>)

(<r> ^request-id <id>)

-->

(<c> ^response <r>)

Binding Input and Output: Agent Developer View

**1: (svs.command <c>
 <c> ^parameter one
 ^parameter two)**

2: process command, create response

**3: (<c> ^response <r>
 <r> ^property a
 ^property b)**

Subgoal Integration: Chunking and GDS

- ▶ If an imagery structure is *supposed* to be a result, it can be created on a superstate's **svs** WME
 - ▶ Chunking automatically captures these
- ▶ If intermediate imagery steps are used during a subgoal, chunking cannot capture those
 - ▶ Chunks get created to build the structures on output-link, but the structures are immediately removed by the cleanup rule since there is no state for them
 - ▶ This could be fixed, but not without modifying chunking
 - ▶ Chunking should *never* completely remove non-symbolic steps from reasoning, though
 - ▶ Solution: disable chunking in these subgoals
- ▶ GDS also causes non-symbolic subgoal processing to be handled differently than symbolic
 - ▶ In some cases, local o-supported imagery structures created based on local non-symbolic reasoning can cause the GDS to remove the goal
 - ▶ No solution, but not a common occurrence