

# Reinforcement Learning in Infinite Mario



Shiwali Mohan and John E. Laird

`shiwali@umich.edu`

`laird@umich.edu`

The Soar Group at University of Michigan

# Research Question

---

- In a constrained cognitive architecture, does
    - describing the world using symbolic, object-oriented representations,
    - hierarchical task decomposition and learning
    - including internal goals and rewards in the design of the agent
- result in better reinforcement learning in a complex task?
- higher average reward
  - faster convergence

# Outline

---

- Domain
- Challenges
- Reinforcement Learning
- Propositional Agent
- Symbolic, Object-Oriented Representation
- Hierarchical Reinforcement Learning
- Design
- Nuggets and Coal

# Domain

## □ Infinite Mario

- Side scrolling game
- Gain points by collecting coins, killing monsters

## □ Domain developed in RL-Glue


- Reinforcement Learning Competition 2009

## □ State Observations

- Visual Scene – 16 x 22 tiles, 13 different types
  - Episode is of arbitrary length
- Monsters – can be of different types, speed etc

## □ Actions

- Typical Nintendo Controls
  - Step right, left or stay
  - Jump
  - Speed toggle



## □ Reward


- +100 on reaching the finish line
- +1.00 for collecting coin, killing monster
- -10.00 for termination of the game
- -0.01 every step

## □ Sample Agent

- Heuristic policy
- Remembers sequence of actions

# Domain

- Learning computationally expensive
  - Episode with 300 steps has ~5000 tiles of 13 different types
  - Use 'good' state representations
- Partial Observations
  - Only a part of the game instance is visible at a time.
  - Assume that only local conditions matter, MDP assumption
- Large, continuous, growing state space
  - Position, speed of objects (monster) are real numbers
  - Many different objects
  - Value function augmentation, good state representation
- Highly dynamic environment
  - High degree of relative movement
  - Despite available input data, predicting behavior is hard
  - Learn from experience




- Learning should generalize across different instances, levels of the game.
  - Representations that are specific to a particular instance of game cannot be used.
- Abundant information
  - Lot of extracted, derived data
  - Learn what is important from instructor

# Reinforcement Learning

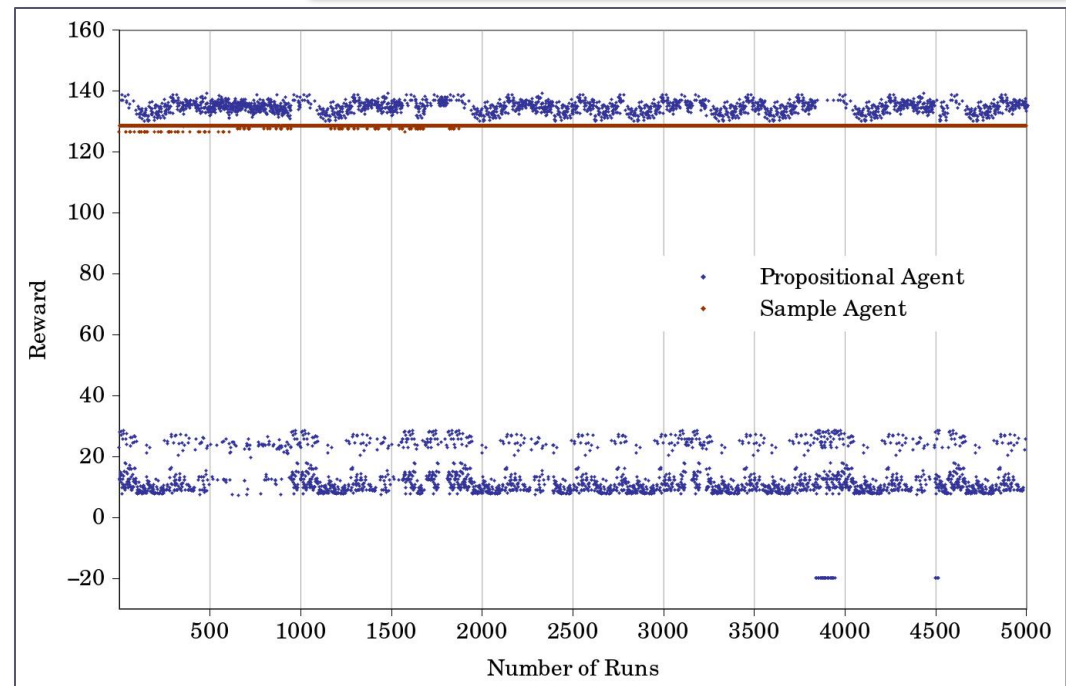

---

- Reinforcement Learning
  - Acquire domain knowledge from experience, online learning
  
- Formally, the basic reinforcement learning model
  - a set of environment states  $S$ ;
  - a set of actions  $A$ ; and
  - a set of scalar "rewards"  $R$ .
  
- Based on the interactions the RL agent develops a policy
  - Maximizes the reward earned



# Propositional Agent


- Enormous state space
  - Visual Scene –  $16 \times 22$  (352) tiles, of 13 different kinds =  $13^{352}$  states
    - All states do not really occur in the game
- Use very local information
  - $5 \times 3$  tiles around Mario
  - Include monsters that are within this range
- Learning is hard
  - Huge state-space
  - Reward achieved after a long sequence of steps
- Not clear how to provide background knowledge to aid learning
  - Extremely difficult, maybe impossible



# Symbolic, Object-Oriented Representation

(agents 2, 3, 4)

- Extract regular objects from inputs
  - Monsters, coins, question-blocks, platforms, pits
- Associate object with its features
  - *speed, type, position*
- Derive features
  - Relative distances between objects
  - Relative distances of objects from Mario
  - attributes like *'isreachable', 'isthreat'* if a object is close enough and should affect agents behavior
- Describe state
- Provide background knowledge
  - *If attribute 'isreachable' for a platform is set, and there is a coin on it, then set attribute 'isreachable' for the coin.*

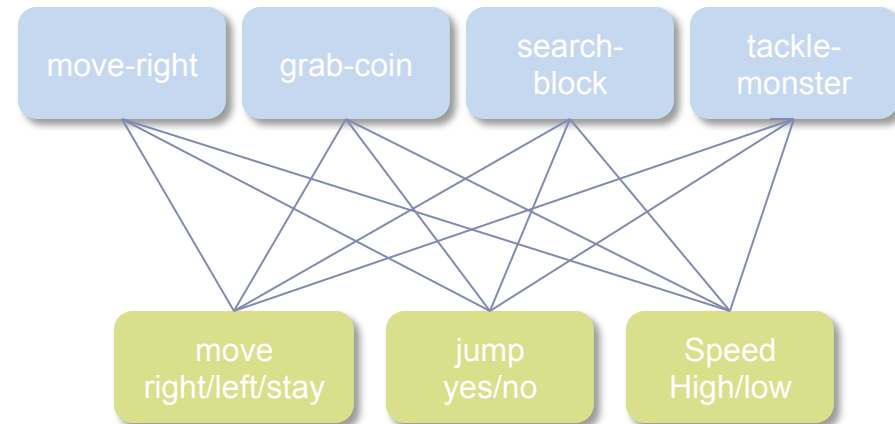


```
state <s> ^name mario-soar      <m1> ^type Goomba
          ^monster <m1>         ^winged no
          ^coin <c1>             ^distx -4.5
          ^coin <c2>             ^disty -2
          ^coin <c3>             ^isthreat no
          ^coin <c4>
          ^coin <c5>             <c1> ^distx 4
          ^coin <c6>             ^disty 1
          ^question-block <q1>   ^isreachable yes
          ^question-block <q2>
```



# Action (Operator) Hierarchy

- GOMS analysis of Mario<sup>1</sup>
  - Predictive of the behavior of human expert
  - Introduced functional-level operators and Keystroke-level
  - Divides the task into smaller tasks
- Two kinds of actions
  - FLOs
    - Functional-level Operators (actions)
    - Abstract macro-actions
    - Sequence of atomic actions
    - With a specific functional goal
  - KLOs
    - Keystroke –level Operators (actions)
    - Atomic actions
    - Move, jump, speed toggle



- Application of Actions
  - Object-Oriented
    - FLOs described for specific objects
    - *tackle-monster* for monsters
  - Control
    - Derived attributes used to control the progression
    - *'isthreat'*, *'isreachable'*

[1] B.E. John and A.H. Vera, "A GOMS analysis of a graphic machine-paced, highly interactive task," Proceedings of the SIGCHI conference on Human factors in computing systems, 1992, pp. 251–258.

# Progression


move-right

grab-coin

search-block

tackle-monster

Trial : 0001  
 Steps : 0007  
 Reward : -0.05  
 Total : -0.05



state: mario-soar  
 available behavior:  
 move-right

state: move-right  
 available actions:  
 atomic-actions

state: mario-soar  
 available behavior:  
 move-right,  
 tackle-monster(m1)

state: tackle-monster  
 available behavior:  
 Atomic-actions

m1 isthreat yes

# Agent 2: Learning at KLO level

---

- Can the agent learn behaviors like grabbing a coin, killing a monster?

- State

- FLO-level: presence, absence of flag attributes like 'isthreat' or 'isreachable'

```
monster <m1> ^isthreat true
```

```
coin <c1> ^isreachable true
```

- KLO-level: features extracted/derived from input

```
monster <m1> ^type Goomba  
            ^winged no  
            ^distx <x>  
            ^disty <y>  
            ...
```

```
coin <c1> ^distx <x>  
         ^disty <y>
```

- Actions

- FLOs : tackle-monster, grab-coin, search-question etc
    - Symbolic preferences used to break a tie
  - KLOs: move{right, left, stay} x jump{yes,no} x speed toggle {on,off}

- Learning


- Given symbolic preferences at the FLO level
    - tackle-monster > move-right
  - Learn the most rewarding sequence of KLOs

- Reward

- As provided by the environment , at KLO level

# Results - Agent 2

- Averaged over 10 trials of 2000 runs each
- Learning algorithm - SARSA
- Learning rate – 0.3
- Discount rate – 0.9
- Exploration policy – Epsilon-greedy
- Epsilon – 0.01
- Reduction-rate – 0.99
- Policy converges at ~1400 episodes
- Average reward earned by converged policy (last 100 runs)=145.97
- Agent performs better than the sample agent.



# Agent 3: Hierarchical Learning

- Can the agent learn preferences between objects as it learns behaviors?
  - Similar to MAXQ-0<sup>1</sup>

- State

- FLO-level: presence, absence of flag attributes like 'isthreat' or 'isreachable'

```
monster <m1> ^isthreat true
           ^distx <x>
           ^disty <y>
```

```
coin <c1> ^isreachable true
         ^distx <x>
         ^disty <y>
```

- KLO-level: features extracted/derived from input

```
monster <m1> ^type Goomba
           ^winged no
           ^distx <x>
           ^disty <y>
           ...
```

```
coin <c1> ^distx <x>
        ^disty <y>
```

- Actions

- FLOs : tackle-monster, grab-coin, search-question etc
  - KLOs: move{right, left, stay} x jump{yes,no} x speed toggle {on,off}

- Learning


- Learn numeric preferences at the FLO level
  - Learn the most rewarding sequence of KLOs

- Reward

- As provided by the environment , at both KLO and FLO level

# Results – Agent 3

- Averaged over 10 trials of 2000 runs each
- Learning algorithm - SARSA
- Learning rate – 0.3
- Discount rate – 0.9
- Exploration policy – Epsilon-greedy
- Epsilon – 0.01
- Reduction-rate – 0.99
- Agent converges to a policy at 400 episodes
- Average reward earned by converged policy (last 100 runs)=144.68
- Agent converges faster than Agent 2 (1400 runs)
- Learns a more specific policy which might be better



# Agent 4: Reward Shaping

---

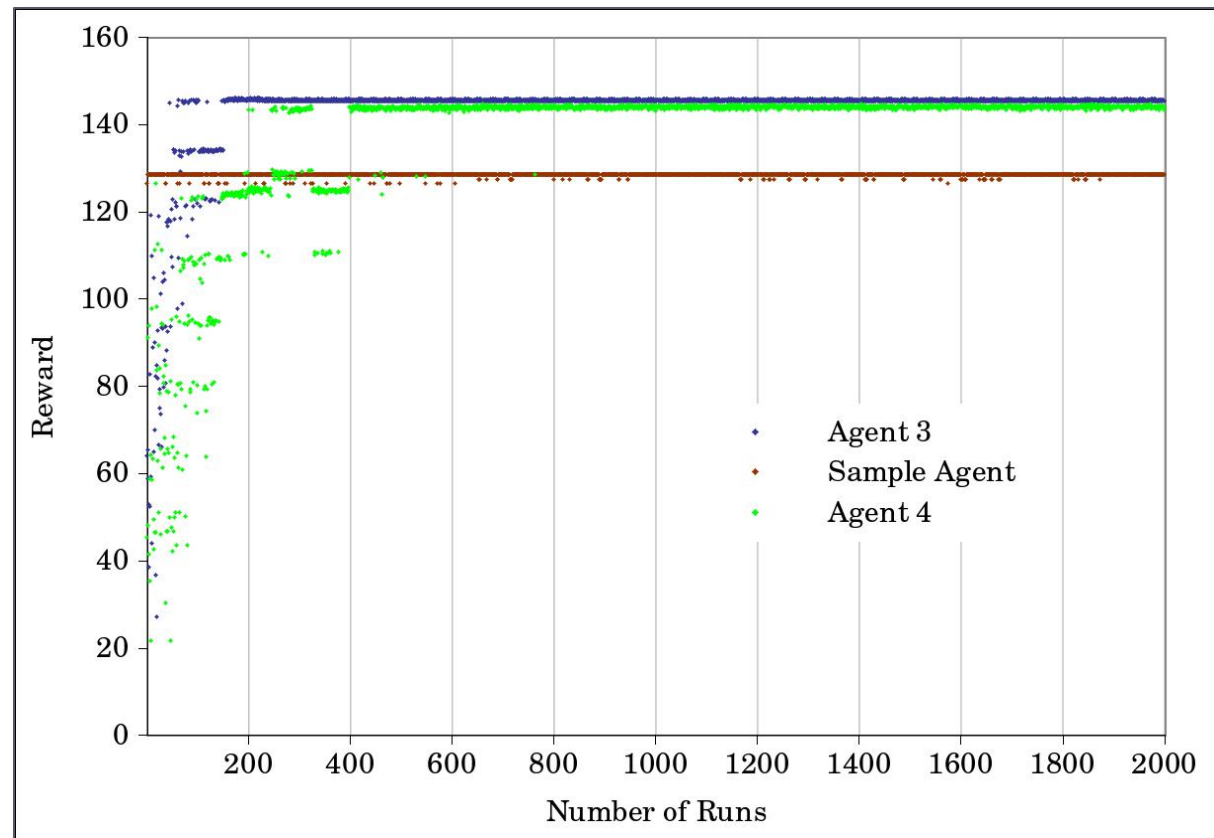
- Can the performance of the agent be improved by introducing internal rewards and goals in the agent design?
  - MAXQ-Q<sup>1</sup>
  - Building in intrinsic goals and rewards<sup>2</sup>
  
- State, Action, Learning
  - As in Agent 3
  
- Reward
  - Agent 3 uses the reward as provided by the environment
    - Agent may get rewarded even if it does not execute the selected FLO correctly,
      - grabbing a coin while tackle-monster is selected
  - Reward the agent at the KLO level only when the goal is achieved
    - +1.00 for correctly executing the selected FIO,
      - killing/avoiding a monster when tackle-monster is selected
    - -0.01 for every step
  - Reward at the FLO level is computed from the reward provided by the environment

[1] T.G. Dietterich, "Hierarchical reinforcement learning with the MAXQ value function decomposition," *Journal of Artificial Intelligence Research*, vol. 13, 2000, pp. 227–303.

[2] S. Singh, R.L. Lewis, A.G. Barto, J. Sorg, and A. Helou, "On Separating Agent Designer Goals from Agent Goals: Breaking the Preferences–Parameters Confound," submitted, 2010.

# Results – Agent 4

- Averaged over 10 trials of 2000 runs each
- Learning algorithm - SARSA
- Learning rate – 0.3
- Discount rate – 0.9
- Exploration policy – Epsilon-greedy
- Epsilon – 0.01
- Reduction-rate – 0.99
- Agent converges to a policy at 200 episodes
- Average reward earned by converged policy (last 100 runs)=145.98
- Average standard deviation (last 100 runs) = 2.0083
- Converges faster than Agent 3 (400 runs), correct intrinsic reward





# Nuggets and Coal

---

- Hierarchical division of the task makes operating in the environment easier
- Hierarchical learning allows the agent to learn policies faster
- Demonstrated that object-oriented representations provide structure to state description
- Symbolic representations allow for easy encoding of the background knowledge
- Encoding intrinsic rewards in an agent helps it learn faster.
- Detailed analytical study of the domain
  - Optimality
  - Good comparison metric
  - Parameter sweeps
  - Transfer Learning
- Learning at one level should help in playing a higher level