

# New SVS Overview

Joseph Xu

Soar Workshop 31

June 2011

# Outline

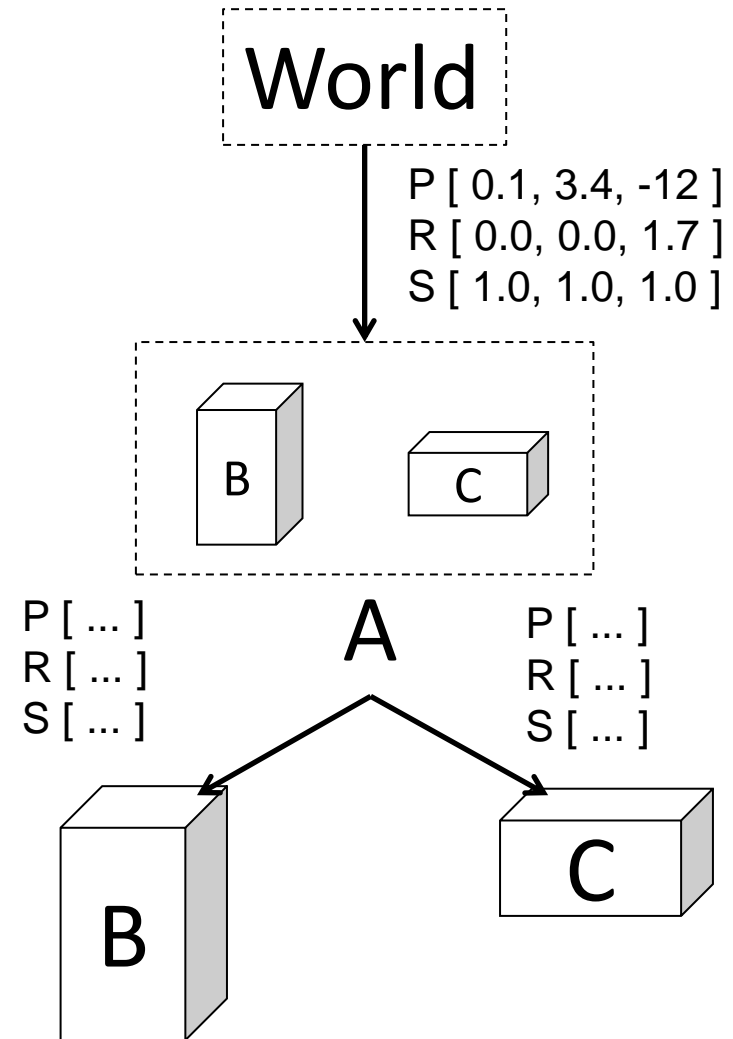
- Functionality in old SVS
- Addition of continuous control
- New implementation
- Conclusions

# The Gist of SVS Past

- Hypothesis: Important properties of continuous environments cannot be completely captured with purely symbolic representations
- Use continuous representation to supplement symbolic reasoning
  - Continuous spatial scene graph
  - Symbolic Soar extracts information by asking questions in the form of spatial predicates
  - Symbolic Soar modifies scene graph using imagery commands

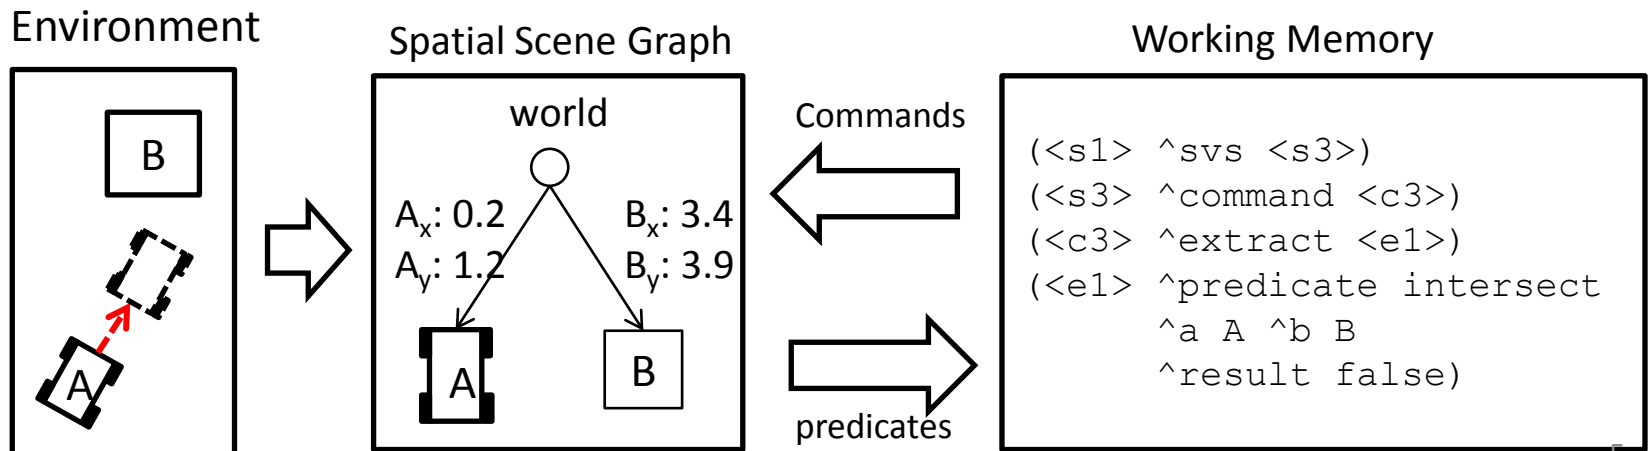
# Spatial Scene Graph

- Set of discrete objects
- Hierarchy of “part-of” relationships
- Each object is transformed relative to parent in terms of position, rotation, and scaling
- Each leaf object has a concrete geometric form



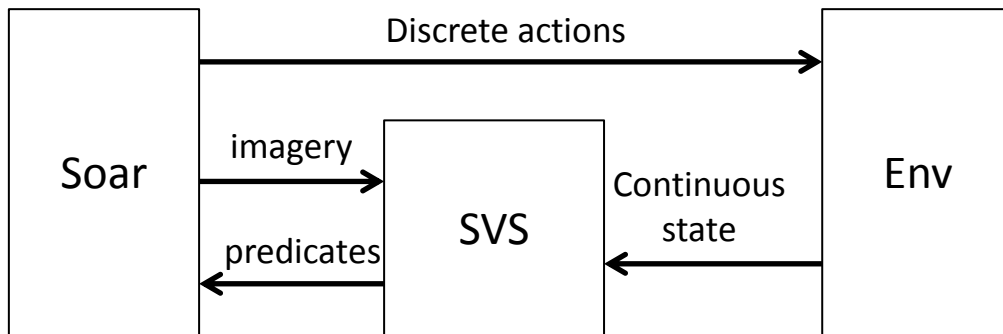
# SVS Commands

- Predicate extraction
  - Ask a question about the current state
    - Is car A intersecting box B?
- Imagery
  - Imagine a change to the state for the purpose of reasoning about it
    - If A is to the right of B, is it to the left of C?
- SVS is only a mechanism, it's not smart
  - Needs knowledge about which predicates to extract and imagery operations to perform



# Applications

- Applied to domains where spatial properties of state are important
- Controlled with given discrete sets of actions
- One step predictions using ad-hoc model learning methods



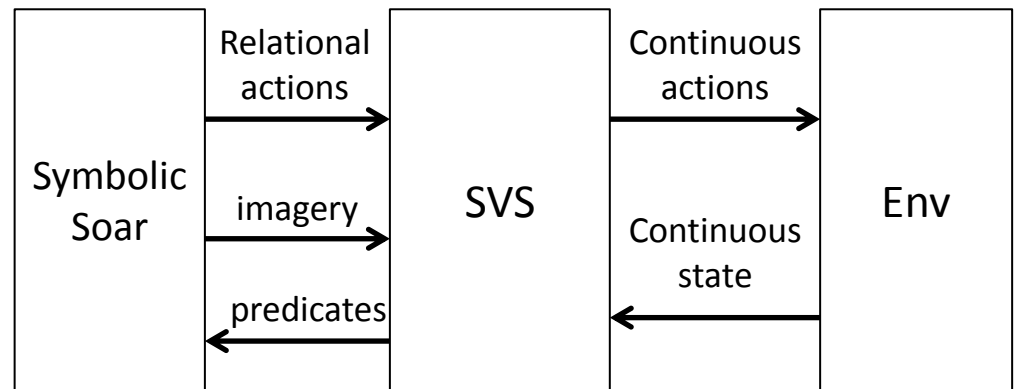
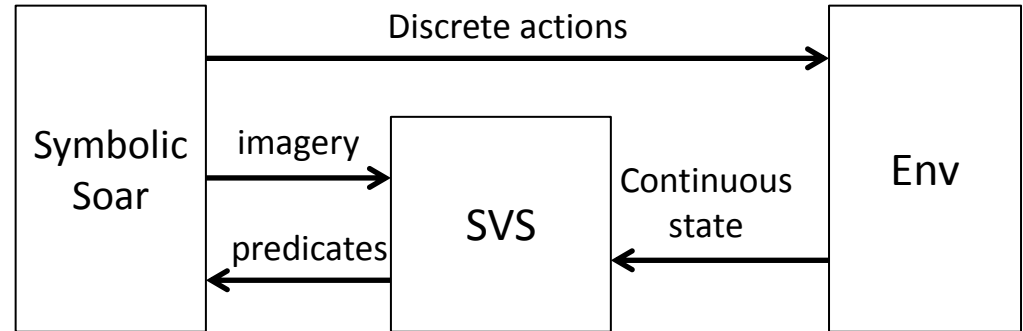
# Continuous Control

- Many real world environments expect continuous control signals from agent
  - Example: robot domain expects left and right motor voltages
- Traditional approach is to hand-code middleware to translate set of discrete actions into continuous output
  - Action discretization is a priori, leading to non-adaptive and non-optimal behavior
  - Not part of the cognitive architecture theory
  - Need new middleware for every new environment

# My Motivation


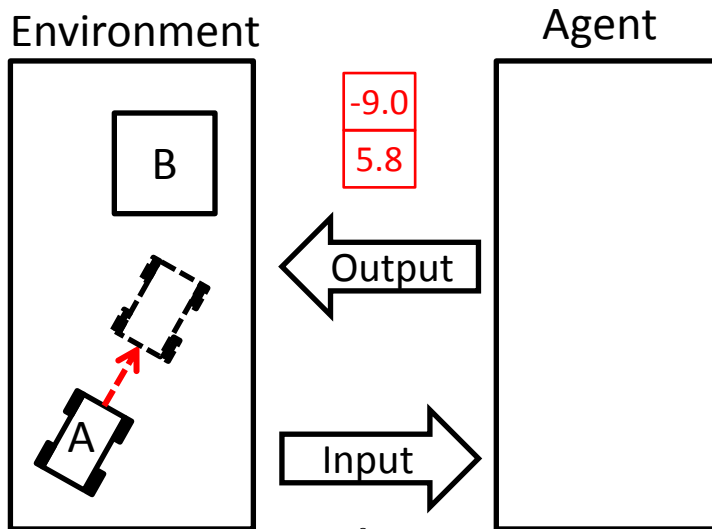
Augment SVS to allow agents to automatically learn continuous control

- Agent autonomously derives a set of relational actions that it can plan over symbolically
- SVS learns how to translate relational actions to continuous output






# Environment Assumptions




- Input to the agent is a scene graph
- Output is fixed-length vector of continuous numbers
  - Agent doesn't know a priori what numbers represent
- Agent runs in lock-step with environment
- Fully observable
- Some noise tolerable

# Relational Actions


- The value of an extracted predicate is the smallest unit of change that's distinguishable to symbolic Soar
- Each potential predicate value change is a *relational action*
  - Combinations of predicates?
- +intersect(A, B)



```
(<s1> ^svs <s3>)  
(<s3> ^command <c3>)  
(<c3> ^extract <e1>)  
(<e1> ^predicate intersect  
^a A ^b B  
^result false)
```



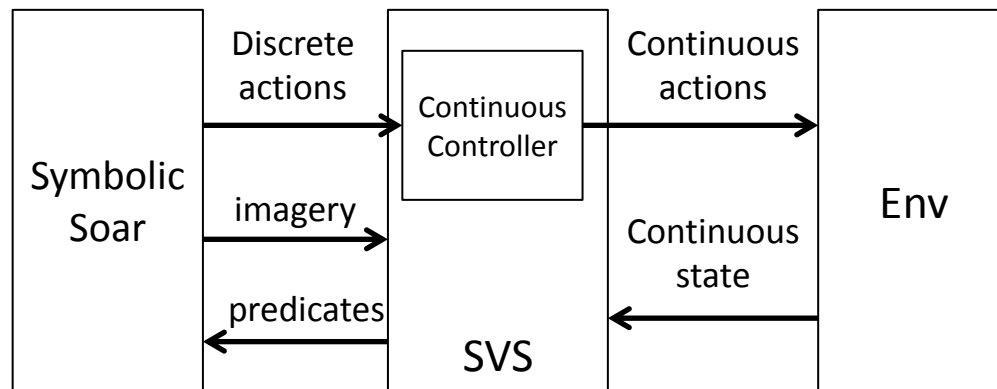
```
(<s1> ^svs <s3>)  
(<s3> ^command <c3>)  
(<c3> ^extract <e1>)  
(<e1> ^predicate intersect  
^a A ^b B  
^result false)
```



```
(<s1> ^svs <s3>)  
(<s3> ^command <c3>)  
(<c3> ^extract <e1>)  
(<e1> ^predicate intersect  
^a A ^b B  
^result true)
```

# Continuous Controller

- Takes a relational action and translates into a multi-step trajectory of continuous-valued outputs to environment  $(u_1, u_2, \dots, u_n)$
- One predicate change takes multiple decision cycles
- May fail to find a trajectory that changes predicate value



# Continuous Planning

- Find the trajectory that will lead to the predicate change the fastest

Continuous Model

$$f(x, u) \rightarrow y$$


Learned (subject of next talk) or hand-coded

Vector of transform values in flattened scene graph

Objective Function

$$g(x) \rightarrow [-\infty, \infty]$$

Given a priori for each predicate



$\text{intersect}(A, B) = \text{true}$

Control Function

$$H_{f,g}(x) \rightarrow u = \text{argmin}_u [g(f(x, u))]$$

Fixed in architecture

GreedySearch( $s$ , *objective*):

for  $u$  in sample(*outputs*):

$\swarrow$  continuous-model( $s$ ,  $u$ )

Downhill  $\searrow$  *objective*( $y$ )

Simplex  $\leftarrow$  *best*:

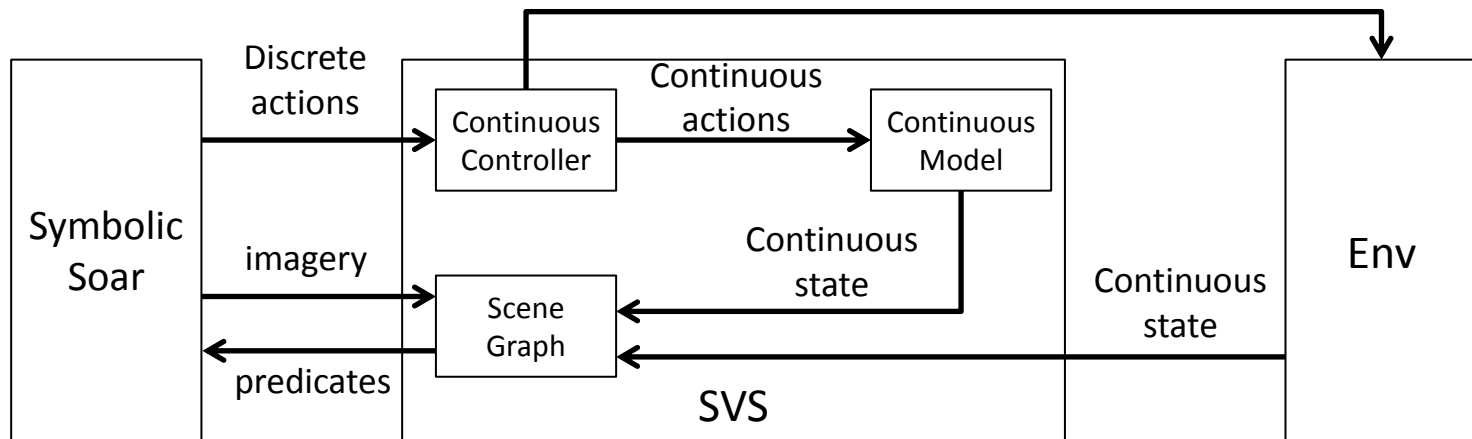
$\text{best} = o$

$\text{best-output} = u$

return *best-output*

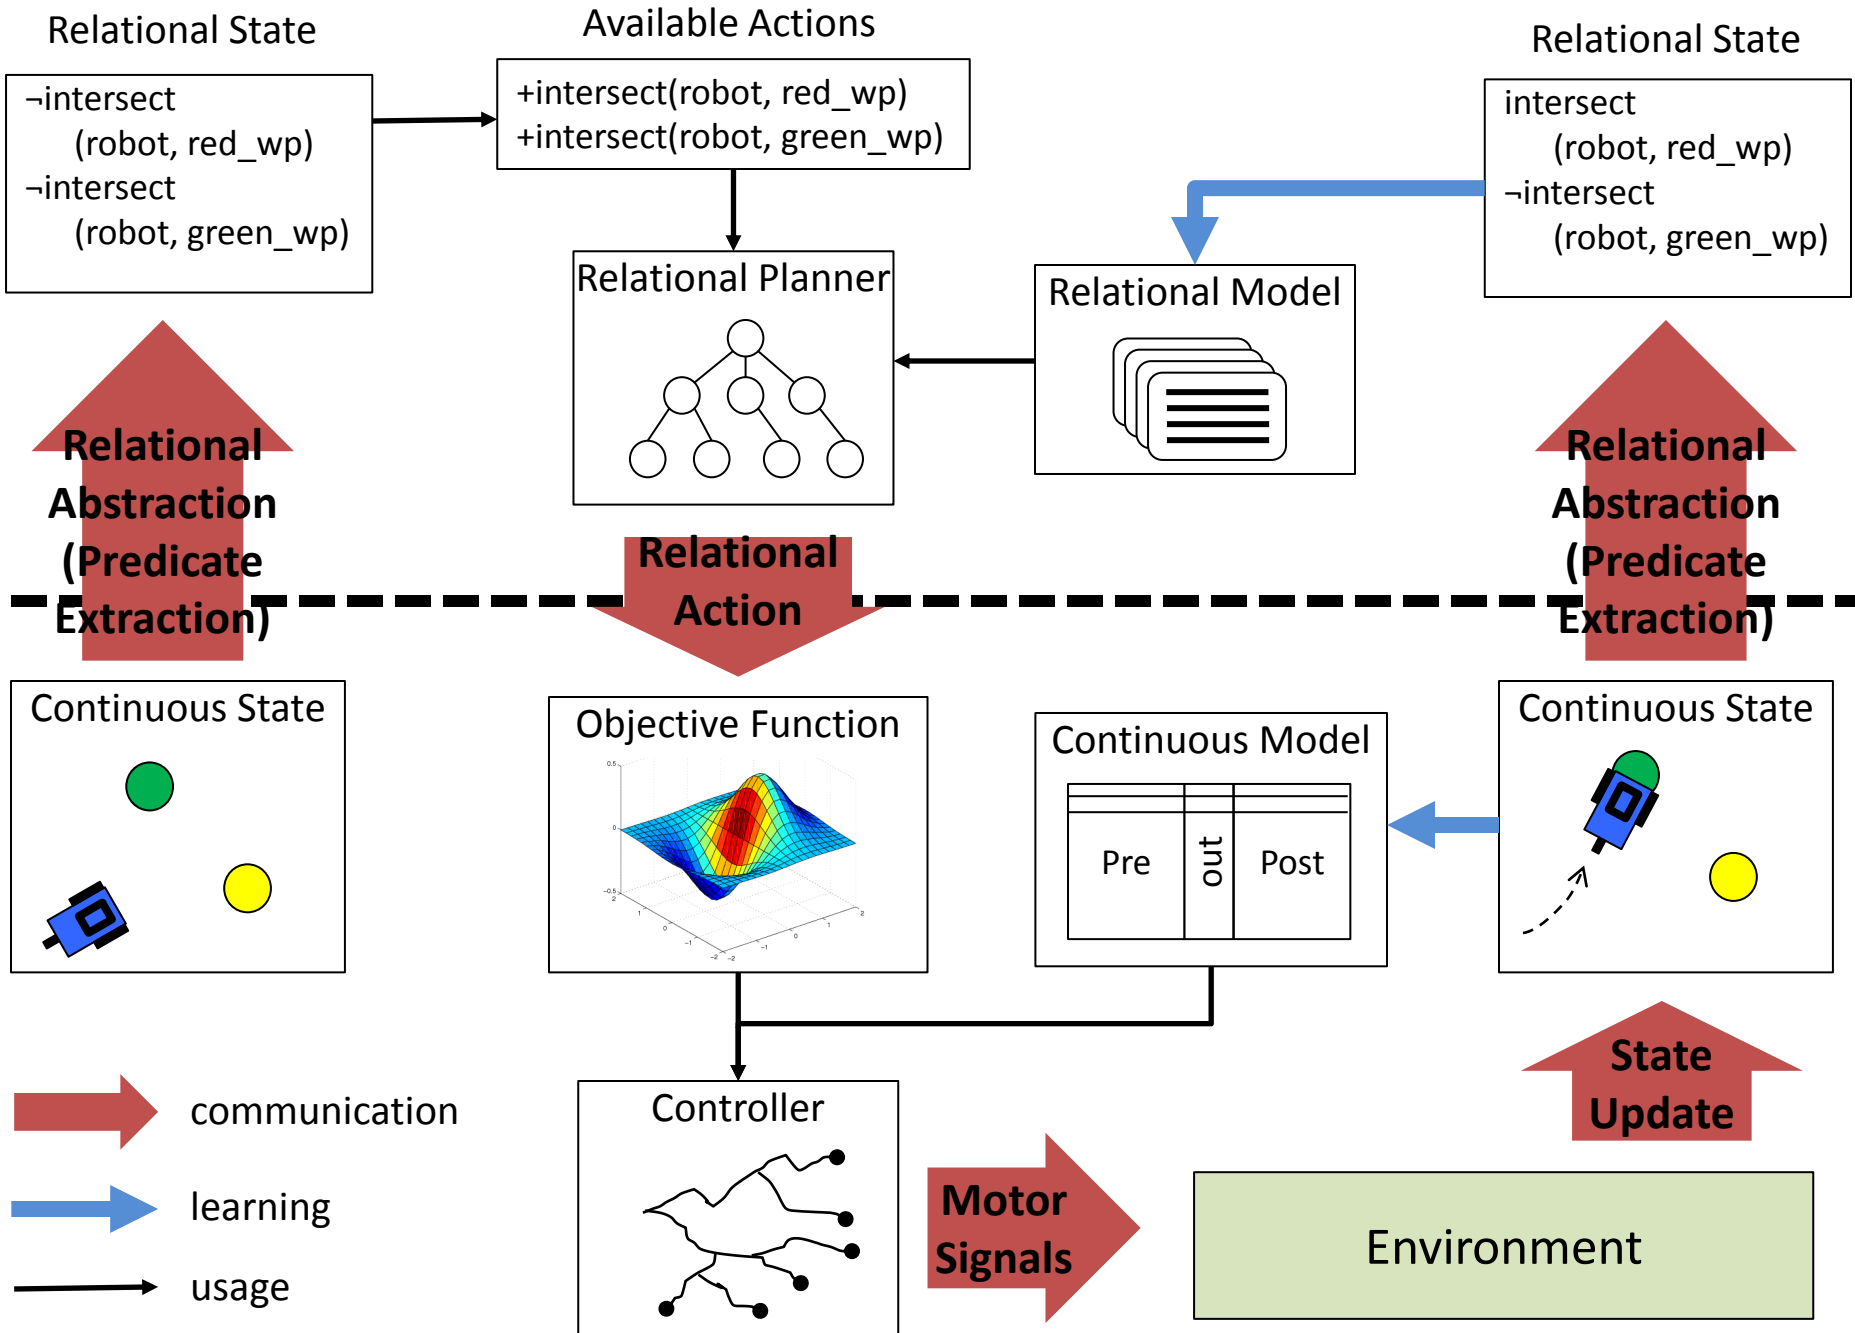
# Imagery with Control

- Agent can perform relational prediction by simulating a trajectory using continuous model on the internal scene graph, instead of sending it out to the environment
- Modified scene graph updates predicates as usual



# Summary

- Sam focused on theory of translating continuous environment state into relational symbolic representations
- I've added theory about where relational actions come from and how to ground them into continuous trajectories in the environment
- Ultimate goal: agents that can learn relational abstractions and plan over them in continuous environments



# New Implementation

- Integrated into kernel rather than communicate on IO link
- Removes dependencies on external libraries (WildMagic, CGAL) to ease installation
- Structured code to be extensible
  - Easy to add new predicates, commands, models
- Main idea: keep it clean enough that future students won't throw it away and start over (happened twice already)



# Nuggets & Coal

## Nuggets

- Adds a task independent mechanism for continuous control
- Usable implementation, soon to be released

## Coal


- Dropped some functionality from previous implementations
- Not optimized, haven't measured performance in non-trivial domains

# Extract Rule

```
sp {cursor-target-intersect
  (state <s> ^superstate nil
    ^svs ( ^command <c>
      ^spatial-scene (
        ^child.id splinter
        ^child.id target)))
```

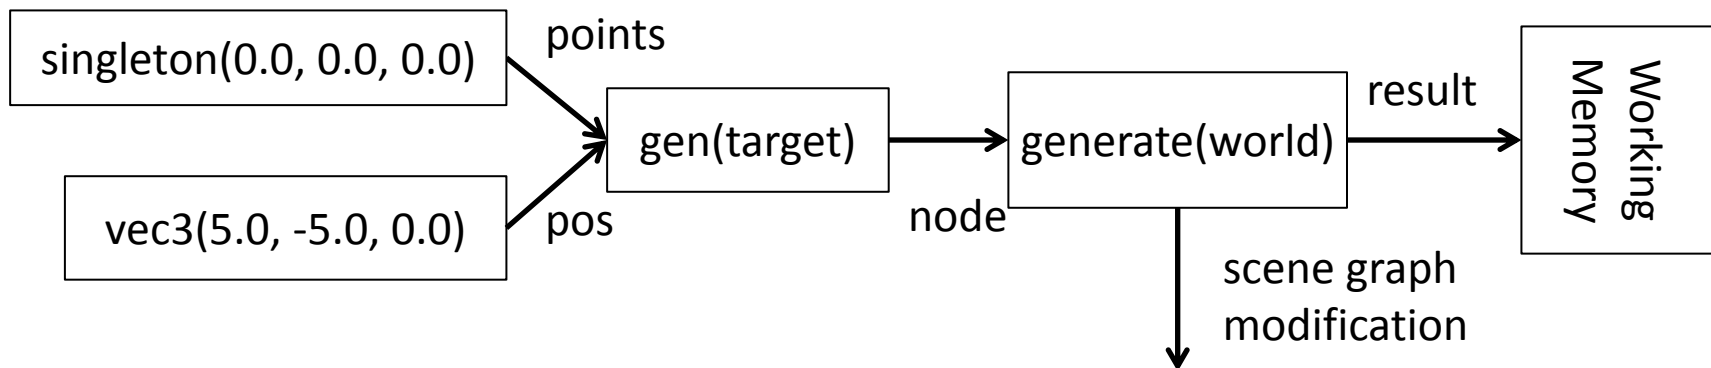
-->

```
(<c> ^extract <e>)
(<e> ^type bbox_intersect ^a <a> ^b <b>)
(<a> ^type bbox ^x <n1>)
(<n1> ^type node ^name splinter)
(<b> ^type bbox ^x <n2>)
(<n2> ^type node ^name target)}
```



# Generate Rule

```
sp {gen-target
  (state <s> ^superstate nil
    ^svs.command <c>)
-->
(<c> ^generate <g>)
(<g> ^parent world
  ^node <n>)
(<n> ^type gen
  ^name target
  ^points.type singleton
  ^pos <pos>)
(<pos> ^type vec3 ^x 5.0 ^y -5.0 ^z 0.0)}
```



# Control Rule

```
sp {seek-target
  (state <s> ^superstate nil
    ^svs ( ^command <c>
      ^spatial-scene ( ^child.id splinter
        ^child.id target)))
-->
(<c> ^control <ctl>)
(<ctl> ^type simplex ^depth 20 ^outputs <out>
  ^objective <obj> ^model model1)
(<out> ^left <left> ^right <right>)
(<left> ^min -1.0 ^max 1.0 ^inc 1.0)
(<right> ^min -1.0 ^max 1.0 ^inc 1.0)
(<obj> ^name euclidean ^a splinter ^b target)}
```