

# Towards an Architecture for Learning with Instruction

Shiwali Mohan and John E. Laird

Computer Science and Engineering  
University of Michigan

June 17th, 2011

# Outline

---

① Introduction


② Instruction Task

③ Agent Design

④ Conclusions


# Motivation

- Long living, general, intelligent agents
  - demonstrate a wide range of behavior
  - on a variety of task
  - knowledge must be added throughout the lifetime of the agent



# Motivation

- Long living, general, intelligent agents
  - demonstrate a wide range of behavior
  - on a variety of task
  - knowledge must be added throughout the lifetime of the agent
- Instruction
  - agent can be *taught* to learn new procedures
  - knowledge/information that can be trusted
  - situated, interactive instruction produces strong human learning<sup>1</sup>
  - semantic, procedural, problem solving etc



<sup>1</sup>Bloom, B. S. (1986). The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring. *Educational Researcher*, 13(6):4-16

## Related Work

- Robo-Soar<sup>1</sup>
  - search guidance in case of an *operator tie*

---

<sup>1</sup>Laird, J. E., Yager, E. S., Hucka, M., and Tuck, C. M. (1991). Robo-Soar: An Integration of External Interaction, Planning, and Learning using Soar. *Robotics and Autonomous Systems*, 8(1-2):113–129

## Related Work

- Robo-Soar<sup>1</sup>
  - search guidance in case of an *operator tie*
- Instructo-Soar<sup>2</sup>
  - Identified properties of tutorial instruction, agent requirements
  - Situated explanation
  - learning through induction and inference, generalize from specific examples

---

<sup>1</sup>Laird, J. E., Yager, E. S., Hucka, M., and Tuck, C. M. (1991). Robo-Soar: An Integration of External Interaction, Planning, and Learning using Soar. *Robotics and Autonomous Systems*, 8(1-2):113–129

<sup>2</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

## Related Work

- Robo-Soar<sup>1</sup>
  - search guidance in case of an *operator tie*
- Instructo-Soar<sup>2</sup>
  - Identified properties of tutorial instruction, agent requirements
  - Situated explanation
  - learning through induction and inference, generalize from specific examples
- Training Personal Robots Using Natural Language Instruction<sup>3</sup>
  - focus on natural language comprehension and *mapping*

---

<sup>1</sup>Laird, J. E., Yager, E. S., Hucka, M., and Tuck, C. M. (1991). Robo-Soar: An Integration of External Interaction, Planning, and Learning using Soar. *Robotics and Autonomous Systems*, 8(1-2):113–129

<sup>2</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

<sup>3</sup>Lauria, S., Bugmann, G., Kyriacou, T., Bos, J., and Klein, A. (2001). Training Personal Robots Using Natural Language Instruction. *Intelligent Systems, IEEE*, 16(5):38–45

## Related Work

- Robo-Soar<sup>1</sup>
  - search guidance in case of an *operator tie*
- Instructo-Soar<sup>2</sup>
  - Identified properties of tutorial instruction, agent requirements
  - Situated explanation
  - learning through induction and inference, generalize from specific examples
- Training Personal Robots Using Natural Language Instruction<sup>3</sup>
  - focus on natural language comprehension and *mapping*
- Task Learning by Instruction: Benefits and Challenges for Intelligent Interactive Systems<sup>4</sup>
  - integration of various learning and reasoning components

---

<sup>1</sup>Laird, J. E., Yager, E. S., Hucka, M., and Tuck, C. M. (1991). Robo-Soar: An Integration of External Interaction, Planning, and Learning using Soar. *Robotics and Autonomous Systems*, 8(1-2):113–129


<sup>2</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

<sup>3</sup>Lauria, S., Bugmann, G., Kyriacou, T., Bos, J., and Klein, A. (2001). Training Personal Robots Using Natural Language Instruction. *Intelligent Systems, IEEE*, 16(5):38–45

<sup>4</sup>Blythe, J., Tandon, P., and Tillu, M. (2007). Task Learning by Instruction: Benefits and Challenges for Intelligent Interactive Systems. *American Association for Artificial Intelligence*



# Domain



**Given** a set of primitive actions:

go-to <area-id>, go-to-room <room-id>,  
 open-door <door-id>, close-door <door-id>,  
 turn-lights-on, turn-lights-off,  
 pick-up <object-id>, put-down <object-id>,

**Learn** *abstract* actions:

patrol [list-of-rooms]  
 turn-all-lights-off  
 get-object <object-id>

# Goals

---

- Identify and characterize different kinds of instructions in a complex environment
  - different kinds of knowledge, different modes of instruction

# Goals

---

- Identify and characterize different kinds of instructions in a complex environment
  - different kinds of knowledge, different modes of instruction
- Identify and characterize requirements on the agent design
  - general set of rules required to support instruction

# Goals

---

- Identify and characterize different kinds of instructions in a complex environment
  - different kinds of knowledge, different modes of instruction
- Identify and characterize requirements on the agent design
  - general set of rules required to support instruction
- Explore the utility of different long-term memories in a complex learning task

# Goals

---

- Identify and characterize different kinds of instructions in a complex environment
  - different kinds of knowledge, different modes of instruction
- Identify and characterize requirements on the agent design
  - general set of rules required to support instruction
- Explore the utility of different long-term memories in a complex learning task
- Develop a comprehensive model of learning that uses different reasoning and learning modules
  - inference, chunking, inductive reasoning, explanation-based generalization

# Goals

---

- Identify and characterize different kinds of instructions in a complex environment
  - different kinds of knowledge, different modes of instruction
- Identify and characterize requirements on the agent design
  - general set of rules required to support instruction
- Explore the utility of different long-term memories in a complex learning task
- Develop a comprehensive model of learning that uses different reasoning and learning modules
  - inference, chunking, inductive reasoning, explanation-based generalization
- Develop evaluation metrics

# Characterizing Instructions: Taxonomy

- Scope
  - Tutorial: instruct as the agent is acting <sup>5</sup>
  - Instruction Manual: instruct before the agent begins acting

---

<sup>5</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

# Characterizing Instructions: Taxonomy

- Scope
  - Tutorial: instruct as the agent is acting <sup>5</sup>
  - Instruction Manual: instruct before the agent begins acting
- Application
  - When commanded: *patrol rooms A B and C*
  - Automated: *when non-friendly person is observed, report back*

---

<sup>5</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324



# Characterizing Instructions: Taxonomy

- Scope
  - Tutorial: instruct as the agent is acting <sup>5</sup>
  - Instruction Manual: instruct before the agent begins acting
- Application
  - When commanded: *patrol rooms A B and C*
  - Automated: *when non-friendly person is observed, report back*
- Context
  - Situated: *go-to-room A, turn-light-on*
  - Hypothetical: *imagine room A, the light-switch is on north wall*

---

<sup>5</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

# Characterizing Instructions: Taxonomy

- Scope
  - Tutorial: instruct as the agent is acting <sup>5</sup>
  - Instruction Manual: instruct before the agent begins acting
- Application
  - When commanded: *patrol rooms A B and C*
  - Automated: *when non-friendly person is observed, report back*
- Context
  - Situated: *go-to-room A, turn-light-on*
  - Hypothetical: *imagine room A, the light-switch is on north wall*
- Communication
  - Agent Initiated: agent does not know how to proceed
  - Instructor Initiated: modify a learned procedure

---

<sup>5</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

# Characterizing Instructions: Taxonomy

- Scope
  - Tutorial: instruct as the agent is acting <sup>5</sup>
  - Instruction Manual: instruct before the agent begins acting
- Application
  - When commanded: *patrol rooms A B and C*
  - Automated: *when non-friendly person is observed, report back*
- Context
  - Situated: *go-to-room A, turn-light-on*
  - Hypothetical: *imagine room A, the light-switch is on north wall*
- Communication
  - Agent Initiated: agent does not know how to proceed
  - Instructor Initiated: modify a learned procedure
- Task Structuring
  - Composite: learn *turn-all-lights-off*
  - Incremental: learn *patrol*, modify *patrol* to include *turning off lights*

---

<sup>5</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

# Instruction Taxonomy

---

- Learning
  - Rote: memorize sequence of instructions
  - Generalized: learn general procedures

# Instruction Taxonomy

---

- Learning
  - Rote: memorize sequence of instructions
  - Generalized: learn general procedures
- Reasoning
  - Simple
  - Inference

# Instruction Taxonomy

- Learning
  - Rote: memorize sequence of instructions
  - Generalized: learn general procedures
- Reasoning
  - Simple
  - Inference
- Knowledge
  - Proposal and goal conditions: *turn-lights-off* when *in a room* and *light is on*
  - Control: *reporting non-friendly persons is more important than picking up objects*
  - Semantic: *objects with color=blue are dangerous*
  - Meta: arguments are sequential, random

# Requirements on Agent Design

- The *Communication* Problem
  - The *Content* Problem
    - What information has to be communicated?
    - Model of the instructor
    - Assumption of shared environment, similar observations

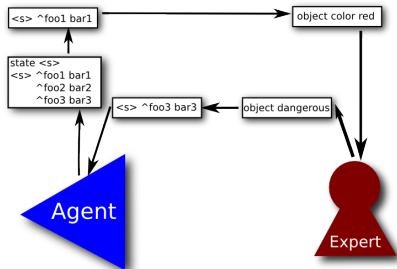
```
<s> ^foo1 bar1
```

```
state <s>  
<s> ^foo1 bar1  
      ^foo2 bar2  
      ^foo3 bar3
```



# Requirements on Agent Design

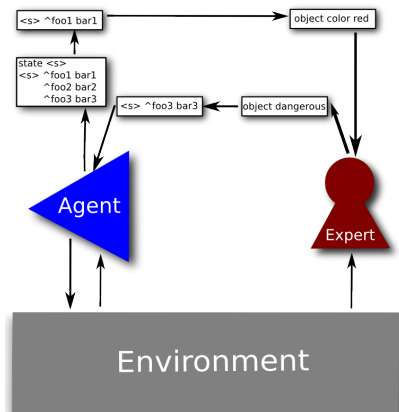
- The *Communication Problem*
  - The *Content Problem*
    - What information has to be communicated?
    - Model of the instructor
    - Assumption of shared environment, similar observations
  - The *Mapping Problem*
    - Common protocol between human and agent
    - Comprehending instructions
    - Generating discourse






# Requirements on Agent Design

- The *Communication* Problem
  - The *Content* Problem
    - What information has to be communicated?
    - Model of the instructor
    - Assumption of shared environment, similar observations
  - The *Mapping* Problem
    - Common protocol between human and agent
    - Comprehending instructions
    - Generating discourse
  - The *Interaction* Problem
    - Maintaining a dialog
    - Integrating communication, learning and acting




# Requirements on Agent Design

- The *Learning Problem*
  - The *Knowledge Problem*
    - Instruction may carry any type of knowledge (semantic, meta, control)
    - Applying knowledge to ongoing task in current context
    - Using one kind of knowledge to learn a different kind (semantic → procedural)



# Requirements on Agent Design

- The *Learning Problem*
  - The *Knowledge Problem*
    - Instruction may carry any type of knowledge (semantic, meta, control)
    - Applying knowledge to ongoing task in current context
    - Using one kind of knowledge to learn a different kind (semantic → procedural)
  - The *Transfer Problem*
    - Learning generally applicable knowledge
    - Transfer to appropriate conditions in future



## Learning from Instruction with Soar: An Example

- Learns abstract action, *patrol*

patrol room 1 room 2 room 3

go-to-room room 1, go-to-room room 2,

go-to-room room 3, go-to-room room 1

## Learning from Instruction with Soar: An Example

- Learns abstract action, *patrol*

patrol room 1 room 2 room 3

go-to-room room 1, go-to-room room 2,

go-to-room room 3, go-to-room room 1

- Instruction type

<b>Scope</b>	<b>Application</b>	<b>Communication</b>	<b>Context</b>
tutorial	automated	agent-initiated	situated
<b>Learning</b>	<b>Reasoning</b>	<b>Knowledge</b>	<b>Task Structuring</b>
rote	-	-	-

# Learning from Instruction with Soar: An Example

- Learns abstract action, *patrol*

patrol room 1 room 2 room 3

go-to-room room 1, go-to-room room 2,

go-to-room room 3, go-to-room room 1

- Instruction type

Scope	Application	Communication	Context
tutorial	automated	agent-initiated	situated
Learning	Reasoning	Knowledge	Task Structuring
rote	-	-	-

- Agent requirements

- The *Mapping Problem*: hand-coded symbols
- The *Content Problem*: pushed to the human
- **The Interaction Problem**: agent-initiated communication
- **The Knowledge Problem**: uses instructions to perform and learn *patrol*
- The *Transfer Problem*: to be investigated later

# Learning from Instruction with Soar: An Example

- Learns abstract action, *patrol*

patrol room 1 room 2 room 3

go-to-room room 1, go-to-room room 2,

go-to-room room 3, go-to-room room 1

- Instruction type

Scope	Application	Communication	Context
tutorial	automated	agent-initiated	situated
Learning	Reasoning	Knowledge	Task Structuring
rote	-	-	-

- Agent requirements


- The *Mapping Problem*: hand-coded symbols
- The *Content Problem*: pushed to the human
- **The Interaction Problem**: agent-initiated communication
- **The Knowledge Problem**: uses instructions to perform and learn *patrol*
- The *Transfer Problem*: to be investigated later

- Soar architecture components

- Procedural, semantic, working memory

# Knowledge Levels


- Procedural Knowledge
  - Immediate application
  - *state-no-change*: implies lack of procedural knowledge for current state






# Knowledge Levels

- Procedural Knowledge
  - Immediate application
  - *state-no-change*: implies lack of procedural knowledge for current state
- Semantic/Episodic Memory
  - Deliberate lookup
  - *retrieve/query failure*



# Knowledge Levels

- Procedural Knowledge
  - Immediate application
  - *state-no-change*: implies lack of procedural knowledge for current state
- Semantic/Episodic Memory
  - Deliberate lookup
  - *retrieve/query failure*
- Human Instructor
  - Deliberate questions
  - Incomplete/incorrect/specific knowledge
    - requires inference
    - generalization
    - *situated explanation* (Huffman and Laird, 1995)<sup>1</sup>



<sup>1</sup>Huffman, S. and Laird, J. (1995). Flexibly Instructable Agents. *Journal of Artificial Intelligence Research*, 3:271–324

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]
- *Wait* for response from the instructor [interaction]

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]
- *Wait* for response from the instructor [interaction]
- *Parse* [mapping]

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]
- *Wait* for response from the instructor [interaction]
- *Parse* [mapping]
- *Assimilate*:
  - create appropriate *structures*
  - store in semantic memory using information stored in *Detect* phase [knowledge]

# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]
- *Wait* for response from the instructor [interaction]
- *Parse* [mapping]
- *Assimilate*:
  - create appropriate *structures*
  - store in semantic memory using information stored in *Detect* phase [knowledge]
- *Apply*:
  - recreate *state stack* [interaction]
  - apply the new piece of information [knowledge]




# The Instruction Cycle

- *Detect* lack of knowledge to proceed further [knowledge]
  - *state-no-change* creates a *lookup-smem* subgoal
  - *retrieval* attempt may lead to a success/failure
  - if *success*, apply the retrieved command in *superstate*
  - if *failure*, store relevant information at the topstate [interaction]
- *Query*:
  - Create *query* using information stored in the *Detect* phase [content]
  - Generate discourse [mapping]
- *Wait* for response from the instructor [interaction]
- *Parse* [mapping]
- *Assimilate*:
  - create appropriate *structures*
  - store in semantic memory using information stored in *Detect* phase [knowledge]
- *Apply*:
  - recreate *state stack* [interaction]
  - apply the new piece of information [knowledge]
- *Explain*: ? [transfer]


## Example Execution

- S1 (robot): ?  
state-no-change




## Example Execution

- S1 (robot): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
failure
- S2 (lookup-memory): store-information




## Example Execution

- S1 (robot): ?  
state-no-change
  - S2 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
  - S1 (robot): ask-instructor
- agent: "robot"  
instructor: "patrol room 1 room 2 room 3"
- S1 (robot): store-instruction




## Example Execution

- S1 (robot): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "robot"  
instructor: "patrol room 1 room 2 room 3"  
  
S1 (robot): store-instruction
- S1 (topstate): ?  
state-no-change




## Example Execution

- S1 (robot): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "robot"  
instructor: "patrol room 1 room 2 room 3"  
  
S1 (robot): store-instruction
- S1 (topstate): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose patrol in S1




## Example Execution

- S1 (robot): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "robot"  
instructor: "patrol room 1 room 2 room 3"  
  
S1 (robot): store-instruction
- S1 (topstate): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose patrol in S1
- S1 (robot): patrol  
operator-no-change



## Example Execution


- S2 (patrol): ?  
state-no-change





## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information




## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor

agent: "patrol"

instructor: "go-to-room 1"

S1 (robot): store-instruction




## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor

agent: "patrol"  
instructor: "go-to-room 1"


S1 (robot): store-instruction

- S1 (topstate): ?  
state-no-change




## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "patrol"  
instructor: "go-to-room 1"  
  
S1 (robot): store-instruction
- S1 (topstate): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose patrol in S1



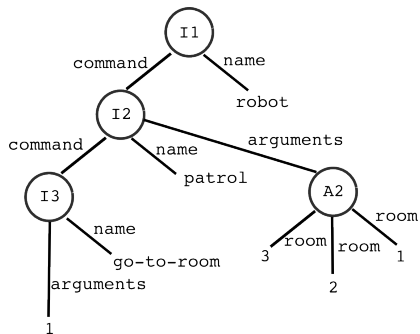
## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "patrol"  
instructor: "go-to-room 1"  
  
S1 (robot): store-instruction
- S1 (topstate): ?  
state-no-change
- S2 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose patrol in S1
- S1 (robot): patrol  
operator-no-change




## Example Execution

- S2 (patrol): ?  
state-no-change




## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose go-to-room 1 (S2)



## Example Execution


- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose go-to-room 1 (S2)
- S2 (patrol): go-to-room 1  
S2 (patrol): ?  
state-no-change






## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose go-to-room 1 (S2)
- S2 (patrol): go-to-room 1  
S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information




## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose go-to-room 1 (S2)
- S2 (patrol): go-to-room 1  
S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "go-to-room 1"  
instructor: "go-to-room 2"  
  
S1 (robot): store-instruction



## Example Execution

- S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
success  
S2 (lookup-memory): propose go-to-room 1 (S2)
- S2 (patrol): go-to-room 1  
S2 (patrol): ?  
state-no-change
- S3 (lookup-memory): retrieve-next  
failure  
S2 (lookup-memory): store-information
- S1 (robot): ask-instructor  
  
agent: "go-to-room 1"  
instructor: "go-to-room 2"  
  
S1 (robot): store-instruction
- ...



# Conclusions

---

- Limitations
  - very limited in scope
  - learns by rote
  - agent initiated communication only

# Conclusions

---

- Limitations
  - very limited in scope
  - learns by rote
  - agent initiated communication only
- Future Work
  - Understanding and solving the *transfer* problem for this instruction set
  - Using episodic memory
  - More detailed investigation of various kinds of instruction

# Nuggets and Coal

---

- Nuggets
  - Have a proof of concept of how instruction can be used for acting
  - Learns 'patrol' and can learn other commands in a similar fashion
- Coal
  - Cannot learn general conditions yet
  - Not a good idea of what the goal of 'patrol' is