

Learning Hierarchical Goal-Oriented Tasks from Situated Interactive Instruction

Shiwali Mohan, John E. Laird

Computer Science and Engineering
University of Michigan, Ann Arbor


June 18th, 2014

Interactive Task Learning




Interactive Task Learning

characteristics



relational goal
structure




hierarchical
decomposition


Interactive Task Learning



characteristics



relational goal
structure



hierarchical
decomposition


acquire

- what?
- how?
- when?


Interactive Task Learning



characteristics



relational goal
structure



hierarchical
decomposition

acquire

- what?
- how?
- when?


desiderata

- multi-task learning
- assisted transfer
- fast generalization
- distributed initiative


Interactive Task Learning



characteristics



relational goal
structure



hierarchical
decomposition

acquire

- what?
- how?
- when?

desiderata

- multi-task learning
- assisted transfer
- fast generalization
- distributed initiative


approach

- composable, hierarchical representations
- knowledge-rich machine learning - EBL
- demo


Interactive Task Learning



characteristics



relational goal
structure



hierarchical
decomposition

desiderata

- multi-task learning
- assisted transfer
- fast generalization
- distributed initiative

acquire

- what?
- how?
- when?

approach

- soar state stack
+ operators
- chunking
+ selection
- demo

Task Representation

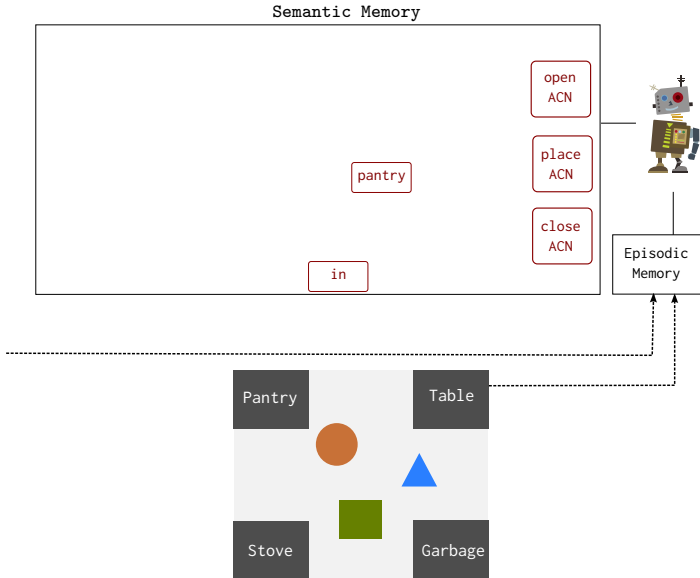
For the task store:

- What?
- parameters
Store the green cylinder.
`store([o], pantry, in([o],pantry))`
 - subtasks
`store: open, move [pick-up, put-down], close`
 - goal
`in(O2,pantry) \wedge closed(pantry)`
-
- How?
- policy
`if [state,task] then execute([subtask])`
 - model
`if [state,task] then [next-state]`
-
- When?
- availability
`if [state] then available(store)`
 - termination
`if [state] then terminate(store)`

Interactive Example Execution

Interaction trace

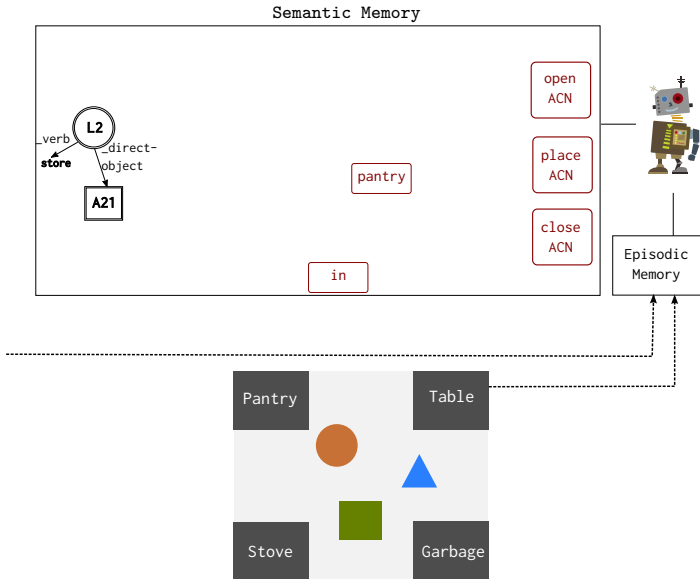
Instructor: Store the green rectangle.



Interactive Example Execution

Interaction trace

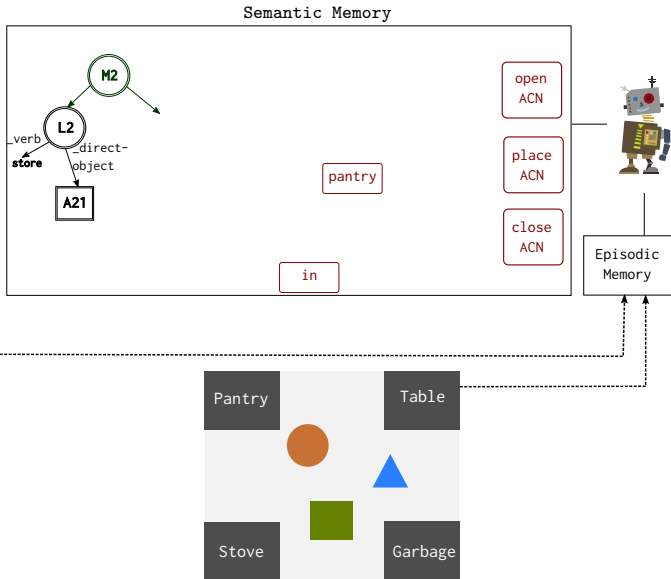
Instructor: Store the green rectangle.



Interactive Example Execution

Interaction trace

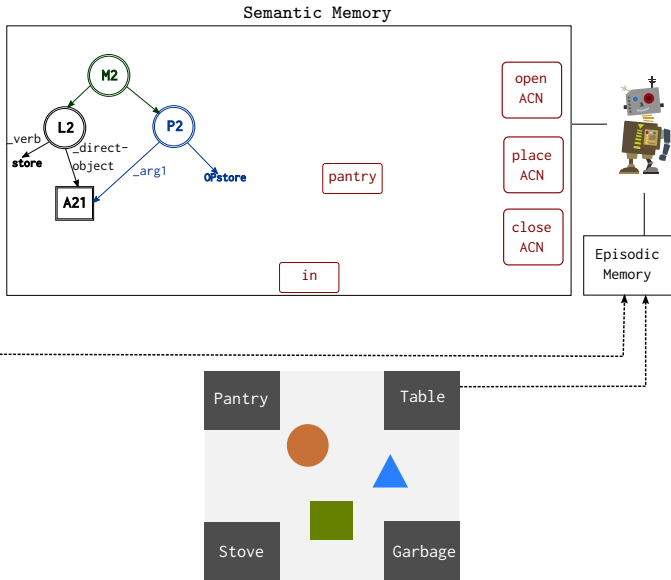
Instructor: Store the green rectangle.



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

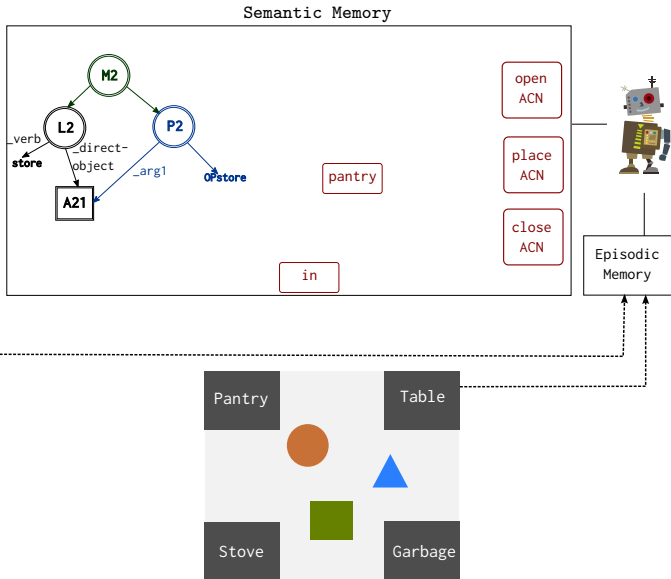


Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?




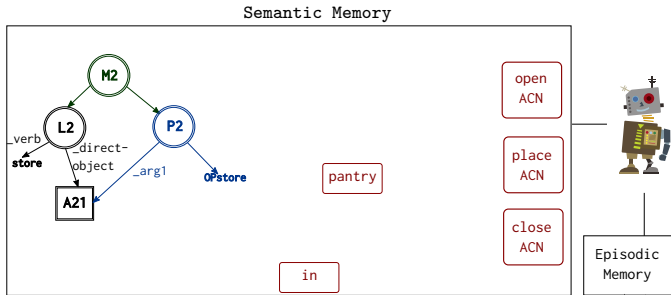
Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.




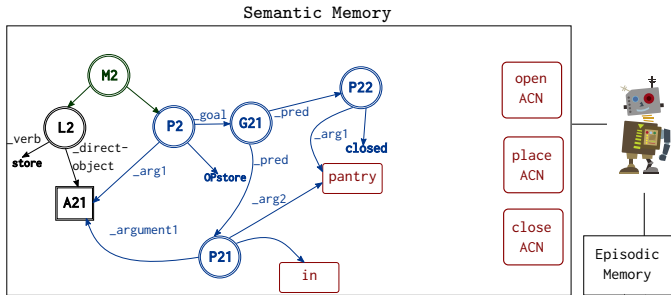
Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.



Interactive Example Execution



Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?



Interactive Example Execution

Interaction trace



Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

Instructor: Open the pantry.



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.


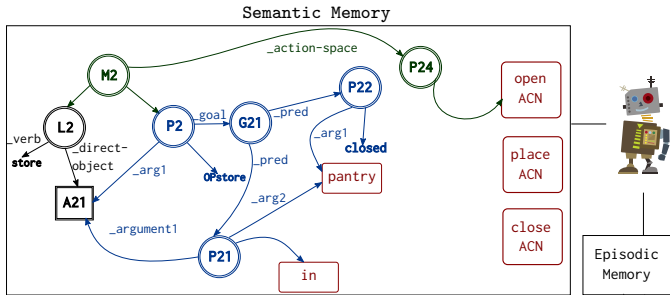
Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

Instructor: Open the pantry.

Agent: Which action should I take?



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?


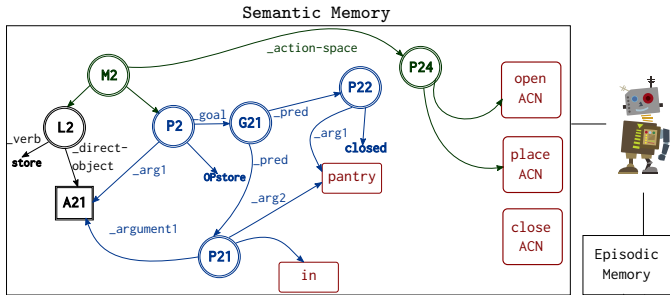
Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

Instructor: Open the pantry.

Agent: Which action should I take?

Instructor: Place the object in the pantry.



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.


Agent: Which action should I take?

Instructor: Open the pantry.

Agent: Which action should I take?

Instructor: Place the object in the pantry.

Agent: OK.



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

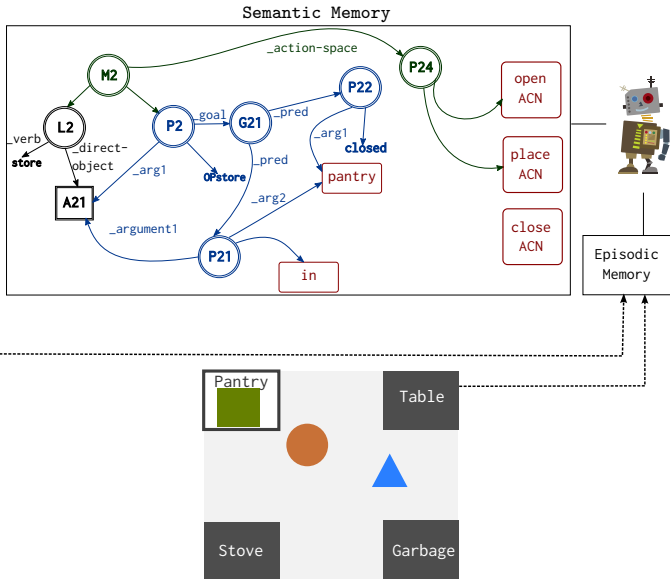
Instructor: Open the pantry.

Agent: Which action should I take?

Instructor: Place the object in the pantry.

Agent: OK.

Agent: explores action space



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

Instructor: Open the pantry.

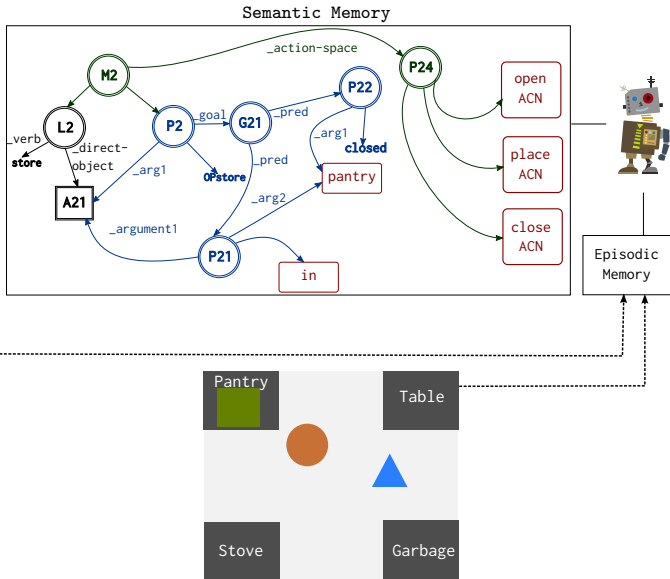
Agent: Which action should I take?

Instructor: Place the object in the pantry.

Agent: OK.

Agent: explores action space

Agent: closes the pantry



Interactive Example Execution

Interaction trace

Instructor: Store the green rectangle.

Agent: What is the goal of the action?

Instructor: The goal is the green rectangle in the pantry and the pantry closed.

Agent: Which action should I take?

Instructor: Open the pantry.


Agent: Which action should I take?

Instructor: Place the object in the pantry.

Agent: OK.

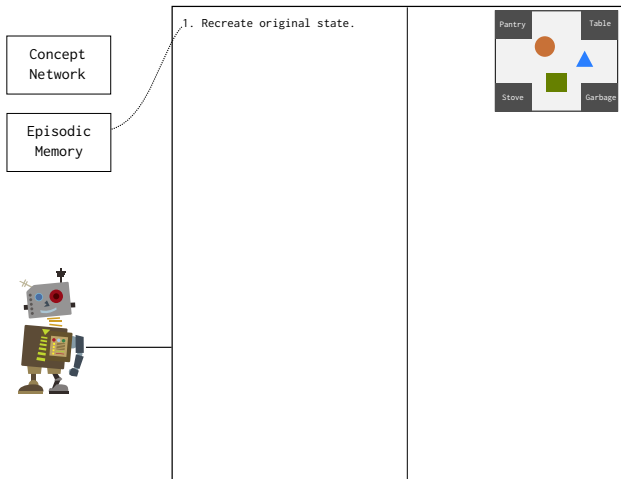
Agent: explores action space

Agent: closes the pantry



Retrospective EBL

Retrospective EBL




Retrospective EBL

Concept
Network

Episodic
Memory

1. Recreate original state.
2. Generate desired state.

Pantry

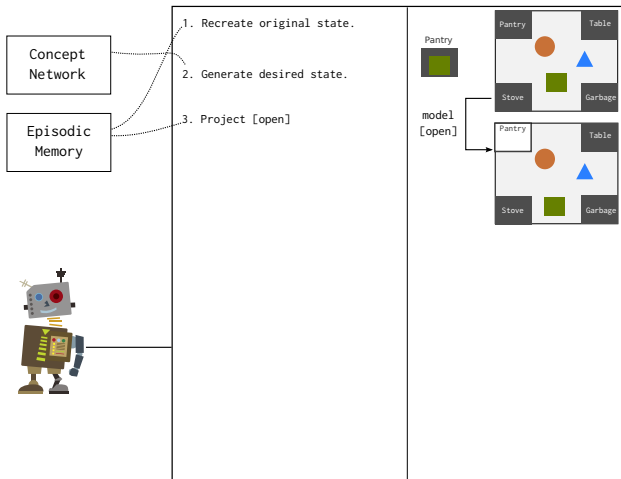


terminate store:
If `store([x])` and
`IN([x], PANTRY)` and
`CLOSED(PANTRY)`

-->

`terminate store[x]`

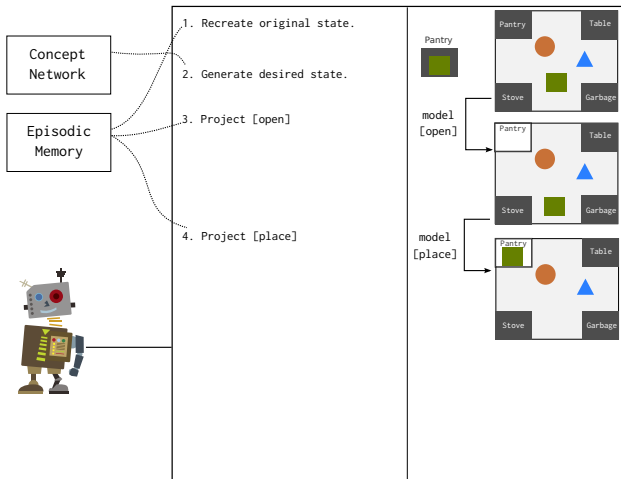
Retrospective EBL



terminate store:
If store ([x]) and
IN ([x], PANTRY) and
CLOSED (PANTRY)

-->
terminate store [x]

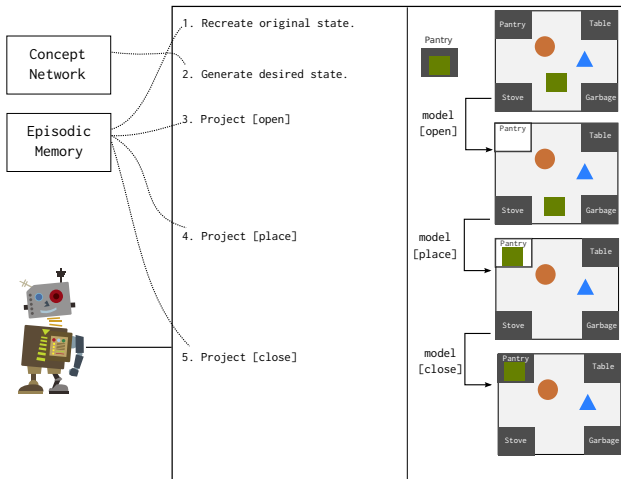
Retrospective EBL



terminate store:
If store ([x]) and
IN ([x], PANTRY) and
CLOSED (PANTRY)

-->
terminate store [x]


Retrospective EBL



terminate store:
If store([x]) and
IN([x], PANTRY) and
CLOSED(PANTRY)

-->
terminate store[x]

Retrospective EBL



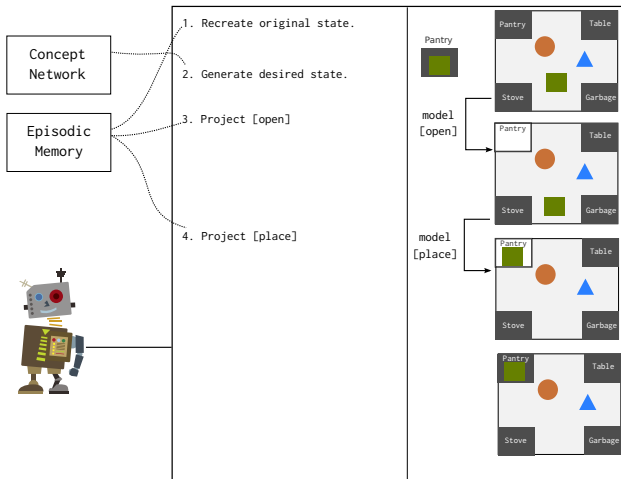
terminate store:
If store([x]) and
IN([x], PANTRY) and
CLOSED(PANTRY)

-->
terminate store[x]

select close:
If store([x]) and
IN([x], PANTRY) and
OPEN(PANTRY)

-->
select close(PANTRY)

Retrospective EBL



```
terminate store:  
If store([x]) and  
IN([x], PANTRY) and  
CLOSED(PANTRY)
```

```
-->  
terminate store[x]
```

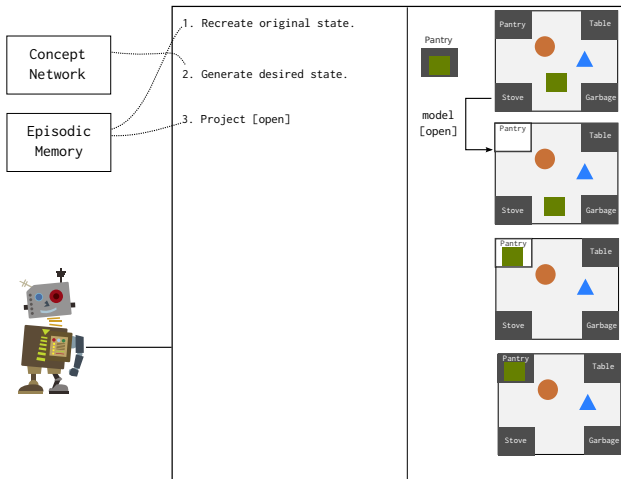
```
select close:  
If store([x]) and  
IN([x], PANTRY) and  
OPEN(PANTRY)
```

```
-->  
select close(PANTRY)
```

```
select place:  
If store([x]) and  
~IN([x], PANTRY) and  
OPEN(PANTRY)
```

```
-->  
select  
place([x], IN, PANTRY)
```

Retrospective EBL



```
terminate store:  
If store([x]) and  
IN([x],PANTRY) and  
CLOSED(PANTRY)
```

```
-->  
terminate store [x]
```

```
select close:  
If store([x]) and  
IN([x],PANTRY) and  
OPEN(PANTRY)
```

```
-->  
select close(PANTRY)
```

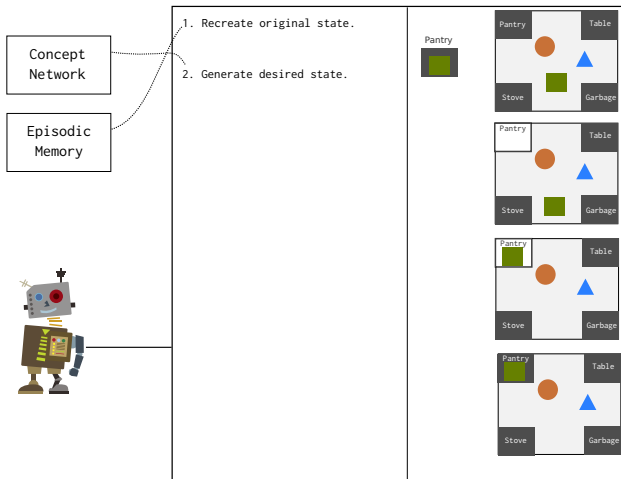
```
select place:  
If store([x]) and  
~IN([x],PANTRY) and  
OPEN(PANTRY)
```

```
-->  
select  
place([x],IN,PANTRY)
```

```
select open:  
If store([x]) and  
~IN([x],PANTRY) and  
CLOSED(PANTRY)
```

```
-->  
select open(PANTRY)
```

Retrospective EBL



terminate store:
If store([x]) and
IN([x], PANTRY) and
CLOSED(PANTRY)

-->
terminate store [x]

select close:
If store([x]) and
IN([x], PANTRY) and
OPEN(PANTRY)

-->
select close(PANTRY)

select place:
If store([x]) and
~IN([x], PANTRY) and
OPEN(PANTRY)

-->
select
place([x], IN, PANTRY)

select open:
If store([x]) and
~IN([x], PANTRY) and
CLOSED(PANTRY)

-->
select open(PANTRY)

available store:
If ~IN([x], PANTRY) or
OPEN(PANTRY)

-->
available store([x])

Multi-task Learning

pick-up & put-down:

`place([x], [rel], [y]), move([x], [y]), discard([x]), store([x])`

functional:

`cook([x]), serve([x])`

organizational:

`stack-3([x], [y], [z]), stack-4([x], [y], [z], [w])`


Generality in Learning

Learns general representations of tasks from few (~2-3) instances

Generality in Learning

Learns general representations of tasks from few (~2-3) instances

abstraction




on (purple object, table)

Generality in Learning

Learns general representations of tasks from few (~2-3) instances

abstraction



on (purple object, table)

predicate selection

select open:

If `store(O1)` and `-IN(O1,PANTRY)` and
`CLOSED(PANTRY)` and ~~`CLOSED(STOVE)`~~ and
~~`OFF(STOVE)`~~ and ~~`-ON(O2,STOVE)`~~ and ...


-->

`select open(PANTRY)`

Generality in Learning




Learns general representations of tasks from few (~2-3) instances

abstraction



on (purple object, table)

causal analysis



predicate selection

select open:

If store(O1) and -IN(O1,PANTRY) and
CLOSED(PANTRY) and ~~CLOSED(STOVE)~~ and
~~OFF(STOVE)~~ and -ON(O2,STOVE) and ...


-->

select open(PANTRY)

Generality in Learning


Learns general representations of tasks from few (~2-3) instances

abstraction



on (purple object, table)

causal analysis



predicate selection

select open:

If `store(O1)` and `-IN(O1,PANTRY)` and
`CLOSED(PANTRY)` and `CLOSED(STOVE)` and
`OFF(STOVE)` and `-ON(O2,STOVE)` and ...

-->

`select open(PANTRY)`

variablization

Store the green rectangle.


The goal is the green rectangle in the pantry and the pantry is closed.

Open the pantry.

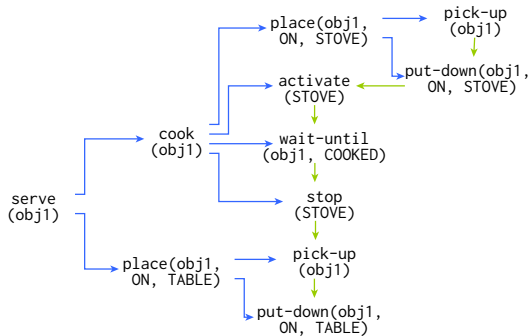
Move the green rectangle to the pantry.

...

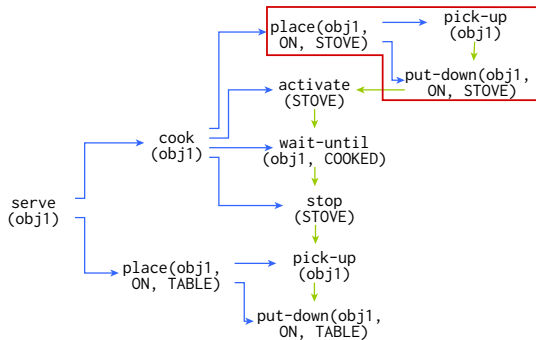
Generality in Learning



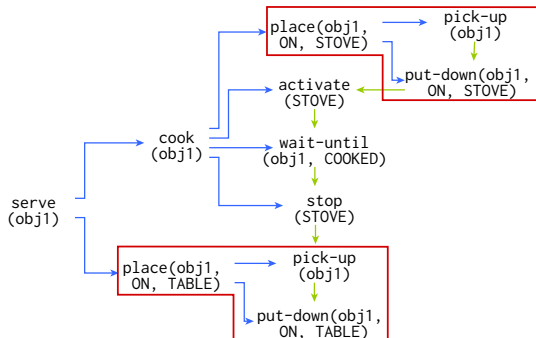
Hierarchical Learning



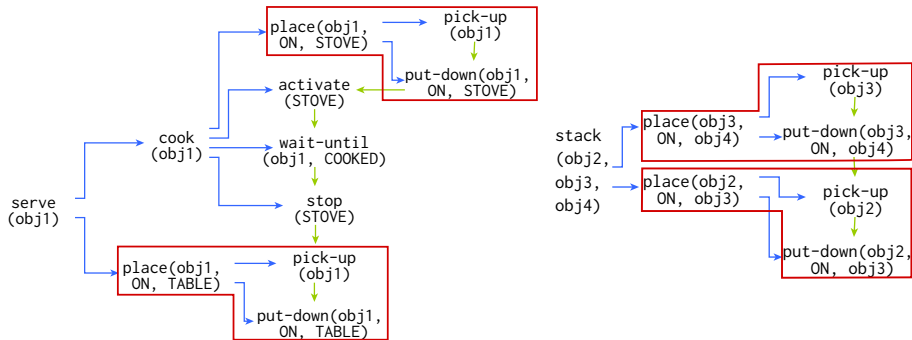
Hierarchical Learning



Hierarchical Learning




Hierarchical Learning



Transfer

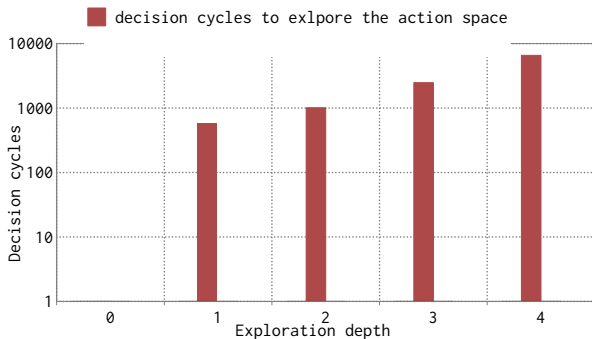
Exploits the common policy space for instruction-aided transfer.



Distributed Initiative

Integrates agent-driven exploration and instruction-guided exploitation


Learning the task *store*.



Distributed Initiative

Integrates agent-driven exploration and instruction-guided exploitation

Learning the task *store*.



Nuggets and Coal

Nuggets

- composable, hierarchical, transferable representation
- multi-task learning
- fast generalization
- distributed initiative of learning
- uses several Soar mechanisms
- Rosie talks!

Coal

- only *achievement* tasks
- not completely robust to instruction errors
- HRI evaluation