




THE UNIVERSITY
of ADELAIDE

Street: A Hardware Production Engine for Soar

Dr. Braden
Phillips

Mostafa
Numan

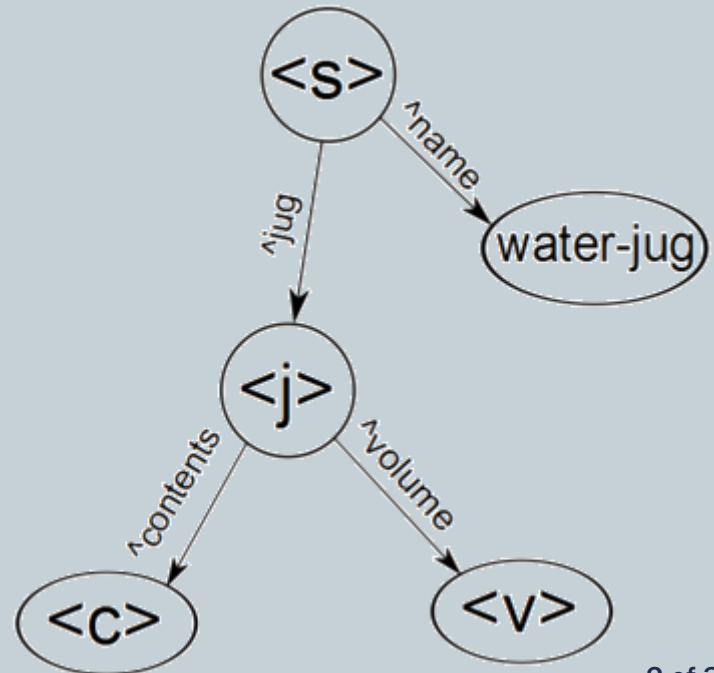
Jesse Frost



Soar, Production Systems

- Working Memory Encoded as WMEs – which can be represented in a graph
- Productions test for subsets of WM that match a specified pattern

```
sp {  
    (state <s> ^name water-jug)  
    (<s> ^jug <j>)  
    (<j> ^contents <c>)  
    (<j> ^volume <v>)  
->  
    (<j> ^empty (- <v> <c>))  
}
```

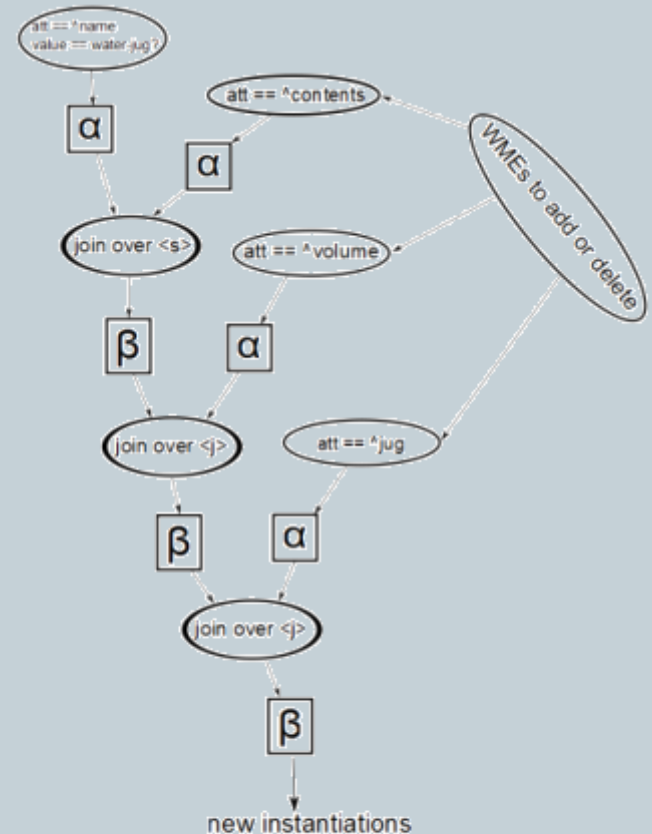


Soar, Production Systems

- How do you find instantiations in working memory?
- Brute force: search every subset all the time
- Rete: track changes to rule satisfaction

Rete

- Rete exploits temporal redundancy
- Stores candidate subsets of working memory
- Tracks changes to these subsets over time

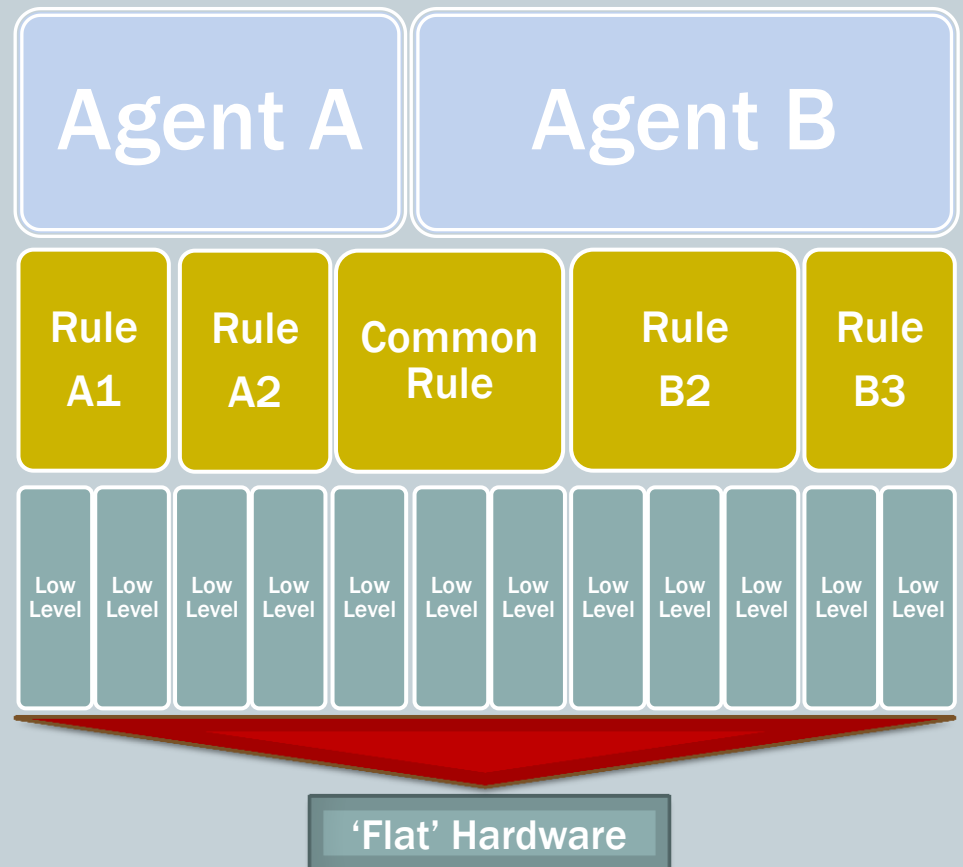


Treat

- Treat is a variation of Rete originally for parallel architectures
- Information in ‘conflict set’ of current instantiations and the beta memories is redundant
- Does not maintain partial joins, just recalculates the information as its needed


Parallelism

- We want a system that allows parallelism at multiple levels
- General Purpose Architectures have bottlenecks



Parallelism

- A hardware platform should be scalable at all levels – need a flat uniform architecture
- Modern Technology makes custom architectures accessible

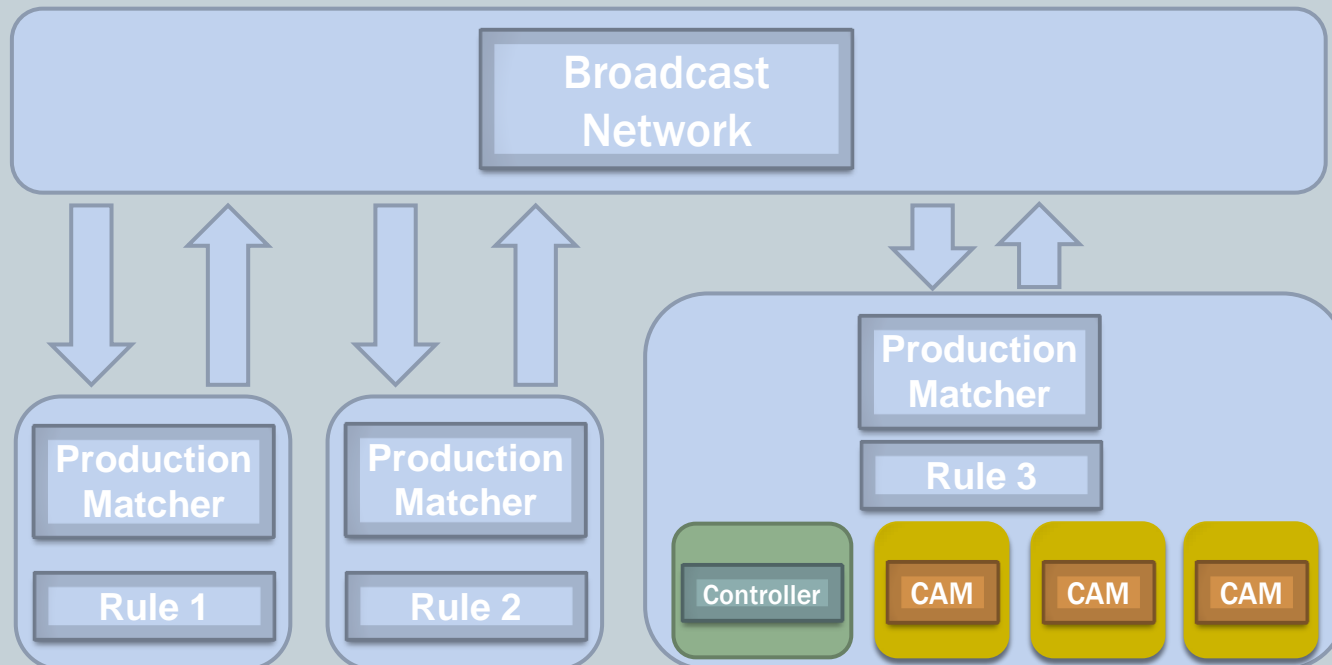


Street

- Street – the name might be from ‘Soar-Treat’
- Not really either
- A hardware platform for executing a parallel production language directly
- Hardware blocks perform the match-act cycle directly (not general purpose processors running a production interpreter)


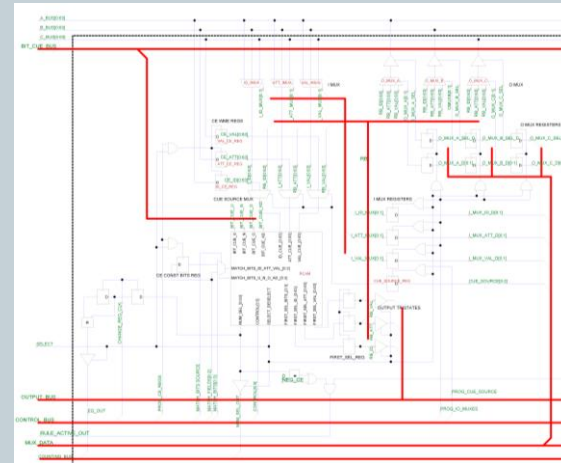
Street

- Use ideas from Rete/Treat
- Subsets of working memory kept in specialised CAMs



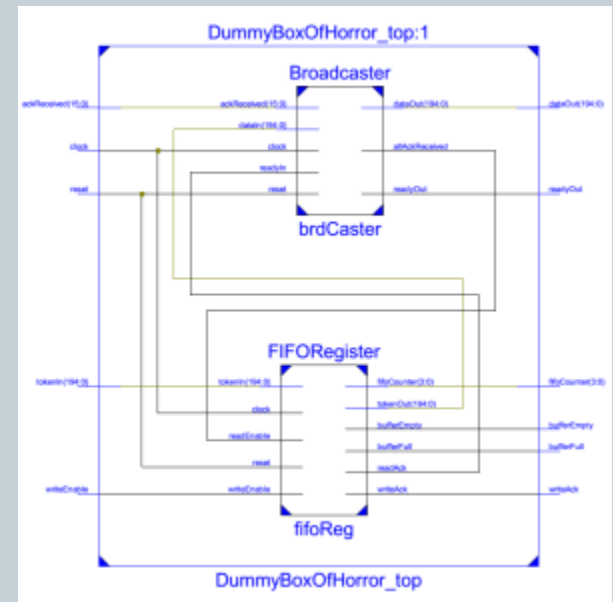
Production Matcher

- Subsets of working memory stored in collections of CAMs
- CAMs designed for single cycle execution of the operations we need
- Each rule is assigned one CAM for each condition element



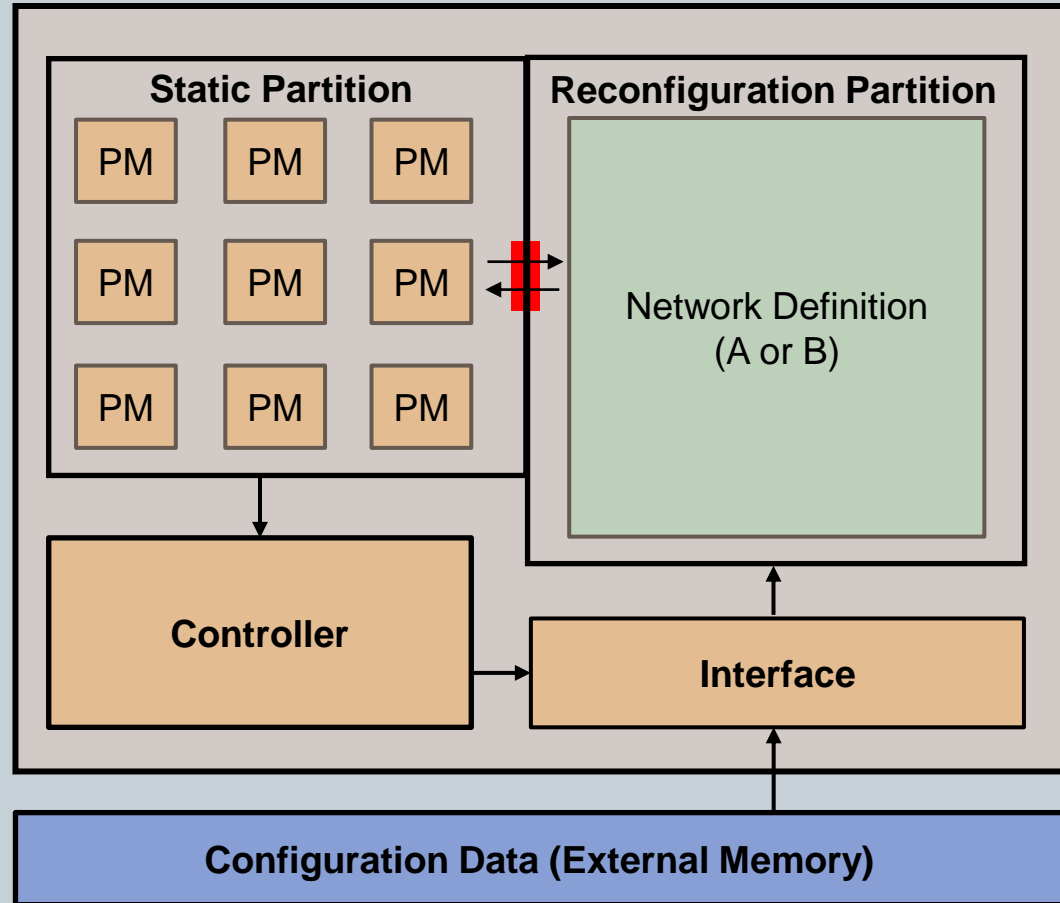
Broadcast Network

- Broadcasts WM changes and synchronises production matchers
- Small broadcasters can be built from FIFO registers
- Gives the architecture a two-section design, one containing matchers and one containing the broadcaster.

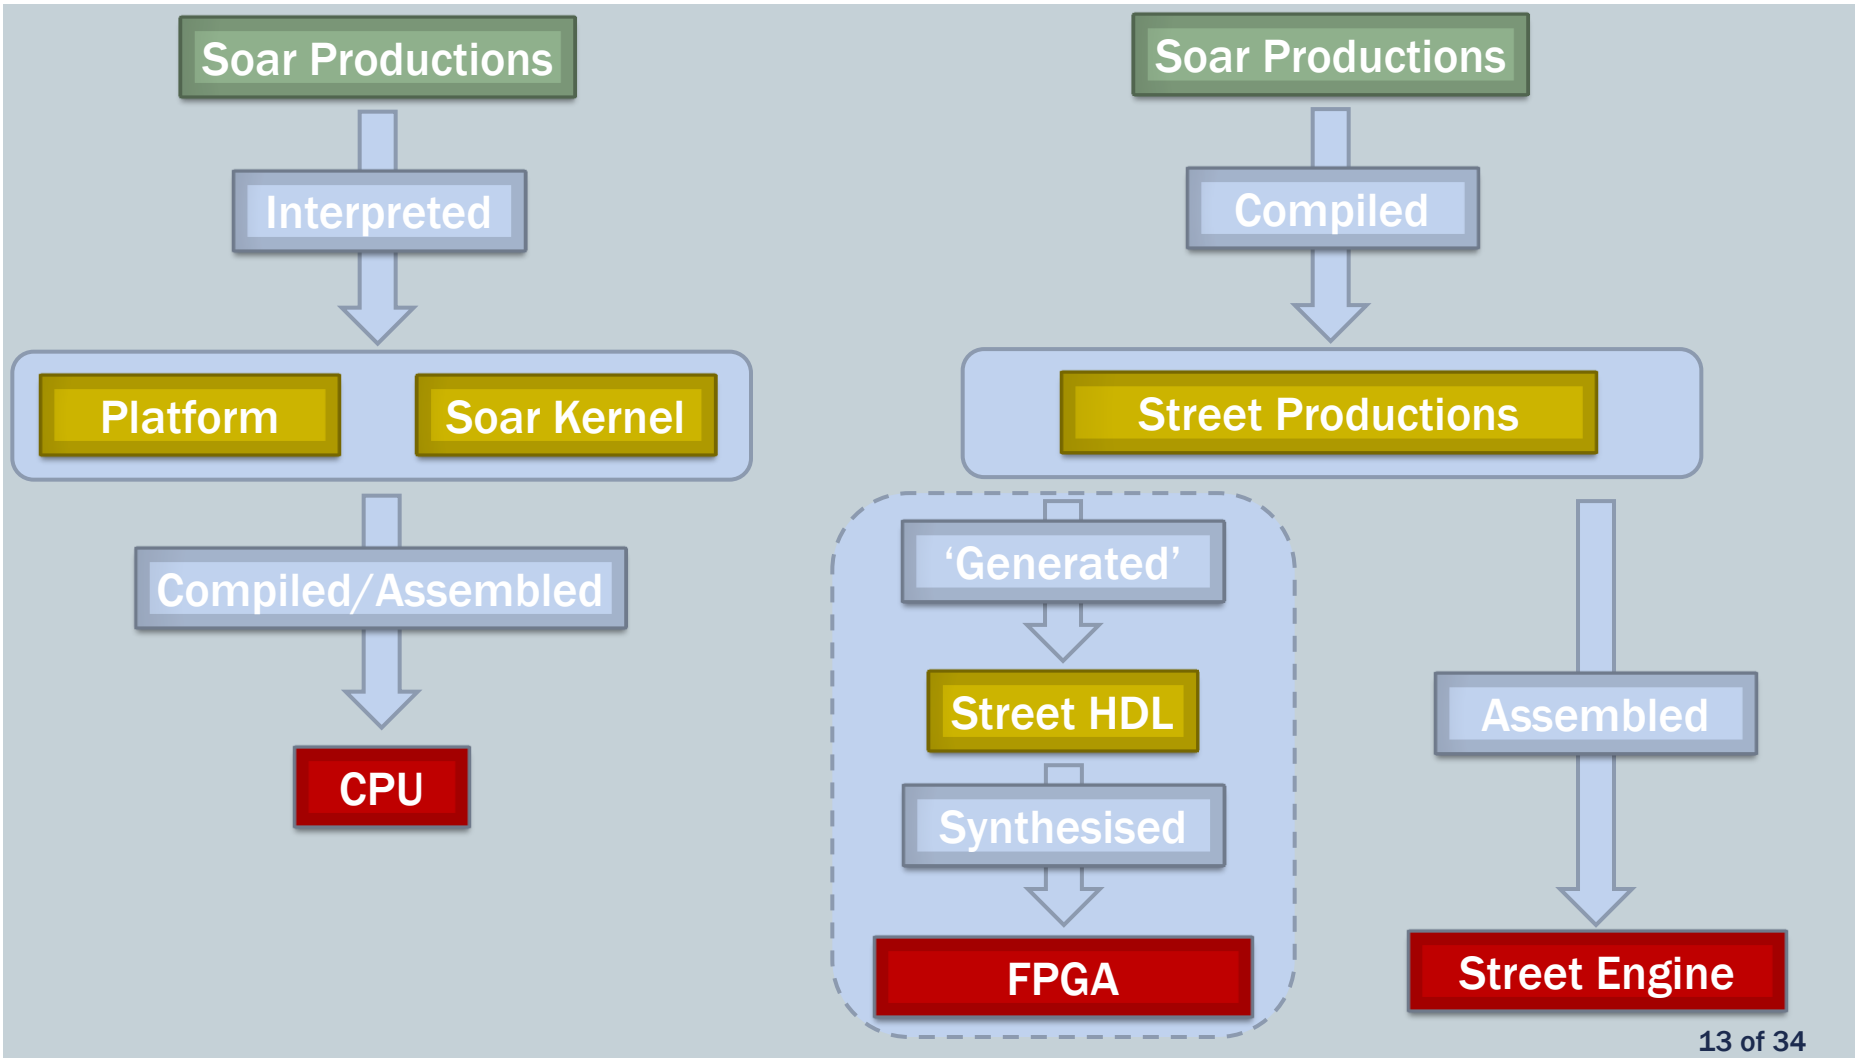


Network-on-Chip

- Network-on-Chip for large systems
- Network definition can be reconfigured dynamically
- Broadcast Network is a potential bottleneck



Street



Street

- Work in progress
- Sorry, no 'Soar-on-a-Chip' just yet
- Preliminary results suggest we could fit several hundred productions on a large commercial FPGA
- Working on proof-of-concept
- Synchronisation and some more complex behaviours for Soar require more rules than the same agent in Soar
- Broadcast Network needs analysis

Vision

- **Application Specific Device – Similar to existing FPGA devices but with CAM blocks rather than conventional LUTs**
- **Very fast cycle times – orders of magnitude faster than software -**
- **As many components of Soar implemented in simple production language as possible – RISC vs CISC**
- **A shift of the programming paradigm from serial rule firing to massively parallel rule firing**

Questions

?