

Mining Episodic Memory

Nate Derbinsky



Research Statement

Analyze the contents of episodic memory across a set of tasks/agents for useful patterns

Goals

- Inform architectural additions/enhancements
- Inform agent design

Initial Directions*

- Episodic summarization & forgetting
- Episodic-to-semantic consolidation


* *Offline, ignoring issues of performance*



Direction 1

Basic idea: event recognition


- Analyze the WME stream and identify groupings that correspond, along some metric, to distinct "events"



Episodic Summarization/Forgetting

Over time, temporally local episodes reduce in specificity

- Coalesce into representative "events"
- Statistical noise pruned/summarized



Potential Outcomes

- Improved EpMem performance
 - Fewer episodes = stabilized query time
 - *Hopefully* equal/better quality
 - Hypothesis: detail useful only in short-term
- Comparison to human cognition
 - Literature: "episodic memories are encoded without great detail"
 - Proposal: progression from high-to-low detail



Candidate Metrics/Signals

- Mechanism performance
 - More WME changes = slower EpMem
 - Serves as a pressure to prune where possible
- Feature/episode...
 - Similarity ala graph structure, temporal distance, feature-level variance (\sim TF/IDF), latent events (\sim LSA)
 - Activation
 - Episode: storage recency, retrieval base-level
- Agent queries
 - Cue = possible signal of details worth keeping
- **Goal.** Content-aware compression



Initial Experimental Framework

1. Gather multi-task episodic databases
 - Sanity check: big switch across unrelated in-the-head agents (counting, water-jug, etc.)
2. Attempt to recover task sequence
 - *Unsure how to evaluate summarization for non-trivial agents*



Candidate Model #1

kMeans in WM space

- Features ala Working Memory Tree
 - Flattens multi-valued attributes, for simplicity
 - Binary feature vector (*bag of words*)
- Greedy search for “optimal” k
 - Distance: {hamming, cosine, hybrid?}
 - Smaller k
 - **Greater compression**
 - Greater within-event variance
 - Greater exemplar variance
 - Larger k
 - Less compression
 - **Less within-event variance**
 - Less exemplar variance



Candidate Model #2

Singular Value Decomposition (SVD)


- Identify features/feature sets that explain the most episode variance
 - Consider value changes = event boundaries
- Smooth within-event feature changes for features that explain little episode variance



Direction 2


Basic idea: online concept learning

- Analyze the WME stream and identify groupings of WMEs that, over time, capture stable distributions of object features, independent of context



Potential Outcomes

- Architectural semantic storage!
 - Useful categorical hierarchies/clusters
 - Summarization of probabilistic distributions



Classes of Candidate Features

1. Individual, over time
 - Variance (e.g. TF/IDF, density estimation)
 - Type (numerical/categorical)
 - Structural location (e.g. input-link, MVAs)

2. Pairwise relations over time
 - Co-variance
 - Structure mapping [w.r.t. existing LTIs]

3. Architectural meta-data
 - Episodic/semantic cues
 - RL reward/update error
 - SVS filtering



Candidate Models

1. Feature co-occurrence matrix
 - Mine for single-level (ala Working-Memory Tree) frequent item sets (e.g. Apriori)
 - Feature variance -> keep vs. discard vs. parameters (e.g. mean/variance)
2. Incremental COBWEB/TRESTLE (MacLellan et al. ACS'15)



Initial Experimental Framework

Issue. What “should” a learned semantic network look like for a domain?

– *I don't have a great answer here...*

1. Synthetic dataset with known priors (e.g. GMM), try to recover
 - Wang: hunting domain
 - Priors on weapon effectiveness w.r.t. prey type
 - Map construction




Potential Interactions

- Episodic category learning
 - In addition to learning concepts, it can be useful to be able to categorize types of events (Doshi, Kira, Wagner ACS'15)
 - Example application: **spontaneous retrieval** of episodes
 - Train a situational classifier, automatically retrieve/signal availability of a similar episode/category exemplar
- Spontaneous retrieval as recommender system
 - Given prior requests, the current context, and **known co-occurrences**, suggest an LTI to supplement the existing context
 - Form of hybrid content/collaborative filtering
- Episodic compression via LTI re-encoding
 - When an episode is retrieved, re-encode subsets of the graph that were STIs with LTIs that have **since been learned**
 - Has the potential to speed/improve retrieval
- Elaborative chunk learning
 - When a concept is learned from perceptual data, automatically augment symbolic interpretation with **learned conceptual LTI**
 - Over time, re-use to learn higher-level features





Evaluation

Nuggets

- Finally looking to episodic memory as a source of useful statistical data
 - Hopefully full of 
- Initial directions lay the groundwork for future architectural mechanisms
 - Episodic summarization, automatic semantic storage
- Potential for beneficial interactions

Coal

- No results yet 
- Huge space of potential signals, metrics, approaches; no  standard



Thank You :)

Questions?

Comments?

Suggestions?

Concerns?

