



# TOWARDS SOLVING VISUAL PUZZLES USING VISUAL KNOWLEDGE REPRESENTATIONS IN SOAR

JAMES BOGGS

UNIVERSITY OF MICHIGAN






# INTRODUCTION

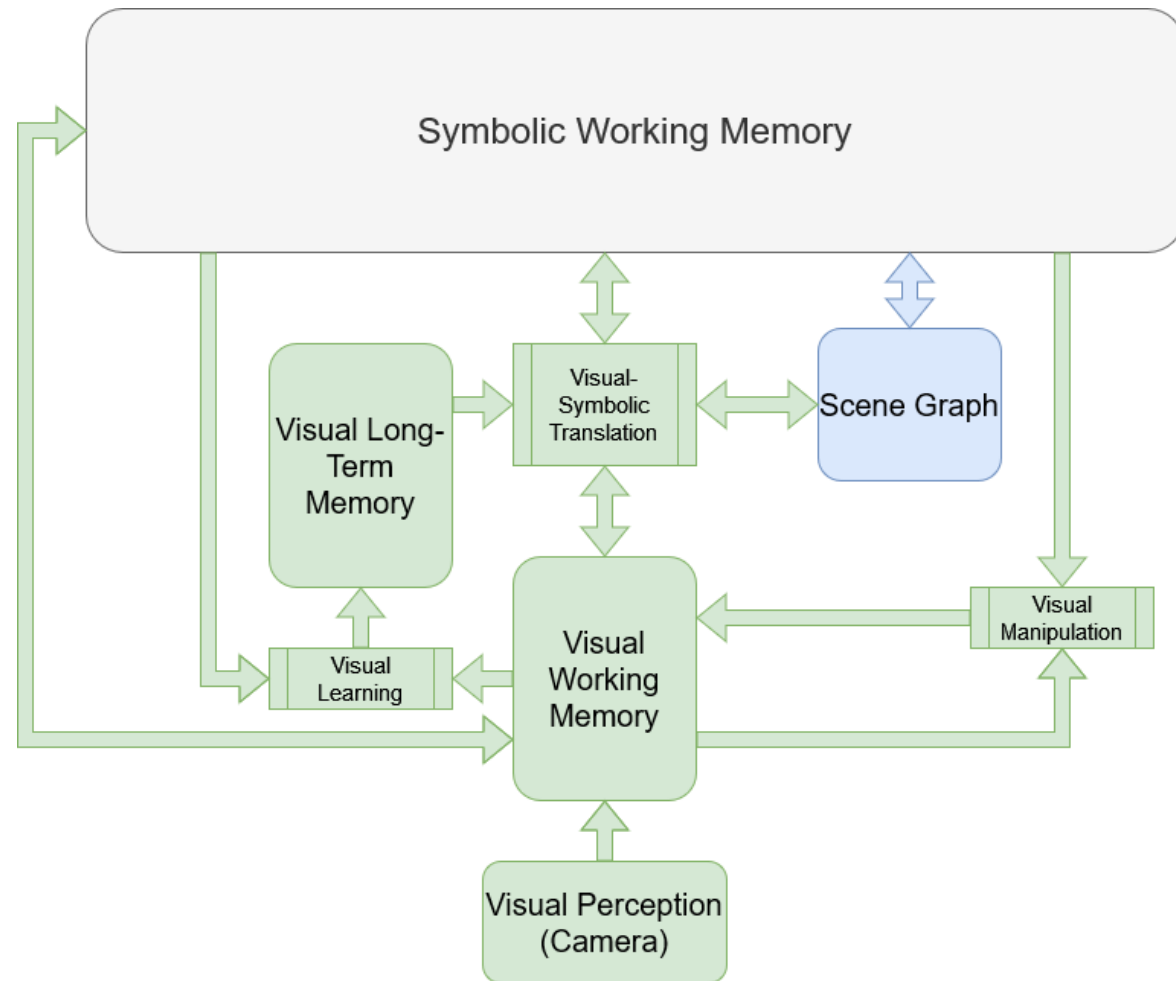
# MOTIVATION

- Many tasks in the real world benefit from or even require visual reasoning



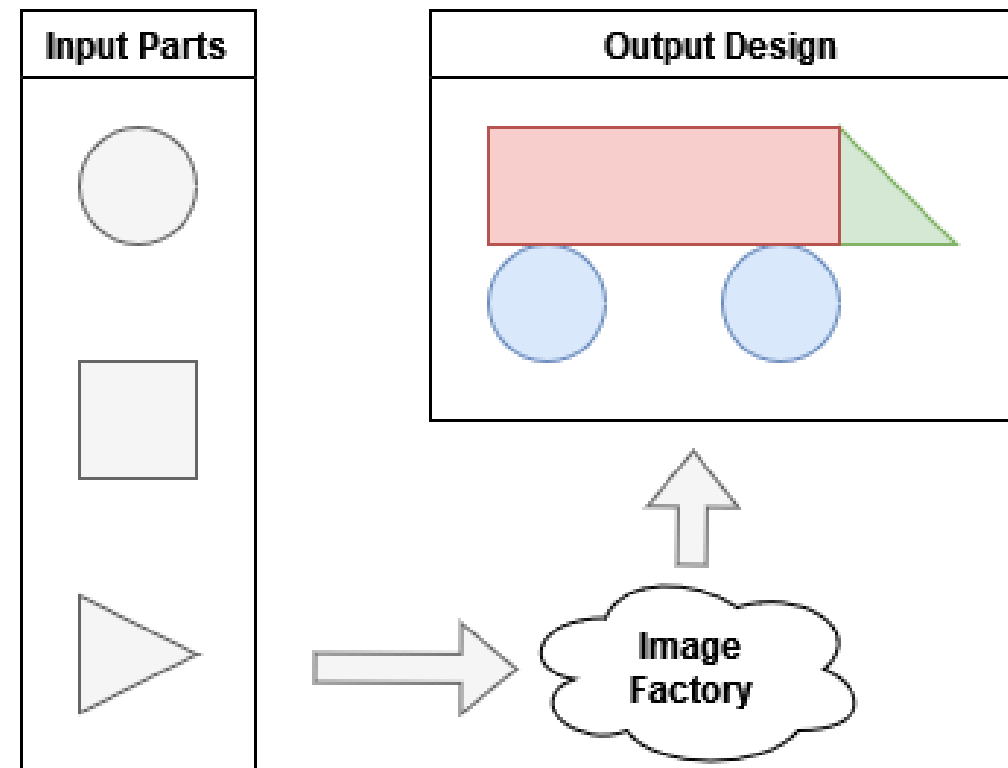
# GENERAL RESEARCH FOCUS

Extend Soar to include visual knowledge and reasoning by creating new visual memories and associated reasoning operations



# OVERVIEW OF CURRENT WORK

- Proof-of-concept agent in simple domain: “Image Factory”
  - Fundamental task: Create output design by modifying and combining input parts
- New visual-symbolic object representation in SVS








# BACKGROUND



# CURRENT SPATIAL-VISUAL SYSTEM

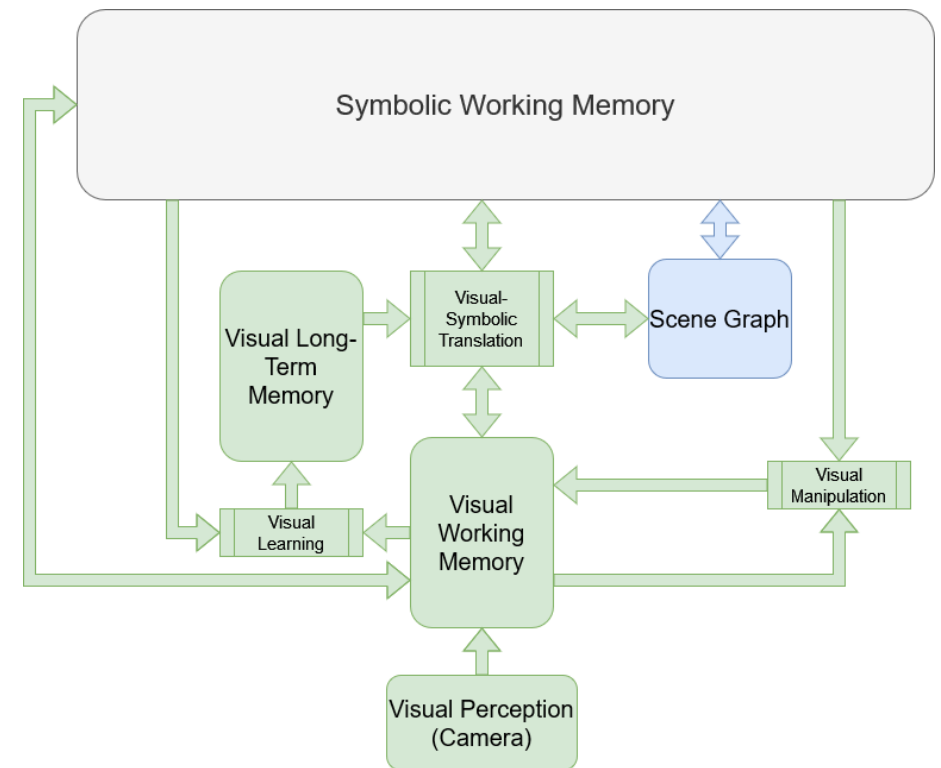


```
(S1 ^svs S3)
(S3 ^command C3 ^spatial-scene S4)
(C3 ^extract E2)
(E2 ^a A1 ^b B1 ^result R7 ^status success ^type intersect)
(A1 ^id car ^status success ^type node)
(B1 ^id pole ^status success ^type node)
(R7 ^record R17)
(R17 ^params P1 ^value false)
(P1 ^a car ^b pole)
```



# STRUCTURE OF SVS2

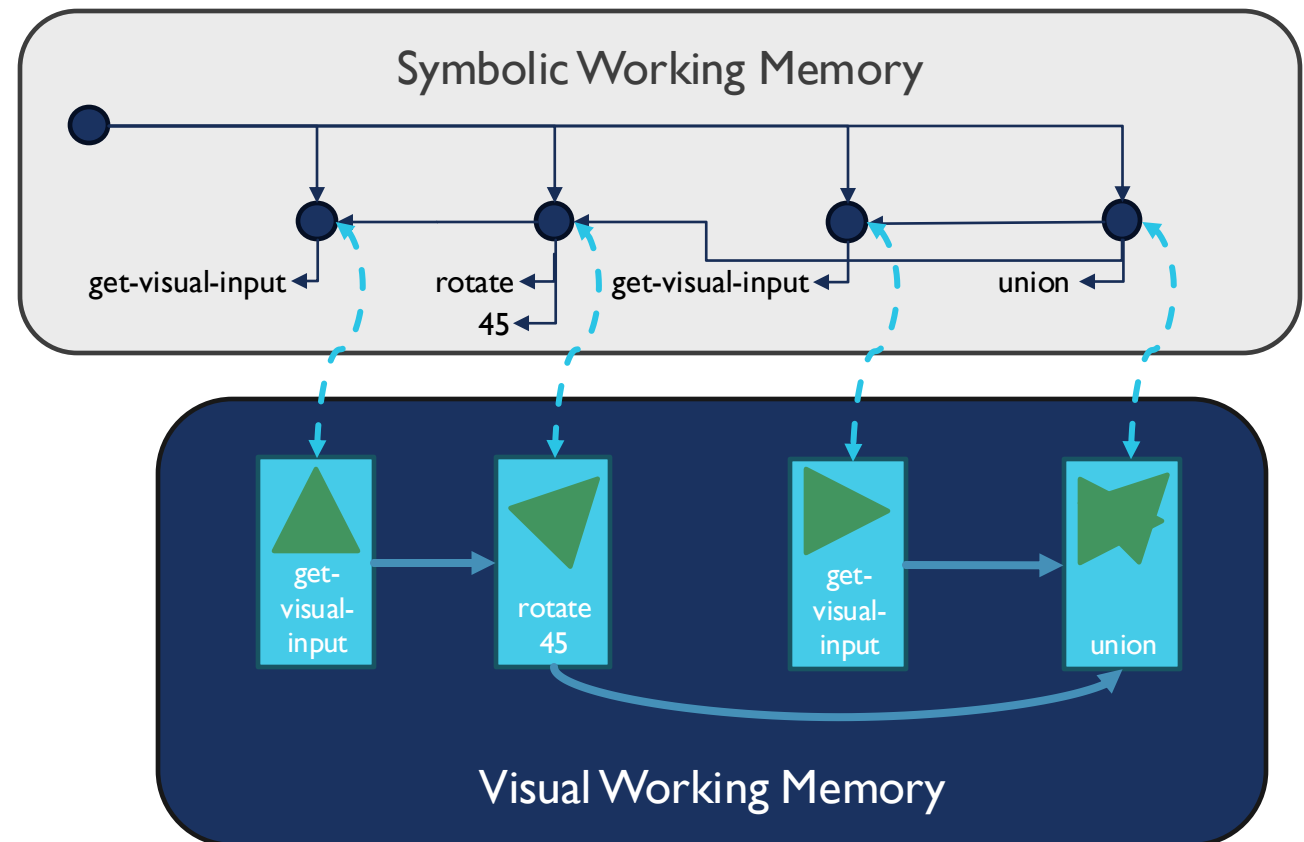
- Three new visual memories:
- Visual Working Memory: holds task-specific visual knowledge and reasoning operations
- Visual Long-Term Memory: holds task-independent, long-term visual knowledge
- Visual Input Buffer: very short term memory storing most recent raw visual percepts





# VISUAL WORKING MEMORY

- Maintains current state of visual reasoning
- Graph of visual operations (VOps) which can generate & manipulate visual knowledge
- VOps + their inputs & outputs connected to symbolic WM



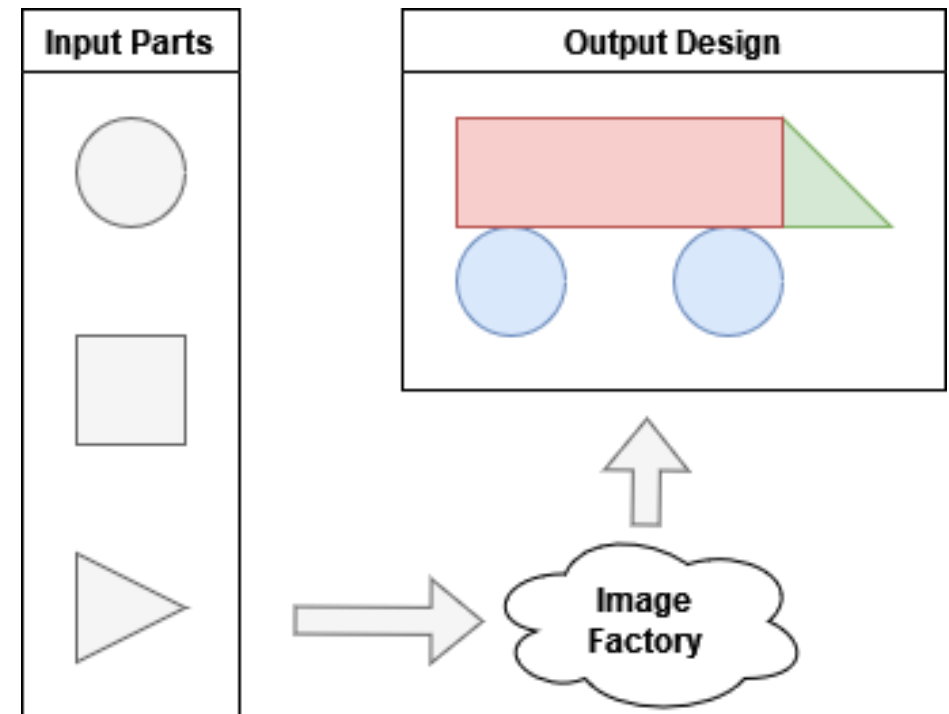


# CURRENT WORK



# IMAGE FACTORY DOMAIN

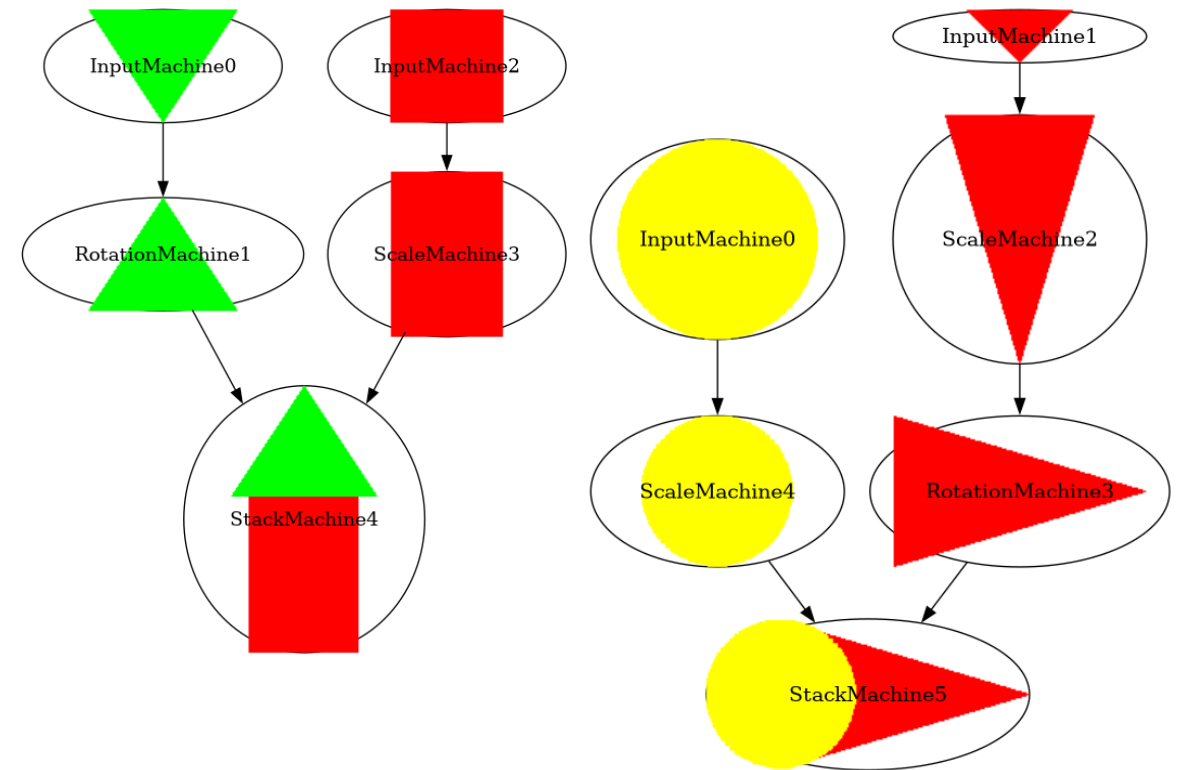
- Objective: given input parts, design a factory which uses machines to assemble the inputs into a specified output product
- Parts & product are both given as images
- Machines can:
  - Generate an input part
  - Scale or rotate an input part
  - Stack two input parts with a relative translation
- Solution should be given as a graph indicating how machines should be connected



# IMAGE FACTORY: SOME SIMPLIFICATIONS

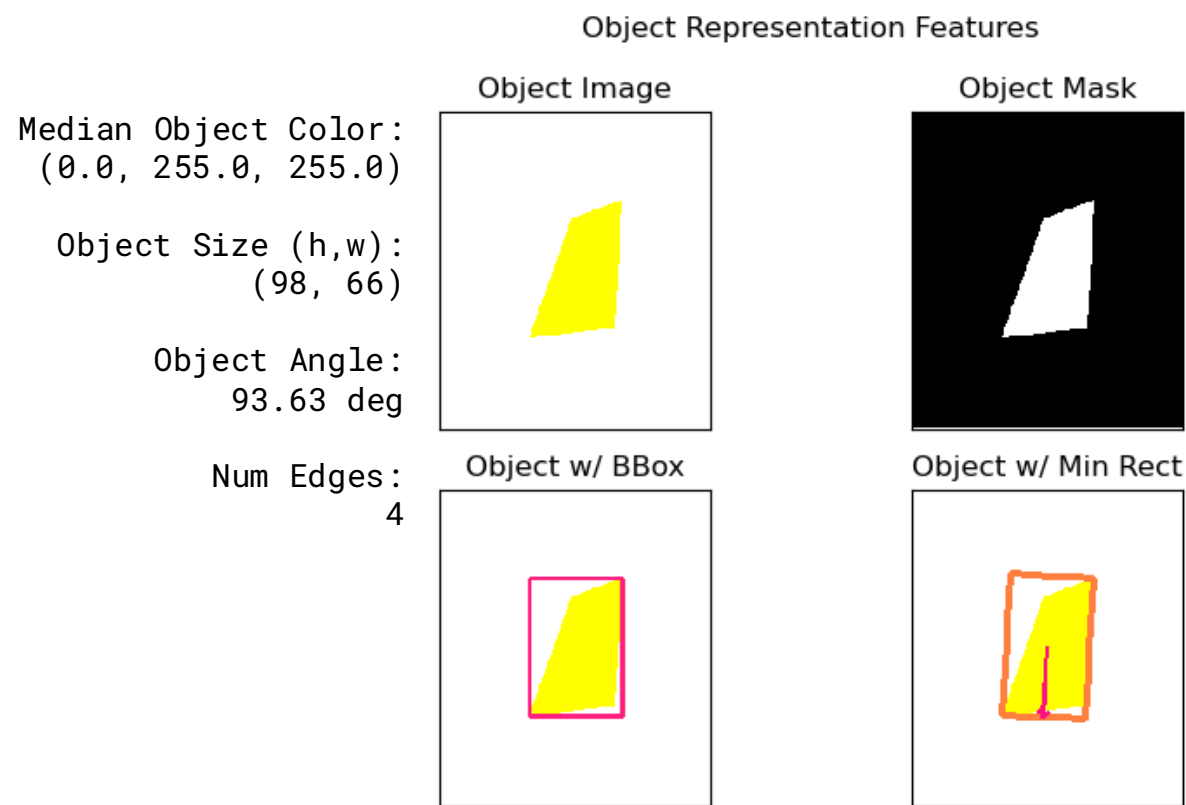
Image Factory domain as described is extremely complex, so some simplifications were made for this work:

- Only two input images
  - Unique shapes and colors
- No color changing machine
- Limited factory size
- Limited values for scaling and rotation






# OBJECT REPRESENTATION

- Fundamental to the domain is *objectness*:
  - there are multiple discrete components which should be treated separately
- Must be able to represent *objects* in addition to images
- Object representation is visual-symbolic: should maintain both visual-spatial & symbolicized features
- Currently, uses hand-crafted feature set, including:
  - Image, mask, height, width, location in original image, color, number of sides, angle of minimal rectangle



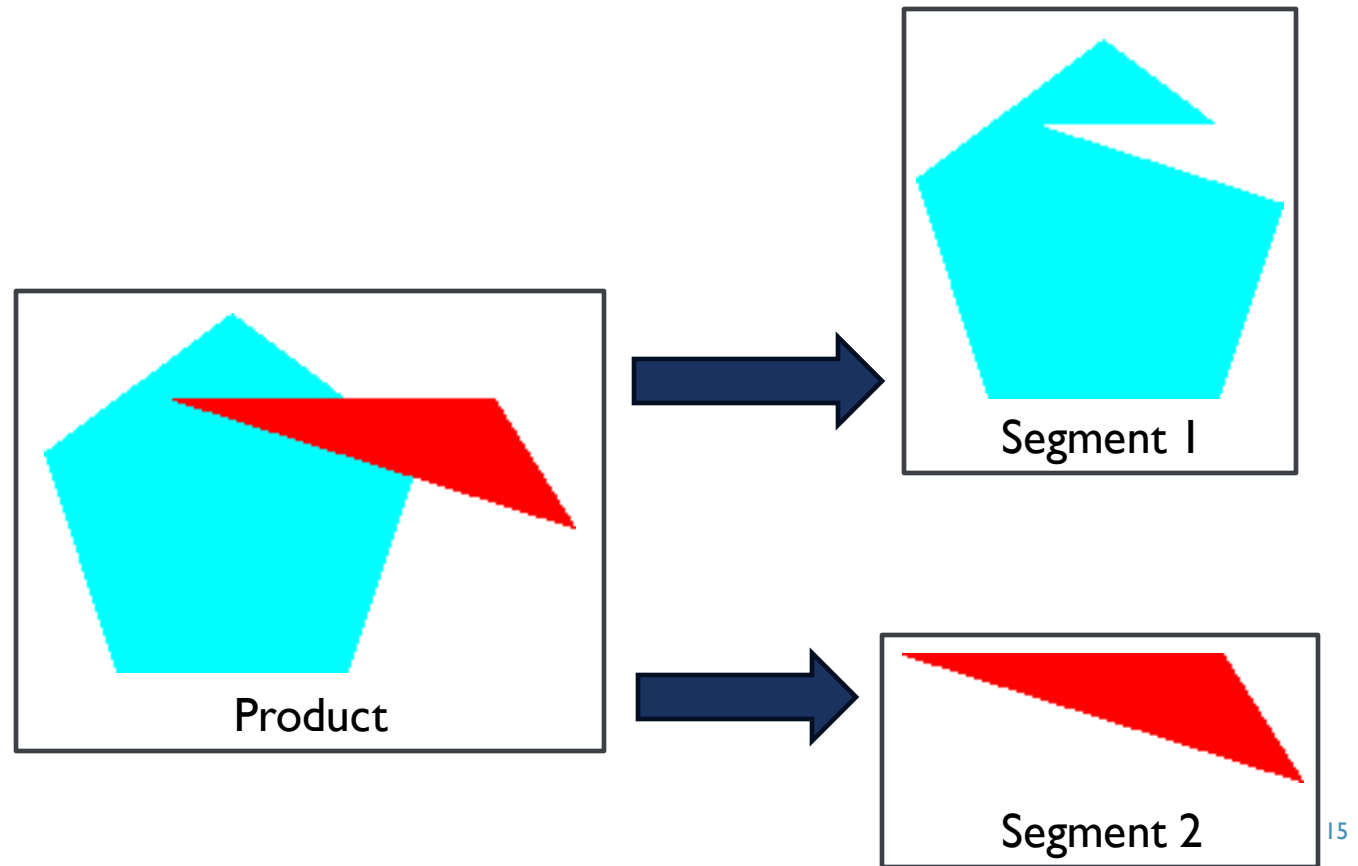
# IMAGE FACTORY AGENT

- Get input parts and desired product




# IMAGE FACTORY AGENT

- Get input parts and desired product
- Segment product image by color to get individual objects



# IMAGE FACTORY AGENT


- Get input parts and desired product
- Segment product image by color
- Since there's no color changing, pair same-colored inputs and segments together






# IMAGE FACTORY AGENT

- Get input parts and desired product
- Segment product image by color
- Pair inputs and segments
- Compare input and segment in a pair to get difference in angles and height & weight ratios



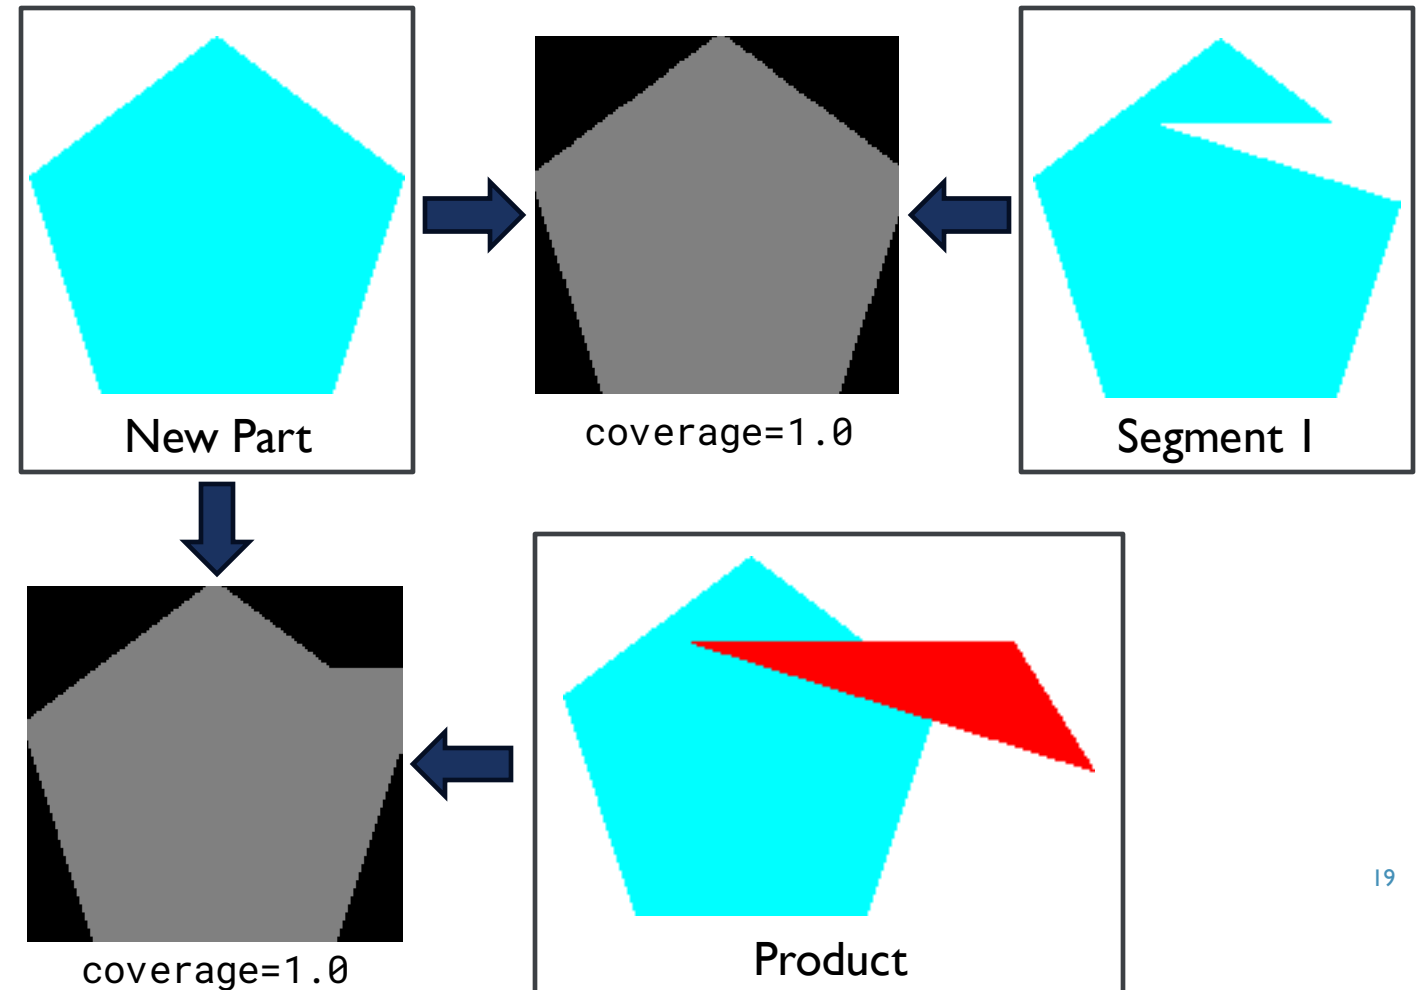
# IMAGE FACTORY AGENT

- Get input parts and desired product
- Segment product image by color
- Pair inputs and segments
- Compare input and segment
- Attempt a scaling or rotation operation on input part, guided by results of comparison




# IMAGE FACTORY AGENT

- Get input parts and desired product
- Segment product image by color
- Pair inputs and segments
- Compare input and segment
- Attempt a scaling or rotation operation
- Compare results with segment and product to detect matches
  - Accounts for and detects obscuration, which it stores to remember which part goes on top in final product



# IMAGE FACTORY AGENT

- Get input parts and desired product
- Segment product image by color
- Pair inputs and segments
- Compare input and segment
- Attempt a scaling or rotation operation
- Detect matches
- Build factory design by following chain of operations which led from an input part to a matching part






# RESULTS & DISCUSSION

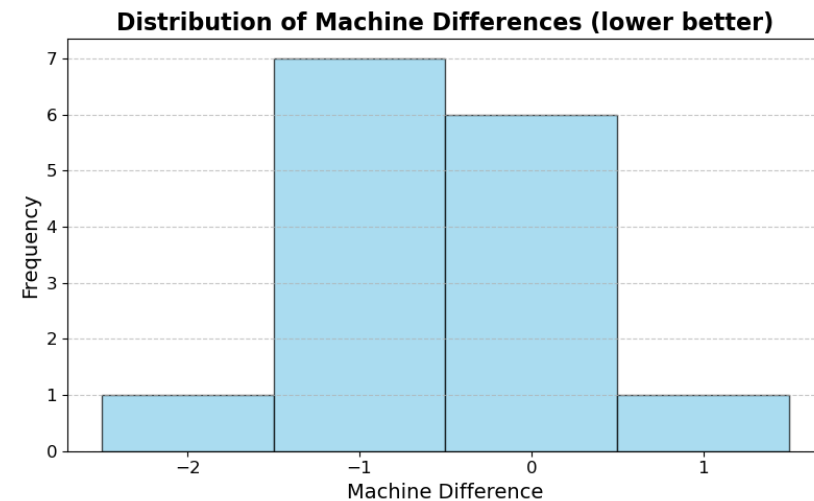
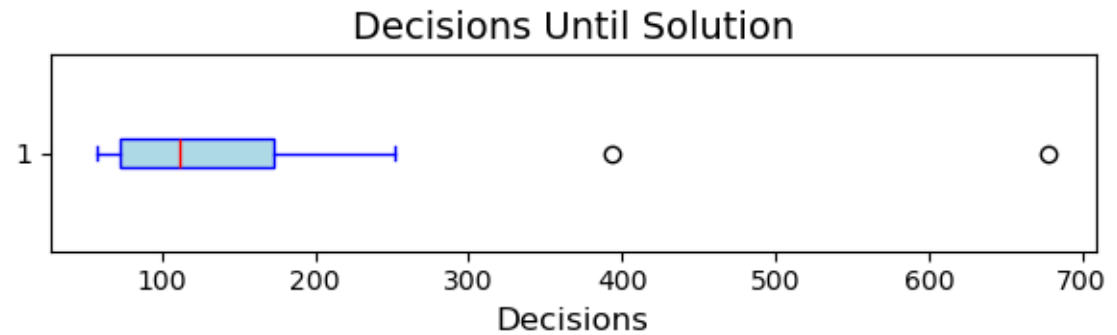
# INITIAL RESULTS

- So far, tested on 16 randomly-generated factories
- 10 agent-generated images were visually identical to goal images
- 5 were within a small margin of error (unnoticeable to humans)
- 1 factory couldn't be completed in reasonable timeframe



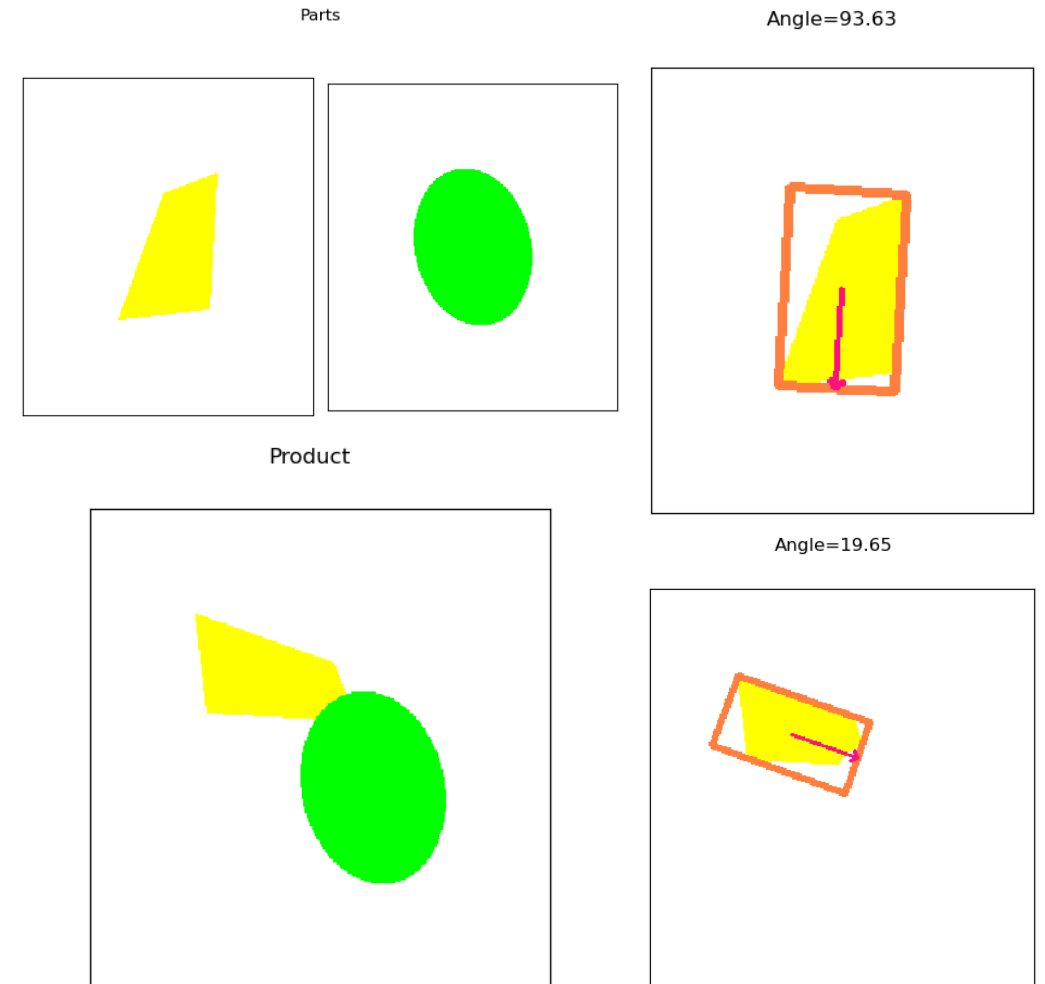
# INITIAL RESULTS

- 15/16 factories were completed in a reasonable amount of time
  - Min decisions: 58
  - Max decisions: 678
  - Longest wall time was 35s
  - Incomplete factory took >2000 decisions
- Agent solutions were almost entirely (14/15) as good or better than original factory
  - Only 1 added an unnecessary machine



# ANALYSIS: FAILURE CASE

- Incomplete factory took >2000 decisions
  - Ended up crashing before finding solution (too much memory usage?)
- Possible reason: factory has an obscured, irregular polygon
- Problem: search is guided by:
  - bounding box height & width ratios for scaling
  - minimal rectangle angle for rotation
- Bounding box ratio thrown off by obscuration
- Angle thrown off by obscuration + irregularity
  - OpenCV minimal rectangles are finicky!
- Result: resorts to exhaustive search







# FUTURE WORK & CONCLUSION



# FUTURE WORK: TWO DIRECTIONS

## Better Soar Code

- Significant room for optimization to prevent slow-down and crashing
- Improve search guidance rules by:
  - adjusting heuristic weights
  - baking in more geometric knowledge
  - probably lots of other little tricks

## Improve Object Representations

- ~~Add more information to hand-crafted representation~~
  - Too painstaking and too fragile!
- Learn object representations via deep learning
  - regress affine transform between input and product part representations
  - predict complete product part from obscured segment

# NUGGETS AND COAL

## Nuggets

- Mostly successful proof-of-concept for visual reasoning with SVS2
- Extends Soar to include images and image manipulation, might be helpful to some
- Hopefully enough to finish my dissertation!?

## Coal

- Hand-crafted object representations are too fragile
- Fragility bleeds over into Soar code: no way to handle noisiness/fuzziness
  - e.g., slight pixel misalignments can throw everything off
- Bonus coal: using PyTorch in Soar is tricky due to namespace collision